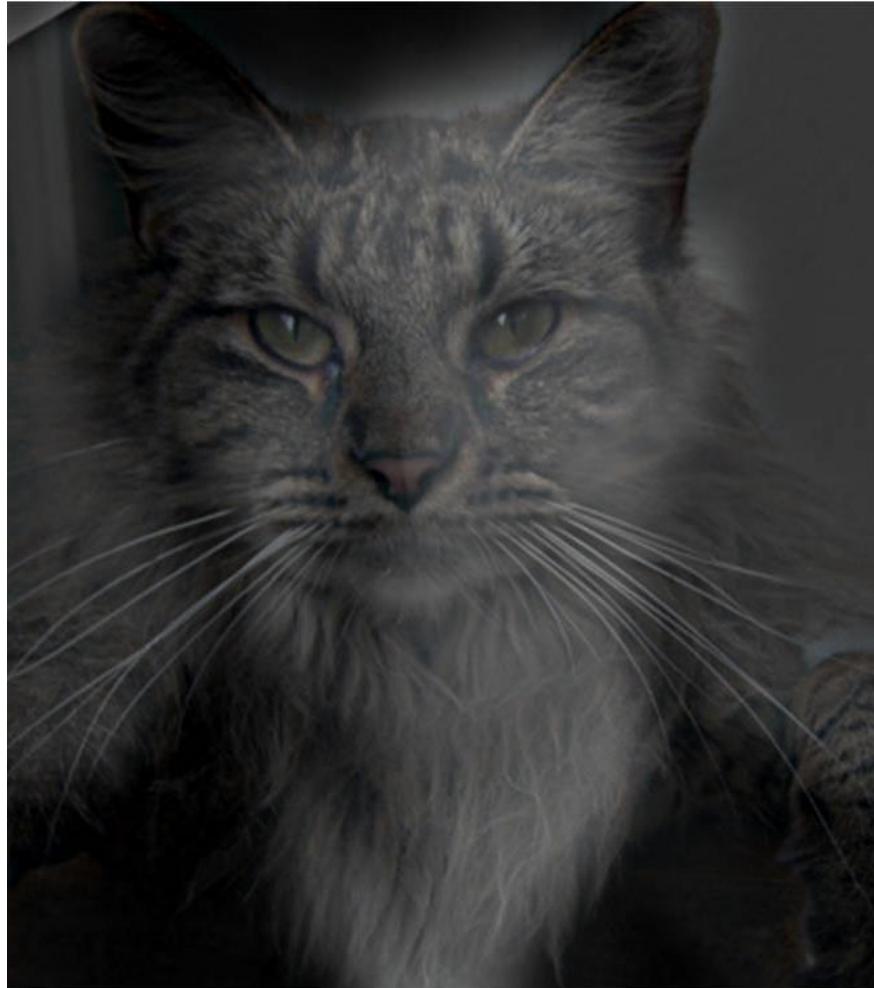


Thinking in Frequency



Computer Vision

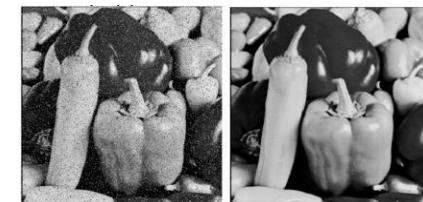
A U G Sankararao, IIIT Sri City

Previous class: Image Filtering

- Linear filtering is sum of dot product at each position
 - Can smooth, sharpen, translate (among many other uses)
- Gaussian filters
 - Low pass filters, separability, variance
- Attend to details:
 - filter size
- Noise models and nonlinear image filters



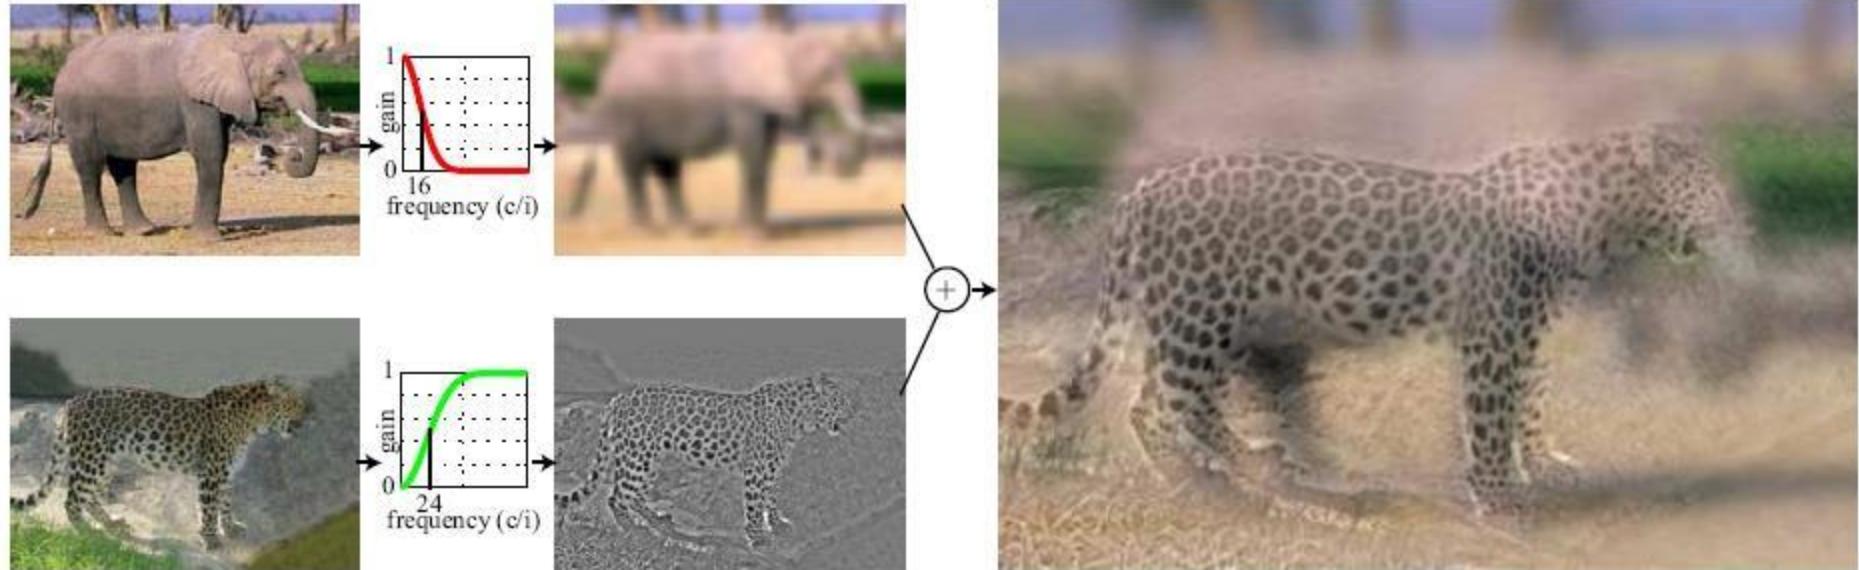
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Today's class

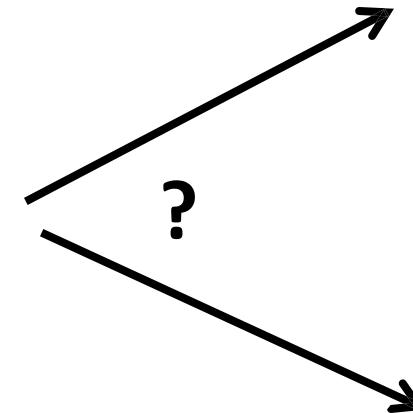
- Fourier transform and frequency domain
- Frequency view of filtering
- Image downsizing and interpolation

Hybrid Images



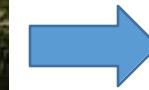
- A. Oliva, A. Torralba, P.G. Schyns,
“Hybrid Images,” SIGGRAPH 2006

Why do we get different, distance-dependent interpretations of hybrid images?



Why does a lower resolution image still make sense to us?

What do we lose?



Thinking in terms of frequency

Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

Any periodic function can be rewritten as a weighted sum of sines and cosines of different frequencies.



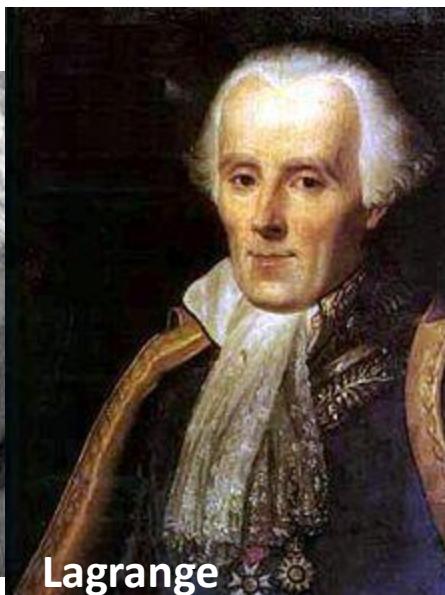
Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

Any periodic function can be rewritten as a weighted sum of sines and cosines of different frequencies.

- Don't believe it?
 - Neither did Lagrange, Laplace, Poisson and other big wigs
 - Not translated into English until 1878!

Laplace



Legendre

Lagrange

Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

Any periodic function can be rewritten as a weighted sum of sines and cosines of different frequencies.

But it's (mostly) true!

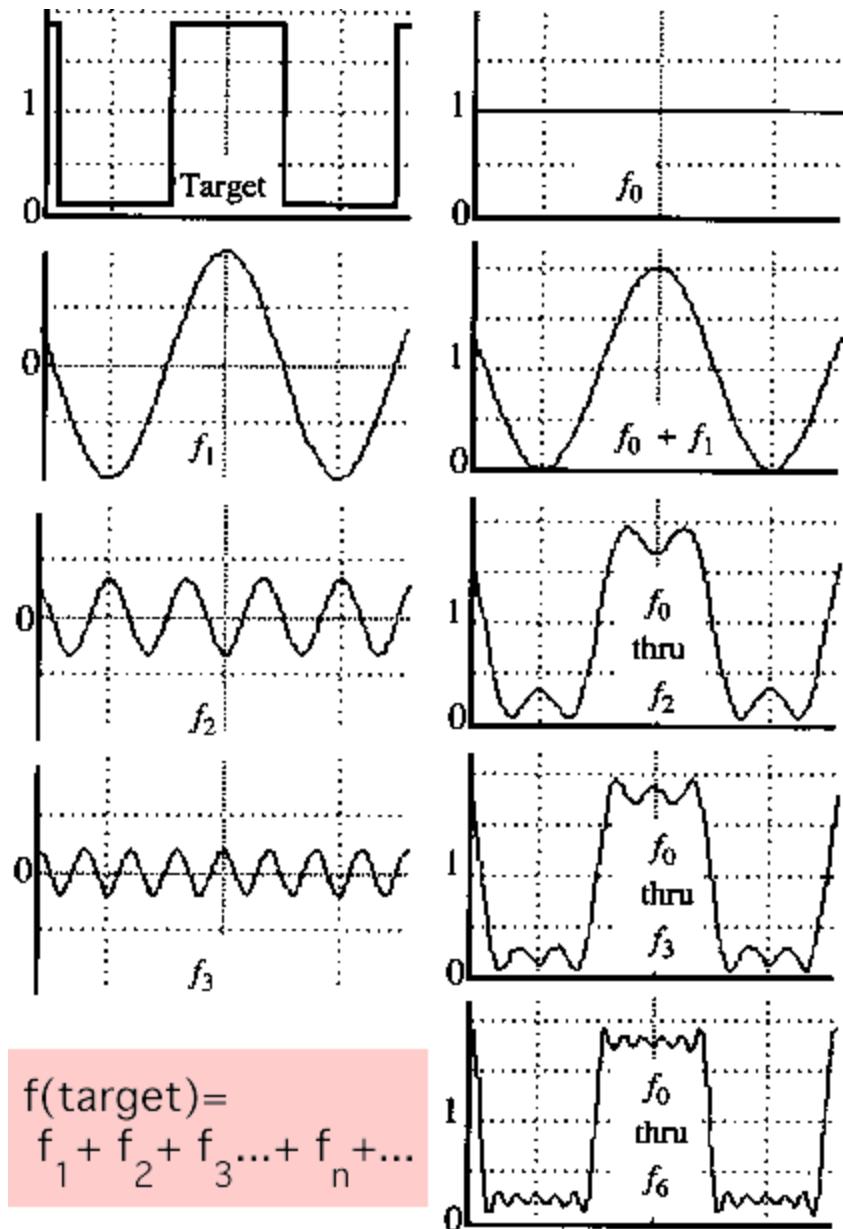
- called Fourier Series

A sum of sines

Our building block:

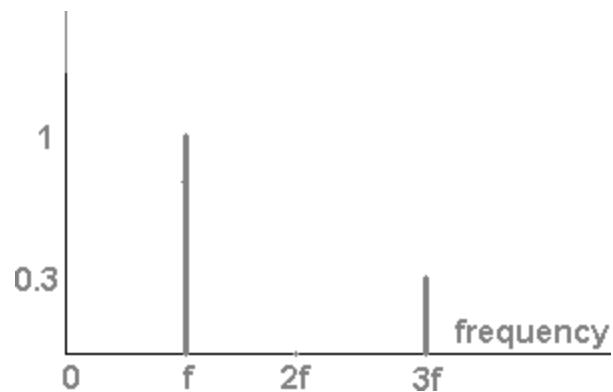
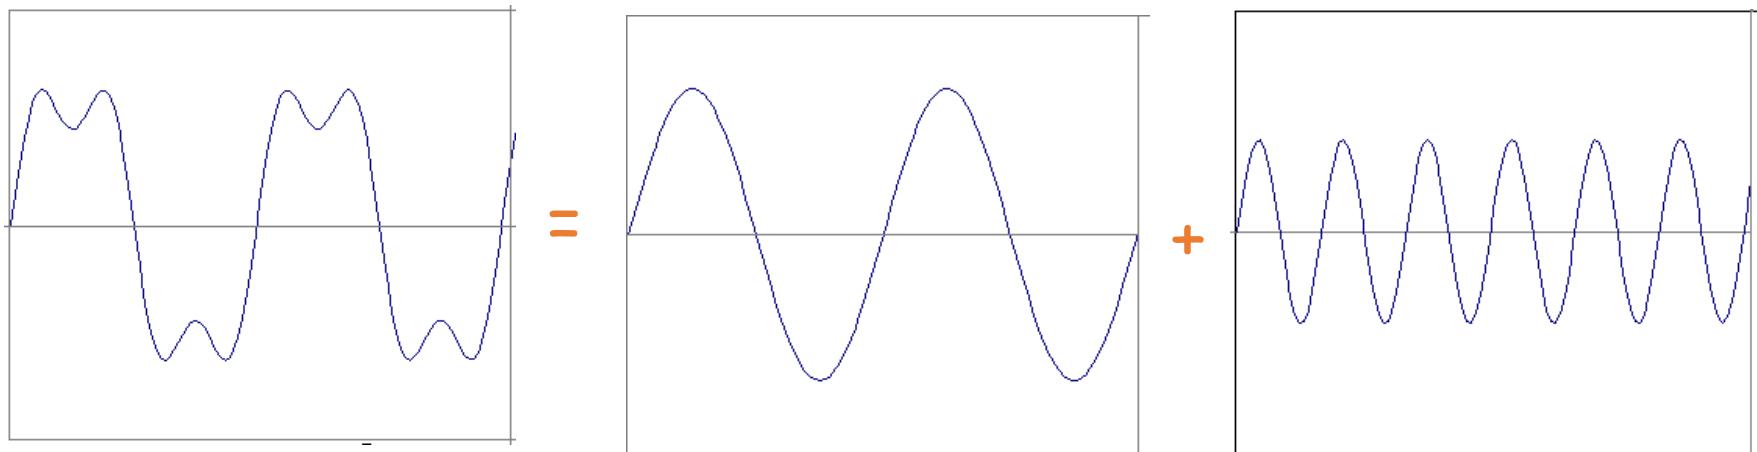
$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal $f(x)$ you want!

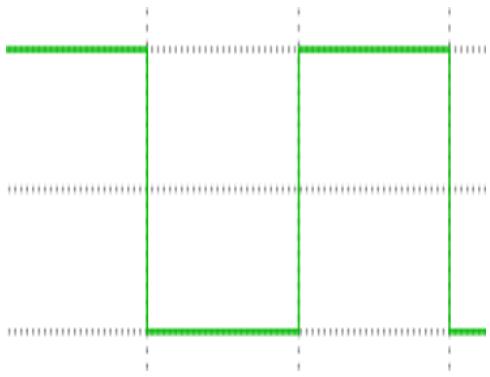


Frequency Spectra

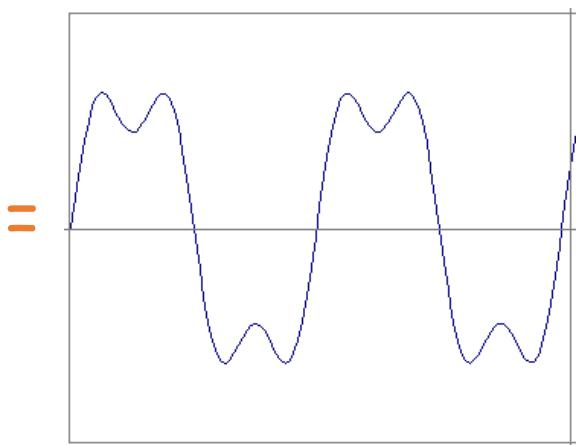
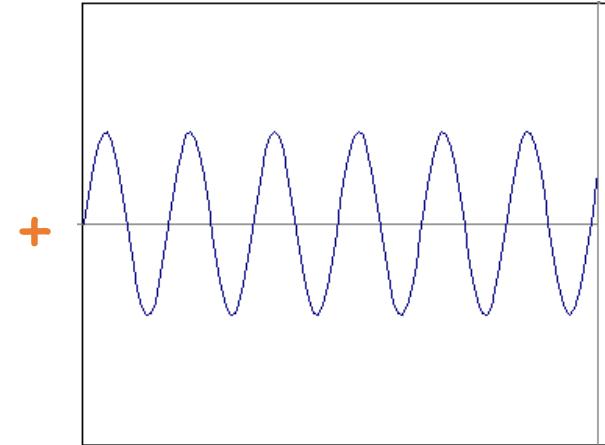
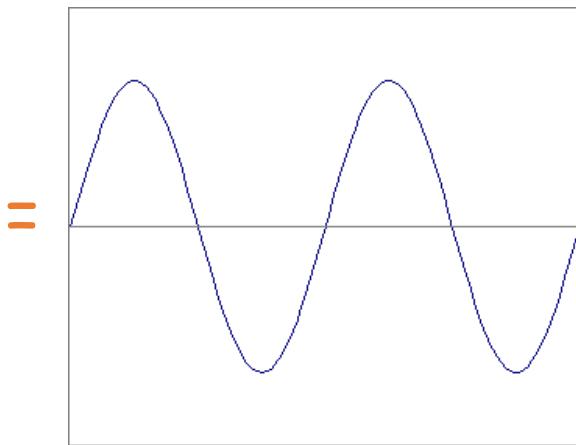
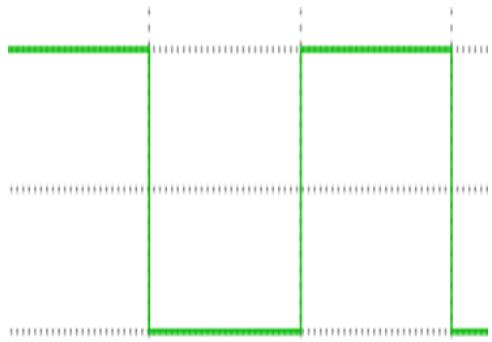
- example : $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$



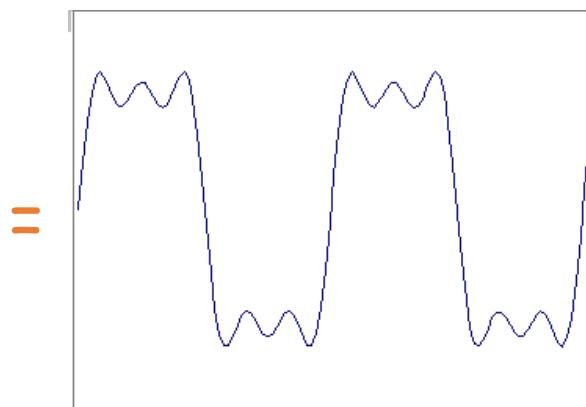
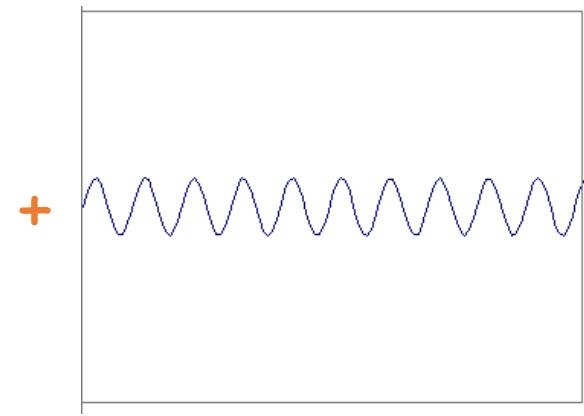
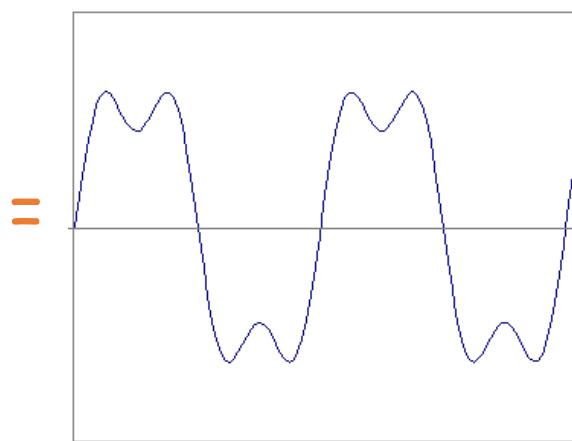
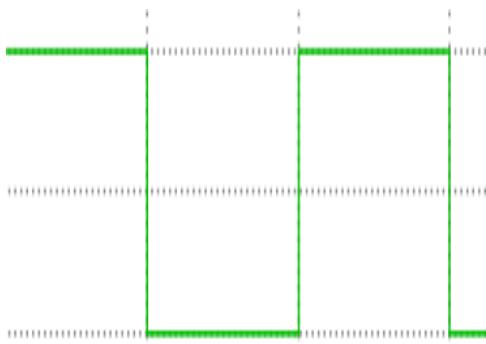
Frequency Spectra



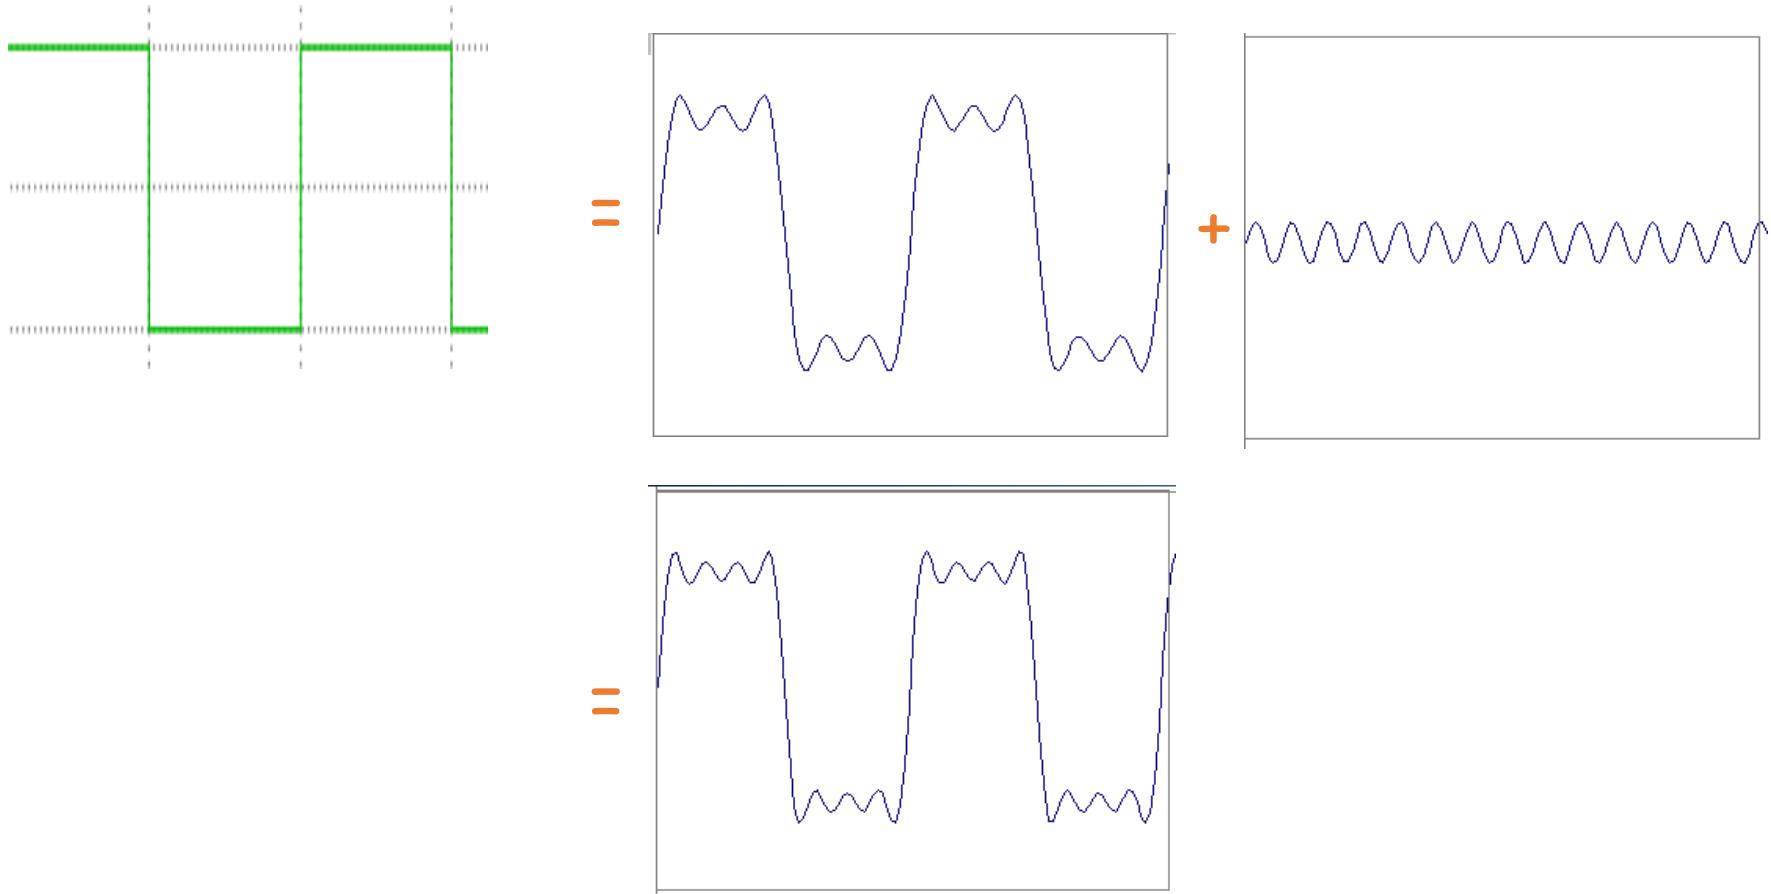
Frequency Spectra



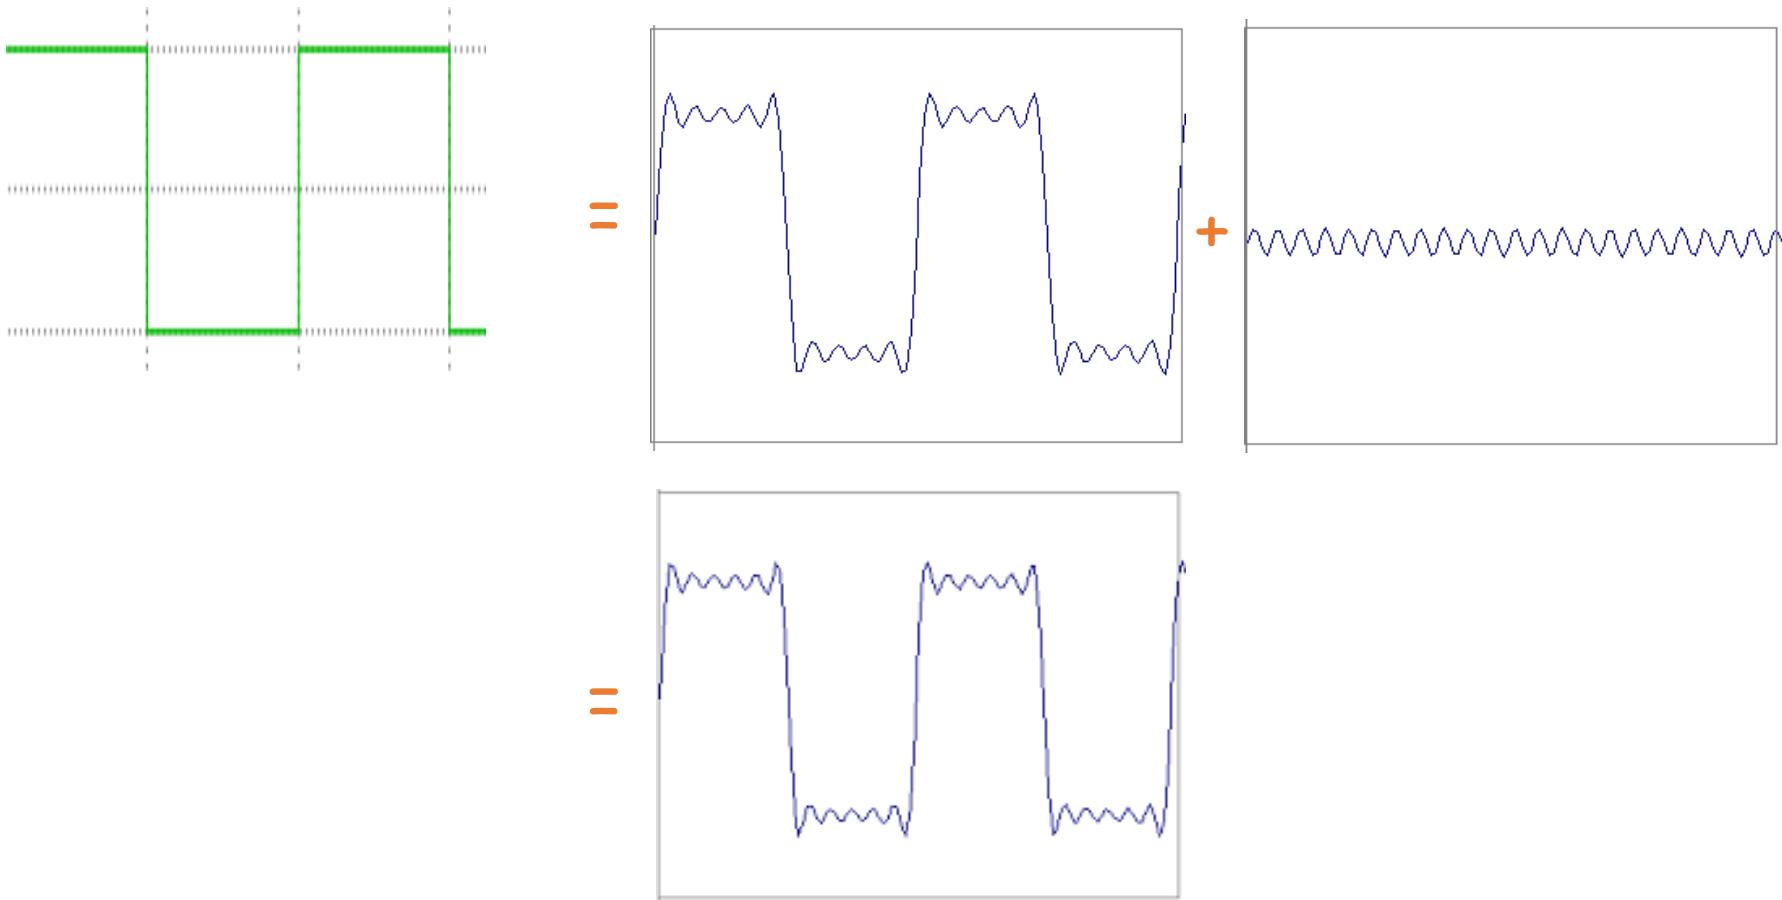
Frequency Spectra



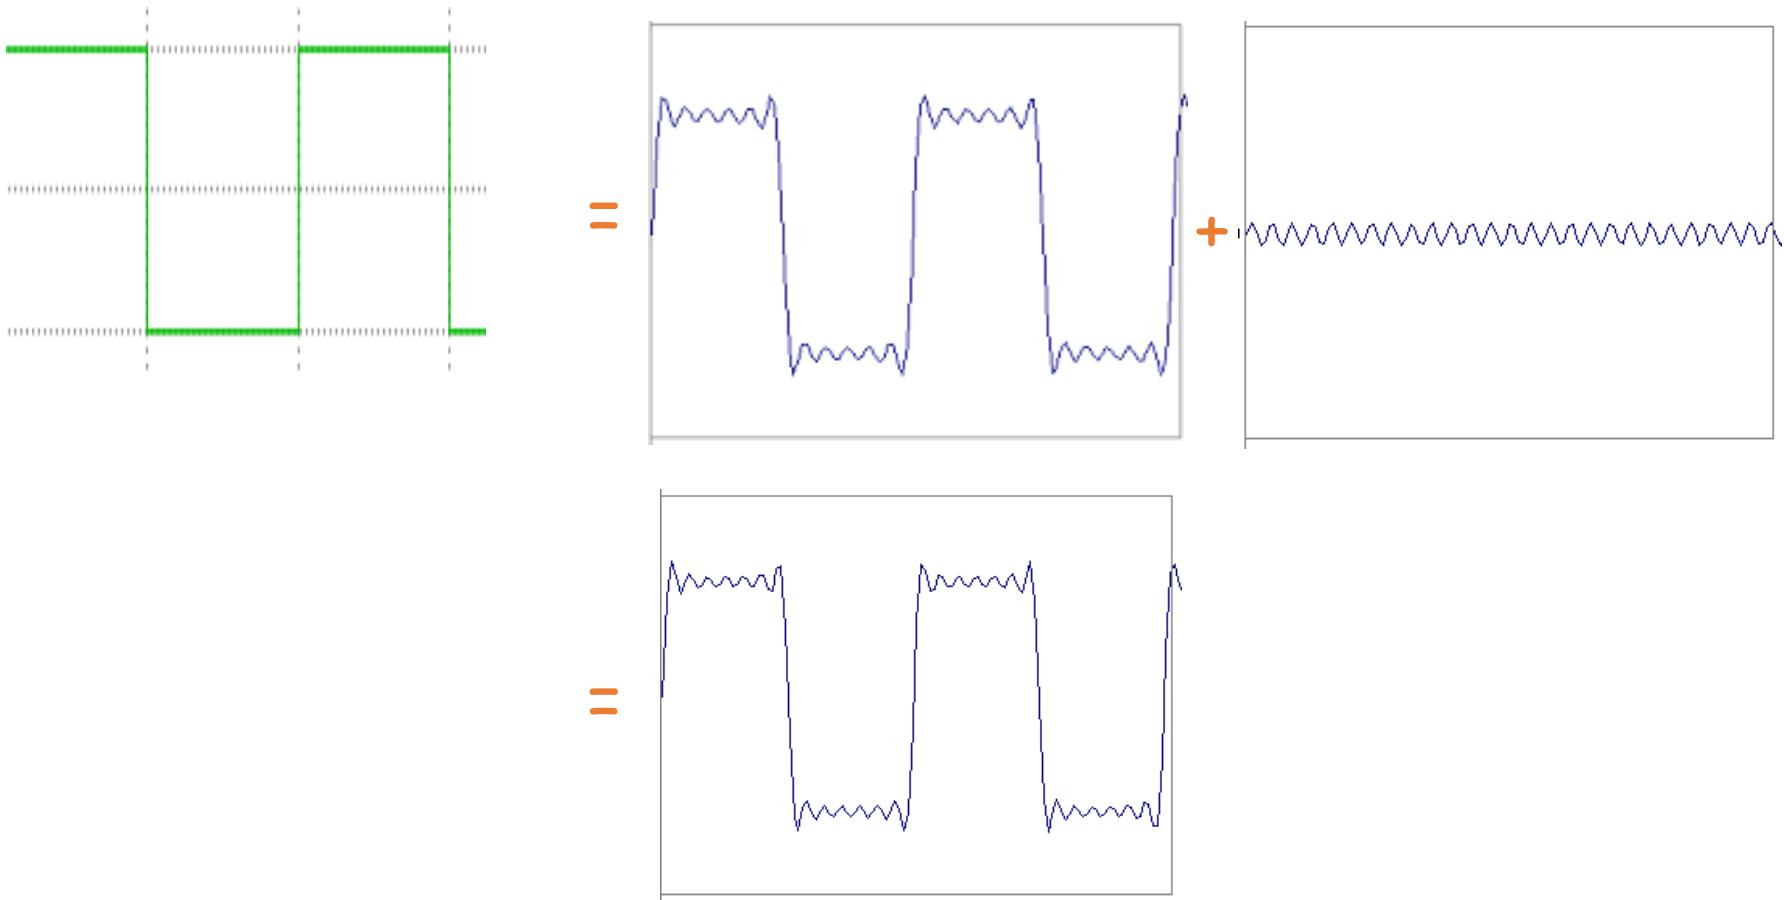
Frequency Spectra



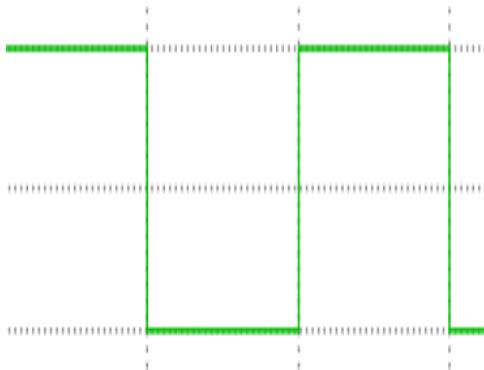
Frequency Spectra



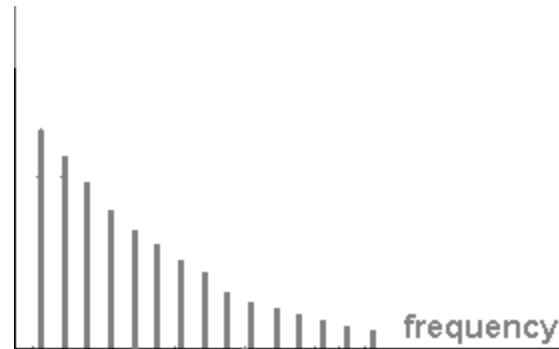
Frequency Spectra



Frequency Spectra



$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



Fourier Transform

We want to understand the frequency ω of our signal. So, let's reparametrize the signal by ω instead of x :



For every ω from 0 to inf, (actually $-\inf$ to \inf), $F(\omega)$ holds the amplitude A and phase ϕ of the corresponding sine

- How can F hold both? Complex number trick!

Fourier Transform

We want to understand the frequency ω of our signal. So, let's reparametrize the signal by ω instead of x :



For every ω from 0 to inf, (actually $-\inf$ to \inf), $F(\omega)$ holds the amplitude A and phase ϕ of the corresponding sine

- How can F hold both? Complex number trick!

$$A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$F(\omega) = R(\omega) + iI(\omega)$$

Even *Odd*

$$\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

Fourier Transform

We want to understand the frequency ω of our signal. So, let's reparametrize the signal by ω instead of x :



For every ω from 0 to inf, (actually $-\inf$ to \inf), $F(\omega)$ holds the amplitude A and phase ϕ of the corresponding sine

- How can F hold both? Complex number trick!

$$A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$F(\omega) = R(\omega) + iI(\omega)$$

Even *Odd*

$$\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

And we can go back:



Fourier Transform

- FT stores the magnitude and phase at each frequency
 - Magnitude encodes how much signal is there at a particular frequency
 - Phase encodes **spatial information** (indirectly)

$$\text{Amplitude: } A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$\text{Phase: } \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

Discrete Fourier Transform

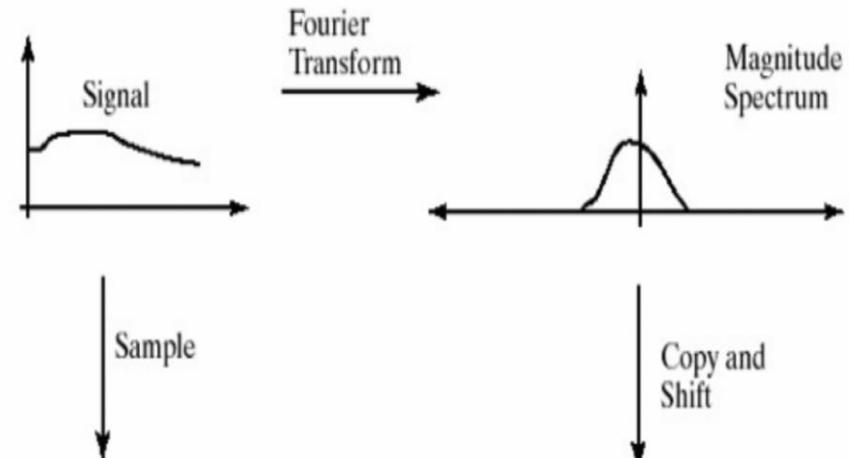
$$F(k) = \sum_{x=0}^{x=N-1} f(x) e^{-i \frac{2\pi k x}{N}} \quad k = -N/2..N/2$$

$$e^{ik} = \cos k + i \sin k \quad i = \sqrt{-1}$$

Continuous vs Discrete FT

- **Continuous Fourier transform (FT):**

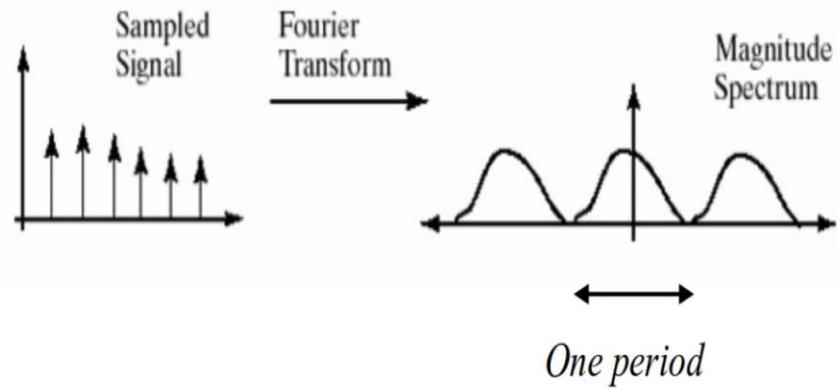
$$H(\omega) = \int_{-\infty}^{\infty} h(x)e^{-j\omega x} dx$$



- **Discrete Fourier Transform (DFT):**

$$H(\omega) = \sum_0^{N-1} h(x)e^{-j\frac{2\pi k}{N}} ;$$

where N is the length of the sampled signal.



Credit: Elgammal, Rutgers University

Discrete Fourier Transform

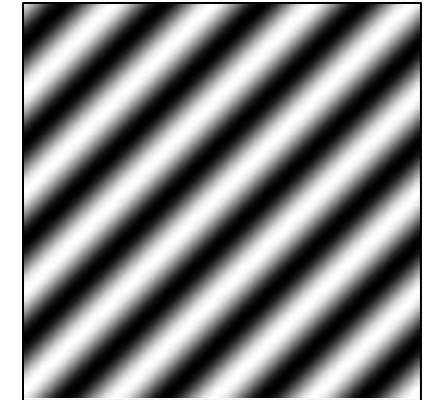
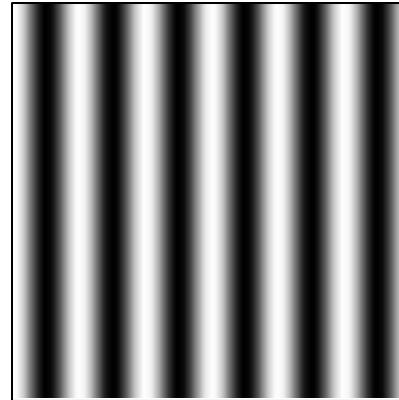
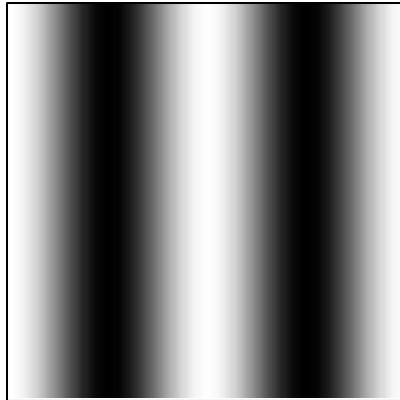
$$F(k) = \sum_{x=0}^{N-1} f(x) e^{-i \frac{2\pi k x}{N}} \quad k = -N/2..N/2$$

$$e^{ik} = \cos k + i \sin k \quad i = \sqrt{-1}$$

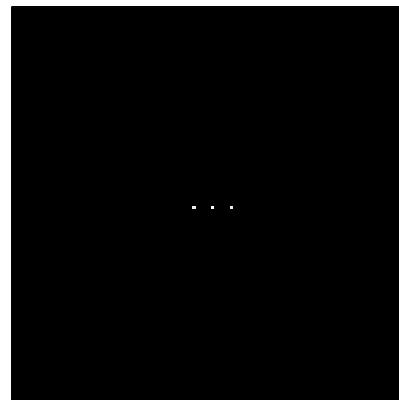
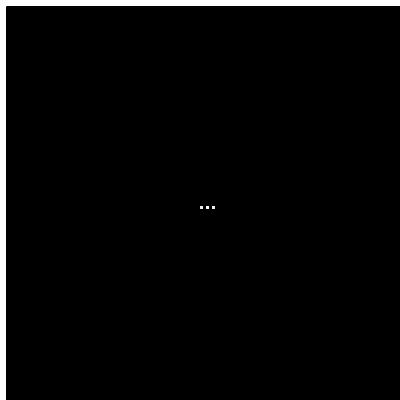
$$F(k_x, k_y) = \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) e^{-i \frac{2\pi (k_x x + k_y y)}{N}}$$

Fourier analysis in images

Intensity
Image

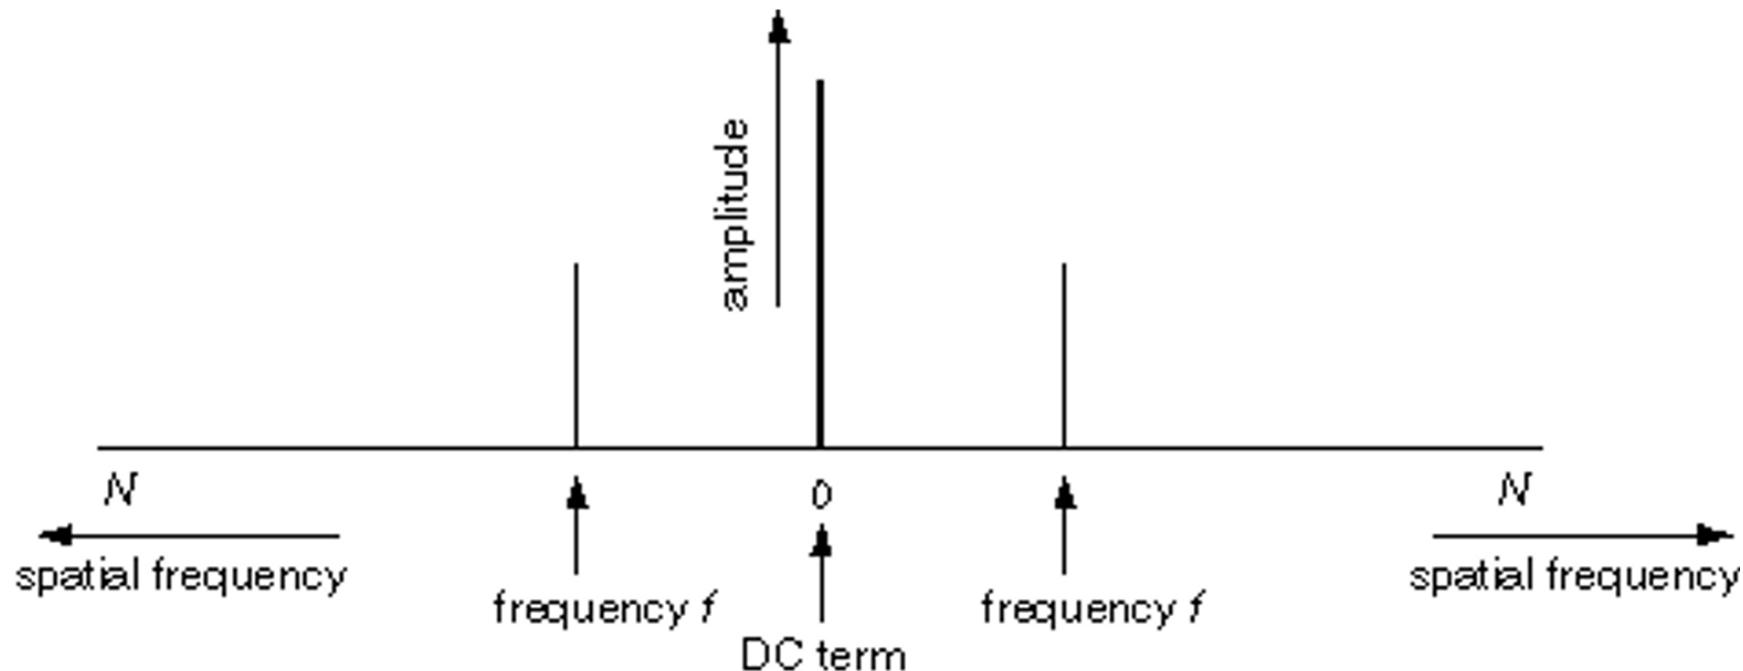


Fourier
Image

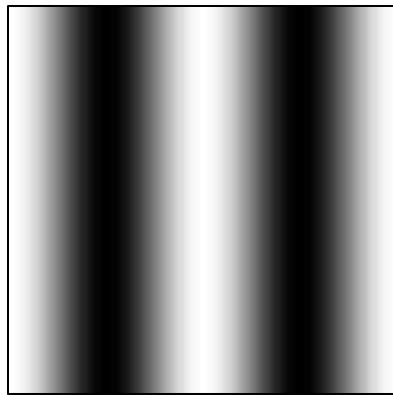


Fourier analysis in images

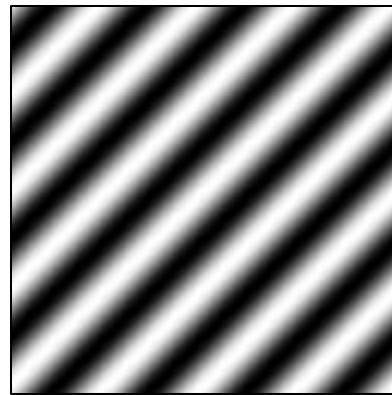
Mirror-image reflections along the axes



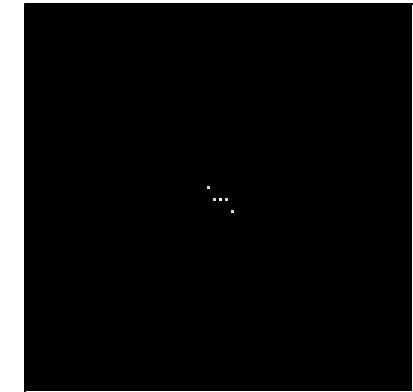
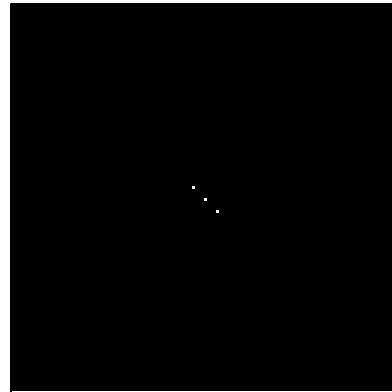
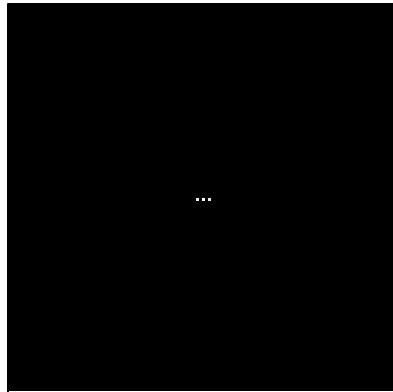
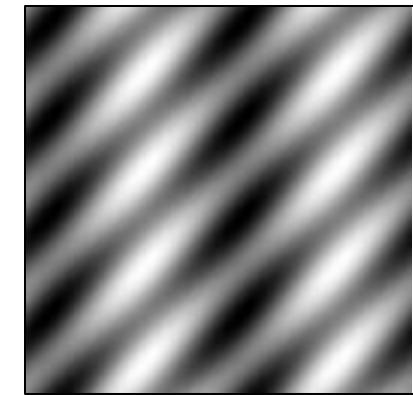
Signals can be composed



+



=



<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>
More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

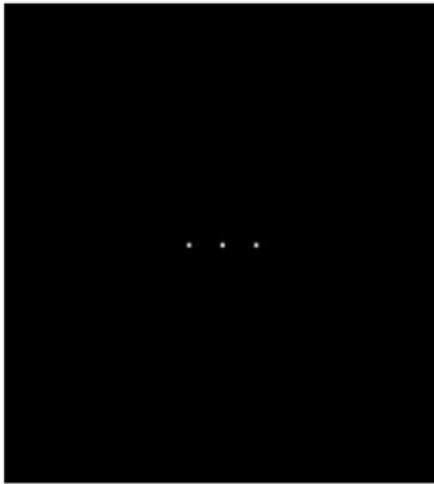
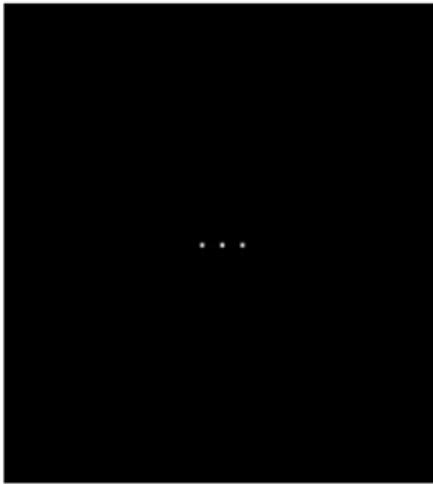
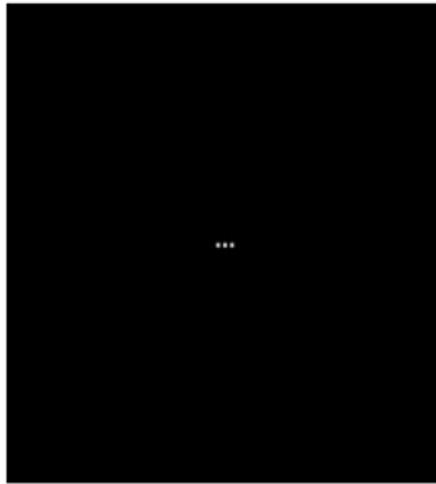
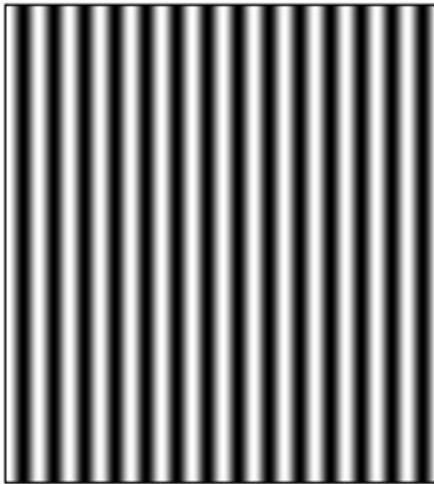
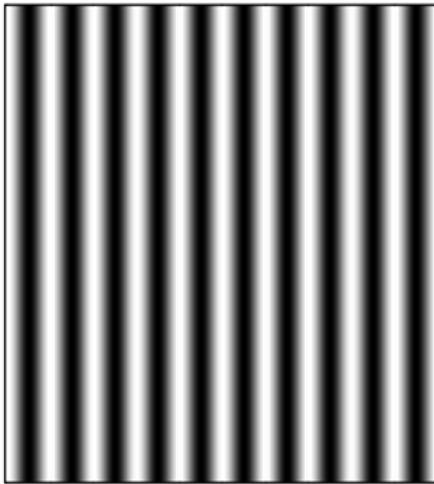
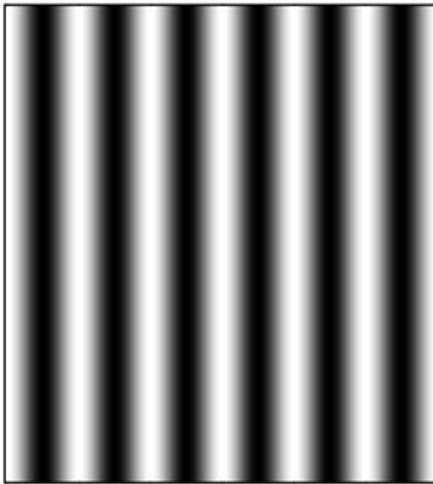
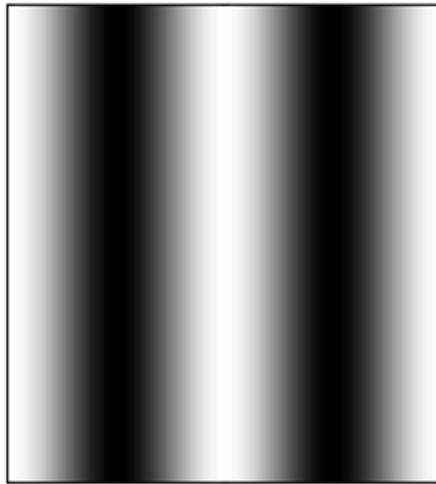
Fourier analysis in images

1

3

5

7



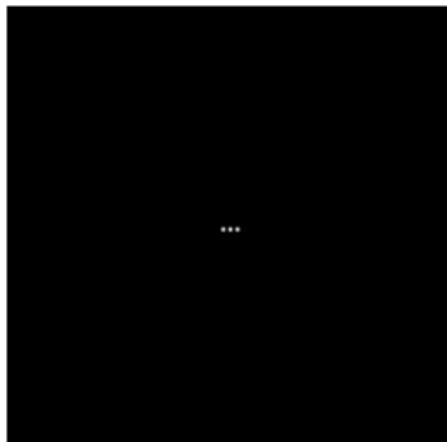
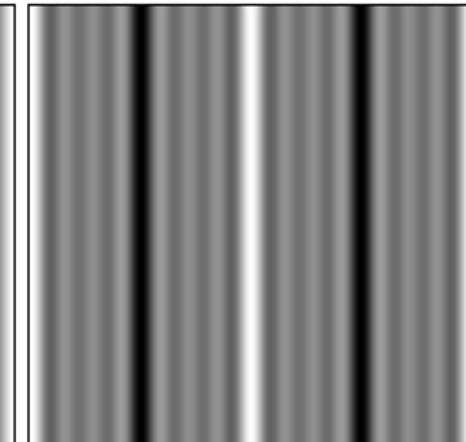
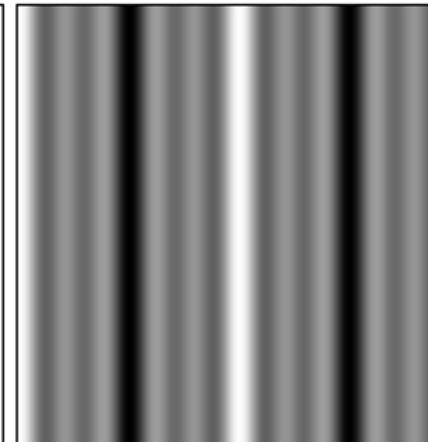
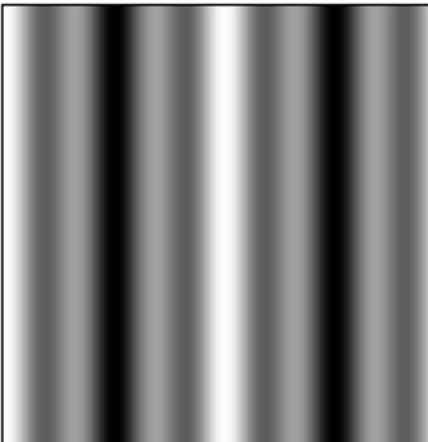
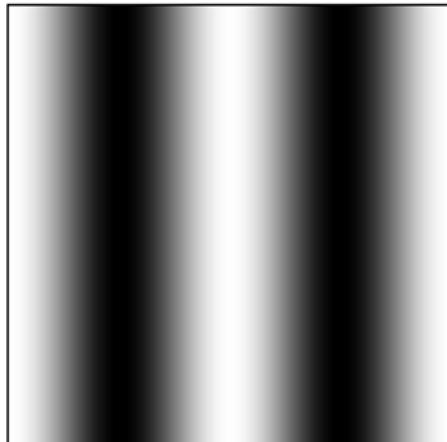
Signals can be composed

1

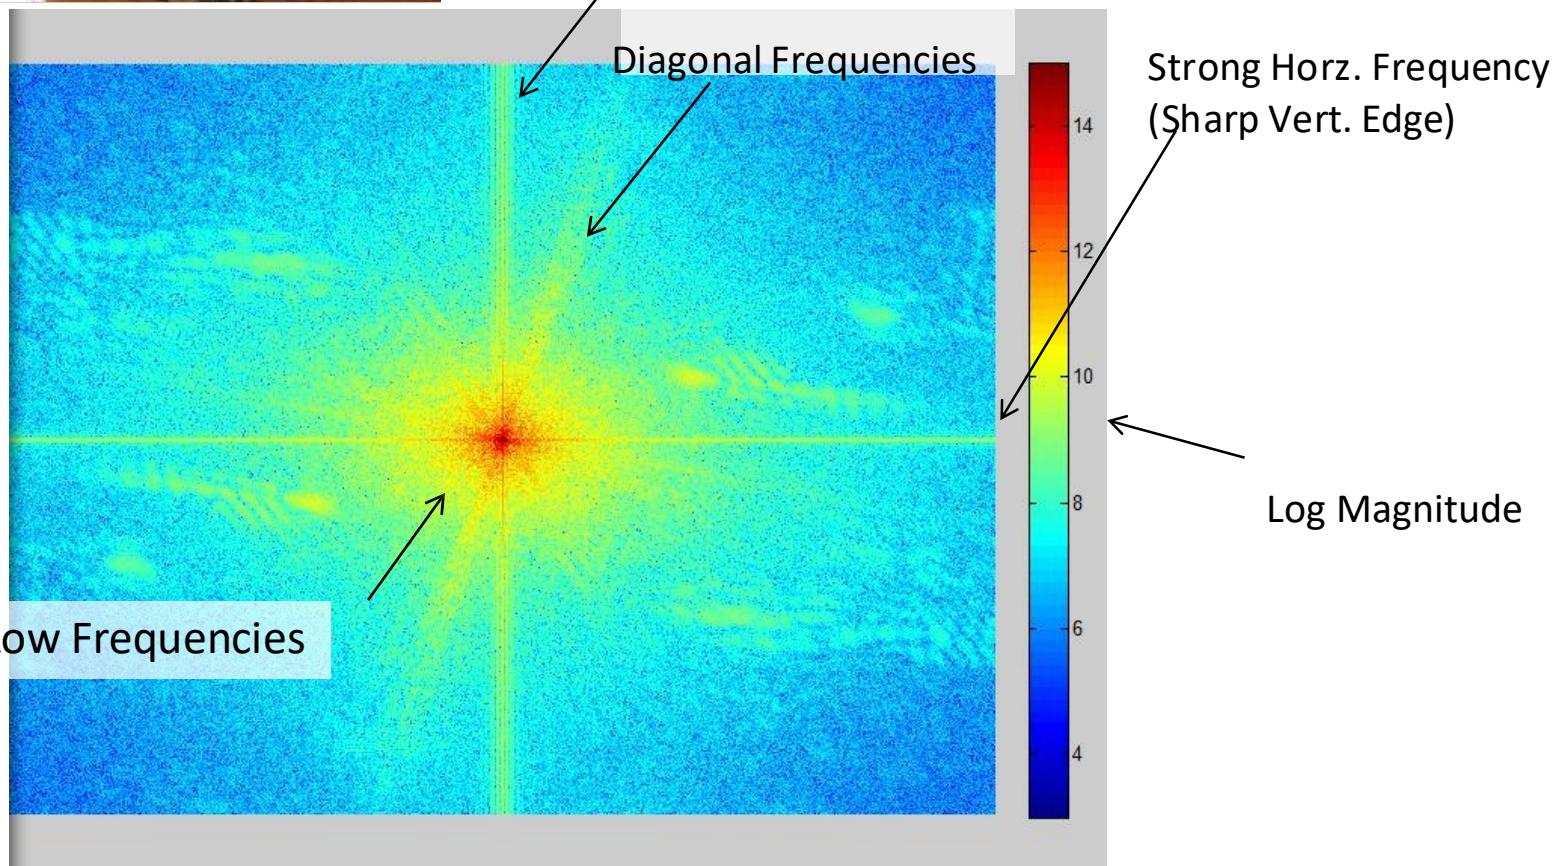
1+3

1+3+5

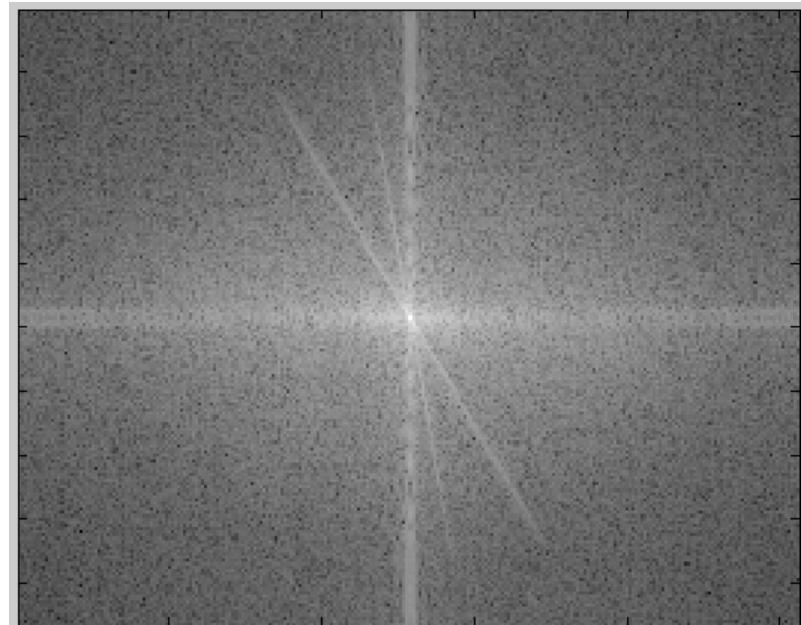
1+3+5+7



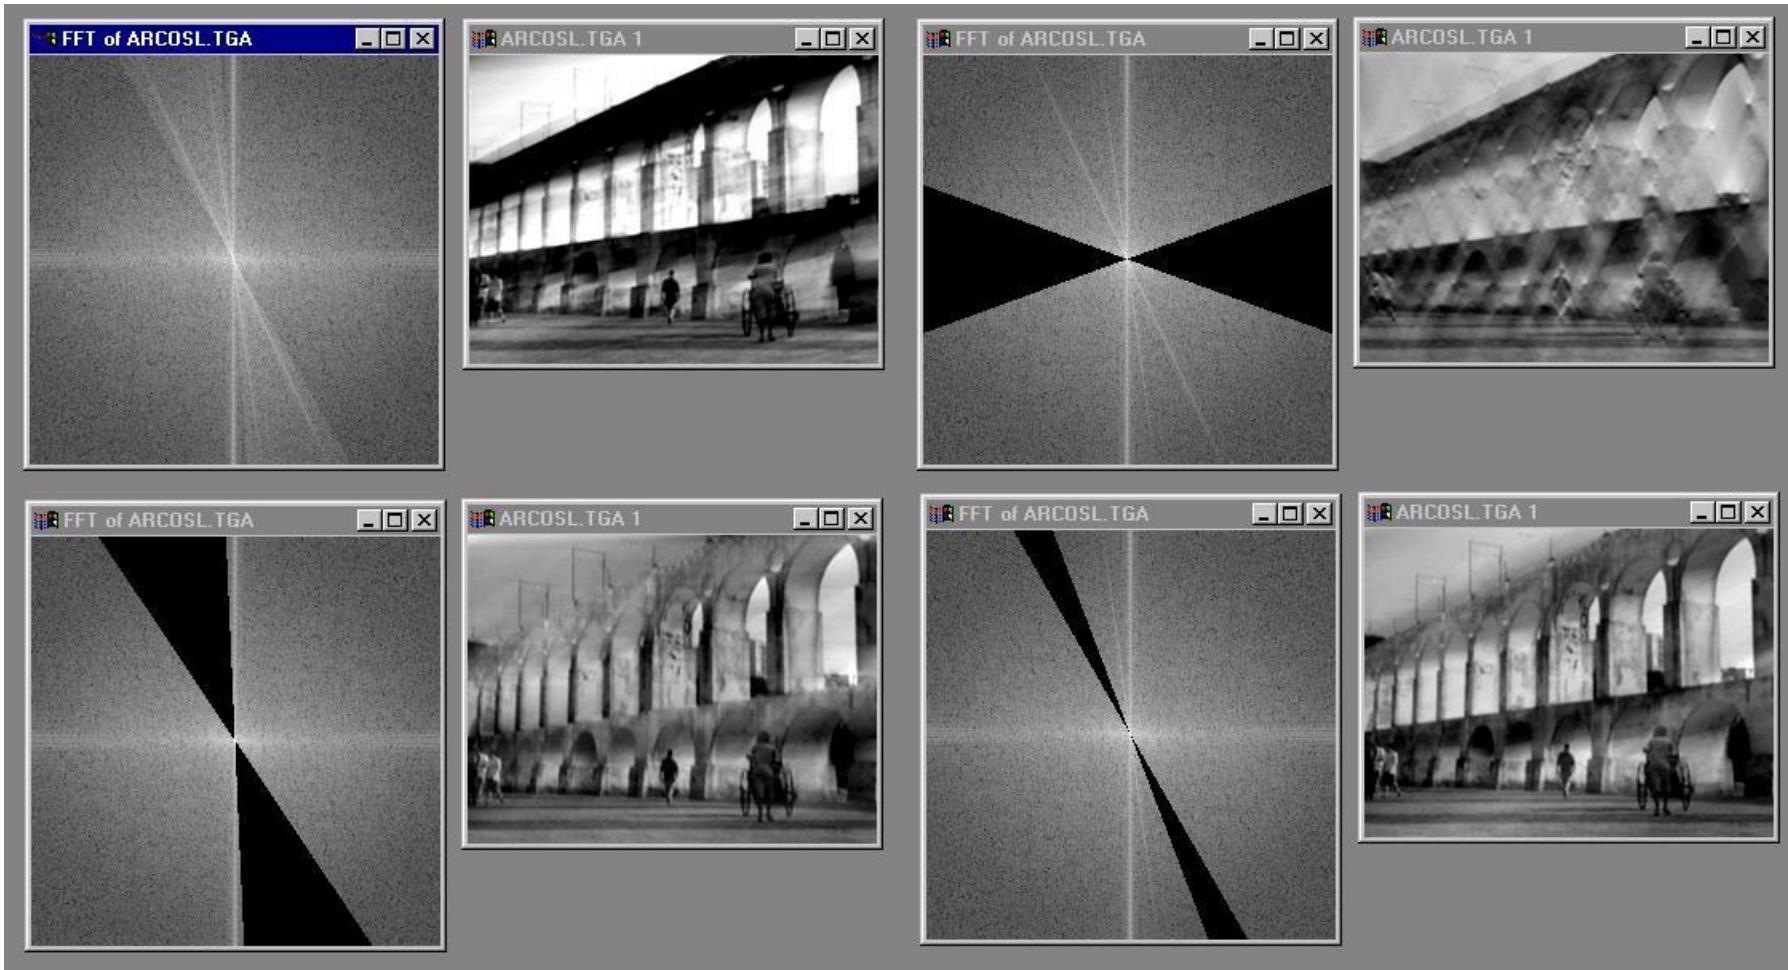
<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>
More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>



Example



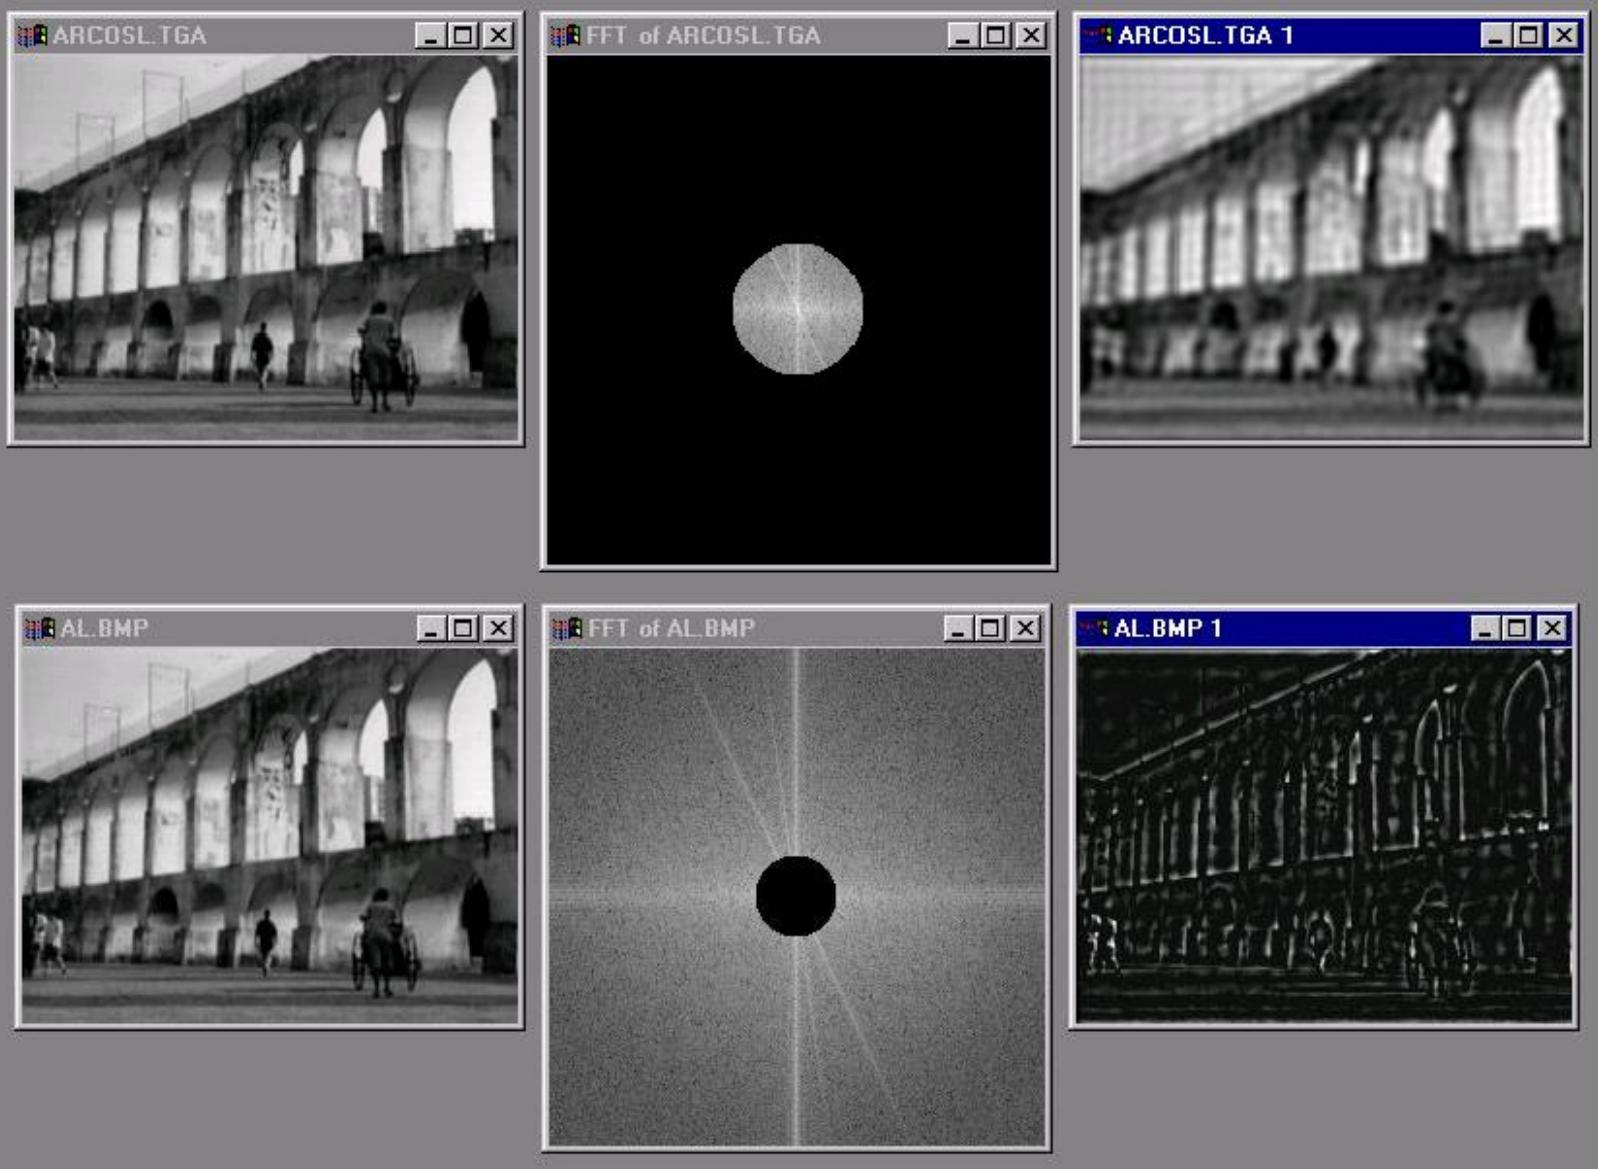
Can change spectrum, then reconstruct



Low and High Pass filters

- Low-Pass Filters: Filters that allow low frequencies to pass through (block high frequencies).
 - Example: Gaussian filter
- High-Pass Filters: Filters that allow high frequencies to pass through (block low frequencies).
 - Example? Edge filter

Low and High Pass filtering



Fast Fourier Transform (FFT)

- Number of arithmetic operations to compute FT of N numbers (i.e., function defined at N points) is proportional to N^2
- Possible to reduce this to $N \log(N)$ using Fast Fourier Transform (FFT)
- FFT is a recursive divide-and-conquer algorithm for computing DFT
- Applications of FFT? Examples: Convolution, correlation

For more, see <https://www.karlsims.com/fft.html>

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

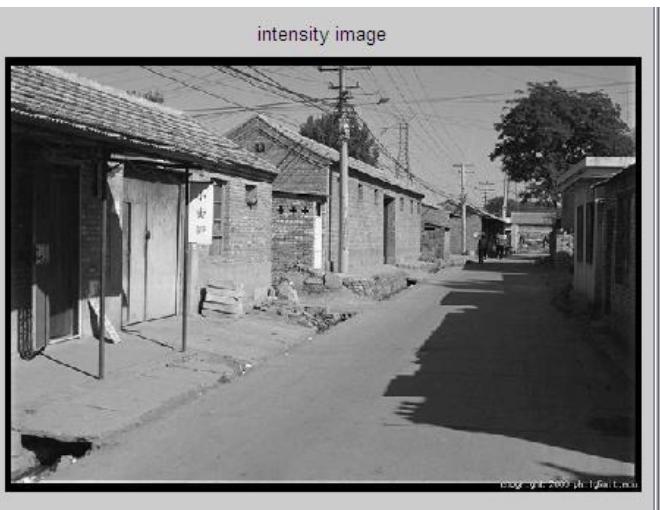
- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

$$\mathcal{F}^{-1}[GH] = \mathcal{F}^{-1}[G] * \mathcal{F}^{-1}[H]$$

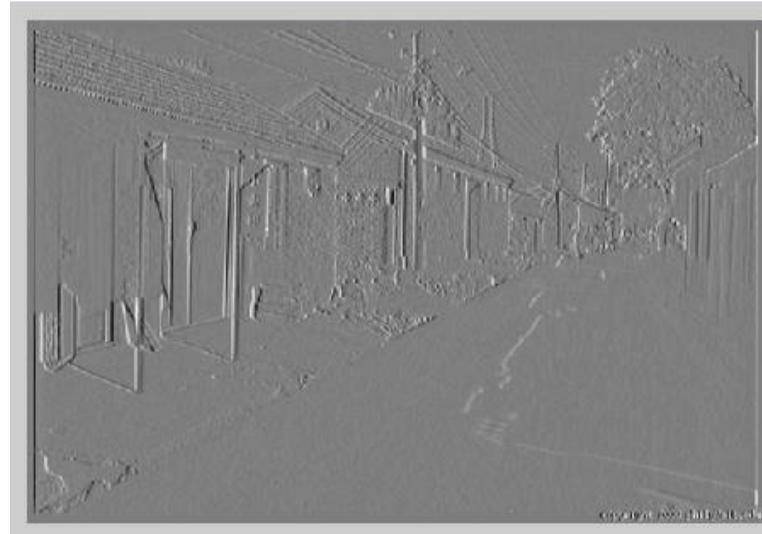
- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

Filtering in spatial domain

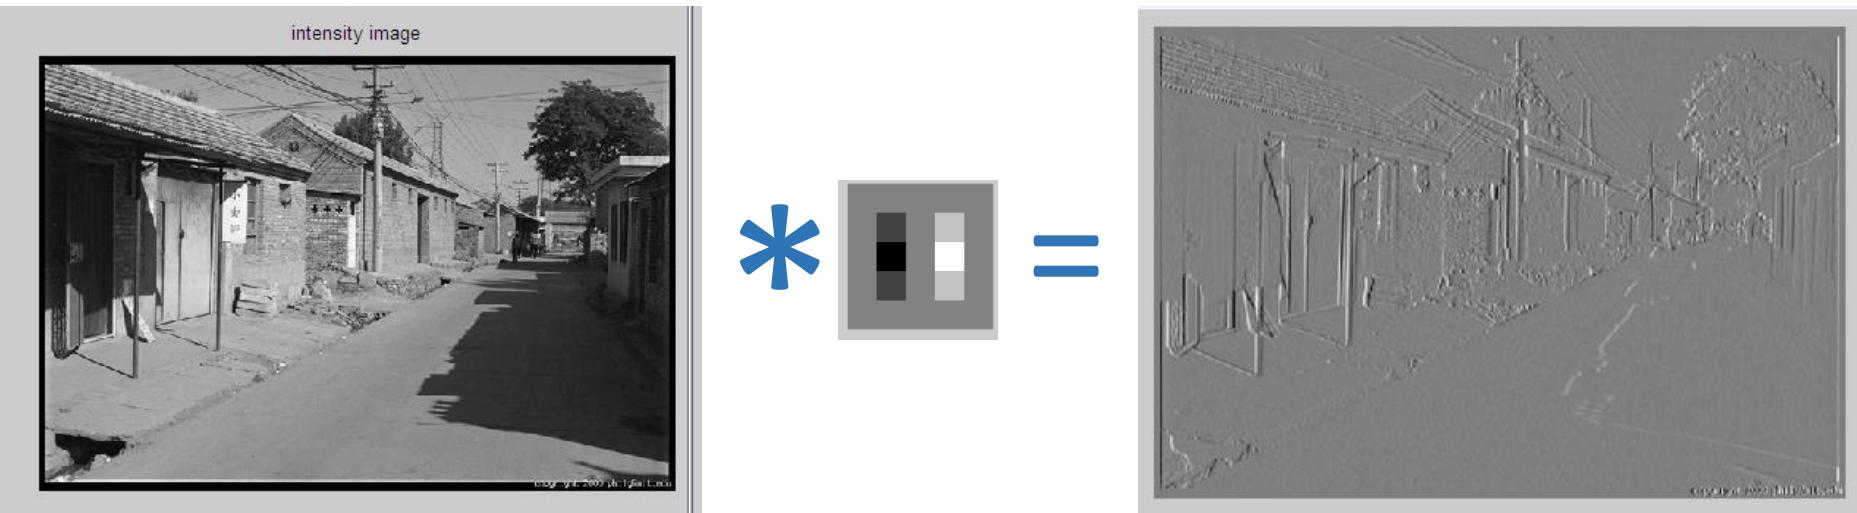
-1	0	1
-2	0	2
-1	0	1



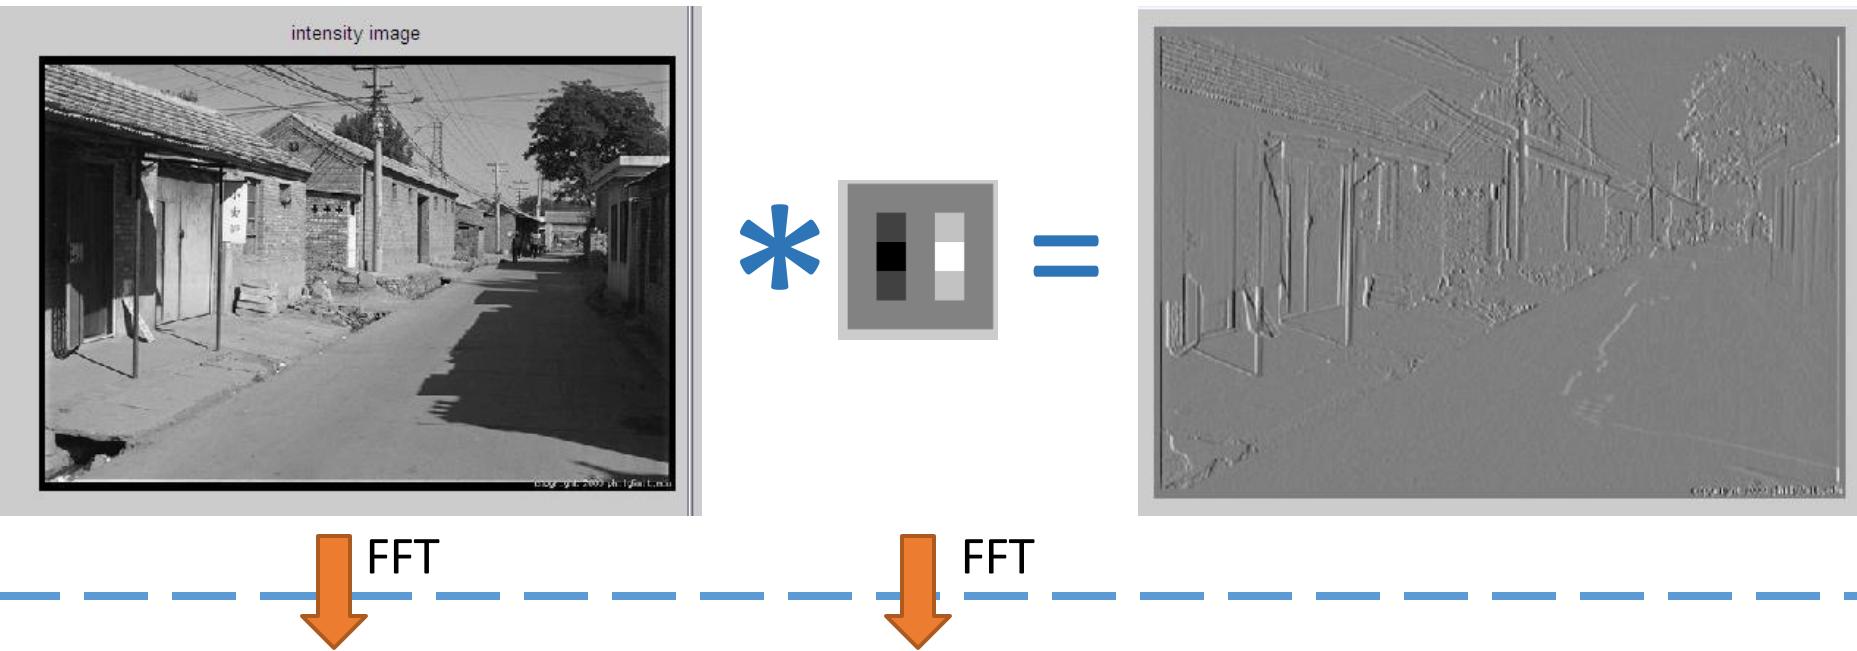
A diagram consisting of three main elements: a blue asterisk (*) on the left, a gray square containing two vertical bars (one black, one white) in the center, and two blue double bars (=) on the right.



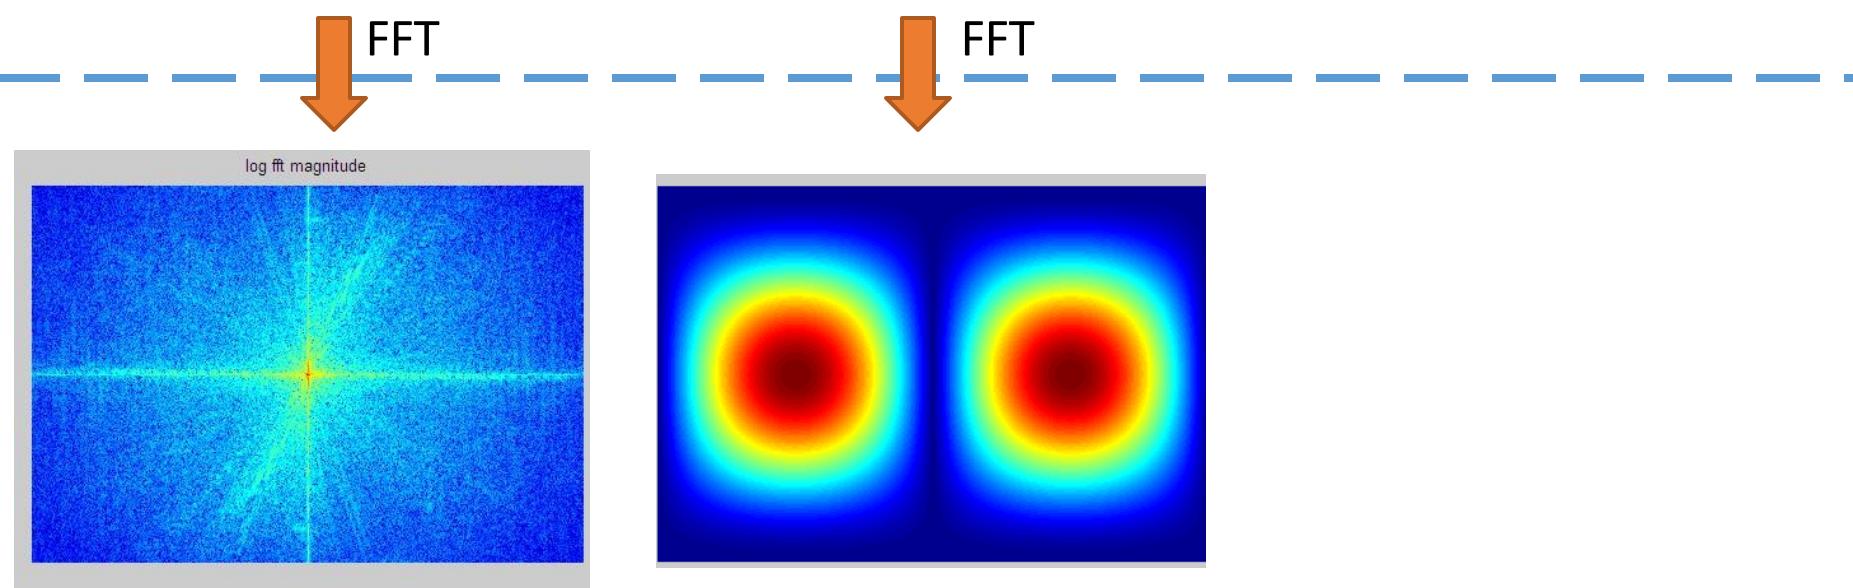
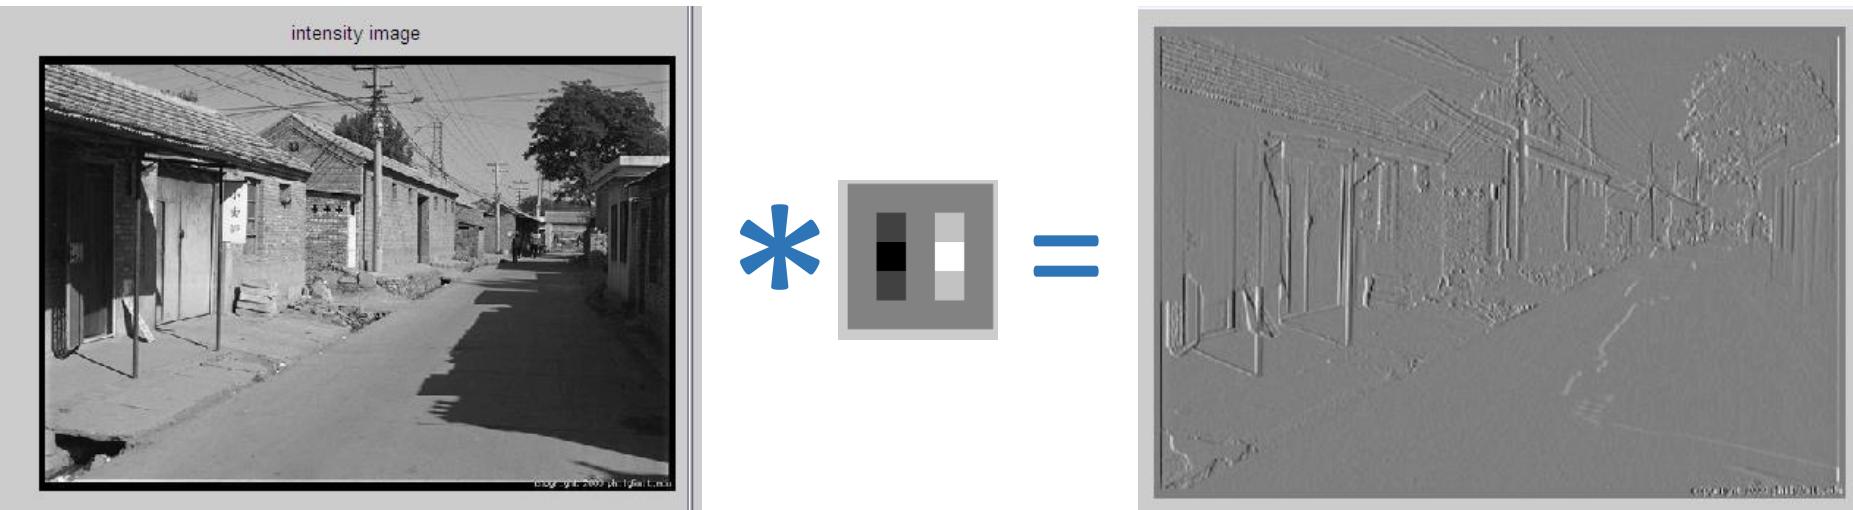
Spatial domain



Spatial domain

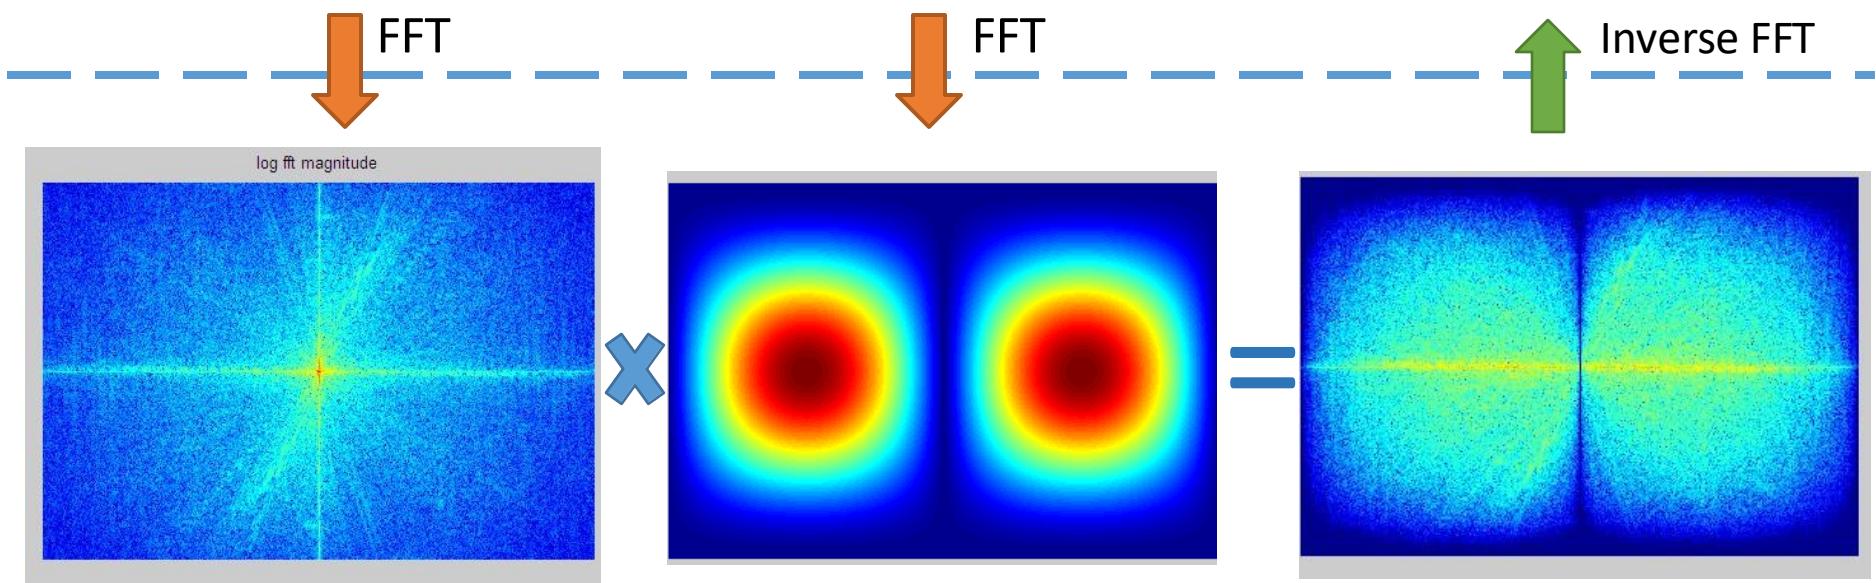
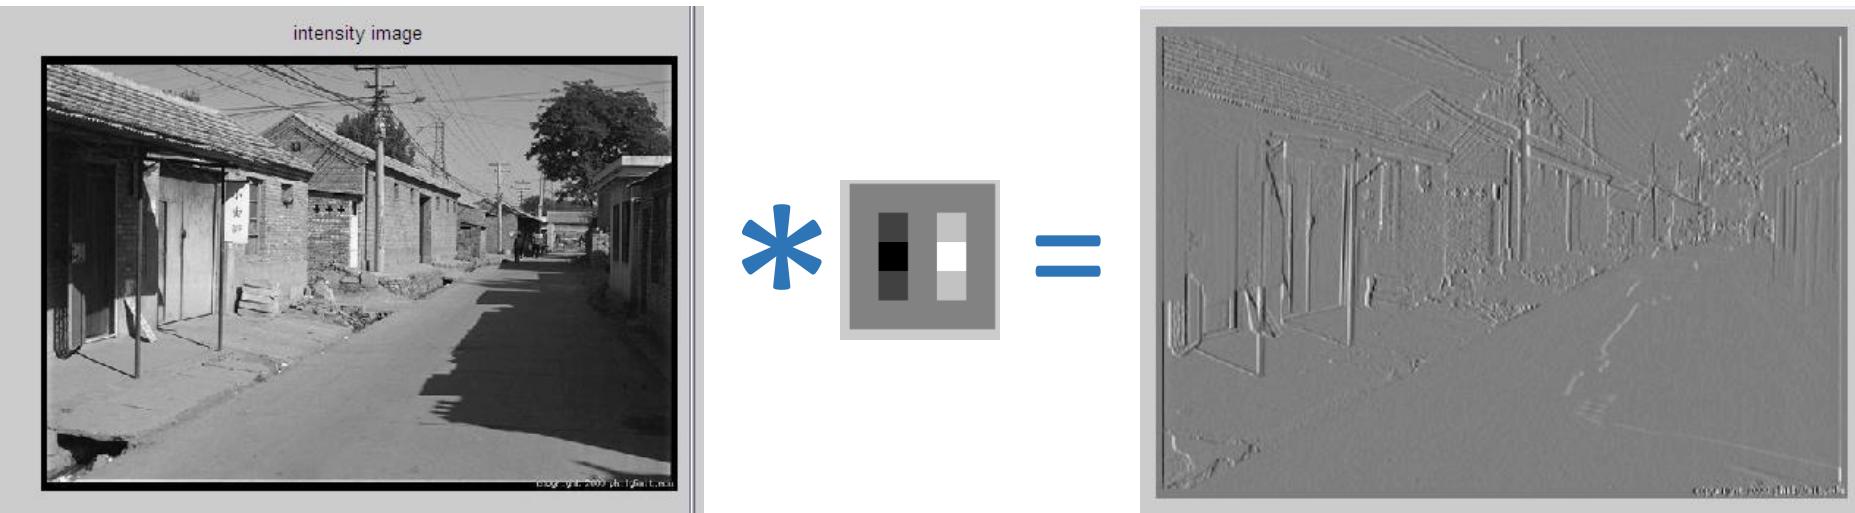


Spatial domain



Frequency domain

Spatial domain



Frequency domain

Questions

Which has more information, the phase or the magnitude?

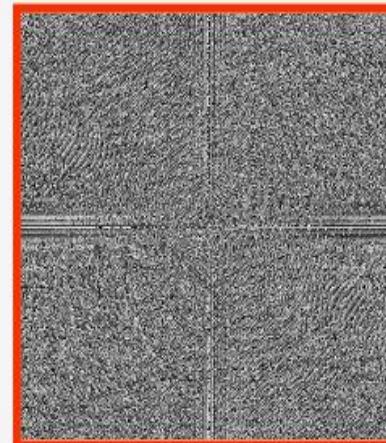
What happens if you take the phase from one image and combine it with the magnitude from another image?

Phase vs. Magnitude

Magnitude

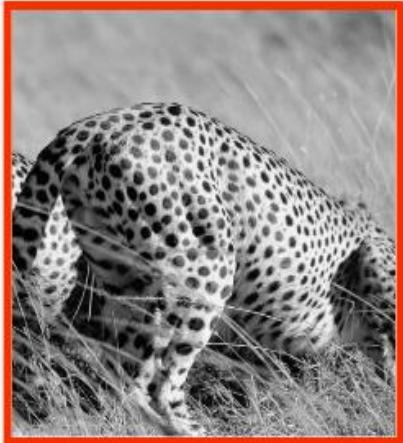


Phase

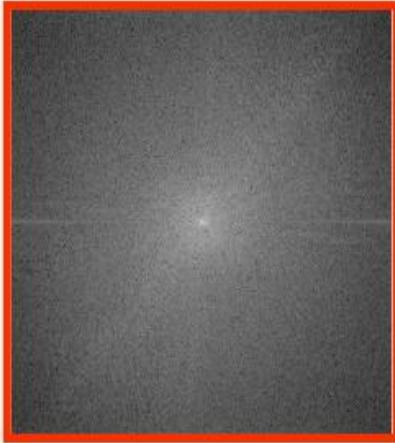


Phase vs. Magnitude

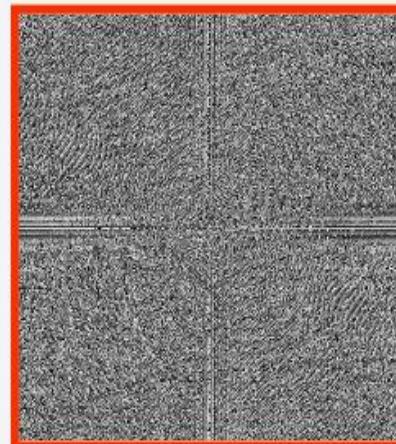
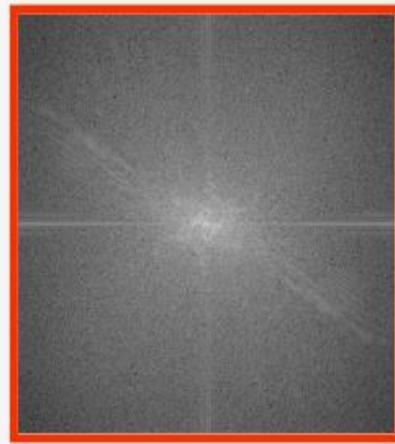
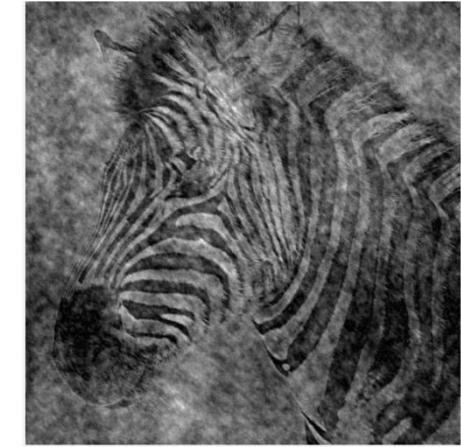
Magnitude



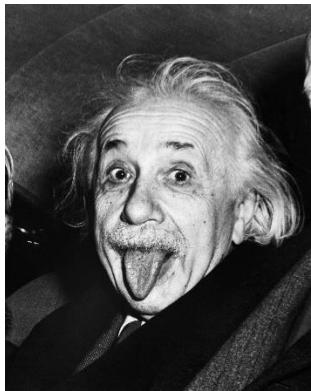
Phase



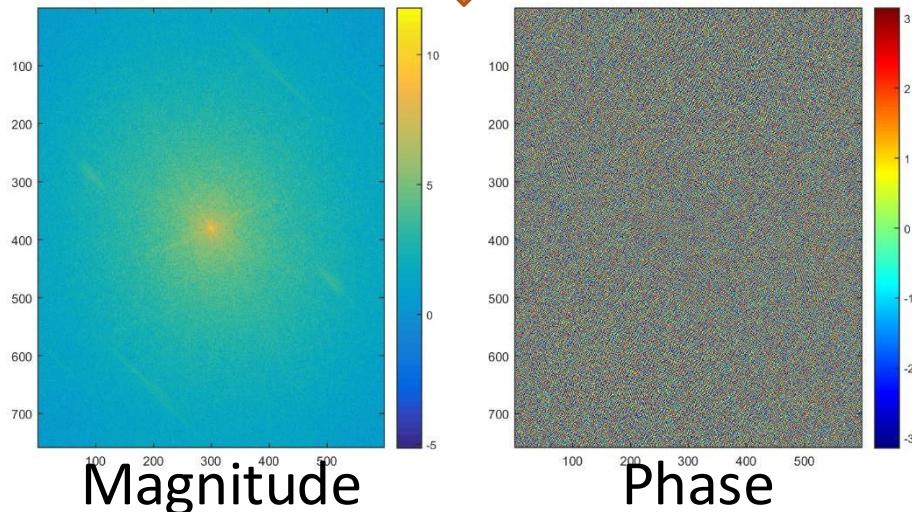
Swap phase and reconstruct



Phase vs. Magnitude



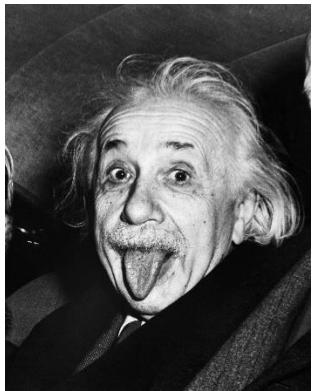
Intensity image



Magnitude

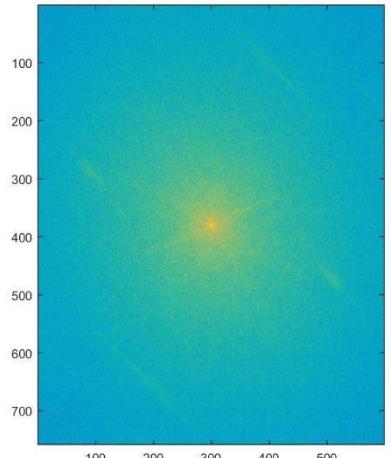
Phase

Phase vs. Magnitude

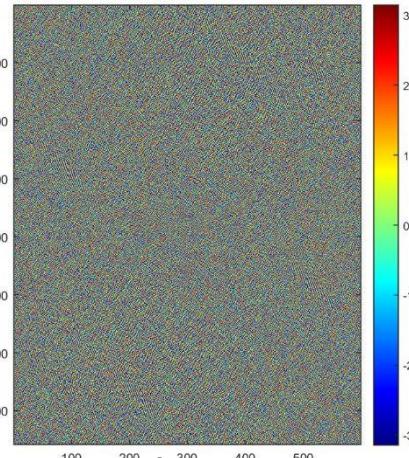


Intensity image

FFT

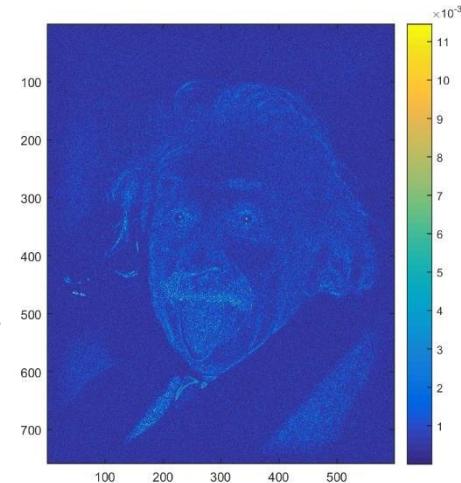


Magnitude

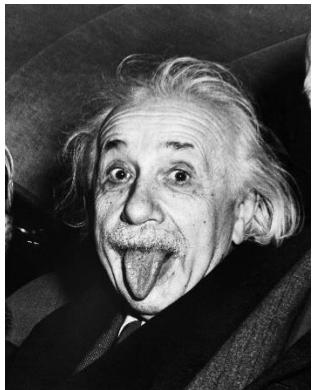


Phase

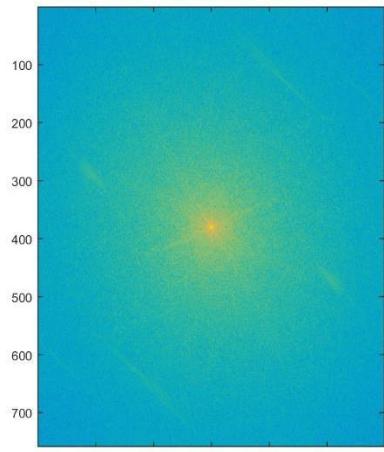
Use random
magnitude Inverse FFT



Phase vs. Magnitude

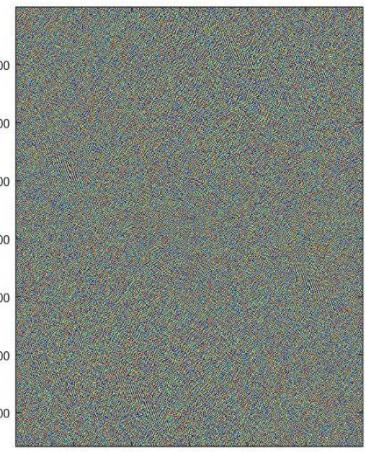


Intensity image



Magnitude

FFT

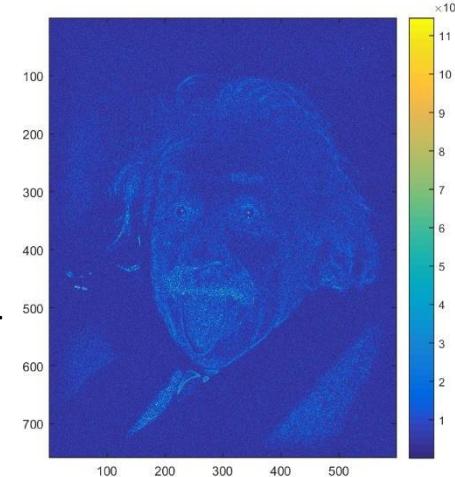


Phase

Use random magnitude



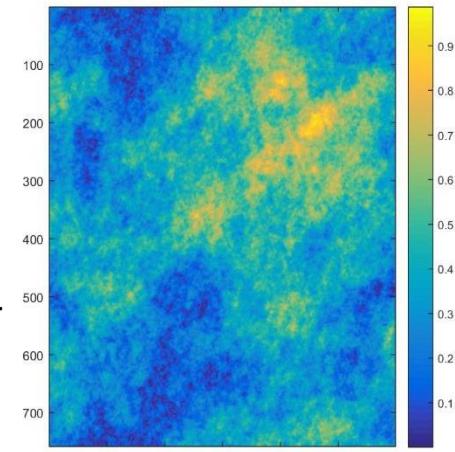
Inverse FFT



Use random phase



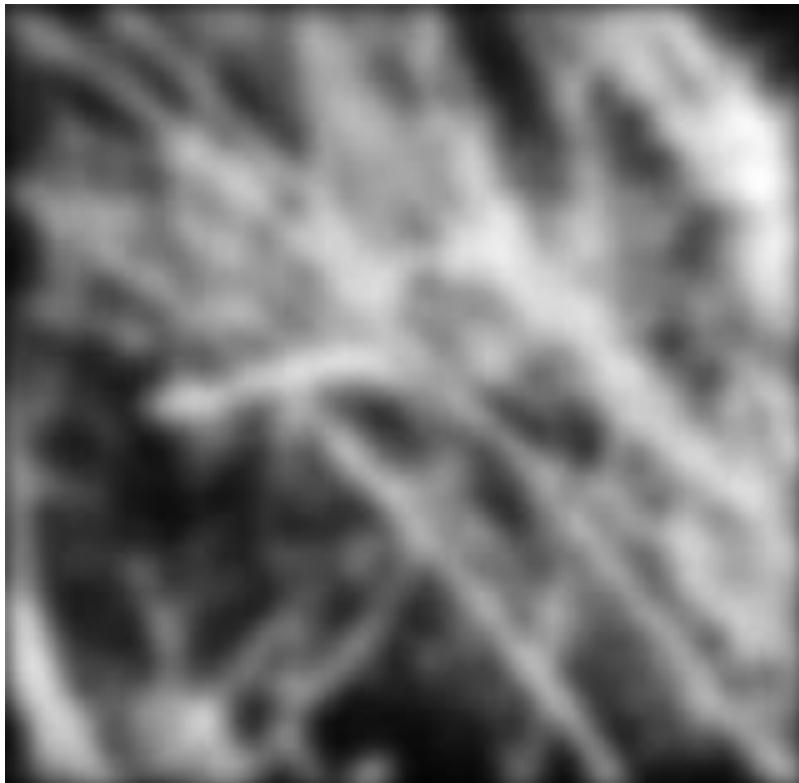
Inverse FFT



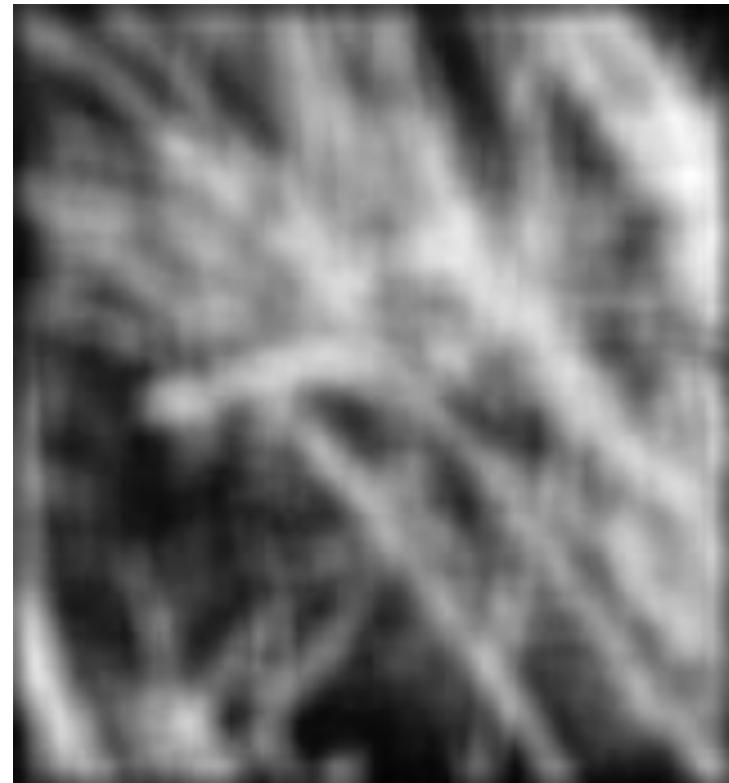
Filtering

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

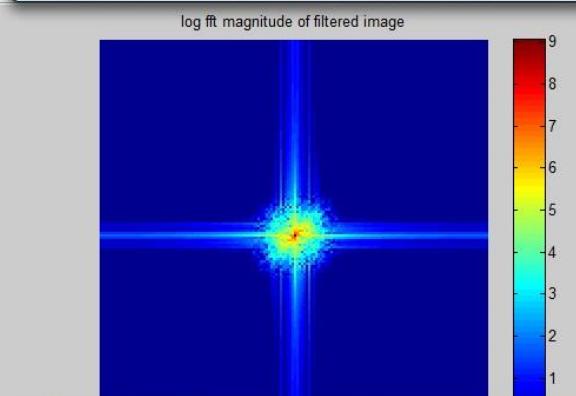
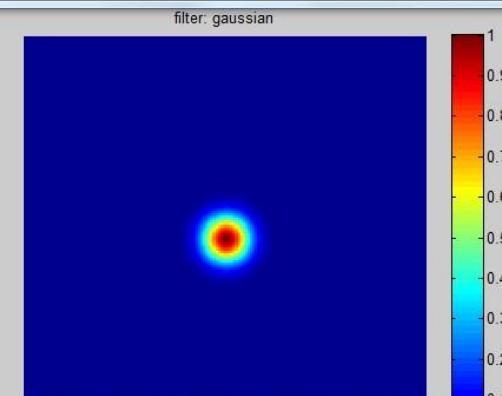
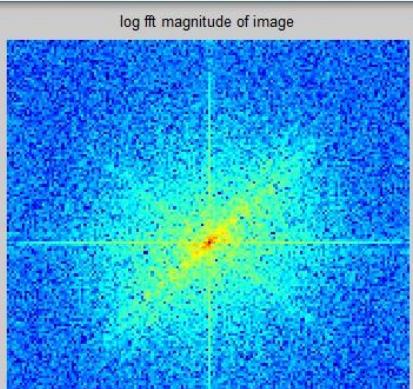
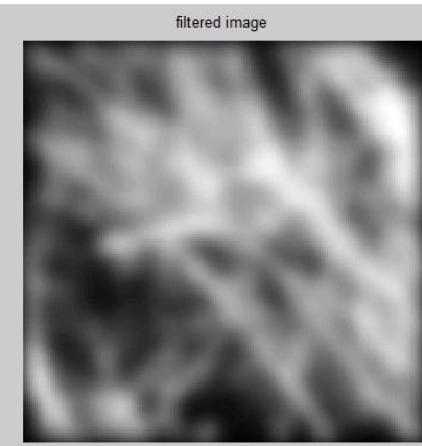
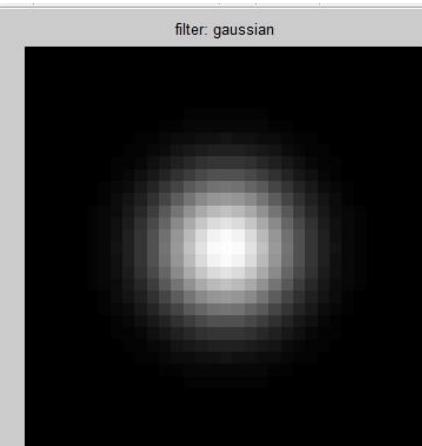
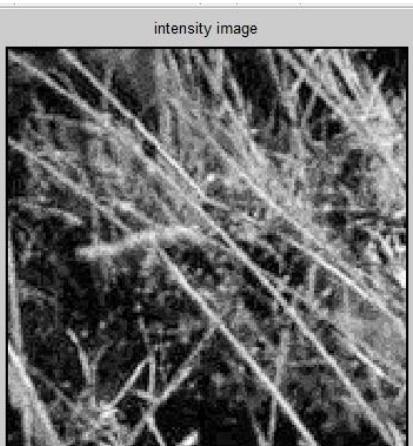
Gaussian



Box filter



Gaussian



Box Filter

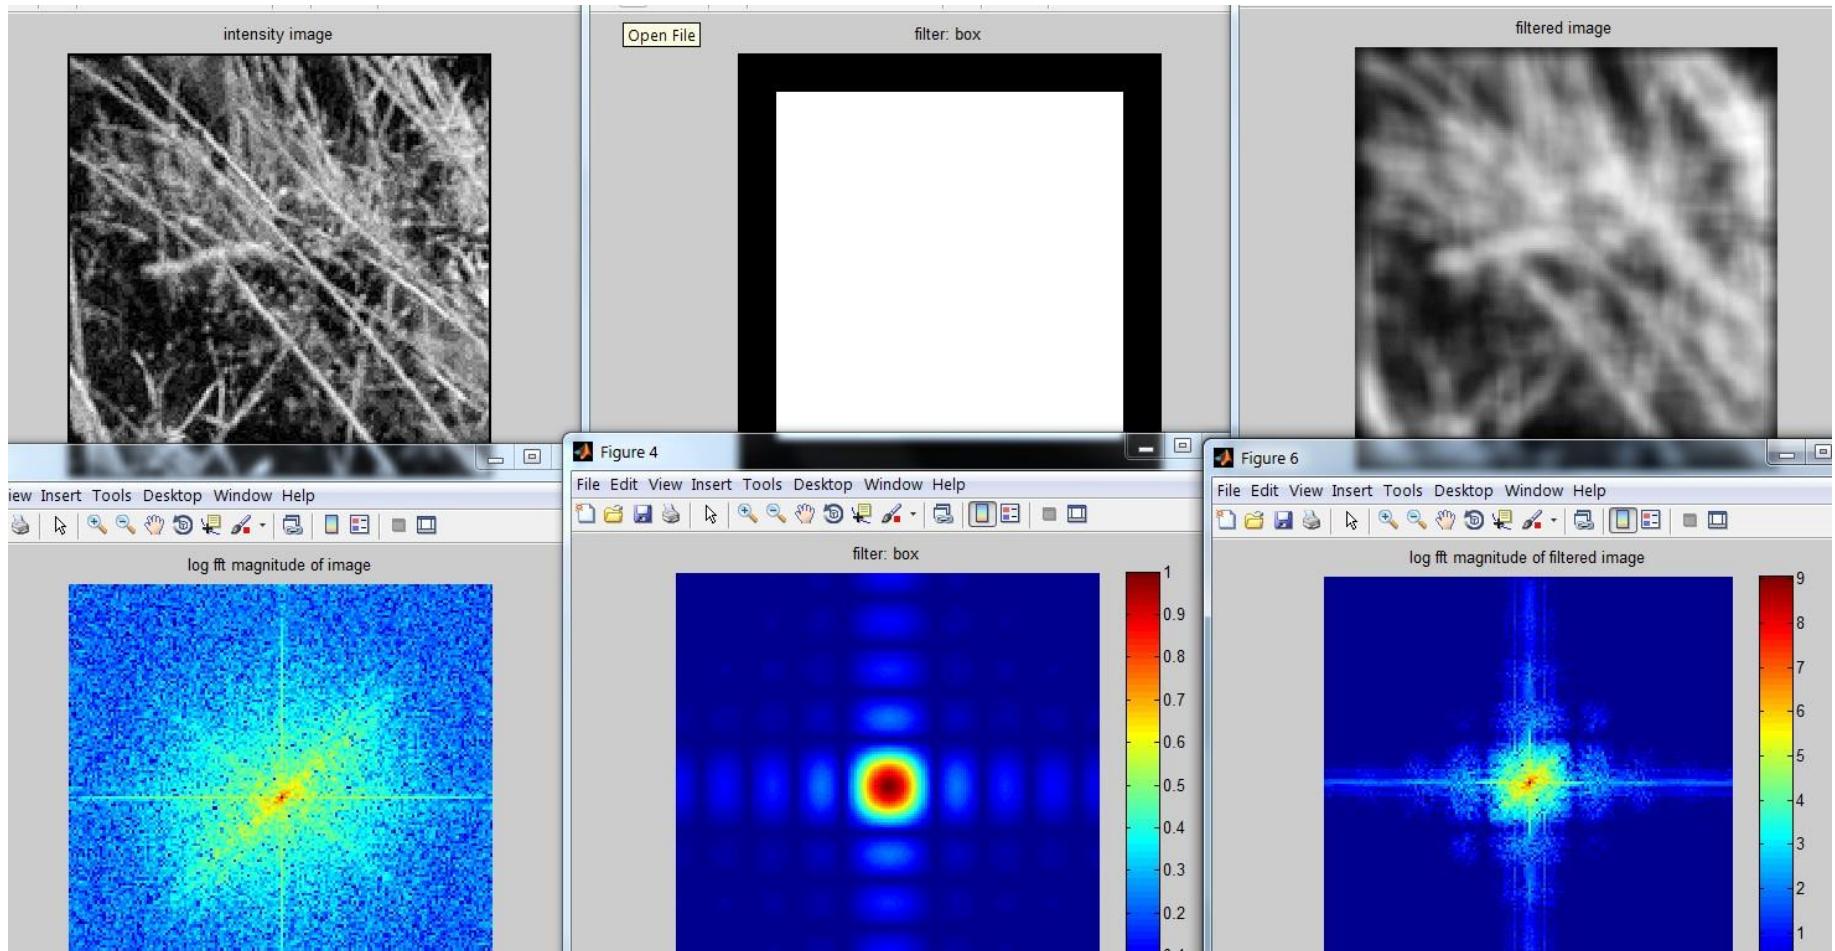


Image half-sizing

This image is too big to fit on the screen. How can we reduce it?

How to generate a half-sized version?

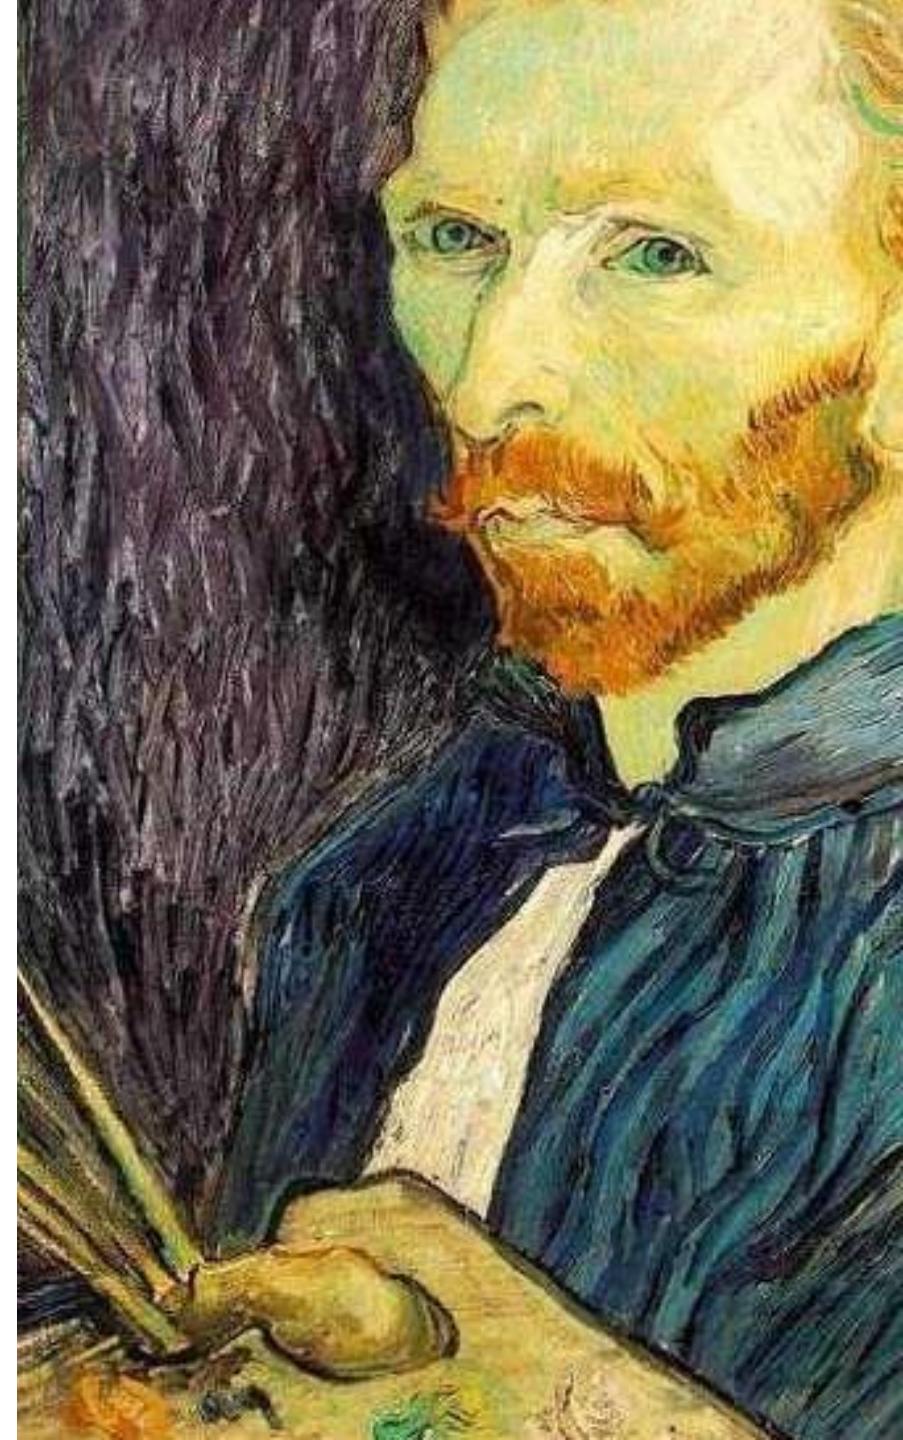
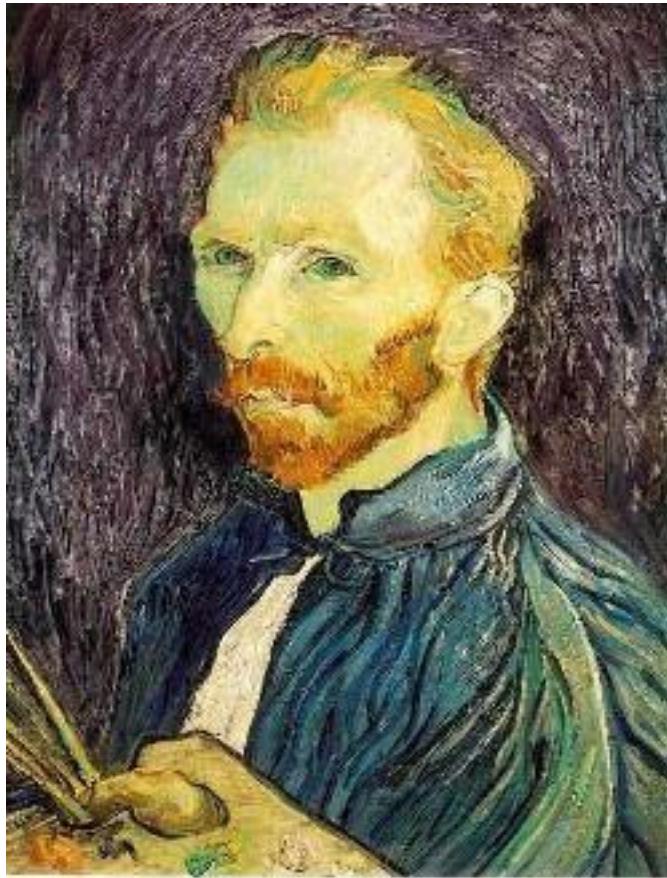


Image sub-sampling



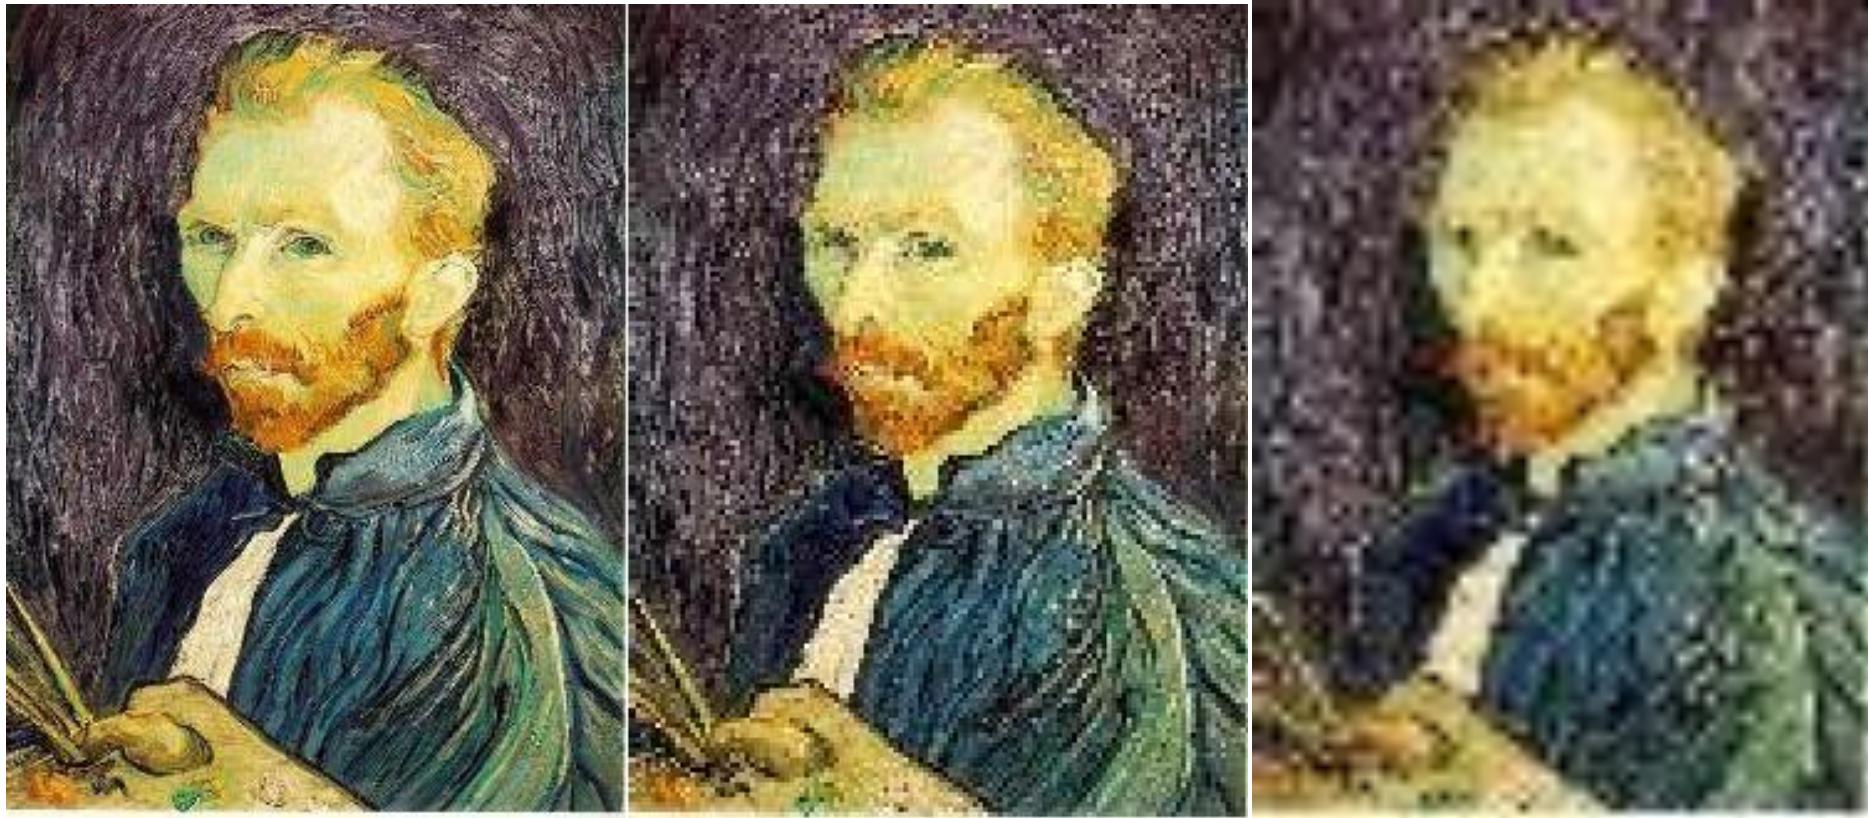
1/4



1/8

Throw away every other row and
column to create a $1/2$ size image
- called *image sub-sampling*

Image sub-sampling



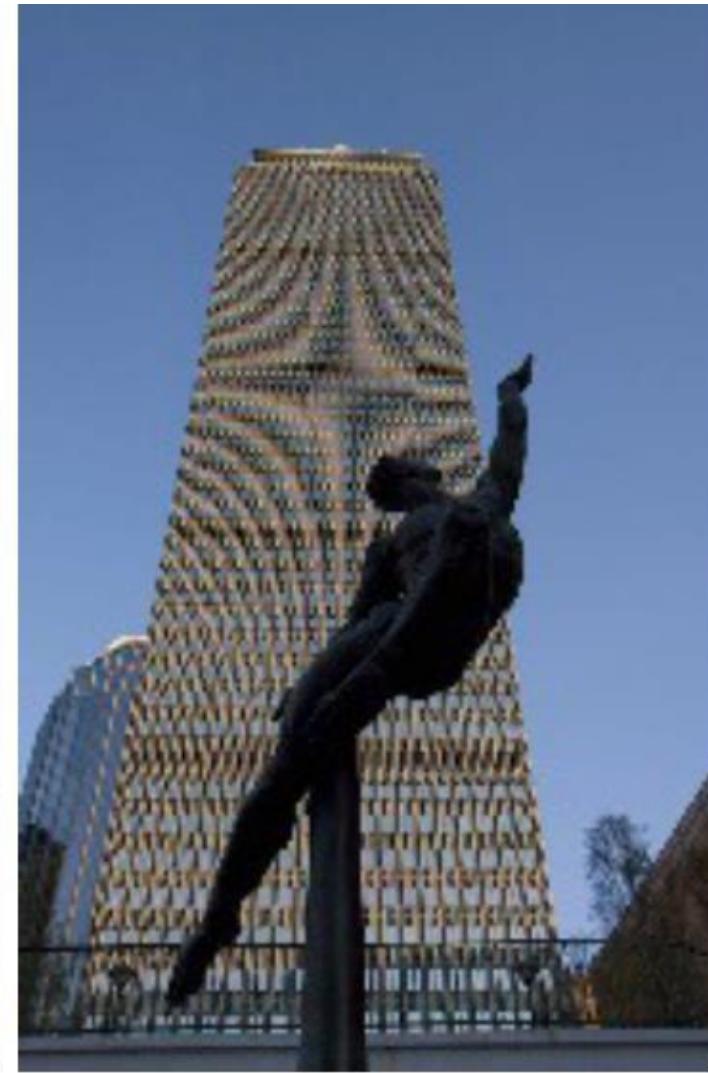
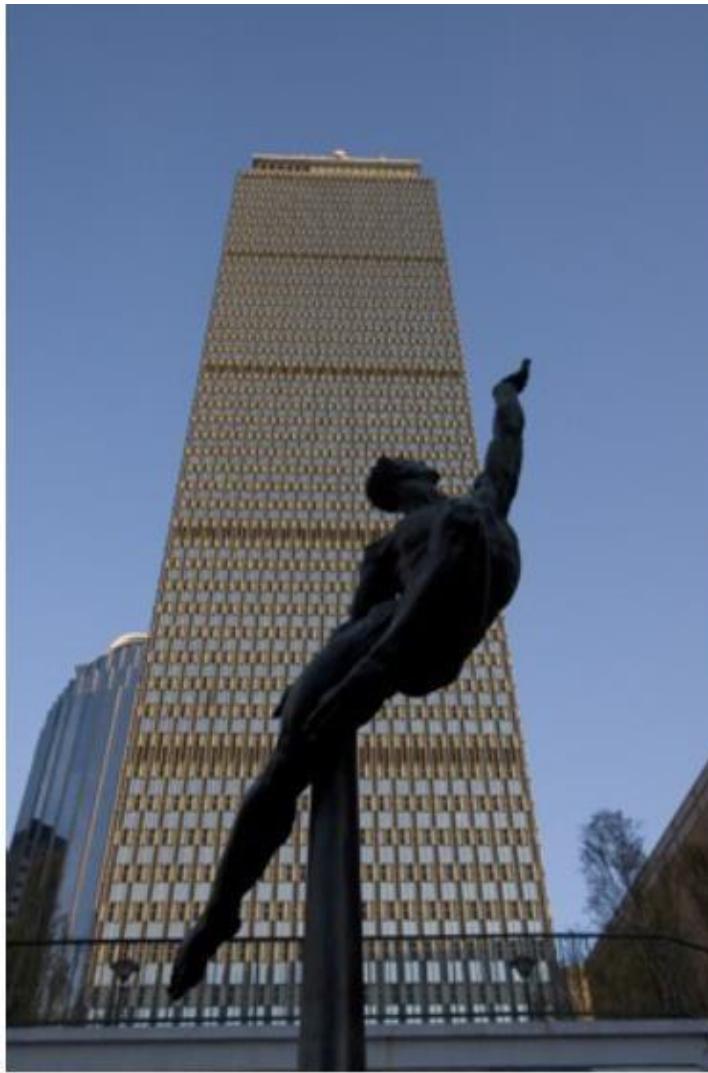
1/2

1/4 (2x zoom)

1/8 (4x zoom)

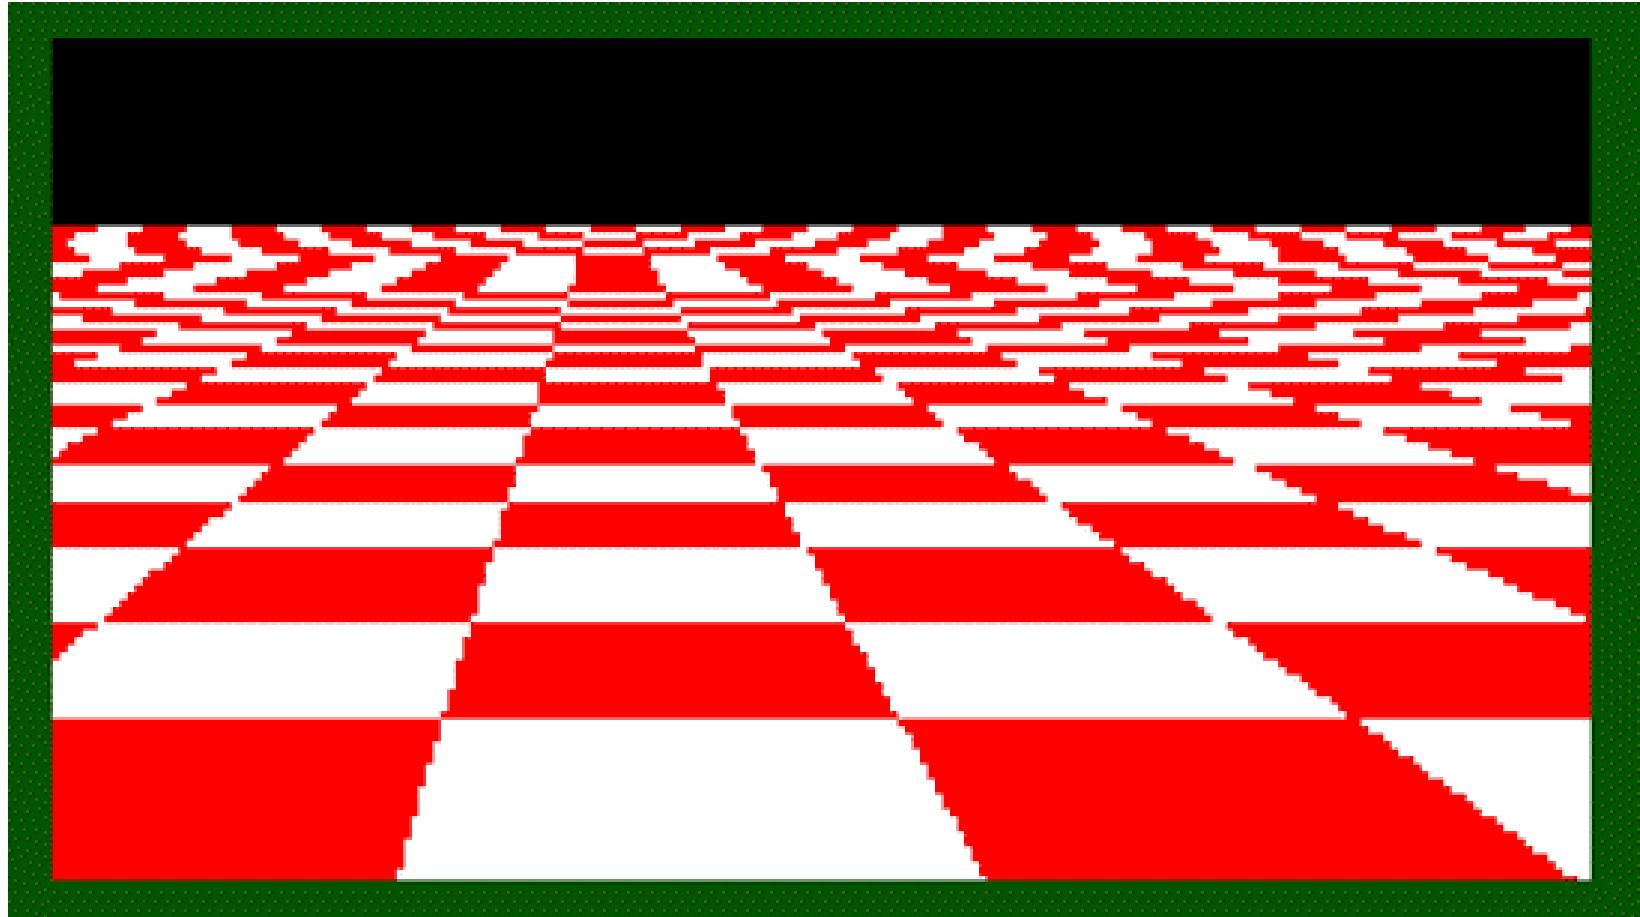
Why does this look so cruddy?
Aliasing! What do we do?

Image sub-sampling



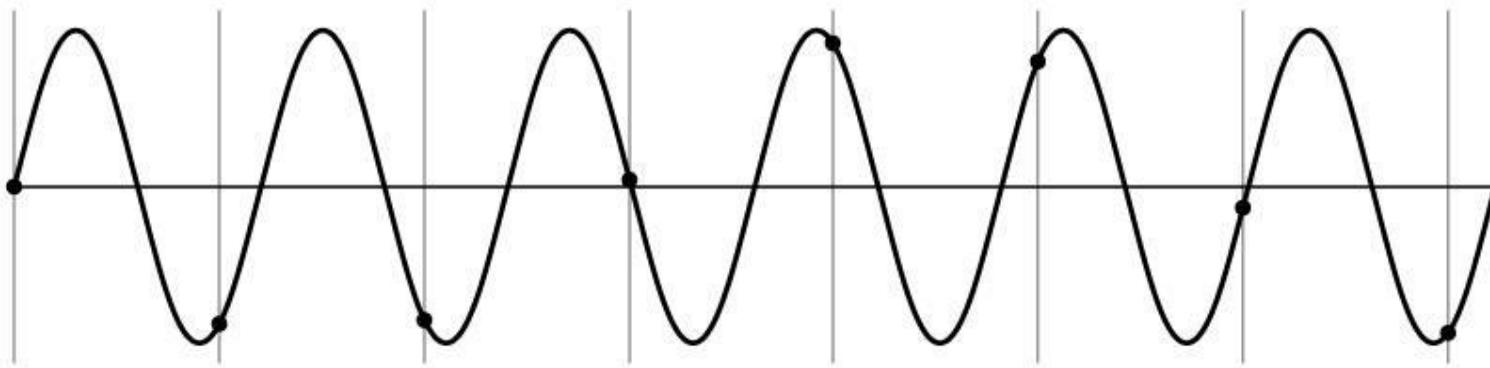
Source: F. Durand

Even worse for synthetic images



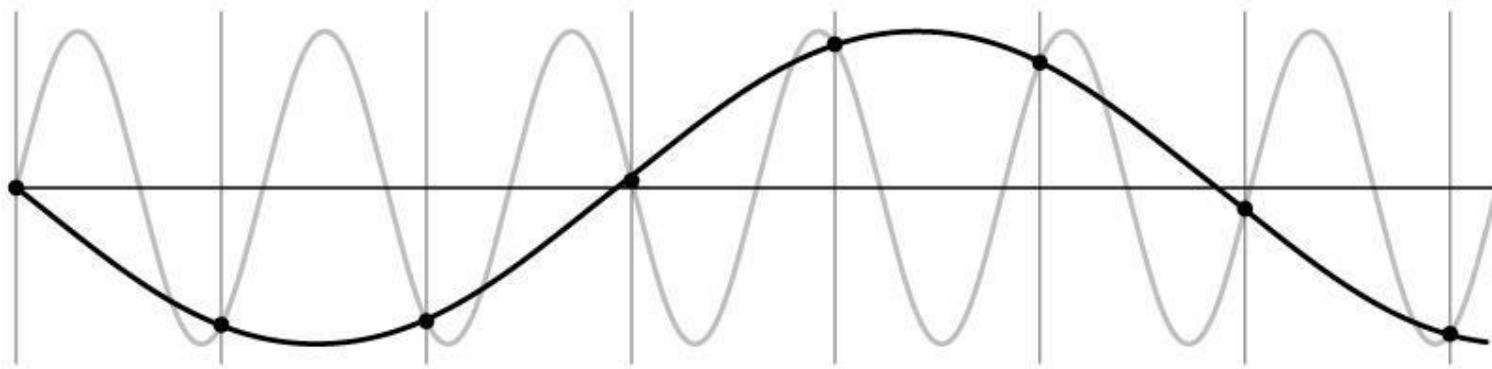
Aliasing problem

- 1D example (sinewave):

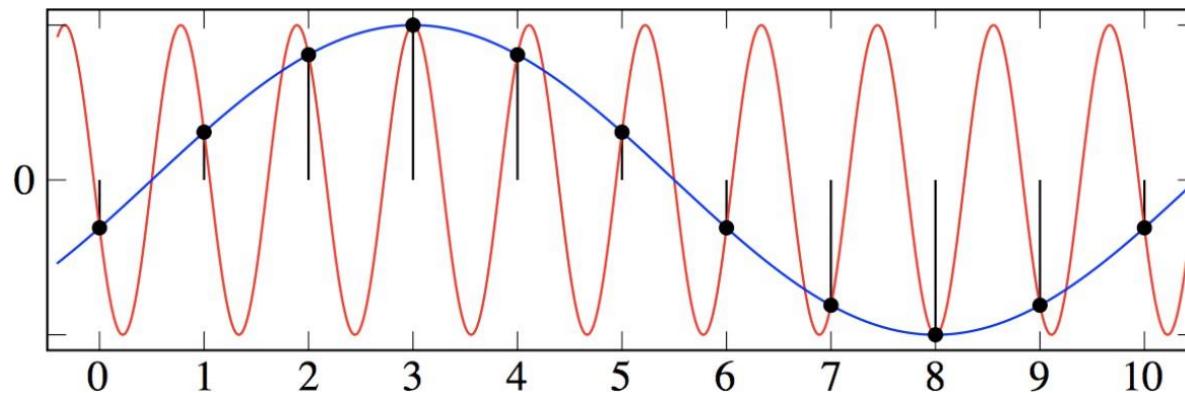


Aliasing problem

- 1D example (sinewave):

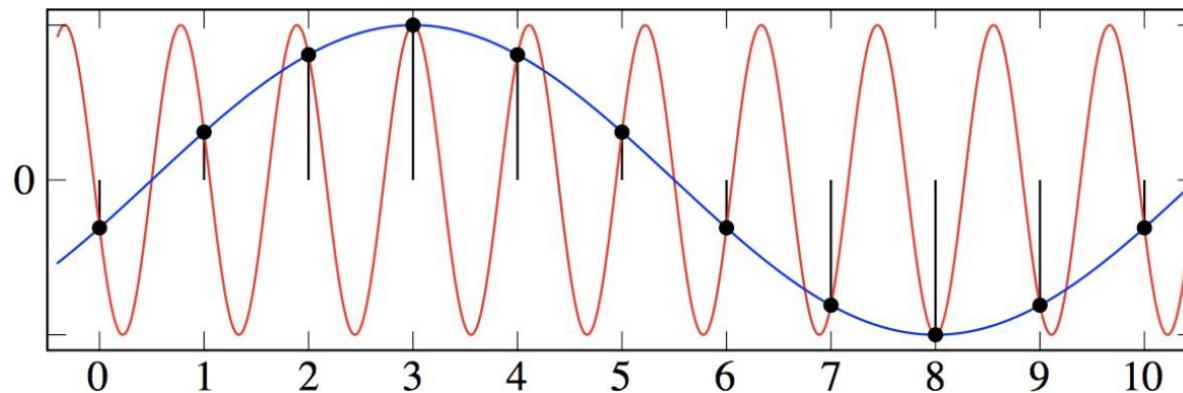


Aliasing



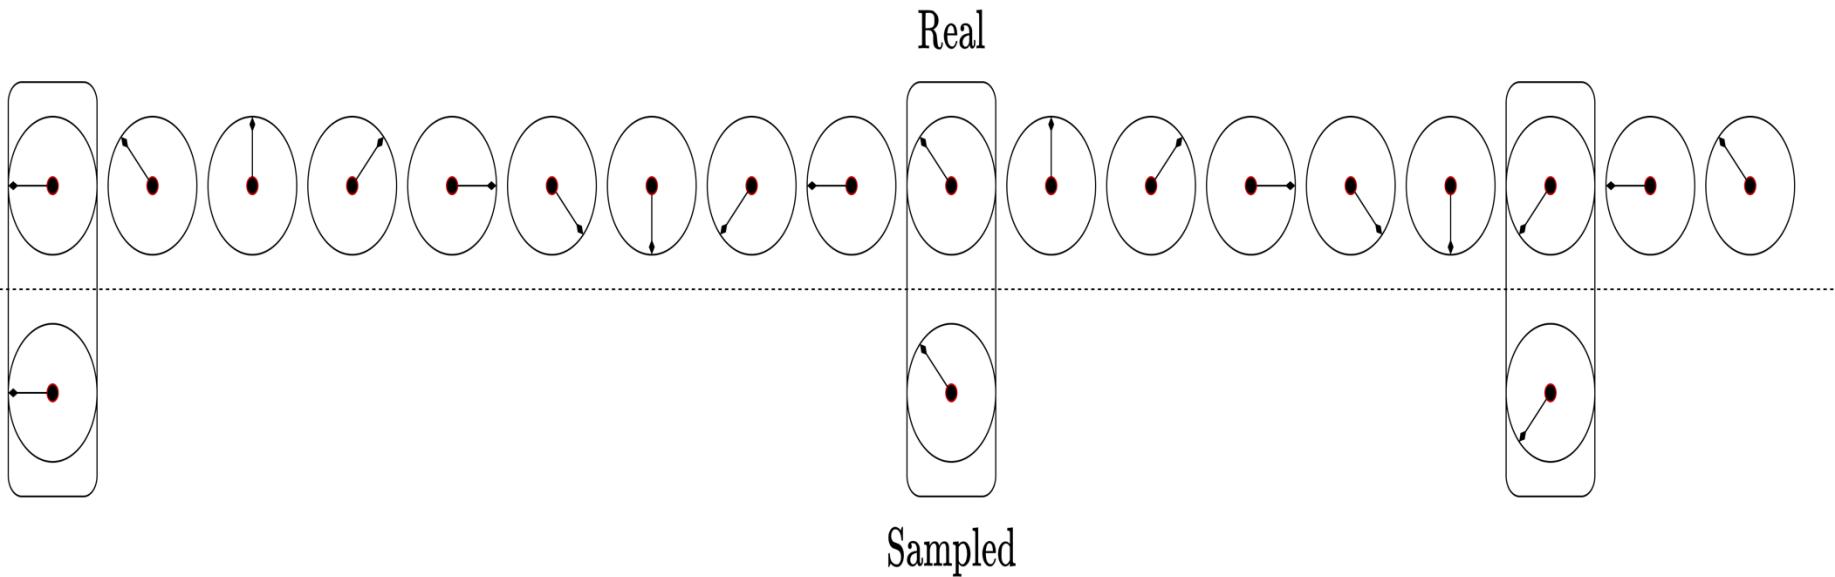
- Occurs when your sampling rate is not high enough to capture the amount of detail in your image
- High frequencies look like low frequencies
- Can give you the wrong signal/image—an *alias*

Aliasing



- Occurs when your sampling rate is not high enough to capture the amount of detail in your image
- Unresolved, high frequencies look like resolved low frequencies
- Can give you the wrong signal/image—an *alias*
- To avoid aliasing:
 - sampling rate $\geq 2 * \text{max frequency in the image}$
 - said another way: \geq two samples per cycle
 - this minimum sampling rate is called the **Nyquist rate**

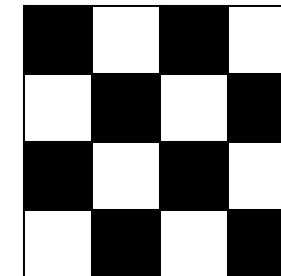
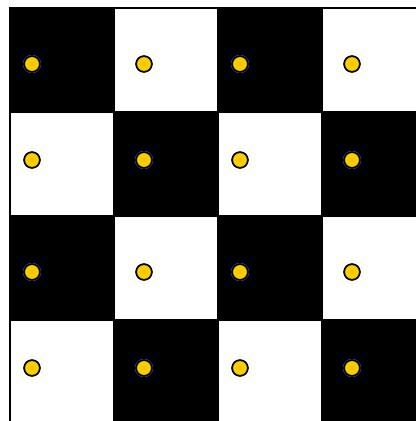
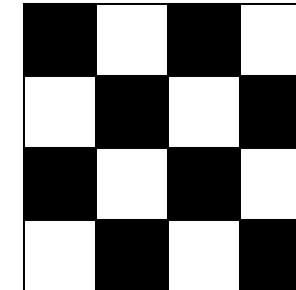
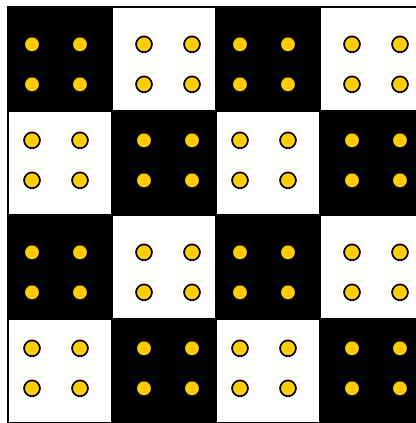
Aliasing: Video



Wagon Wheel effect: Wheels spins in the opposite direction at high speed.

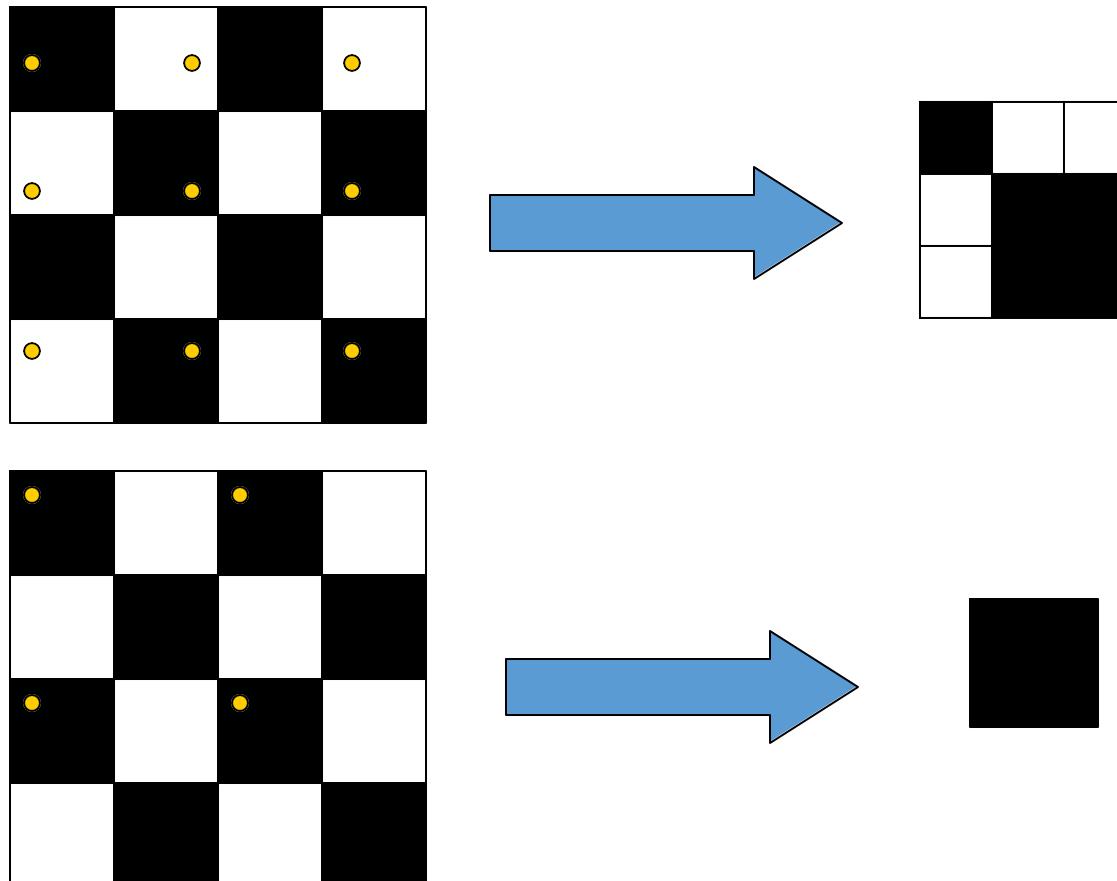
<https://www.youtube.com/watch?v=jHS9JGkEOmA>

Sampling an image



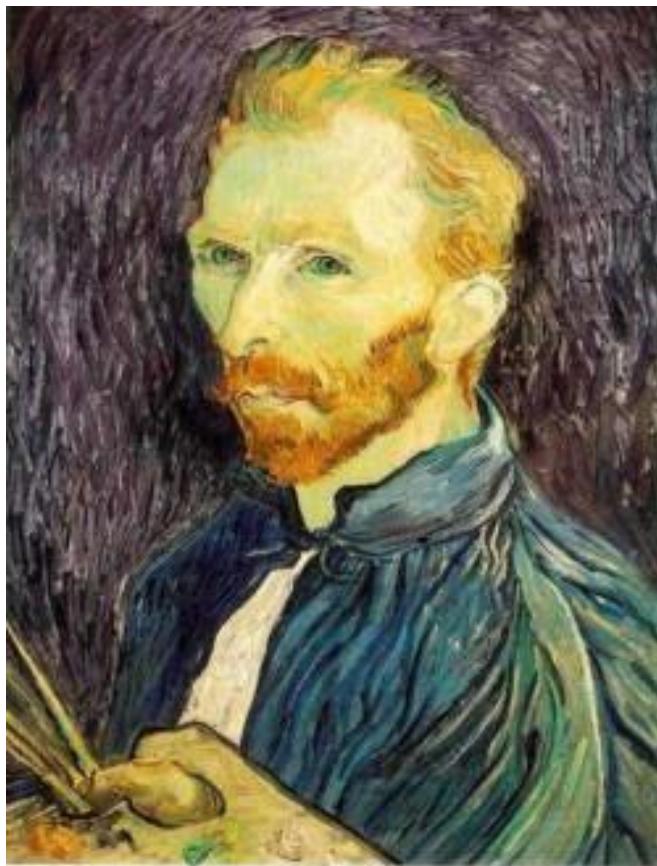
Examples of GOOD sampling

Undersampling

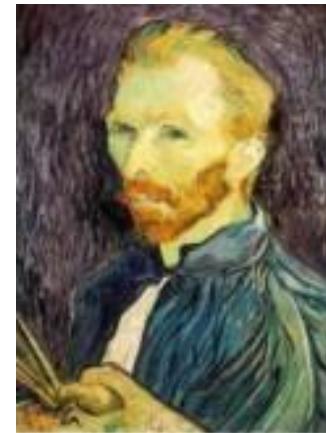


Examples of BAD sampling -> Aliasing

Gaussian (low-pass) pre-filtering



Gaussian 1/2



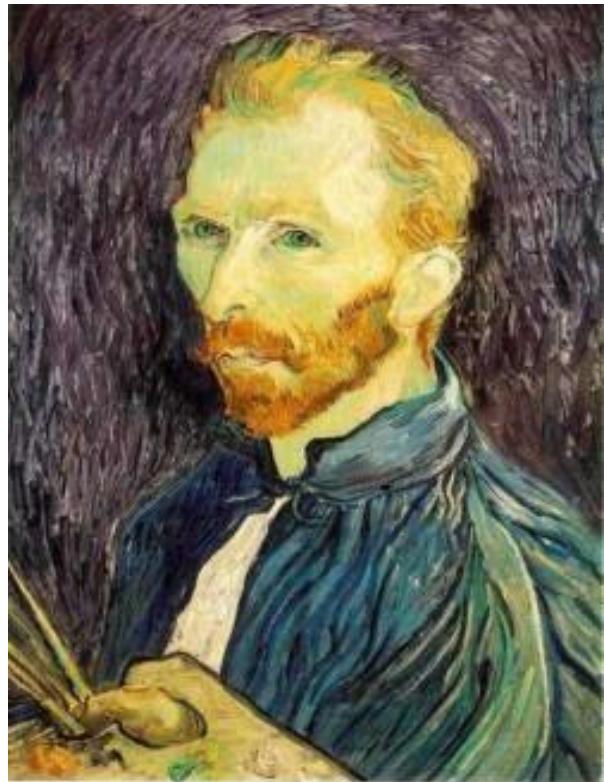
G 1/4



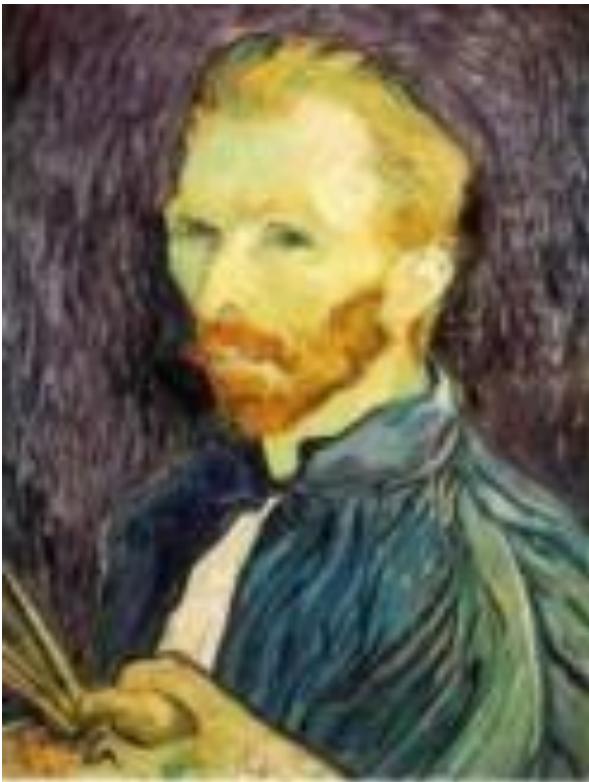
G 1/8

- Solution: filter the image, *then* subsample

Subsampling with Gaussian pre-filtering



Gaussian 1/2



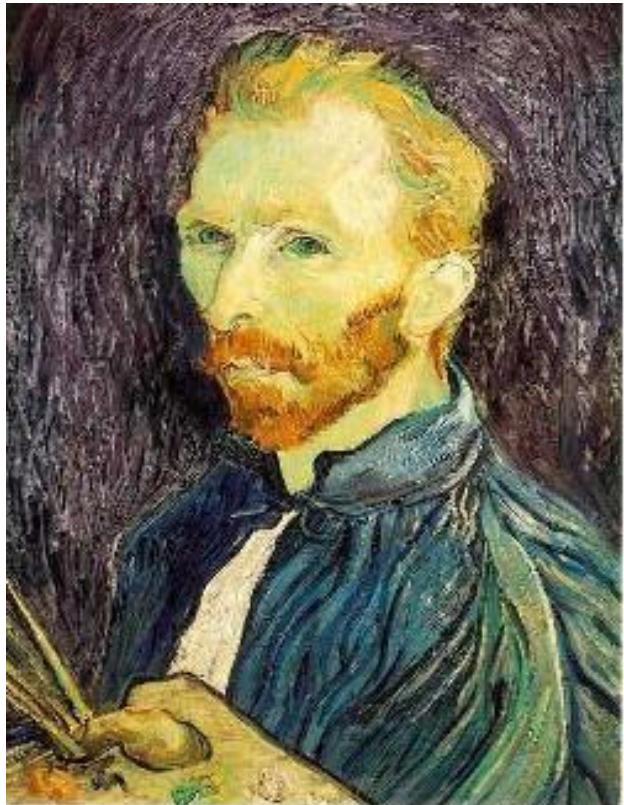
G 1/4



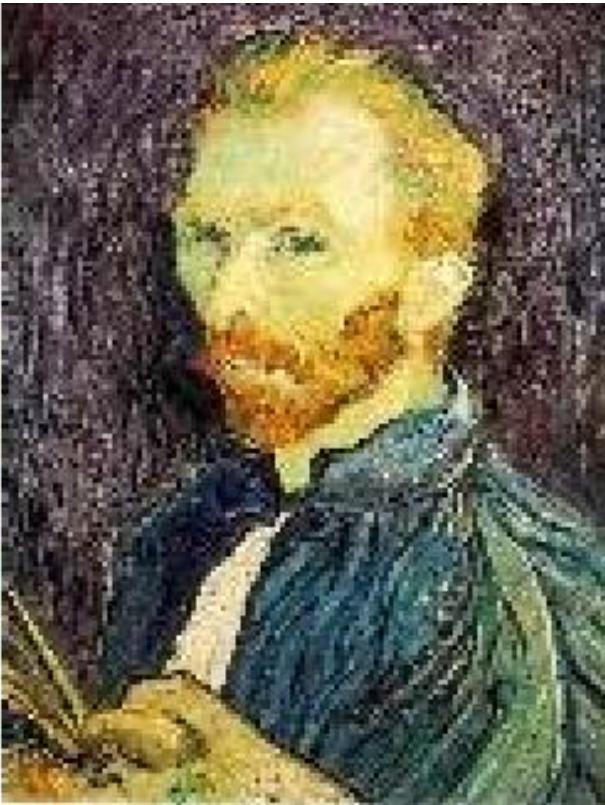
G 1/8

- Solution: filter the image, *then* subsample

Compare with...



1/2



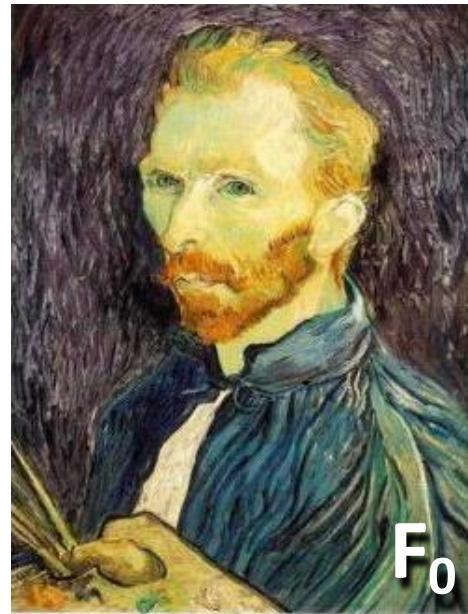
1/4 (2x zoom)



1/8 (4x zoom)

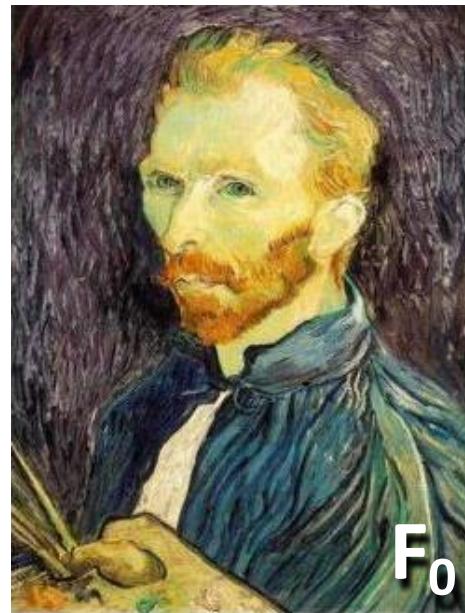
Gaussian pre-filtering

- Solution: filter the image, *then* subsample

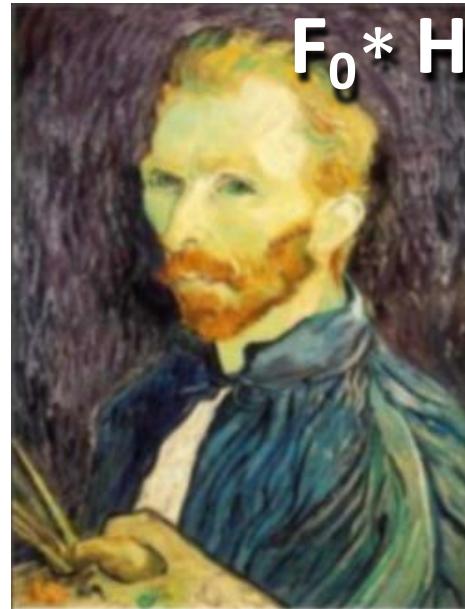


Gaussian pre-filtering

- Solution: filter the image, *then* subsample

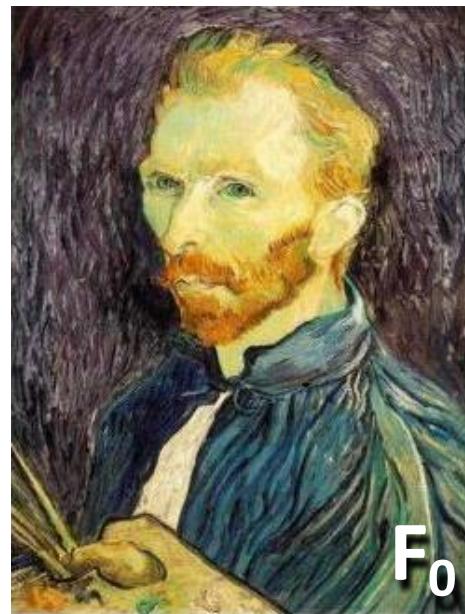
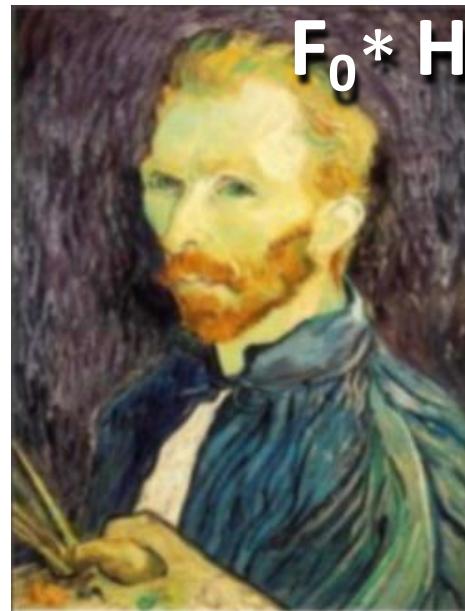


blur



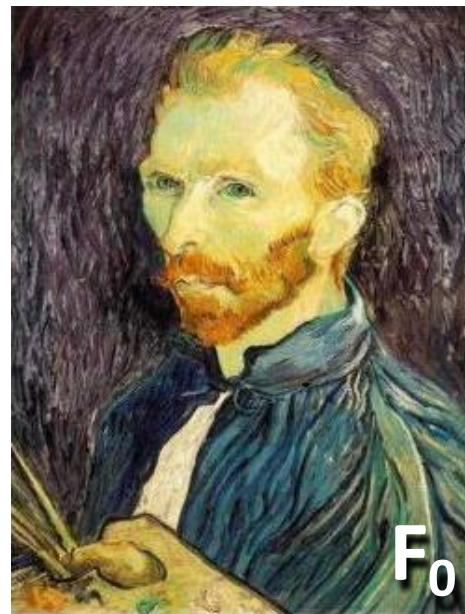
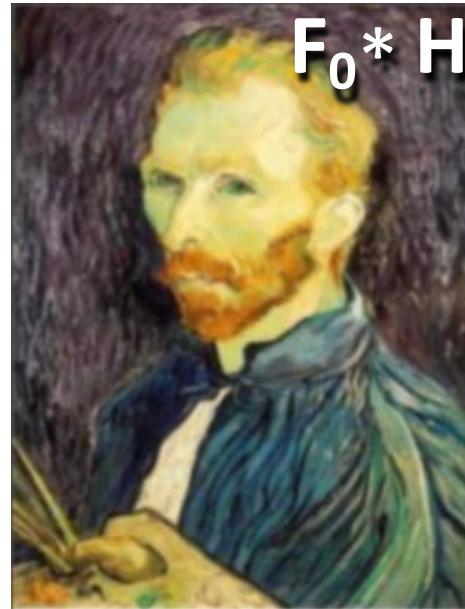
Gaussian pre-filtering

- Solution: filter the image, *then* subsample

 F_0 F_1  $F_0 * H$

Gaussian pre-filtering

- Solution: filter the image, *then* subsample

 F_0 F_1  $F_0 * H$ $F_1 * H$

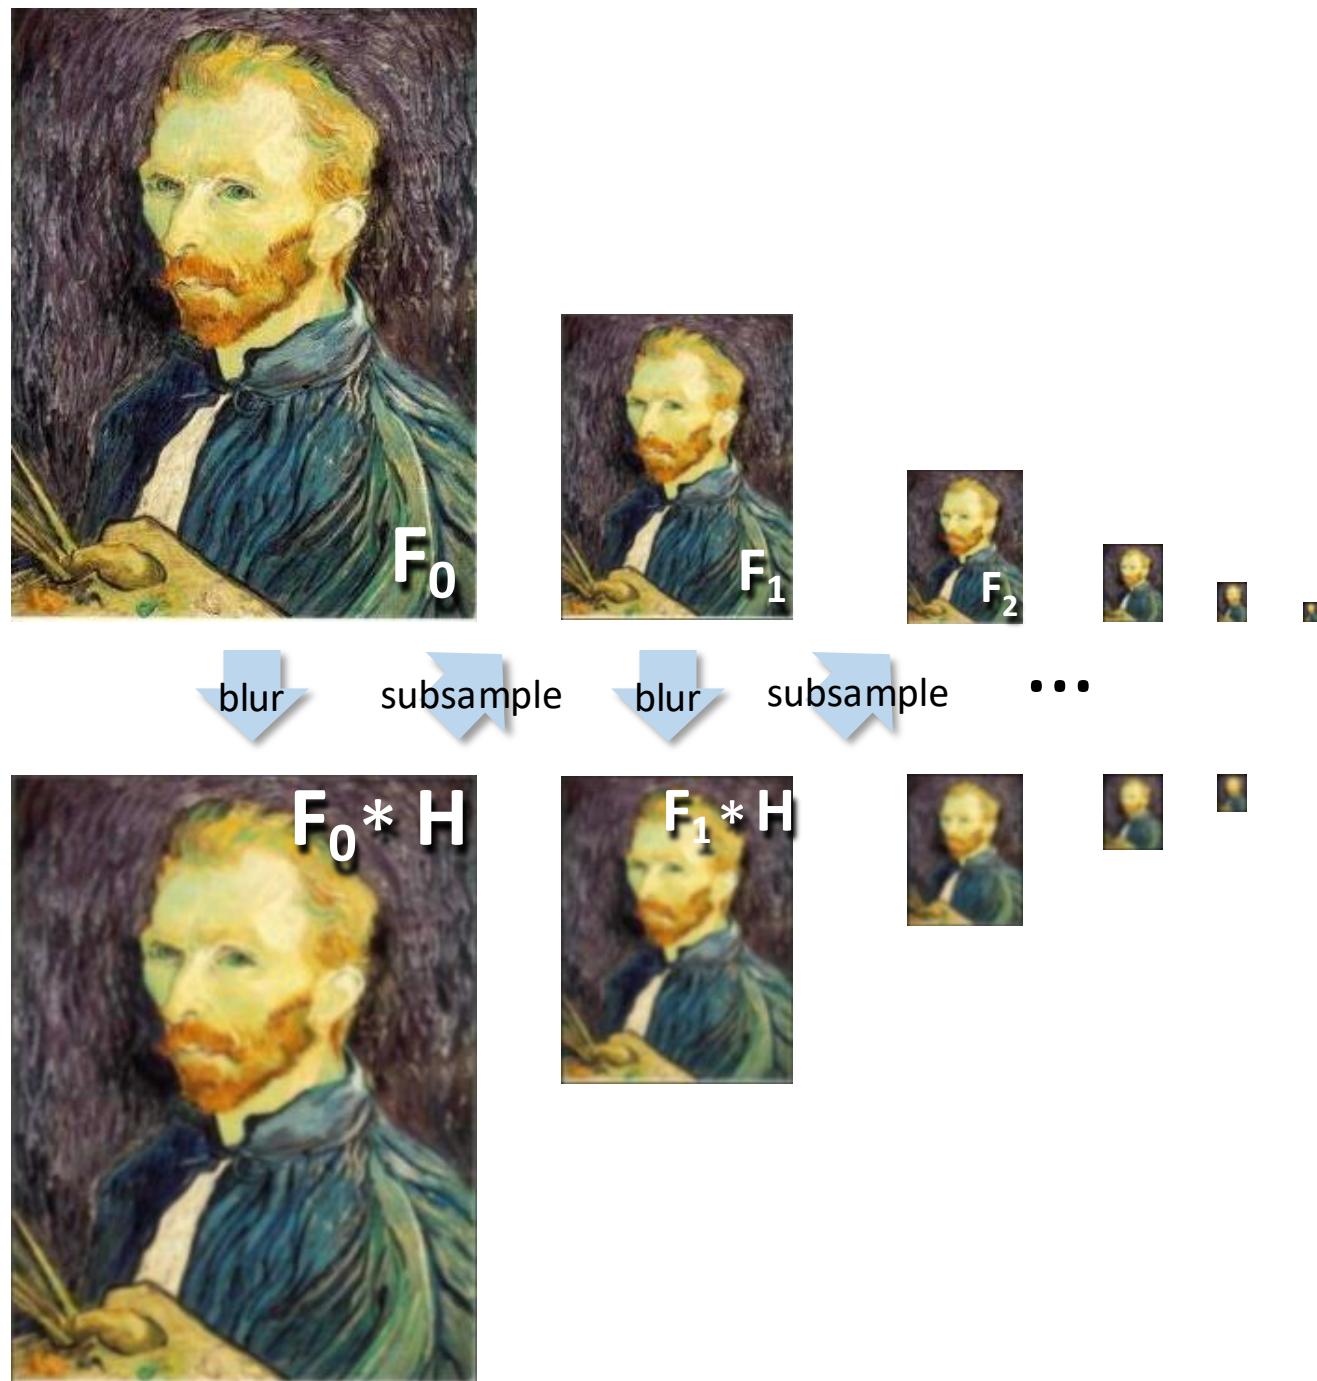
Gaussian pre-filtering

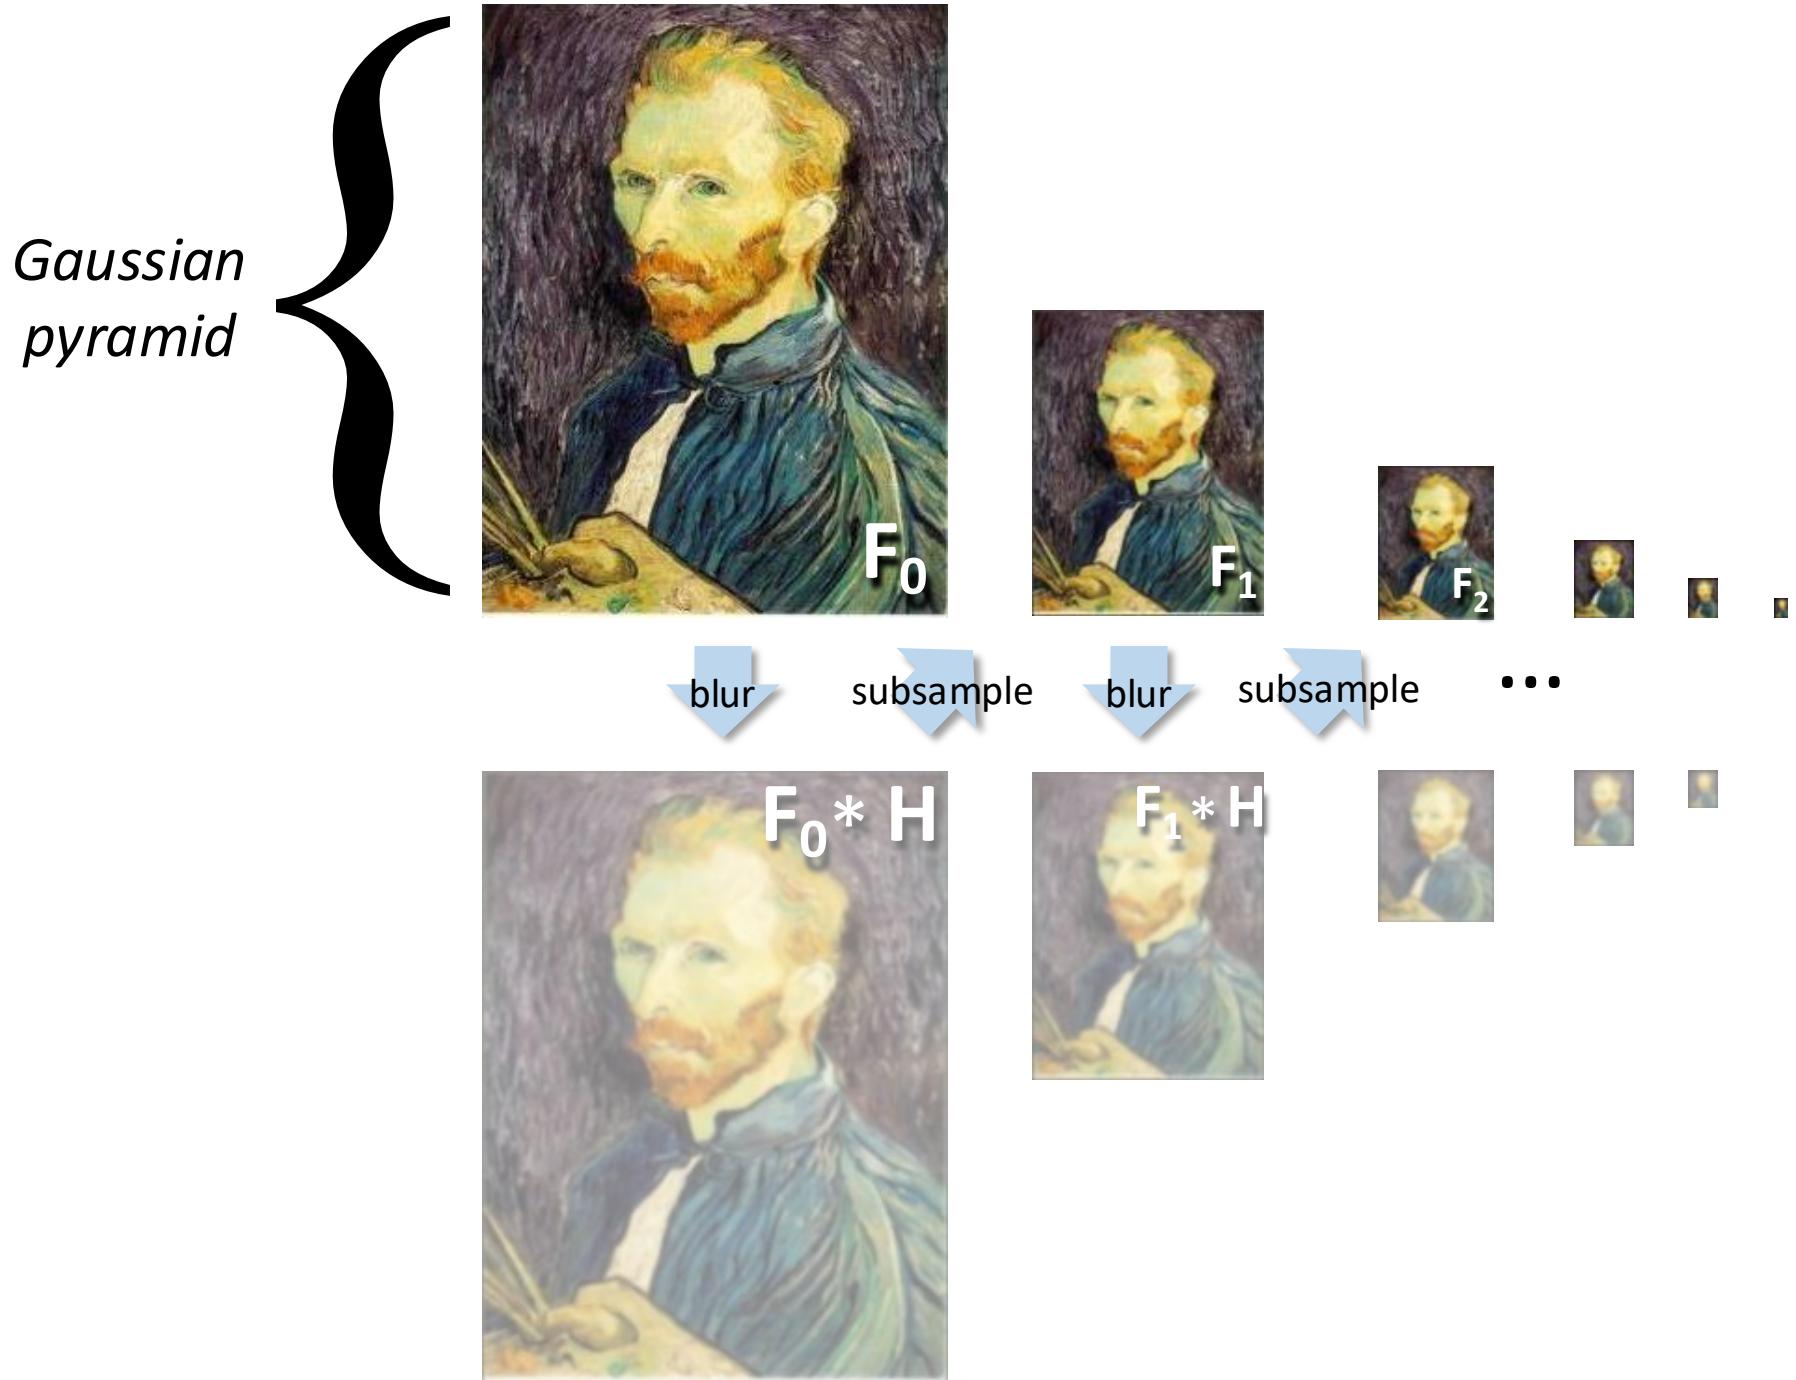
- Solution: filter the image, *then* subsample



Gaussian pre-filtering

- Solution: filter the image, *then* subsample

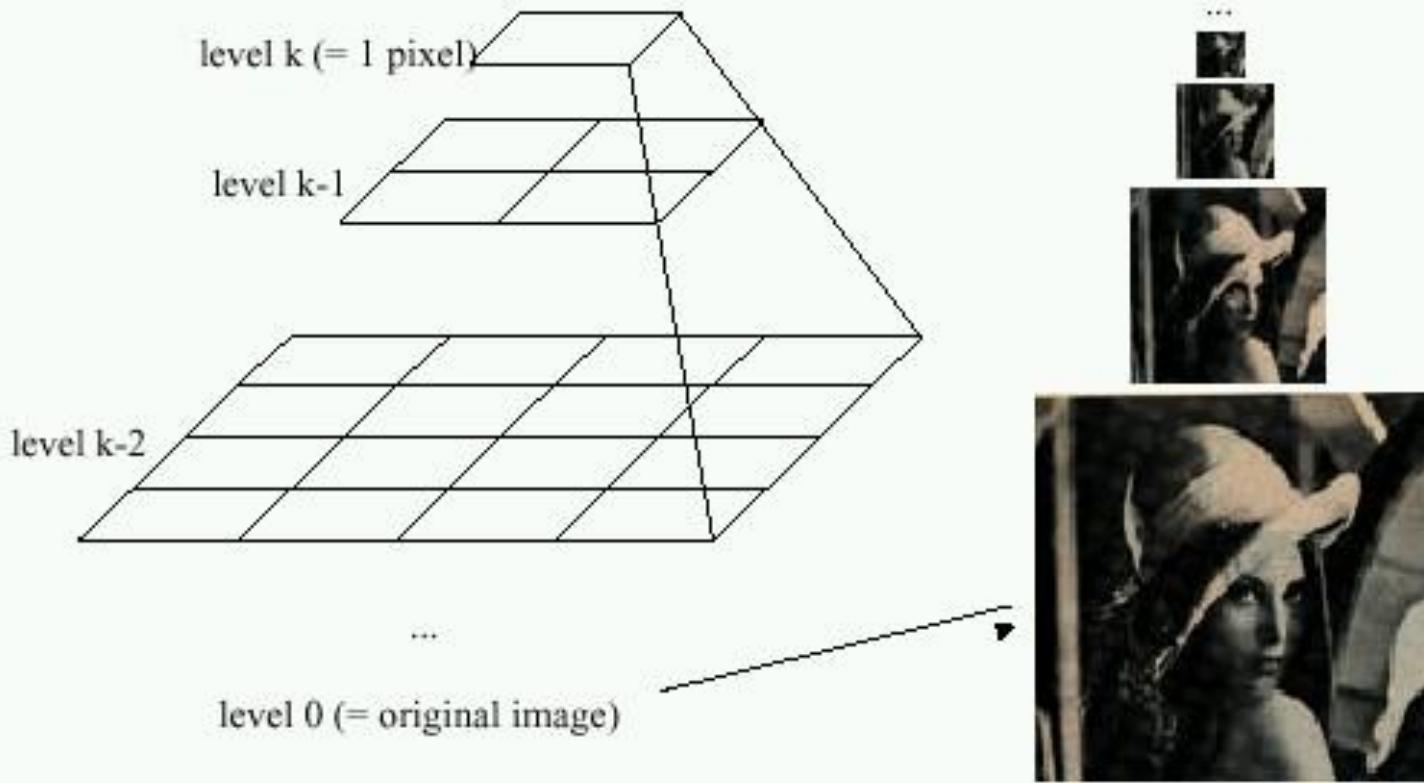




Gaussian pyramids

[Burt and Adelson, 1983]

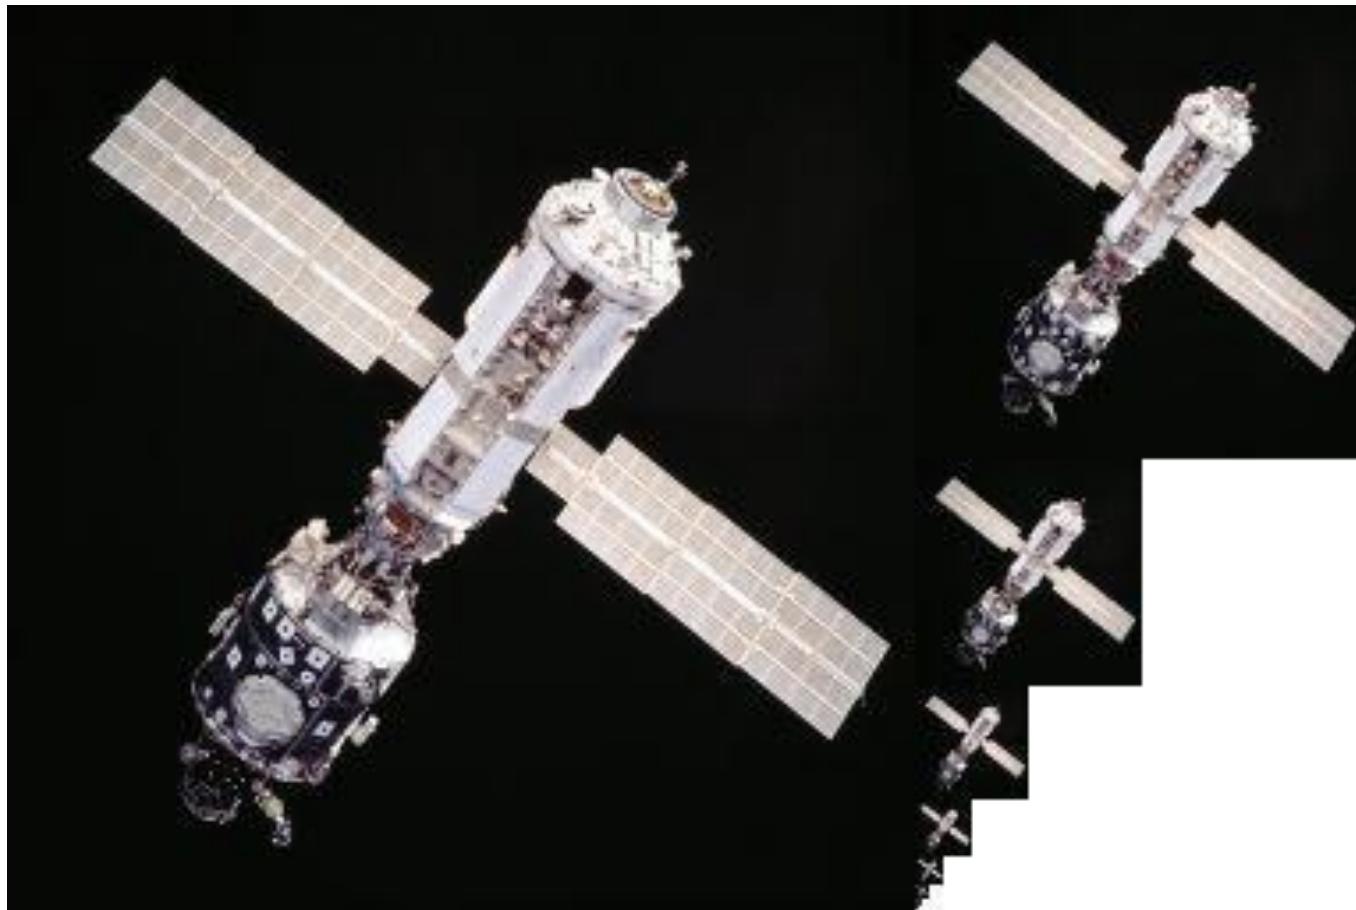
Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N=2^k$)



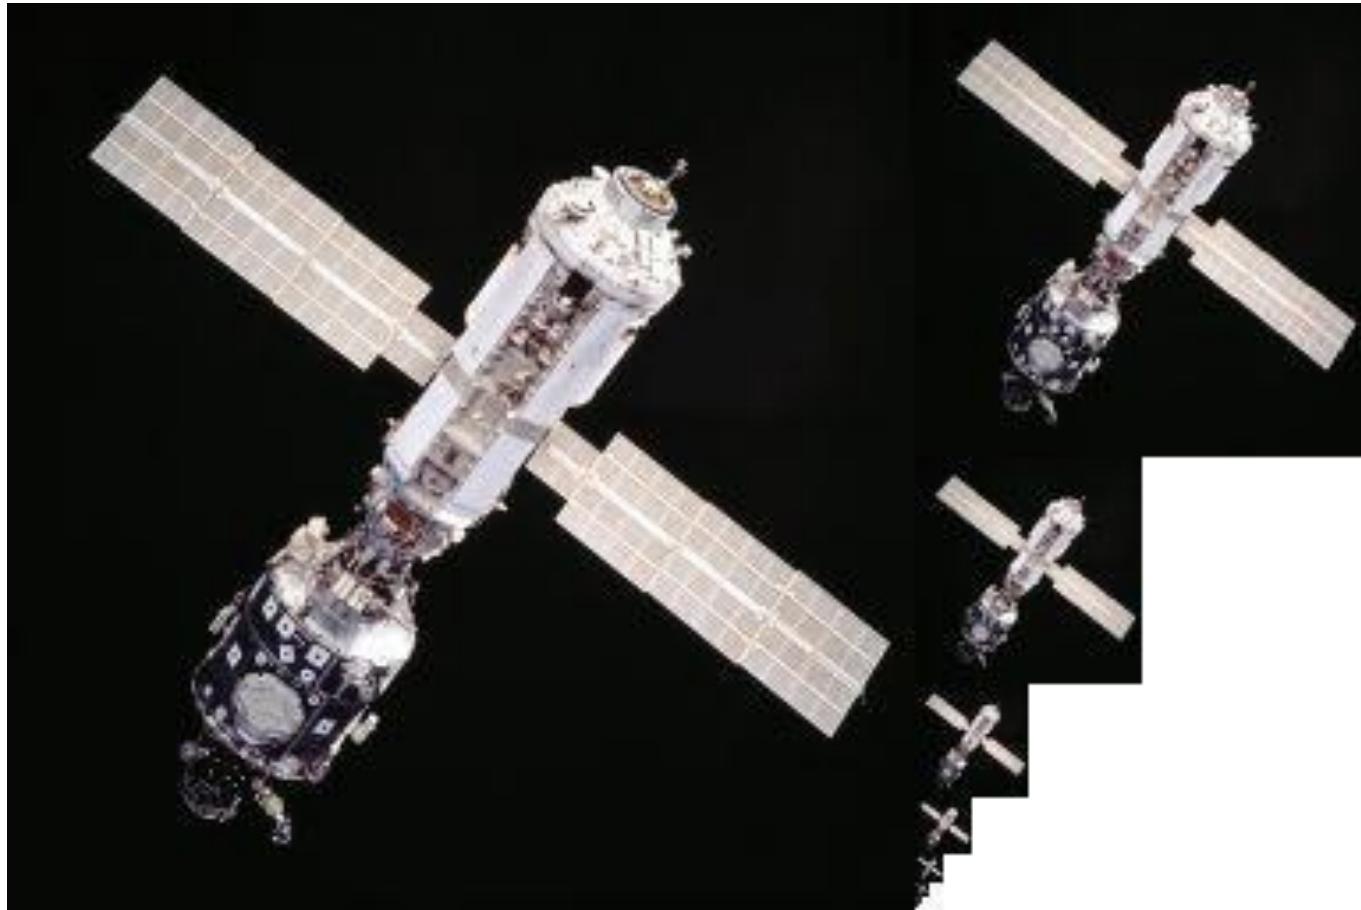
Gaussian Pyramids have all sorts of applications in computer vision

Source: S. Seitz

Gaussian Pyramid



Gaussian Pyramid



- How much space does a Gaussian pyramid take compared to the original image?

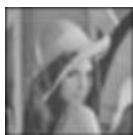
The Laplacian Pyramid

Gaussian Pyramid

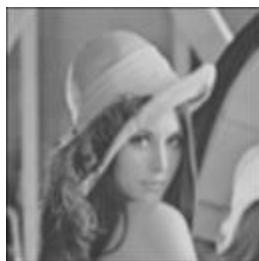
G_n



G_2



G_1

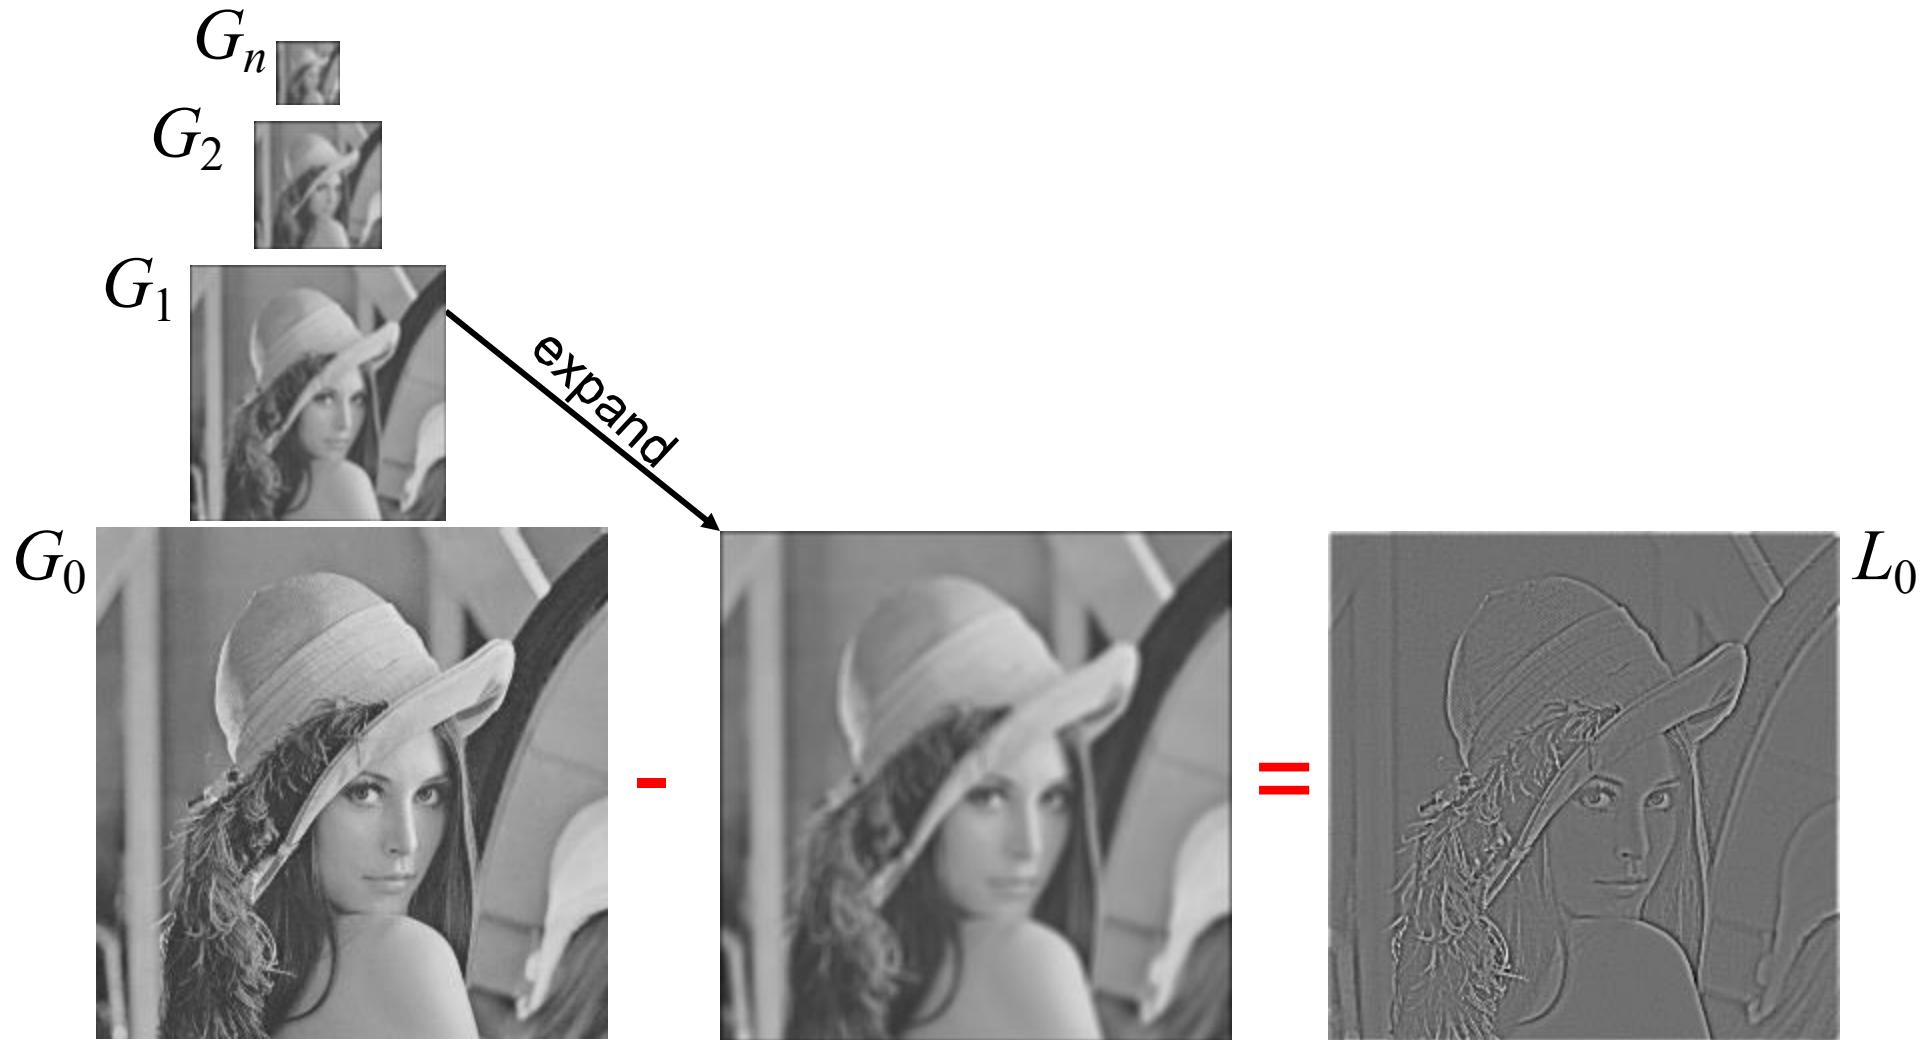


G_0



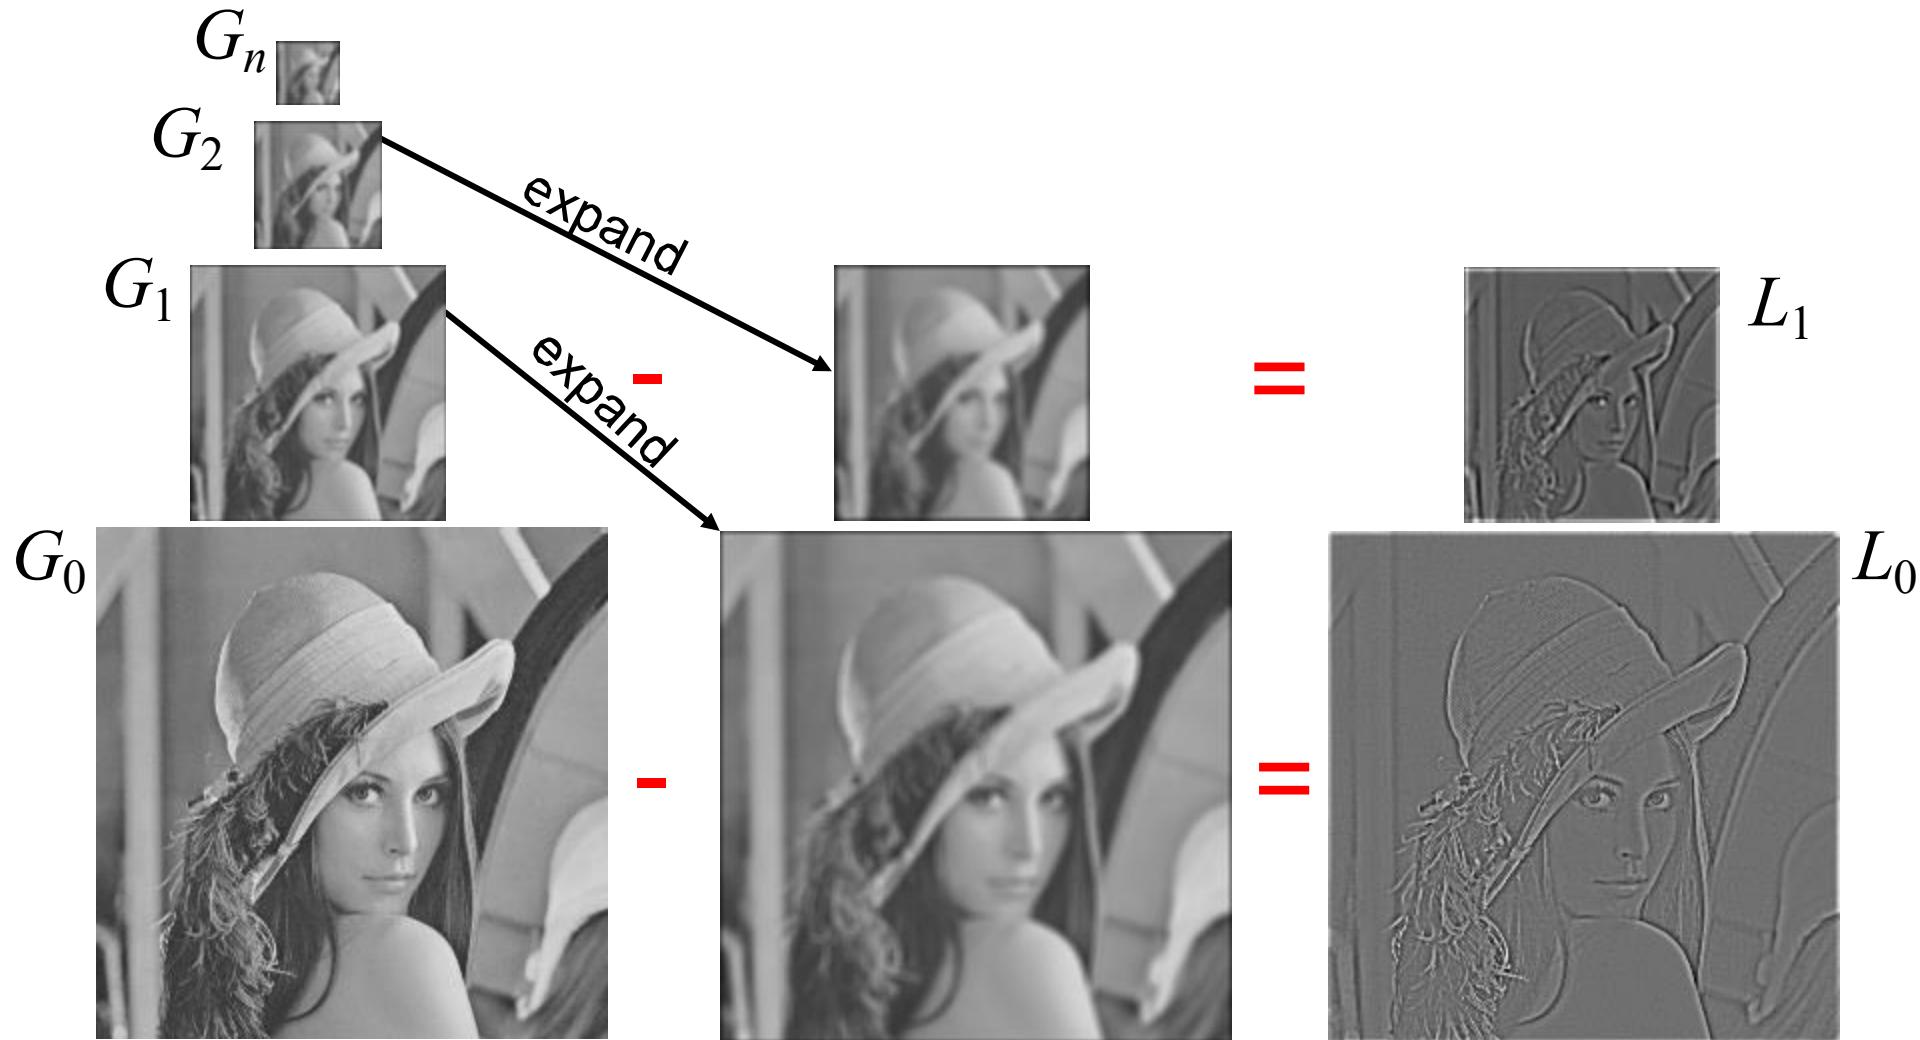
The Laplacian Pyramid

Gaussian Pyramid



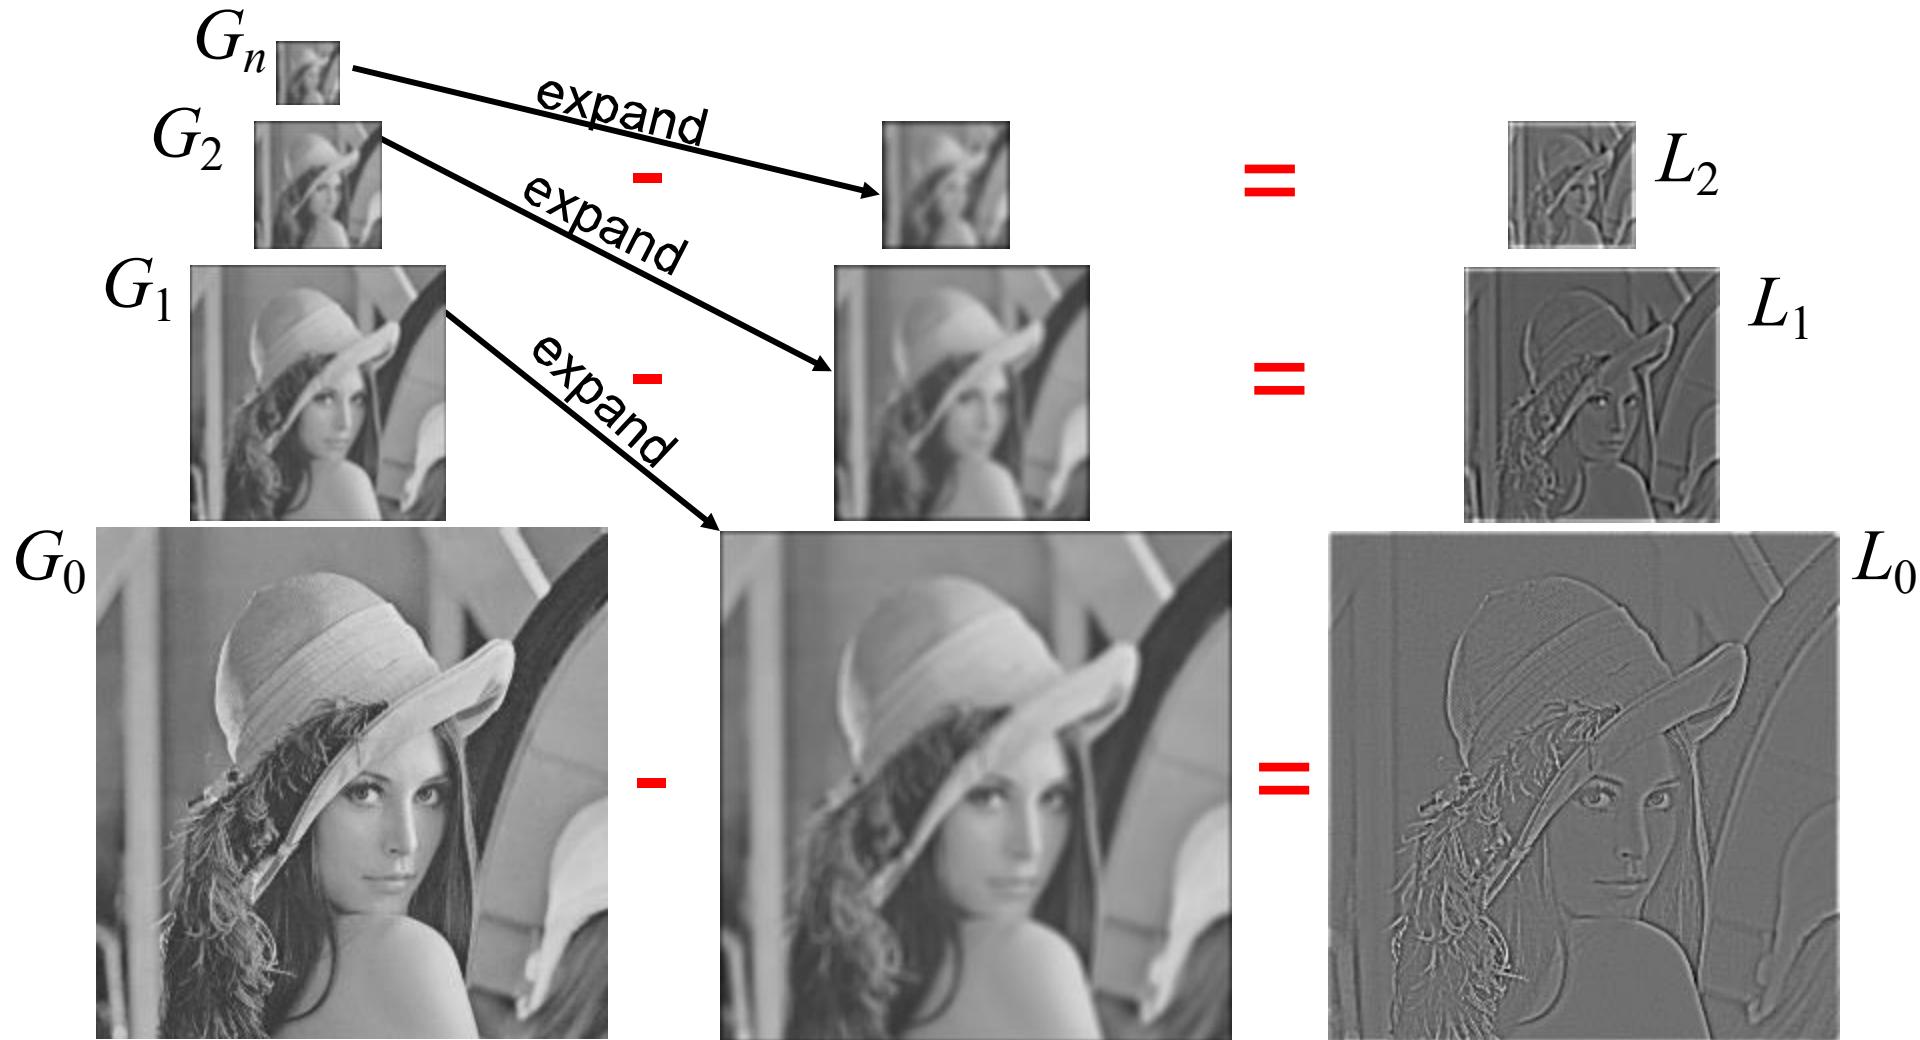
The Laplacian Pyramid

Gaussian Pyramid



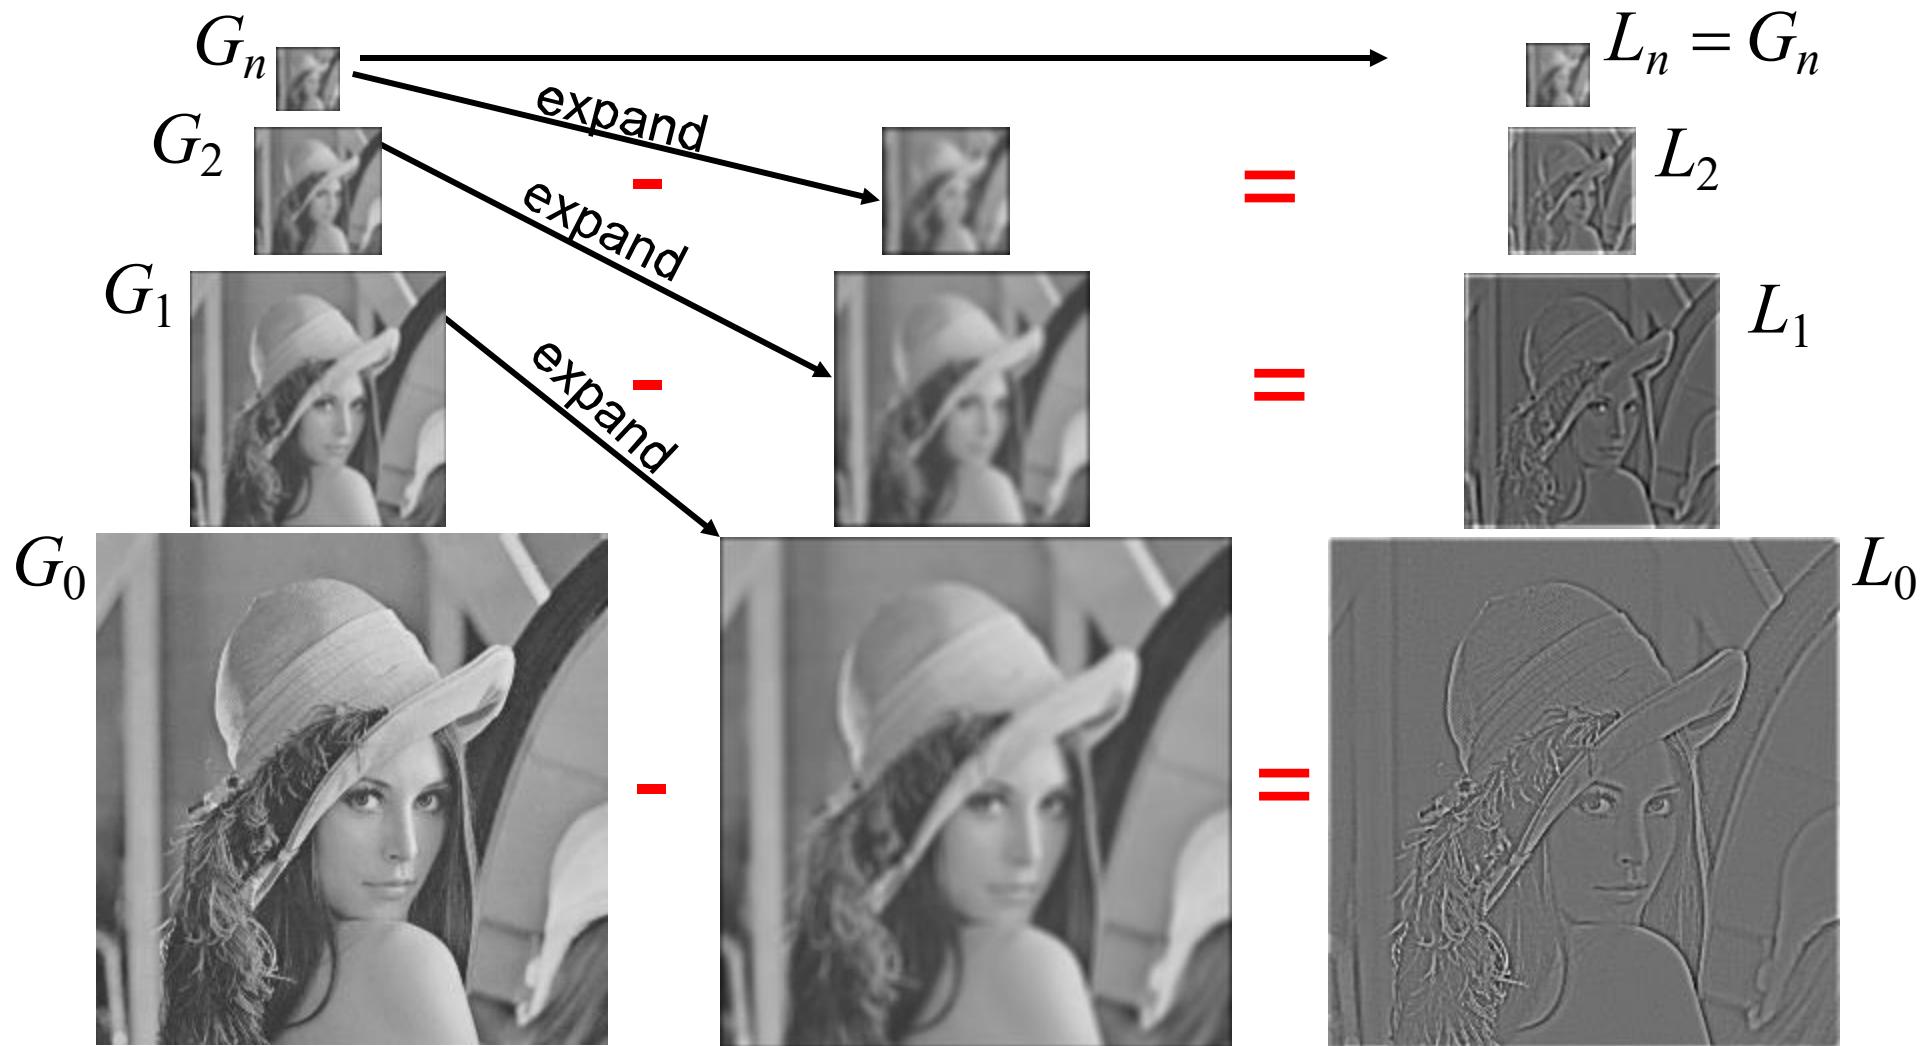
The Laplacian Pyramid

Gaussian Pyramid



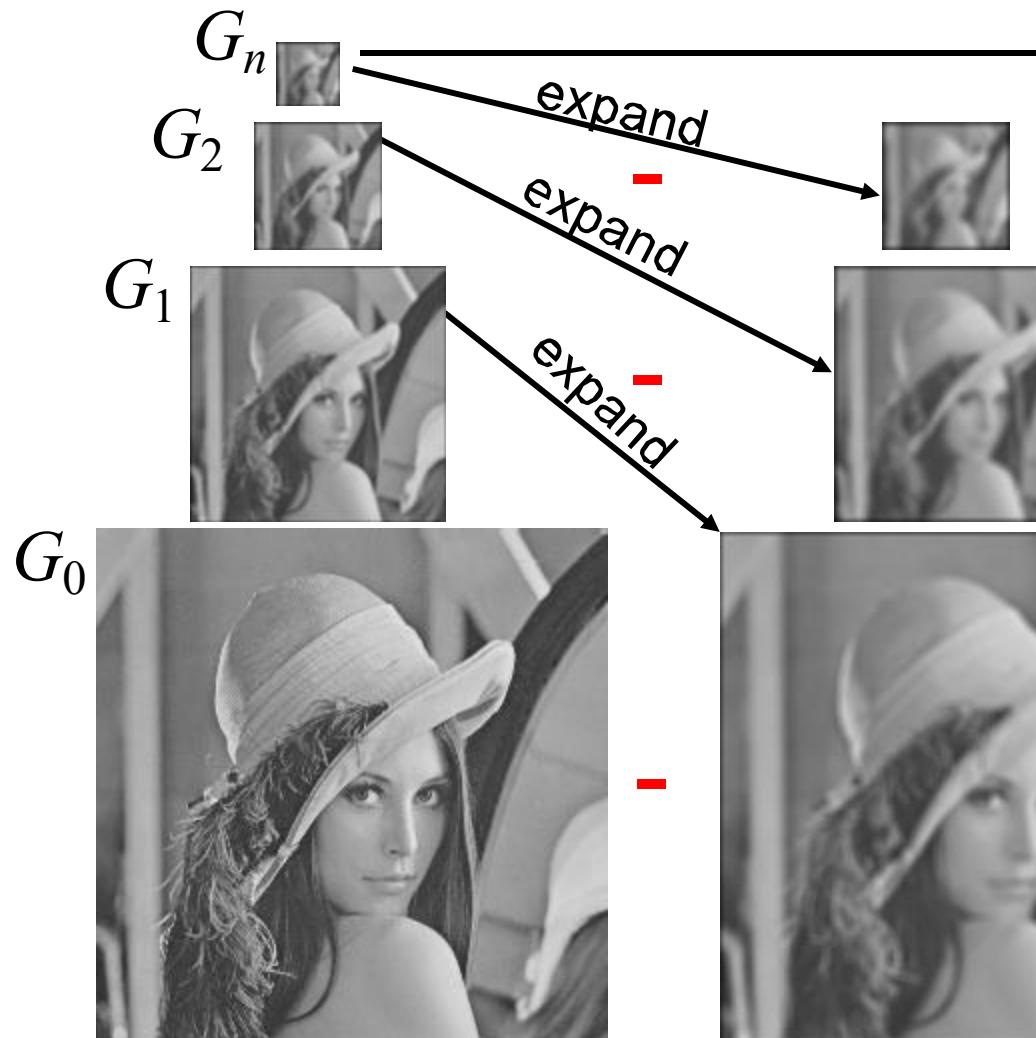
The Laplacian Pyramid

Gaussian Pyramid

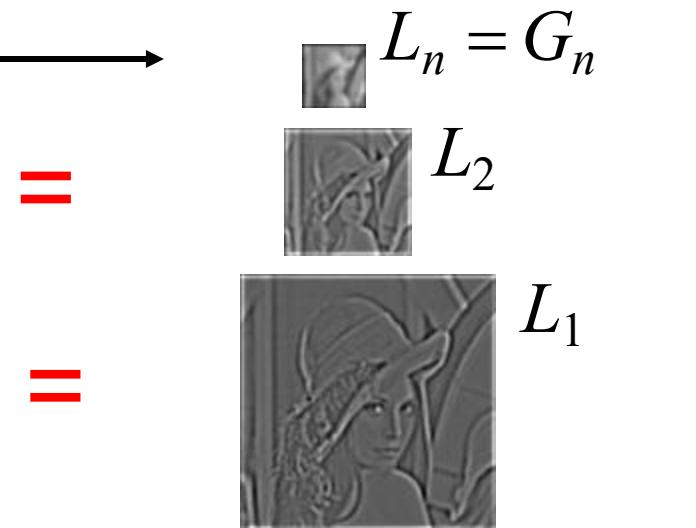


The Laplacian Pyramid

Gaussian Pyramid



Laplacian Pyramid



L_0

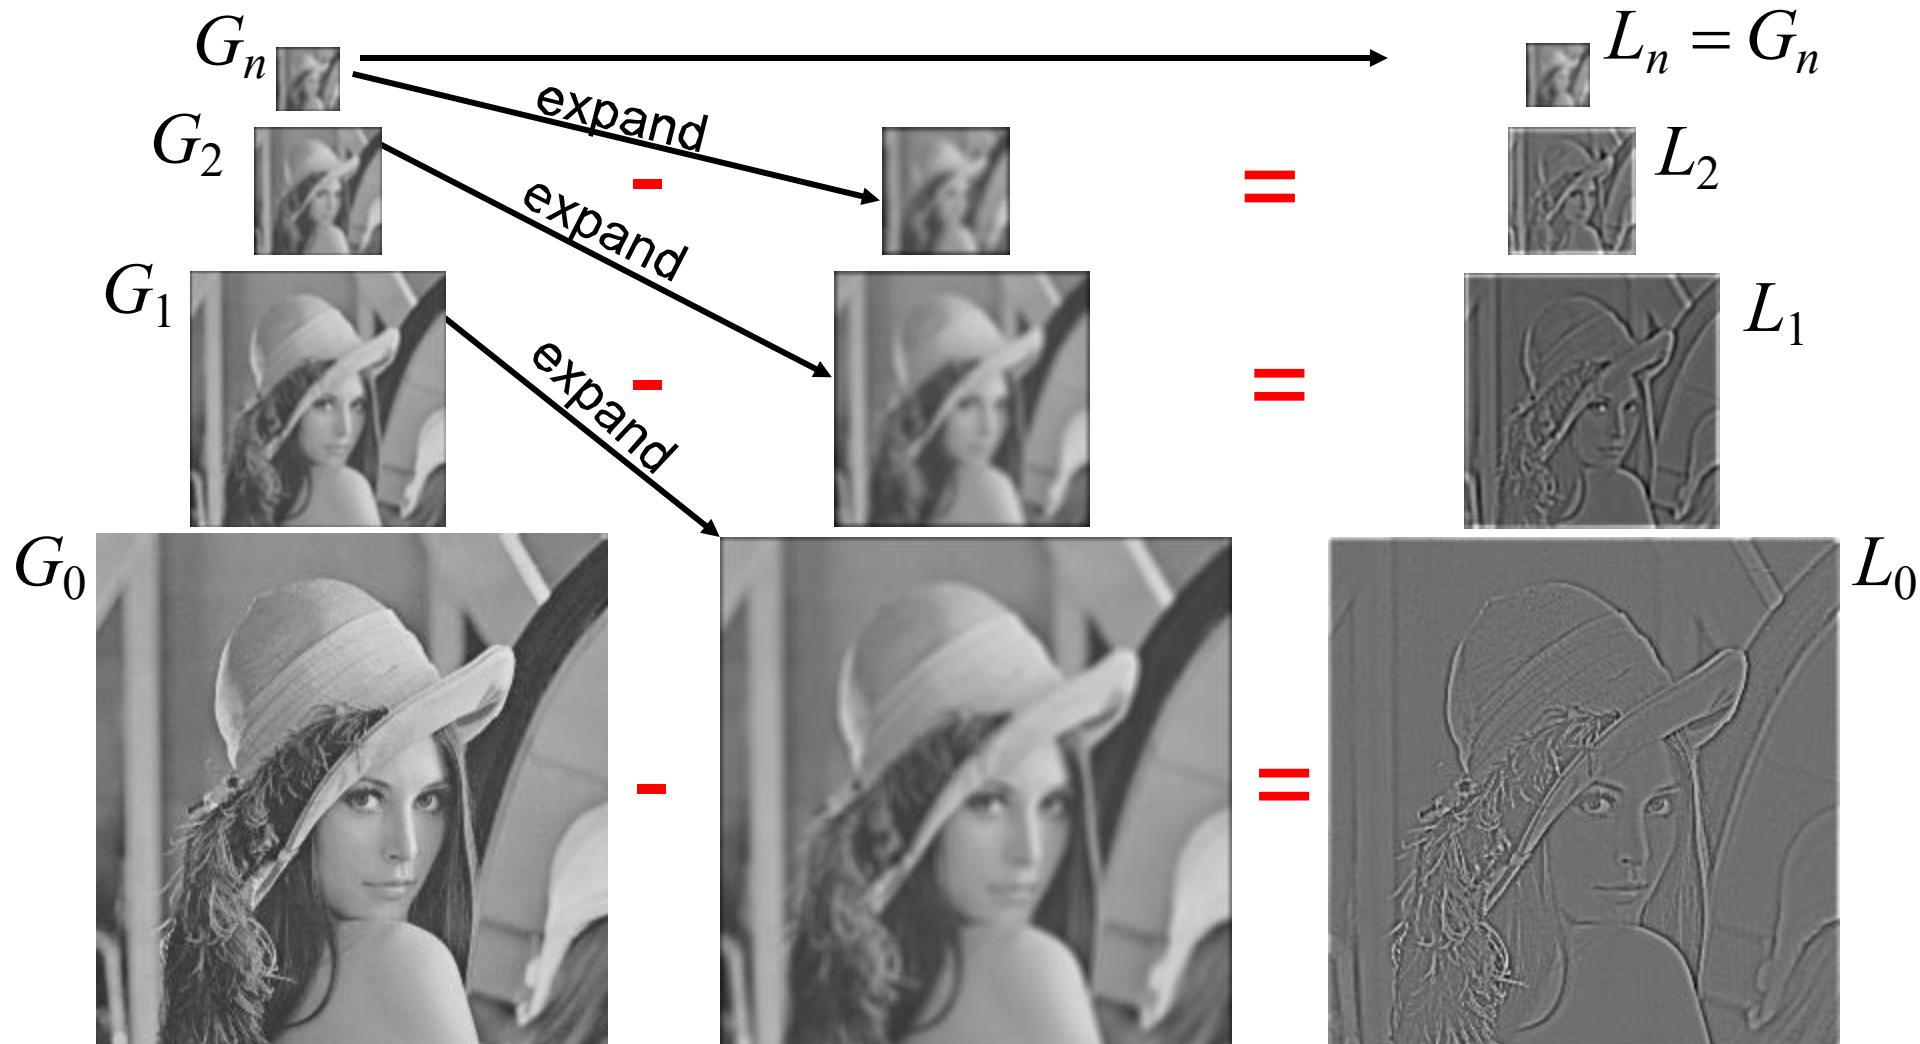
The Laplacian Pyramid

$$L_i = G_i - \text{expand}(G_{i+1})$$

Gaussian Pyramid

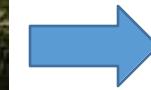
$$G_i = L_i + \text{expand}(G_{i+1})$$

Laplacian Pyramid

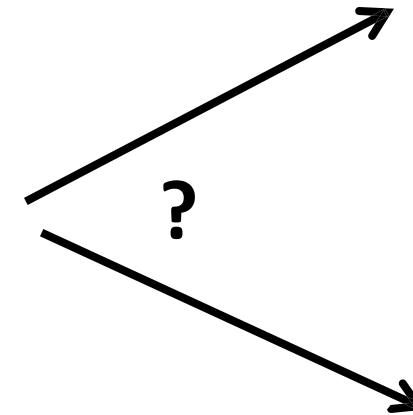


Why does a lower resolution image still make sense to us?

What do we lose?

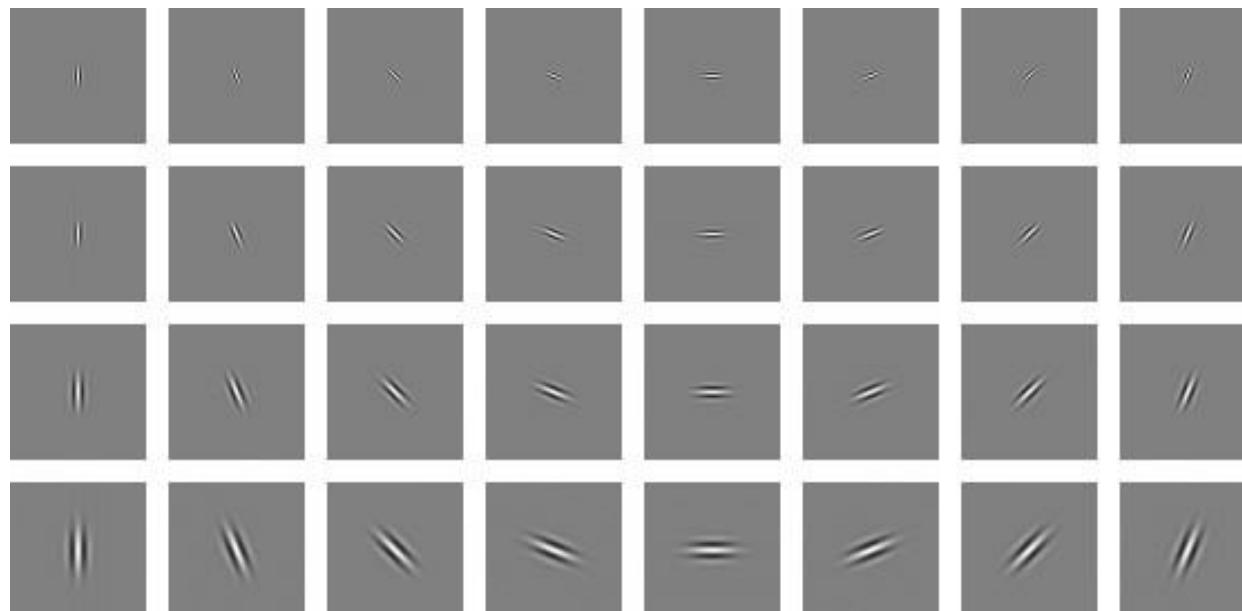


Why do we get different, distance-dependent interpretations of hybrid images?



Clues from Human Perception

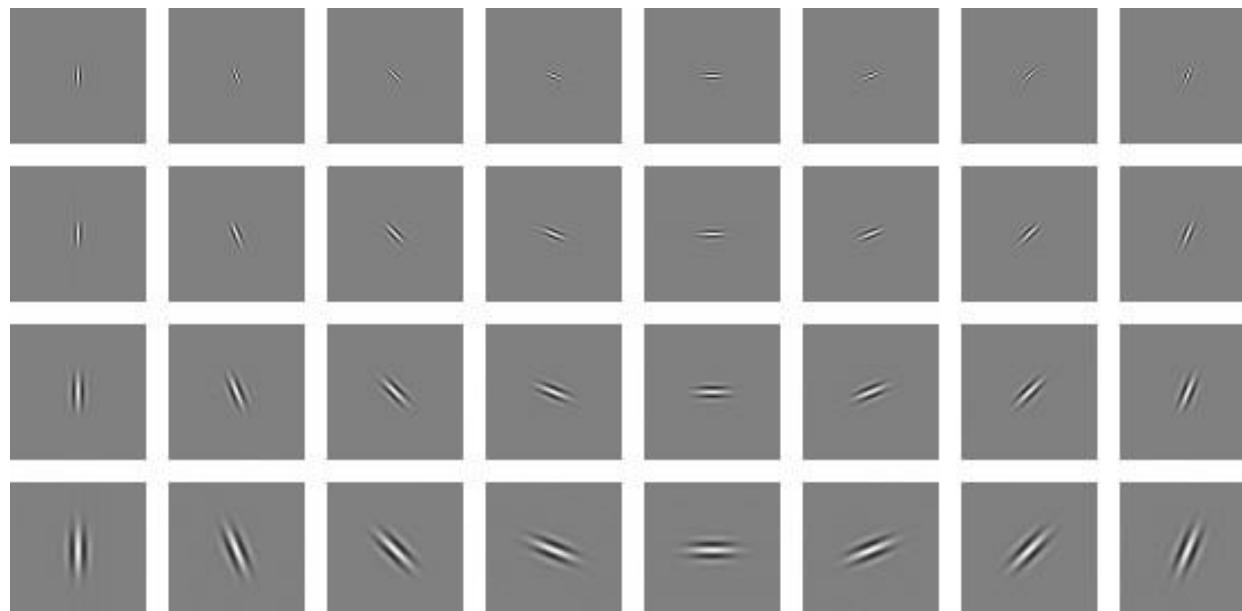
- Early processing in humans filters for various orientations and scales of frequency



Early Visual Processing: Multi-scale edge and blob filters

Clues from Human Perception

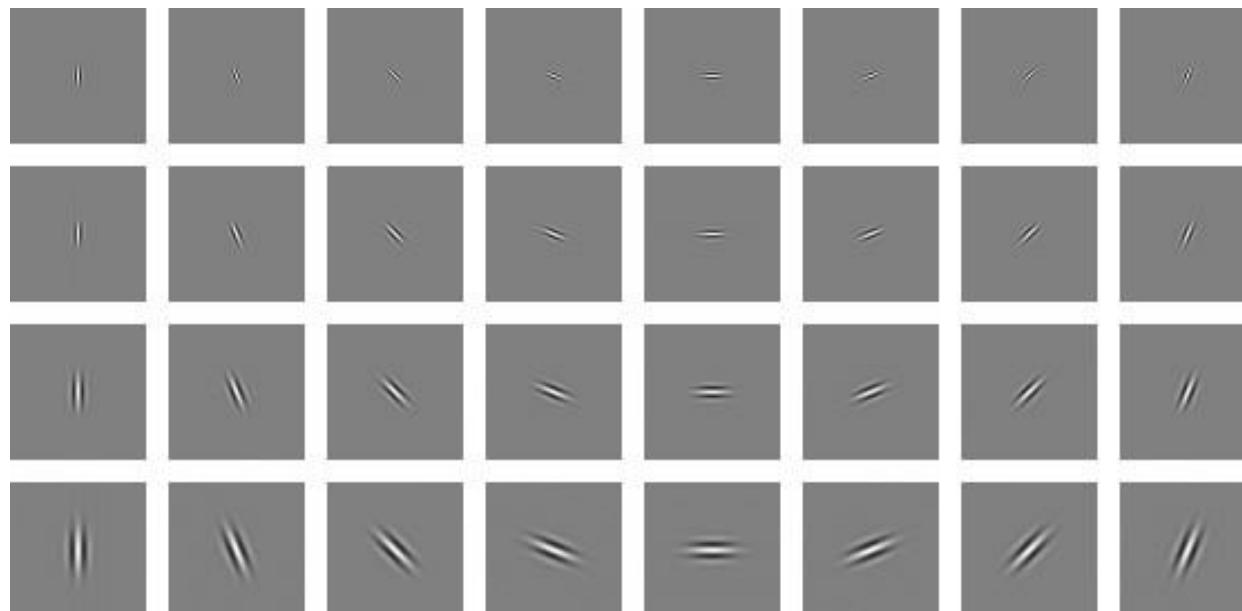
- Early processing in humans filters for various orientations and scales of frequency
- Perceptual cues in the mid-high frequencies dominate perception



Early Visual Processing: Multi-scale edge and blob filters

Clues from Human Perception

- Early processing in humans filters for various orientations and scales of frequency
- Perceptual cues in the mid-high frequencies dominate perception
- When we see an image from far away, we are effectively subsampling it



Early Visual Processing: Multi-scale edge and blob filters

Upsampling

- This image is too small for this screen:
- How can we make it 10 times as big?



Upsampling

- This image is too small for this screen:
- How can we make it 10 times as big?
- Simplest approach:
repeat each row
and column 10 times
- (“Nearest neighbor
interpolation”)



Image interpolation

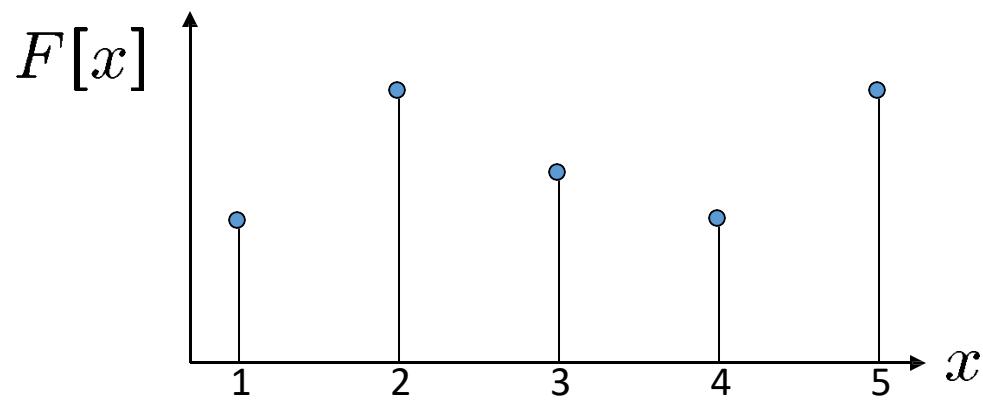
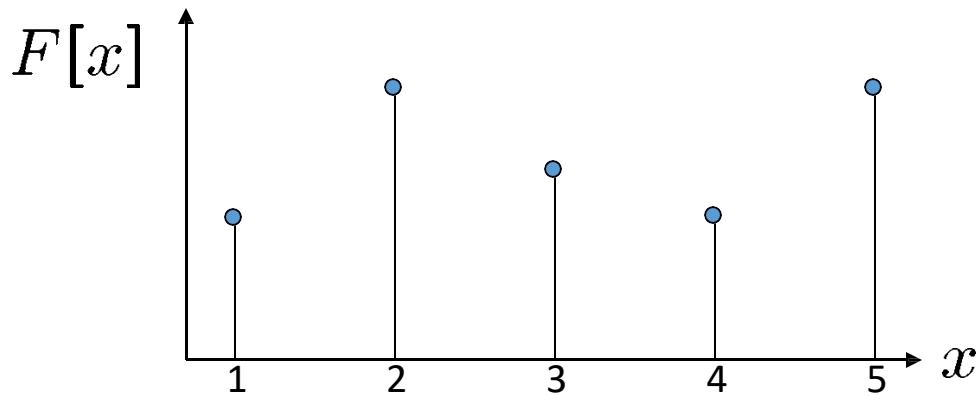
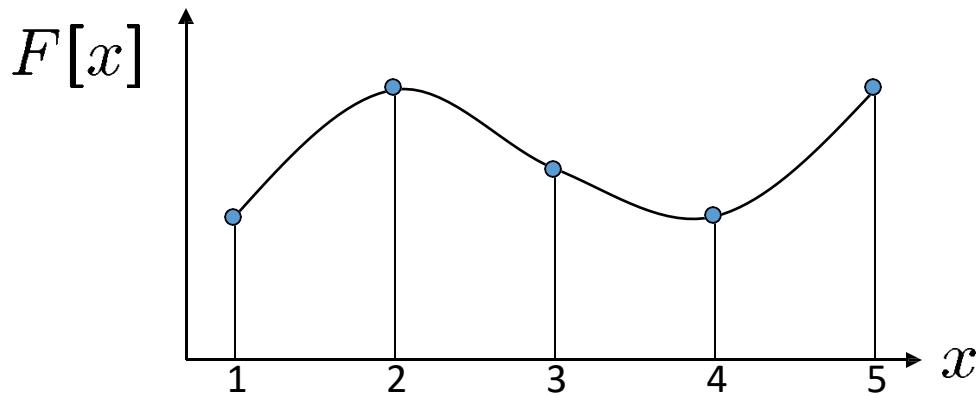


Image interpolation



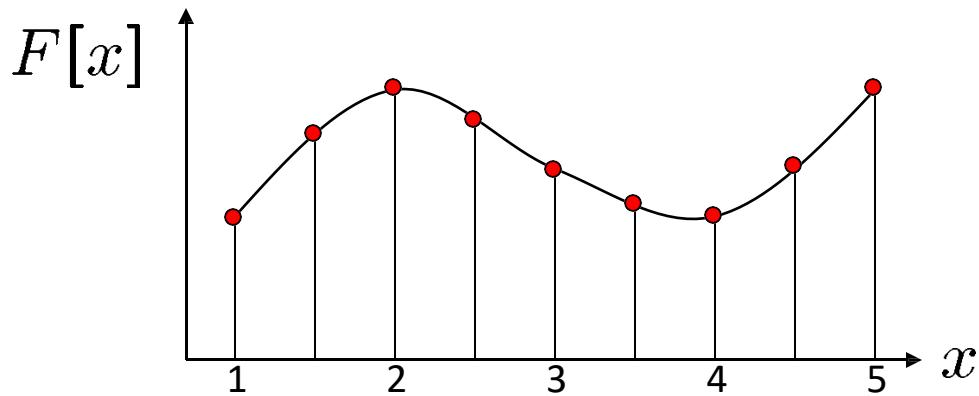
- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

Image interpolation



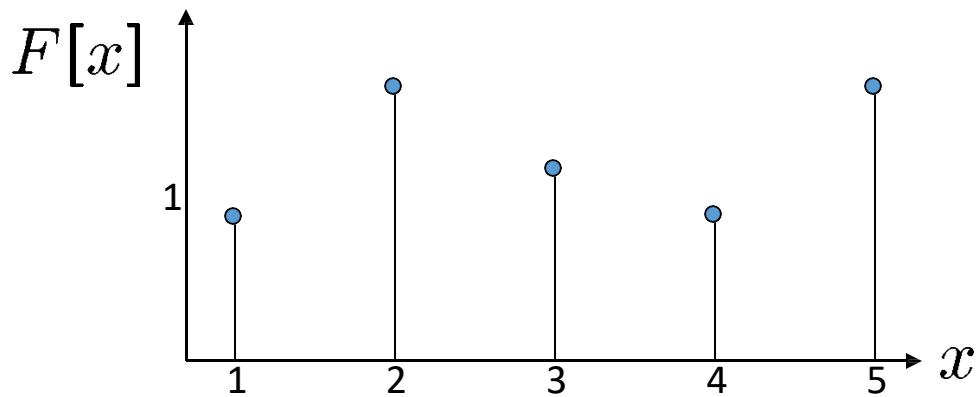
- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

Image interpolation



- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

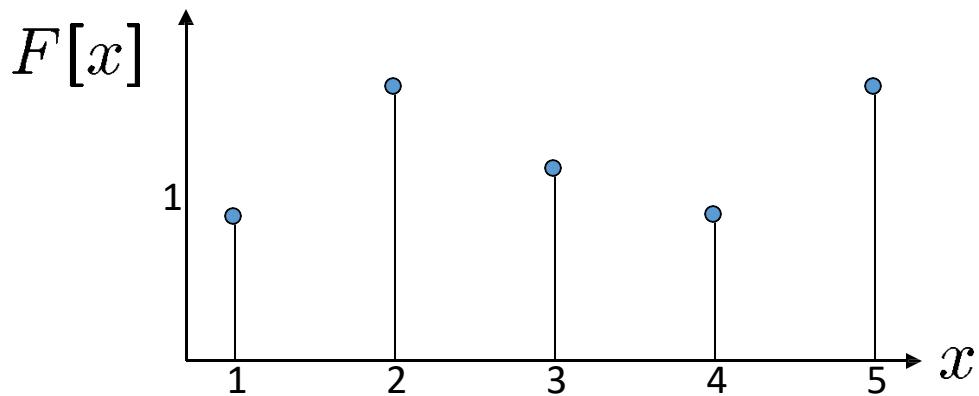
Image interpolation



$d = 1$ in this example

- What if we don't know f ?

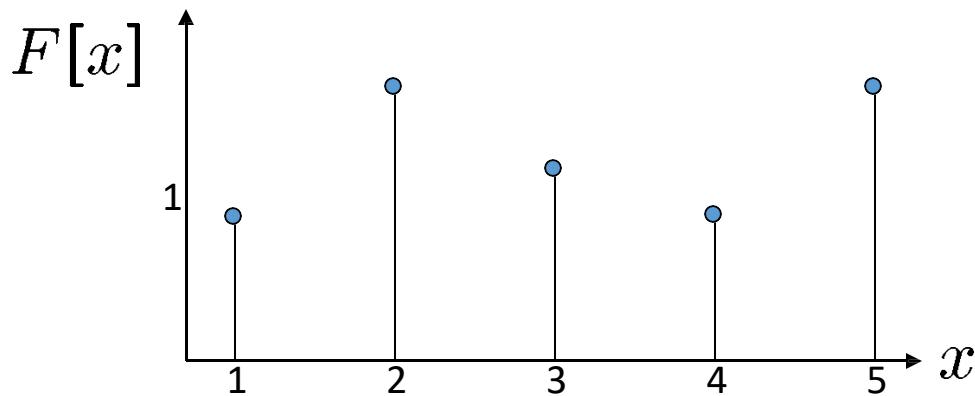
Image interpolation



$d = 1$ in this example

- What if we don't know f ?
 - Guess an approximation: \tilde{f}
 - Can be done in a principled way: filtering

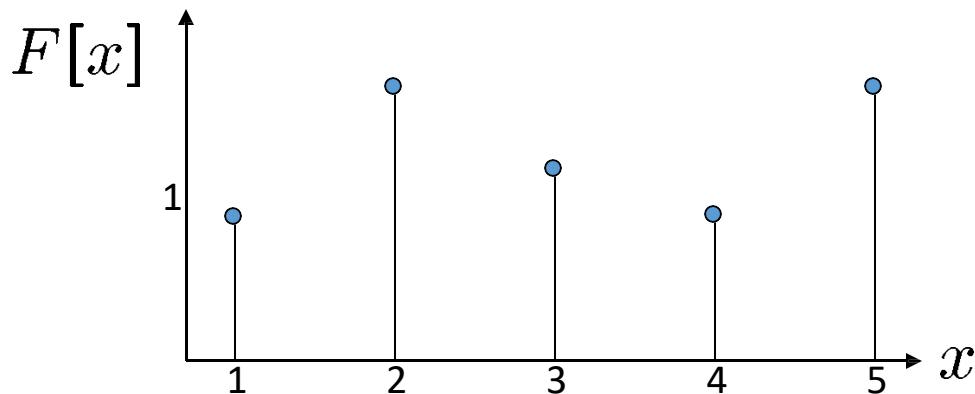
Image interpolation



$d = 1$ in this example

- What if we don't know f ?
 - Guess an approximation: \tilde{f}
 - Can be done in a principled way: filtering
 - Convert F to a continuous function:
$$f_F(x) = F\left(\frac{x}{d}\right) \text{ when } \frac{x}{d} \text{ is an integer, 0 otherwise}$$

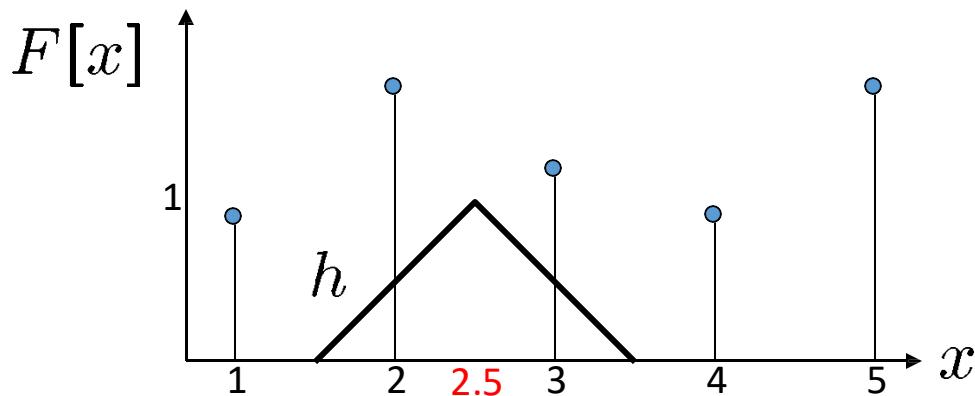
Image interpolation



$d = 1$ in this example

- What if we don't know f ?
 - Guess an approximation: \tilde{f}
 - Can be done in a principled way: filtering
 - Convert F to a continuous function:
$$f_F(x) = F\left(\frac{x}{d}\right) \text{ when } \frac{x}{d} \text{ is an integer, 0 otherwise}$$
 - Reconstruct by convolution with a *reconstruction filter*, h
$$\tilde{f} = h * f_F$$

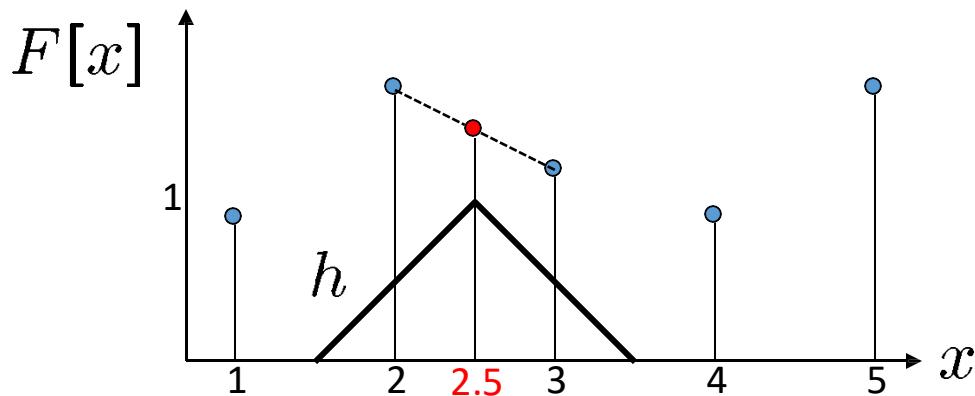
Image interpolation



$d = 1$ in this example

- What if we don't know f ?
 - Guess an approximation: \tilde{f}
 - Can be done in a principled way: filtering
 - Convert F to a continuous function:
$$f_F(x) = F\left(\frac{x}{d}\right) \text{ when } \frac{x}{d} \text{ is an integer, 0 otherwise}$$
 - Reconstruct by convolution with a *reconstruction filter*, h
$$\tilde{f} = h * f_F$$

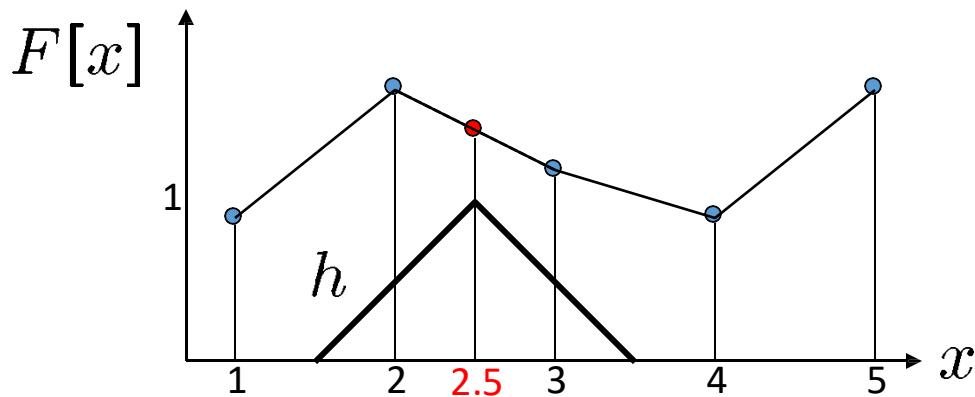
Image interpolation



$d = 1$ in this example

- What if we don't know f ?
 - Guess an approximation: \tilde{f}
 - Can be done in a principled way: filtering
 - Convert F to a continuous function:
$$f_F(x) = F\left(\frac{x}{d}\right) \text{ when } \frac{x}{d} \text{ is an integer, 0 otherwise}$$
 - Reconstruct by convolution with a *reconstruction filter*, h
$$\tilde{f} = h * f_F$$

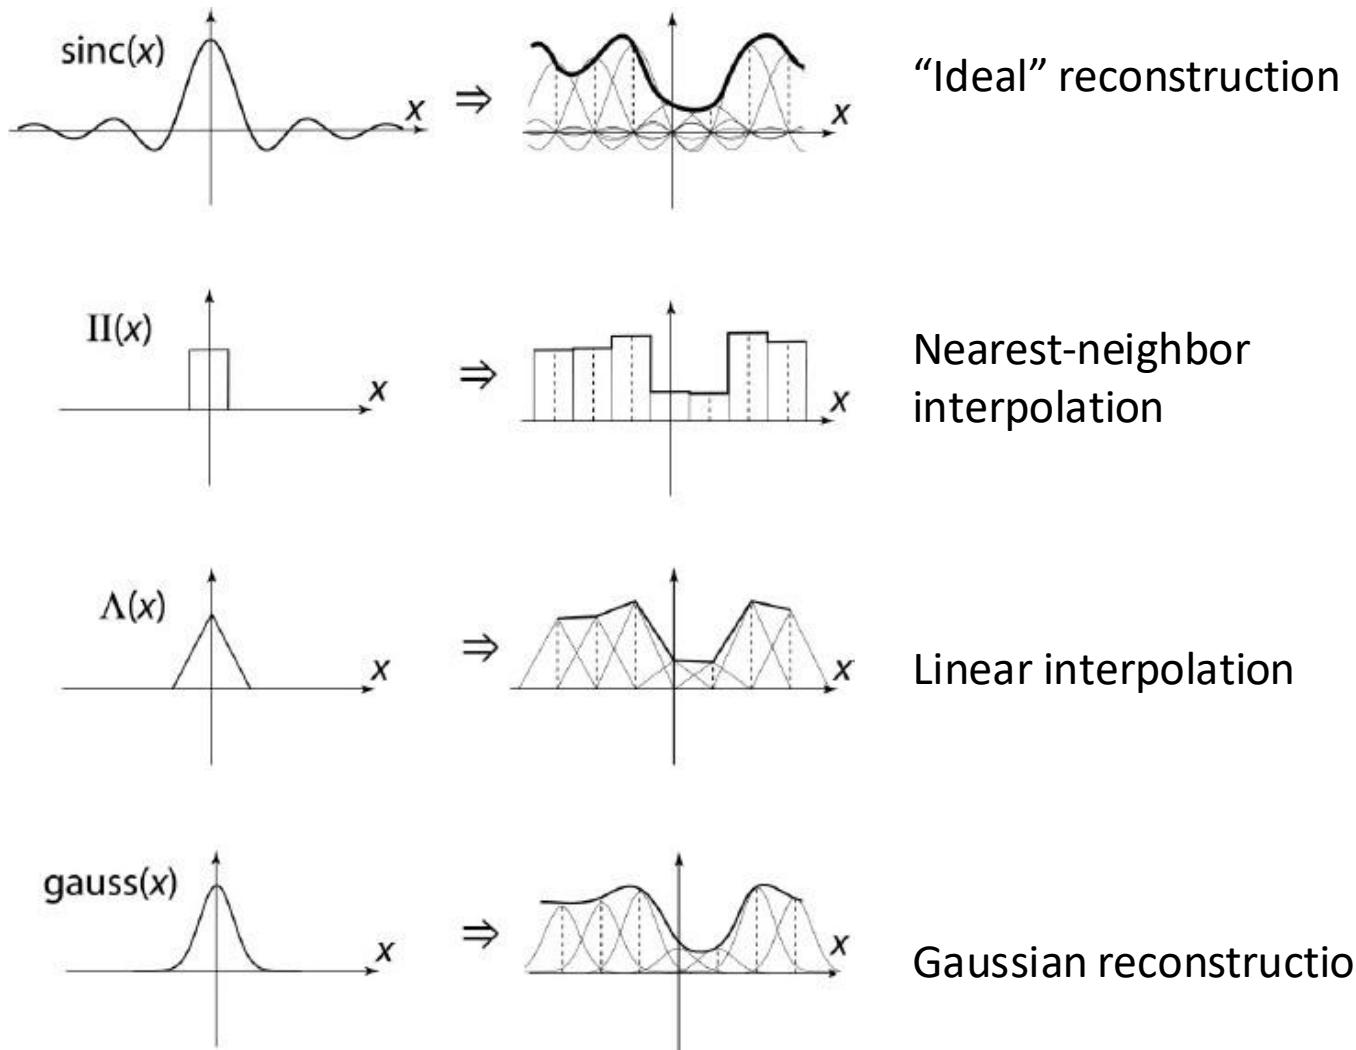
Image interpolation



$d = 1$ in this example

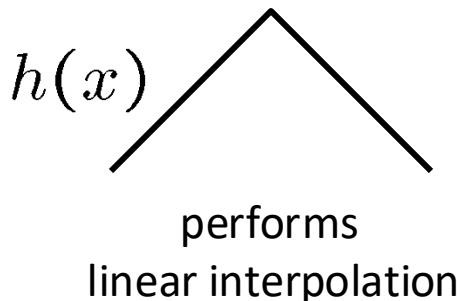
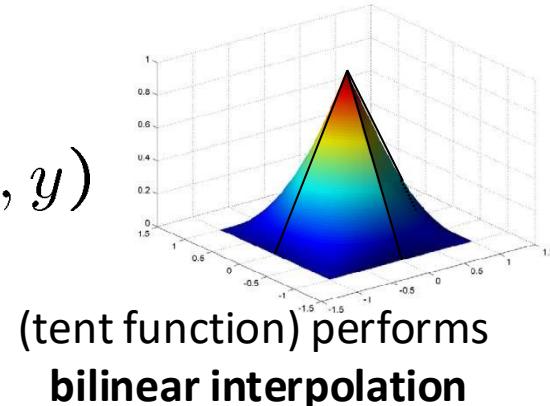
- What if we don't know f ?
 - Guess an approximation: \tilde{f}
 - Can be done in a principled way: filtering
 - Convert F to a continuous function:
$$f_F(x) = F\left(\frac{x}{d}\right) \text{ when } \frac{x}{d} \text{ is an integer, 0 otherwise}$$
 - Reconstruct by convolution with a *reconstruction filter or interpolation kernel*, h
$$\tilde{f} = h * f_F$$

Image interpolation



Reconstruction filters

- What does the 2D version of this hat function look like?


$$h(x, y)$$


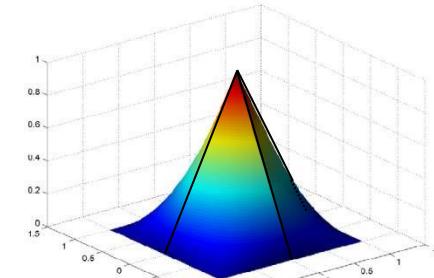
Reconstruction filters

- What does the 2D version of this hat function look like?

$$h(x)$$

$$h(x, y)$$

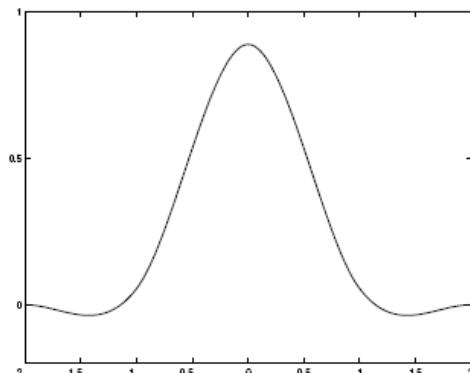
performs
linear interpolation



(tent function) performs
bilinear interpolation

Better filters give better resampled images

- **Bicubic** is common choice



$$r(x) = \begin{cases} \frac{1}{6} \left[(12 - 9B - 6C)|x|^3 + (-18 + 12B + 6C)|x|^2 + (6 - 2B) \right] & |x| < 1 \\ \frac{1}{6} \left[((-B - 6C)|x|^3 + (6B + 30C)|x|^2 + (-12B - 48C)|x| + (8B + 24C) \right] & 1 \leq |x| < 2 \\ 0 & \text{otherwise} \end{cases}$$

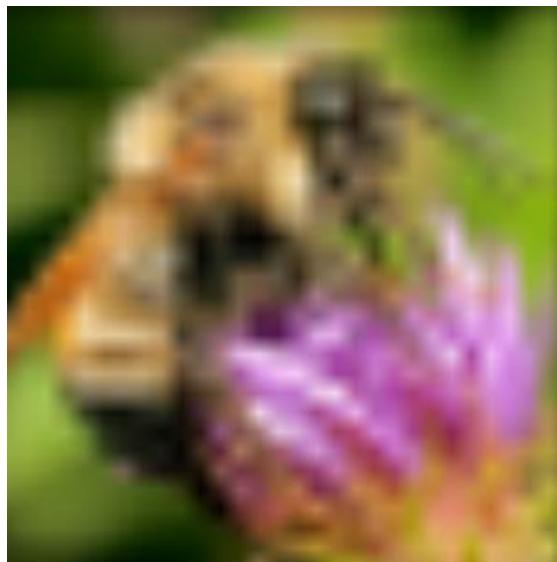
Cubic reconstruction filter

Image interpolation

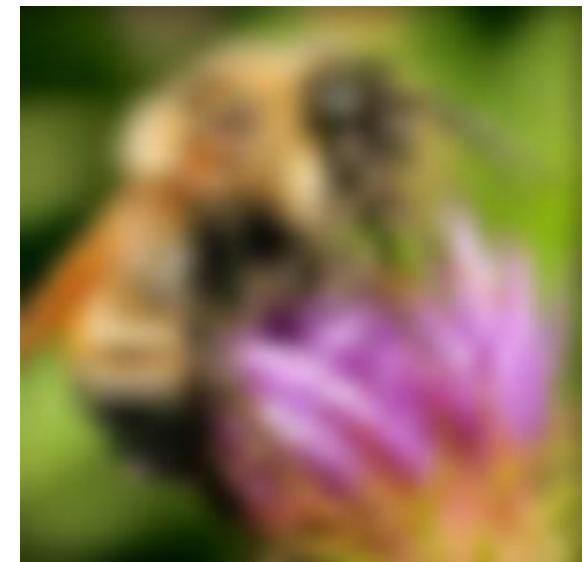
Original image:  x 10



Nearest-neighbor interpolation

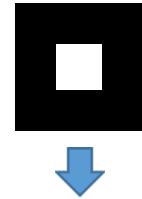


Bilinear interpolation



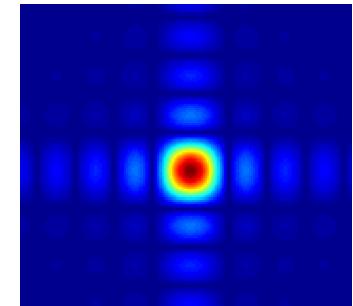
Bicubic interpolation

Things to Remember

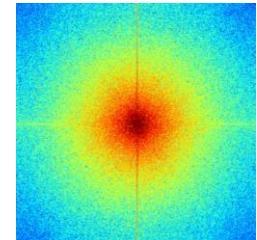


- Sometimes it makes sense to think of images and filtering in the frequency domain

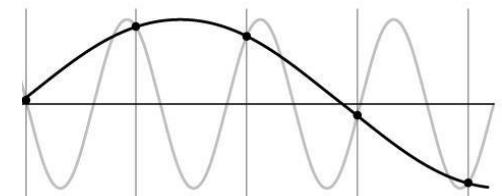
- Fourier analysis



- Can be faster to filter using FFT for large images ($N \log N$ vs. N^2 for auto-correlation)



- Remember to low-pass before sampling



Acknowledgements

- Thanks to the following researchers for making their teaching/research material online
 - Forsyth
 - Steve Seitz
 - Noah Snavely
 - J.B. Huang
 - Derek Hoiem
 - D. Lowe
 - A. Bobick
 - S. Lazebnik
 - K. Grauman
 - R. Zaleski

Thank you

- Next class:
 - Edge Detection

