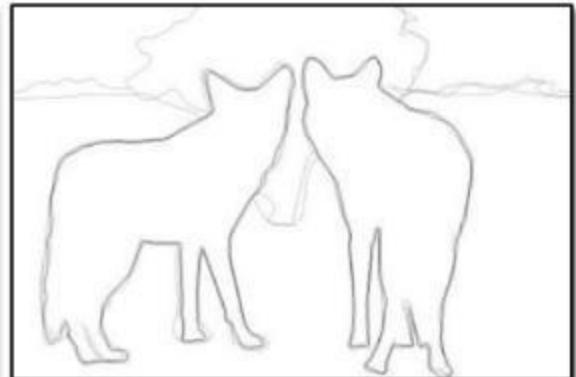
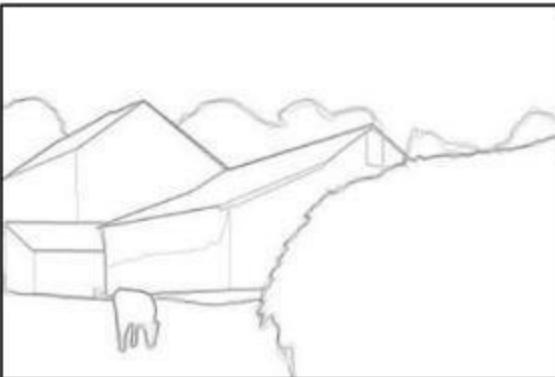
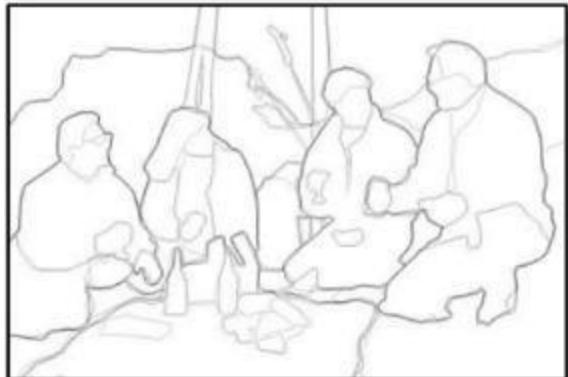


Edge Detection



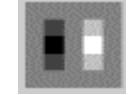
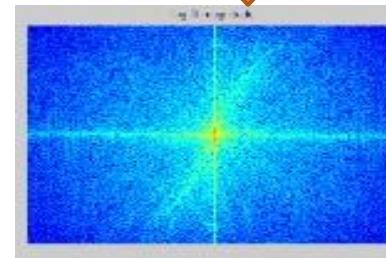
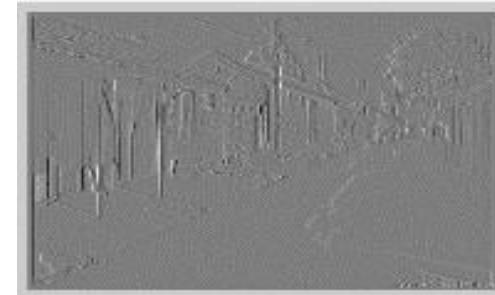
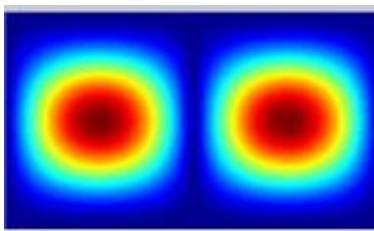
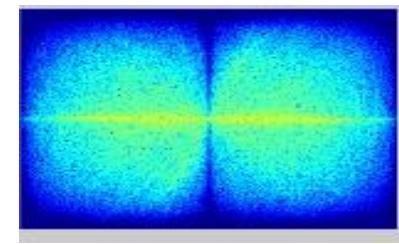
Computer Vision

Adduru U G Sankararao, IIIT Sri City

Previous two classes: Image Filtering

Spatial domain

- Smoothing, sharpening, measuring texture

 $*$  $=$  \times  $=$ 

Frequency domain

- Denoising, sampling

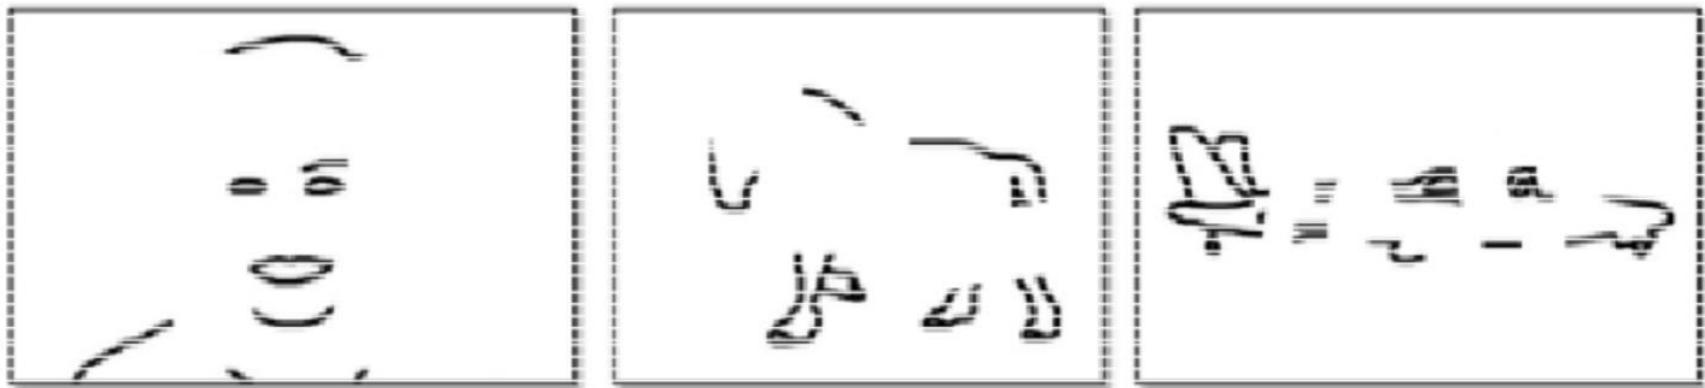
Today's class

- Detecting edges
- Finding straight lines

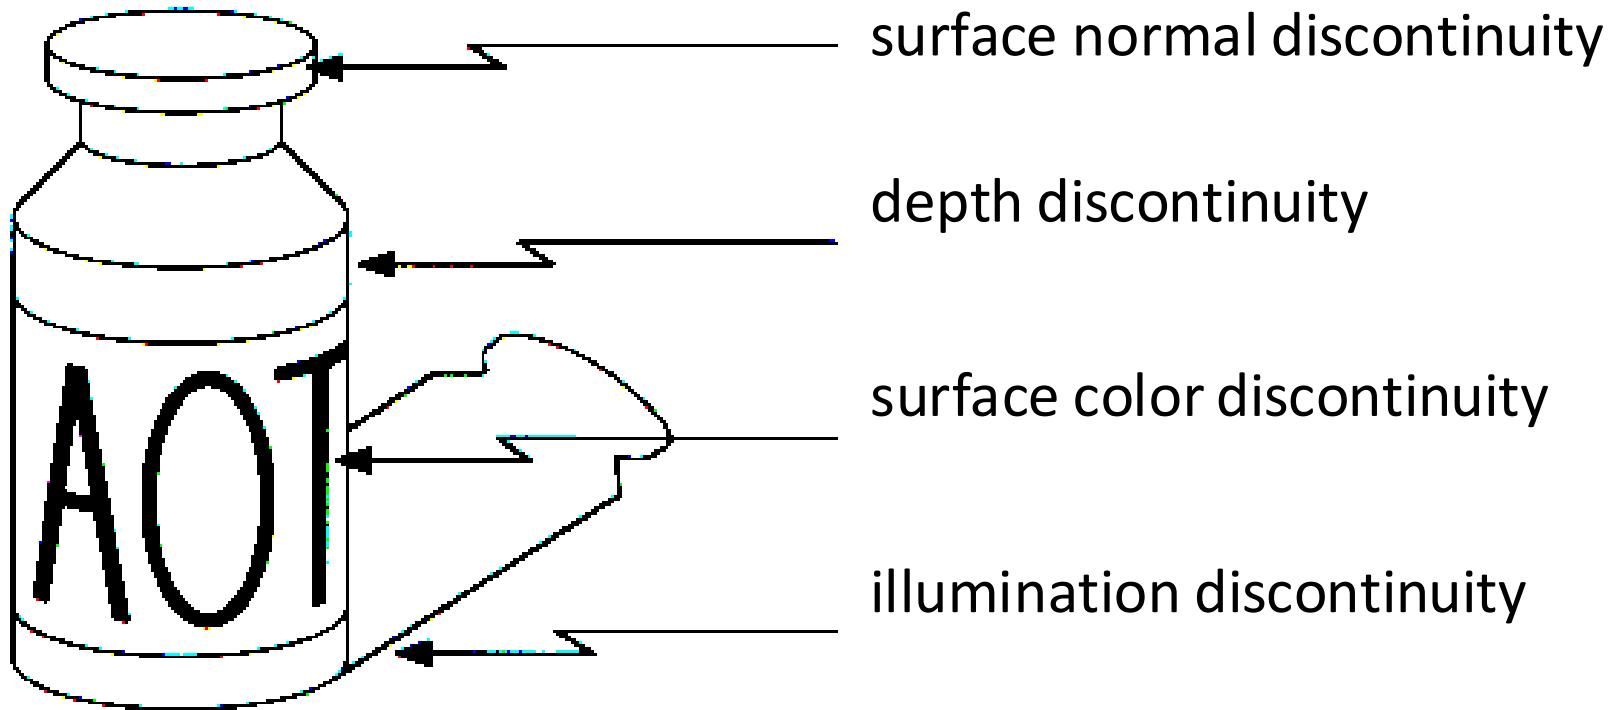


Edge Detection

- Map image from 2D matrix of pixels to a set of curves or line segments or contours => More compact representation than pixels
- Key idea: Look for strong gradients, and then post-process



How edges are formed



Edges are caused by a variety of factors

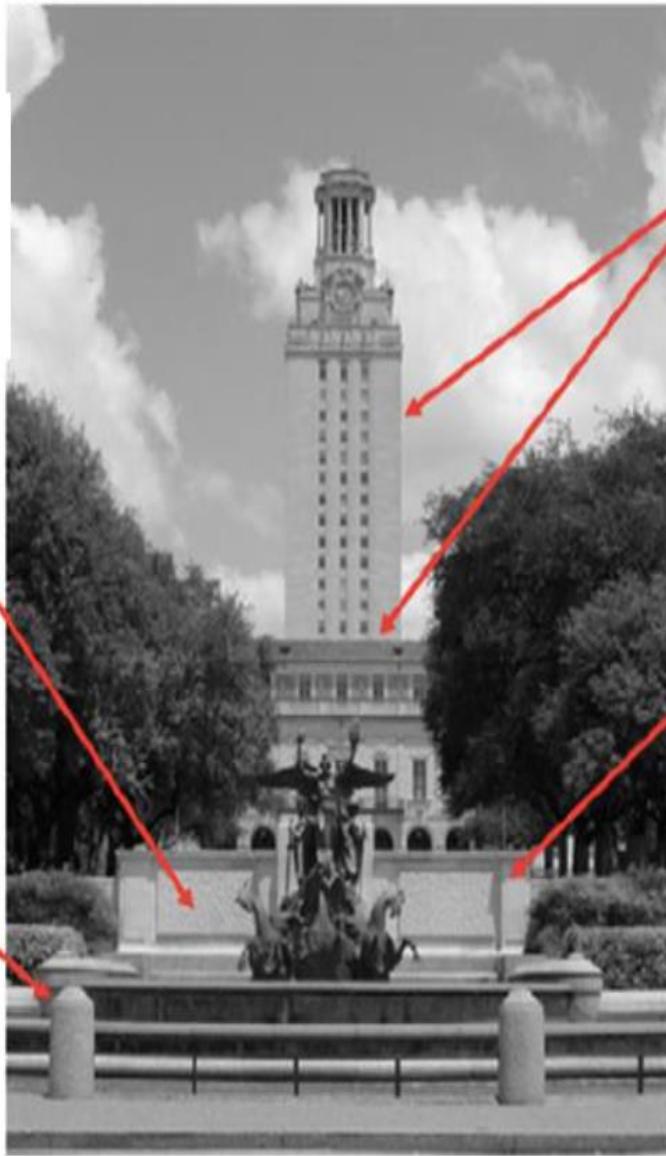
How edges are formed

Surface color/appearance
discontinuity

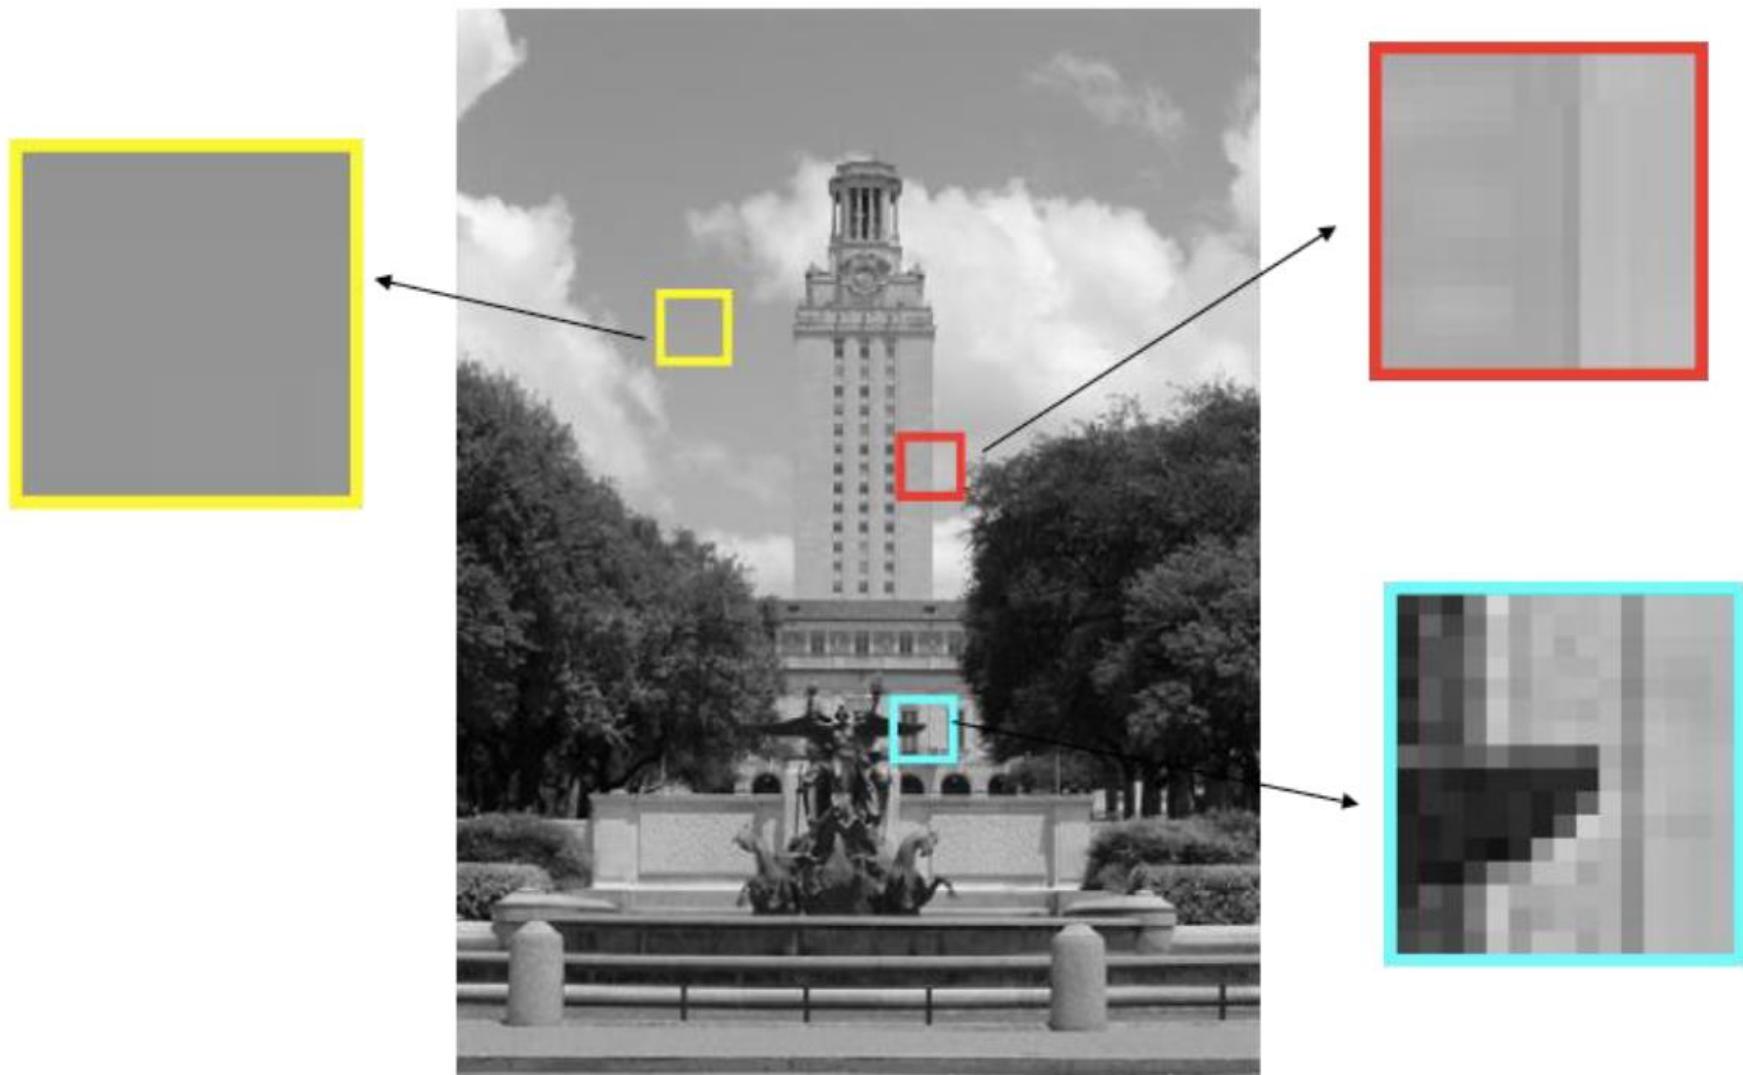
Surface normal
discontinuity

Depth discontinuity

Illumination discontinuity

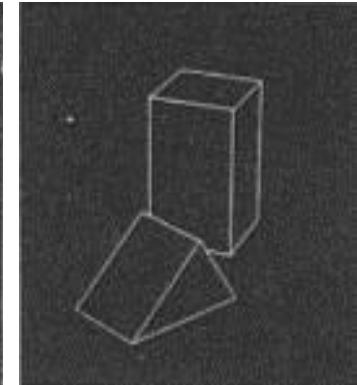
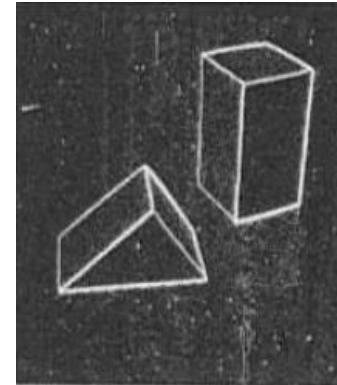
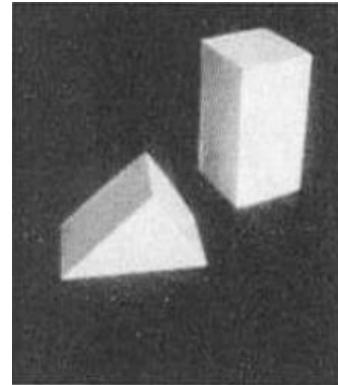


Looking more locally



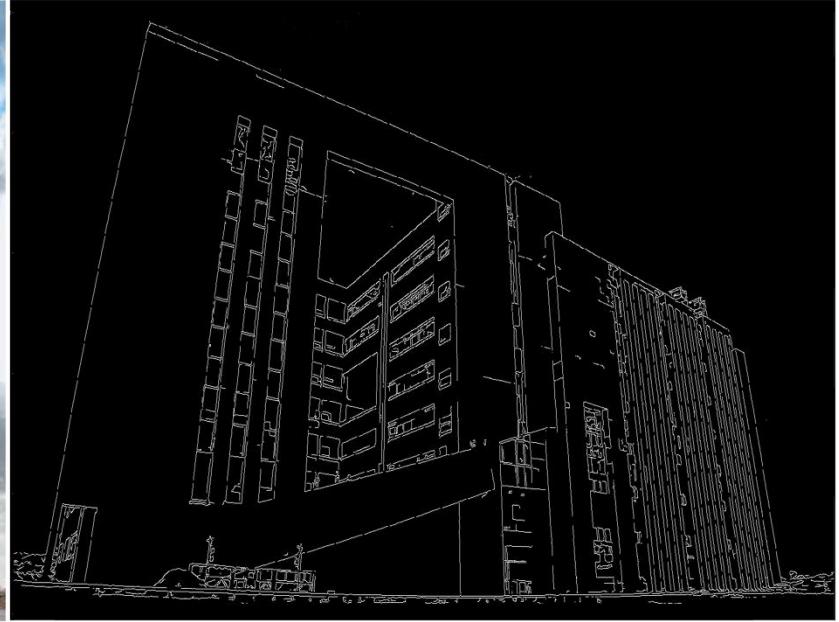
Why are Edges Important?

- Group pixels into objects or parts
- Allow us to track important features (e.g., corners, curves, lines).
- Cues for 3D shape
- Guiding interactive image editing



Why are Edges Important?

- Group pixels into objects or parts
- Allow us to track important features (e.g., corners, curves, lines).
- Cues for 3D shape
- Guiding interactive image editing



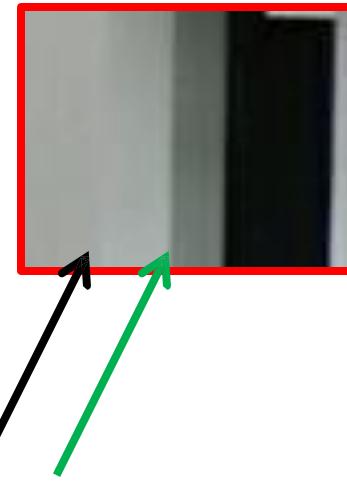
Closeup of edges



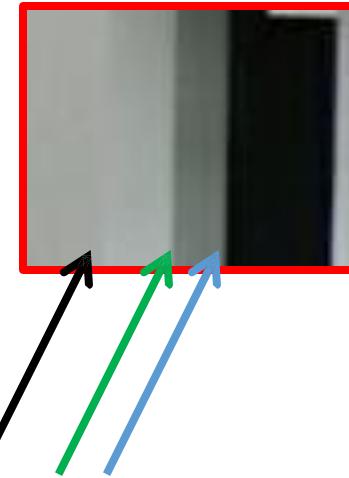
Closeup of edges



Closeup of edges



Closeup of edges



Closeup of edges

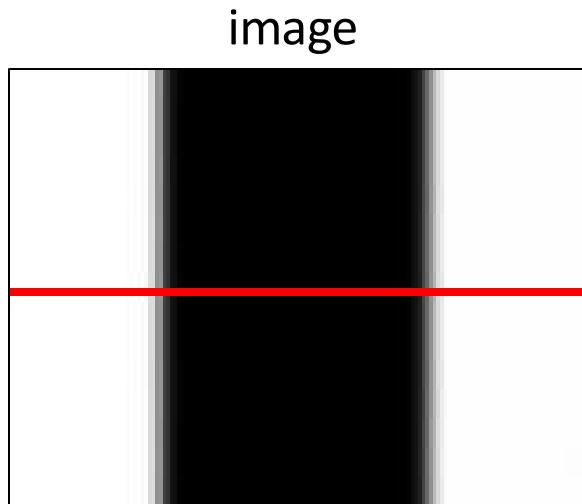


Closeup of edges



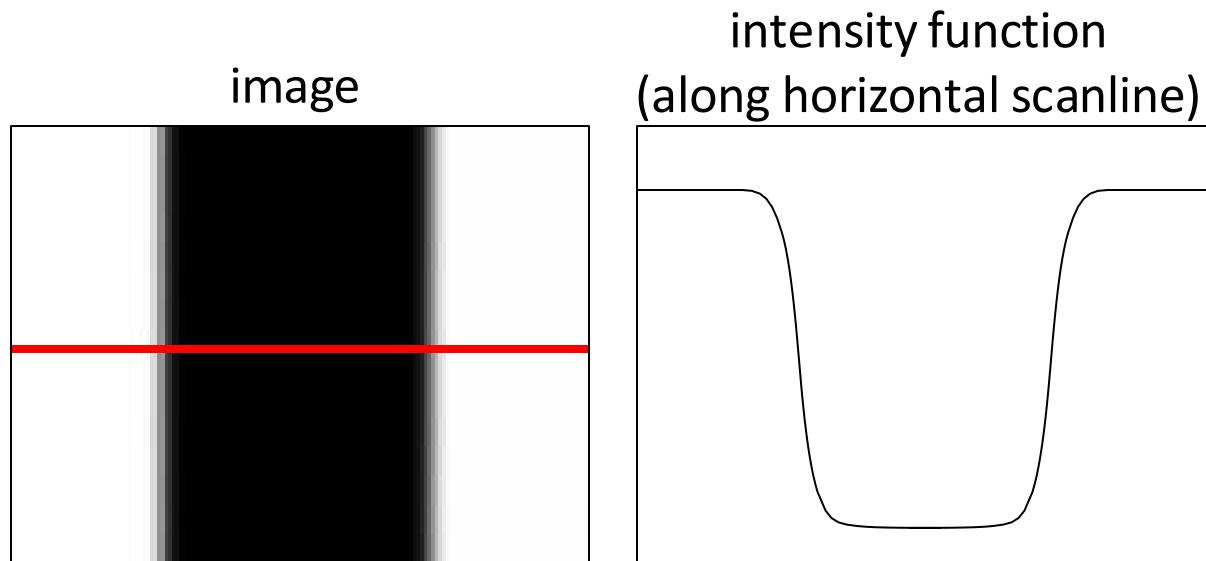
Characterizing edges

- An edge is a place of rapid change in the image intensity function



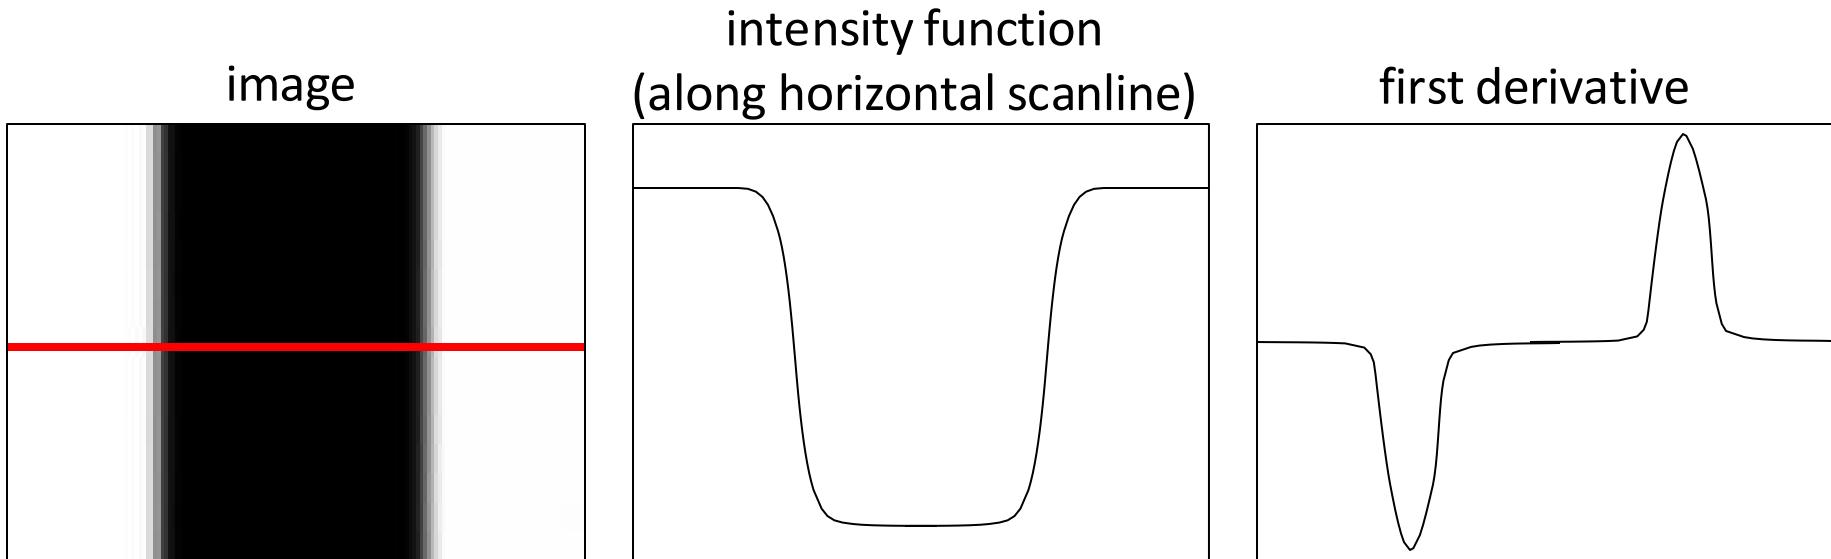
Derivatives & edges

- An edge is a place of rapid change in the image intensity function



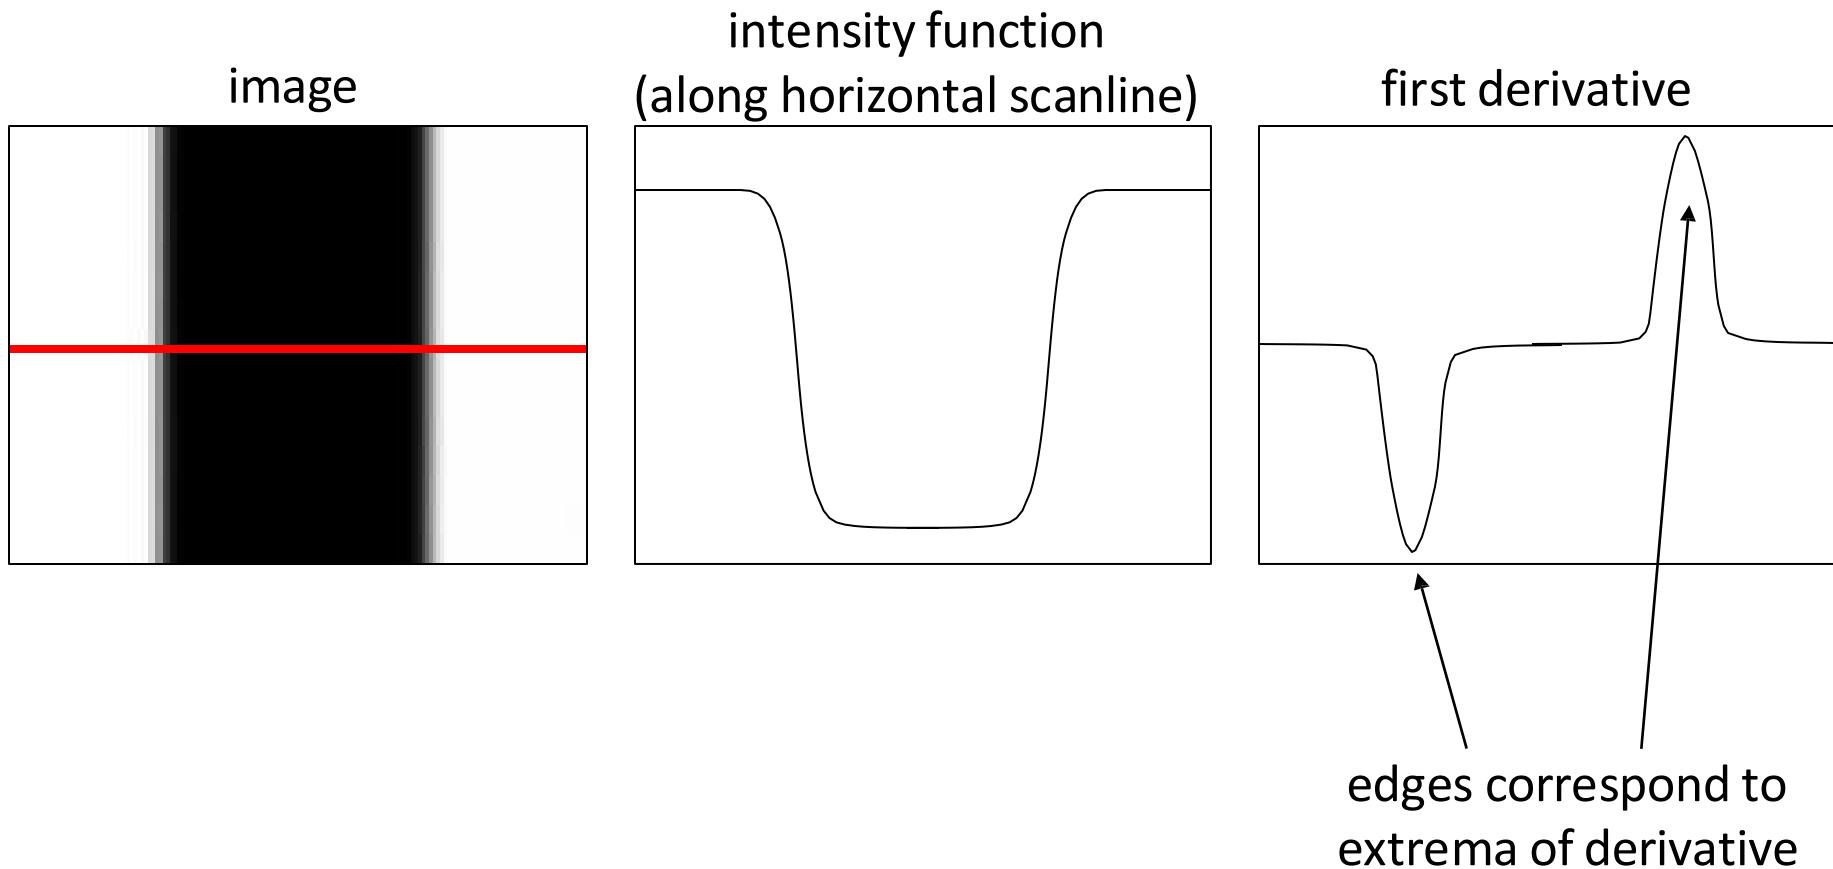
Derivatives & edges

- An edge is a place of rapid change in the image intensity function



Derivatives & edges

- An edge is a place of rapid change in the image intensity function



Derivatives with Convolution

- For 2D function, $f(x, y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement above as convolution, what would be the associated filter?

Derivatives with Convolution

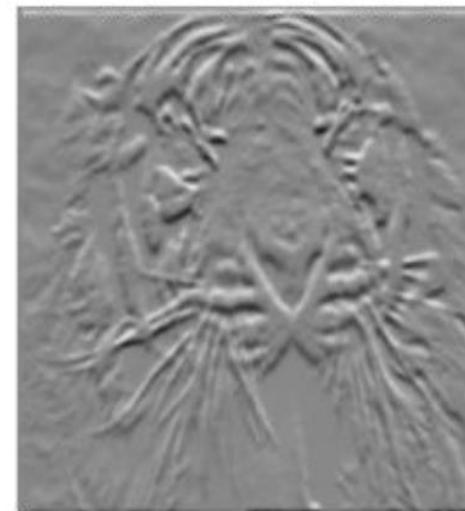


$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

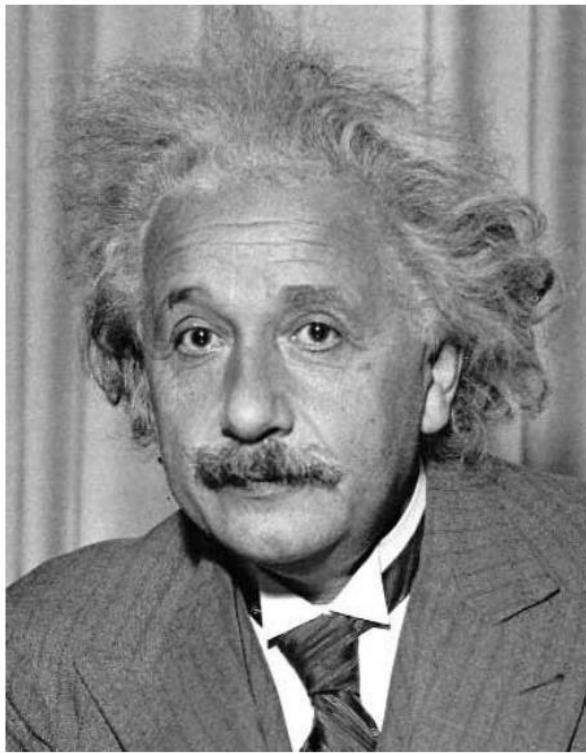


-1
1

or

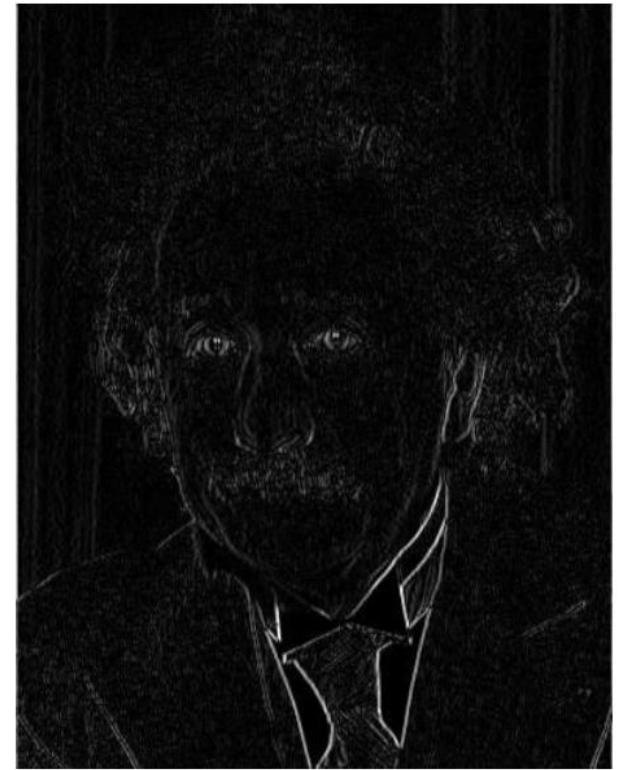
1
-1

Sobel Edge Detection Filter



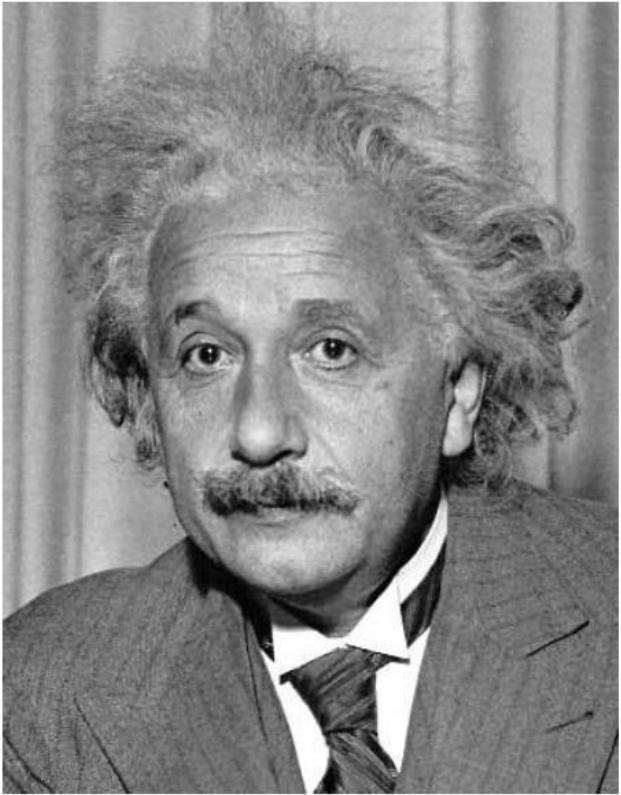
1	0	-1
2	0	-2
1	0	-1

Sobel



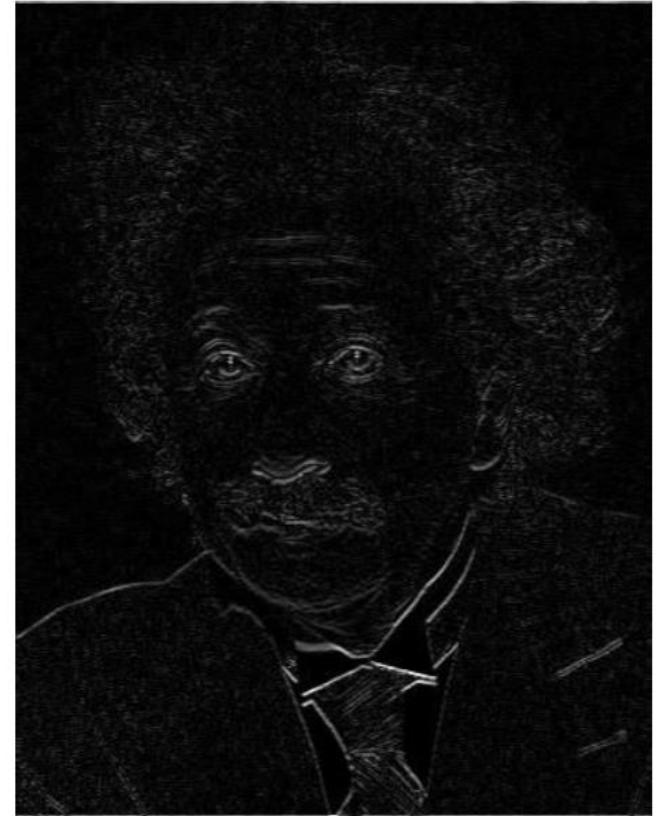
Vertical Edge
(absolute value)

Sobel Edge Detection Filter



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

Finite Difference Filters

Prewitt

M_x
-1 0 1
-1 0 1
-1 0 1

M_y
1 1 1
0 0 0
-1 -1 -1

Sobel

-1 0 1
-2 0 2
-1 0 1

1 2 1
0 0 0
-1 -2 -1

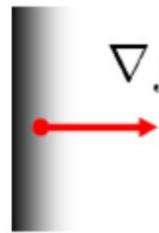
Roberts

0 1
-1 0

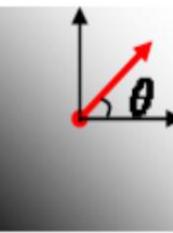
1 0
0 -1

Image Gradients

- The gradient of an image $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of most rapid change in intensity


$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

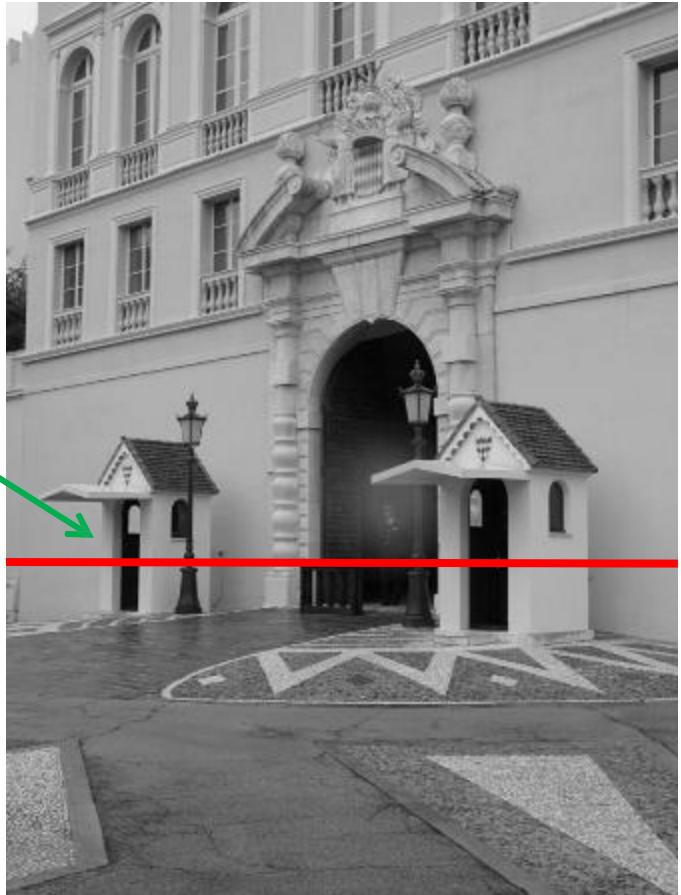

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The **gradient direction** (orientation of edge normal) is given by:

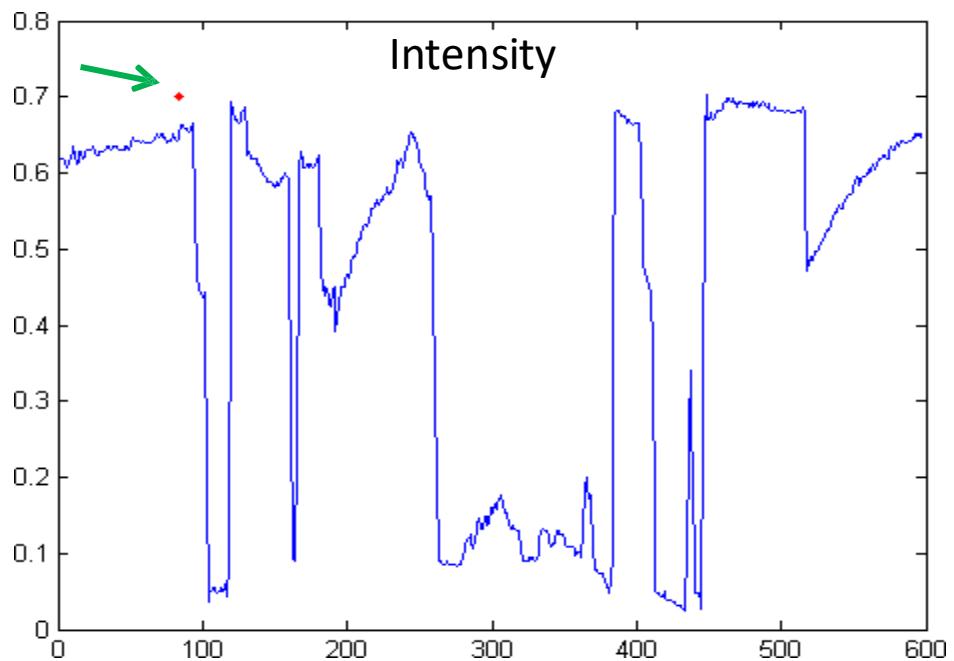
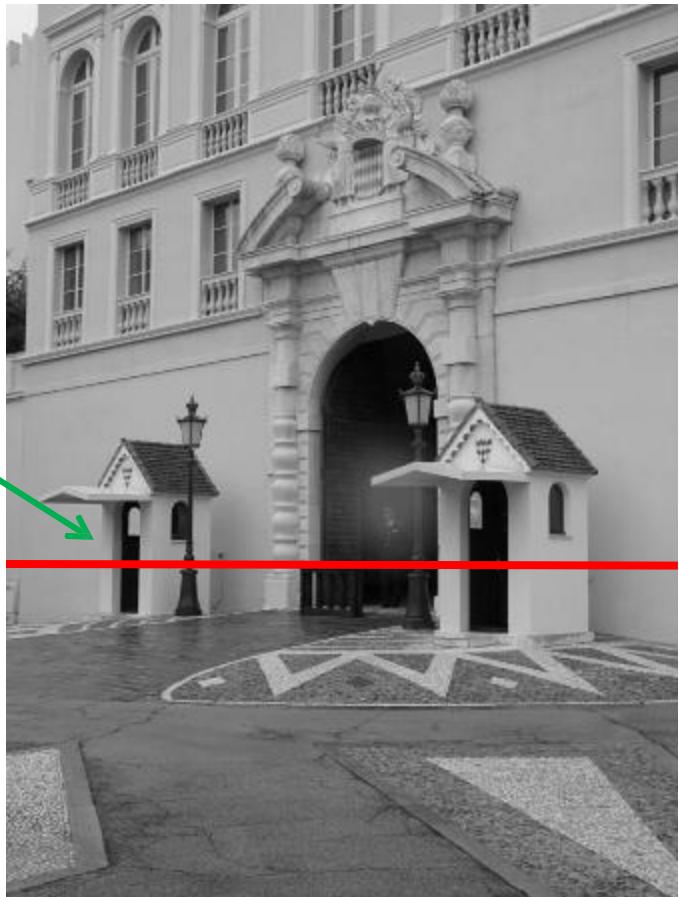
$$\theta = \tan^{-1} \left(\frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}} \right)$$

- The **edge strength** is given by the magnitude $\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$

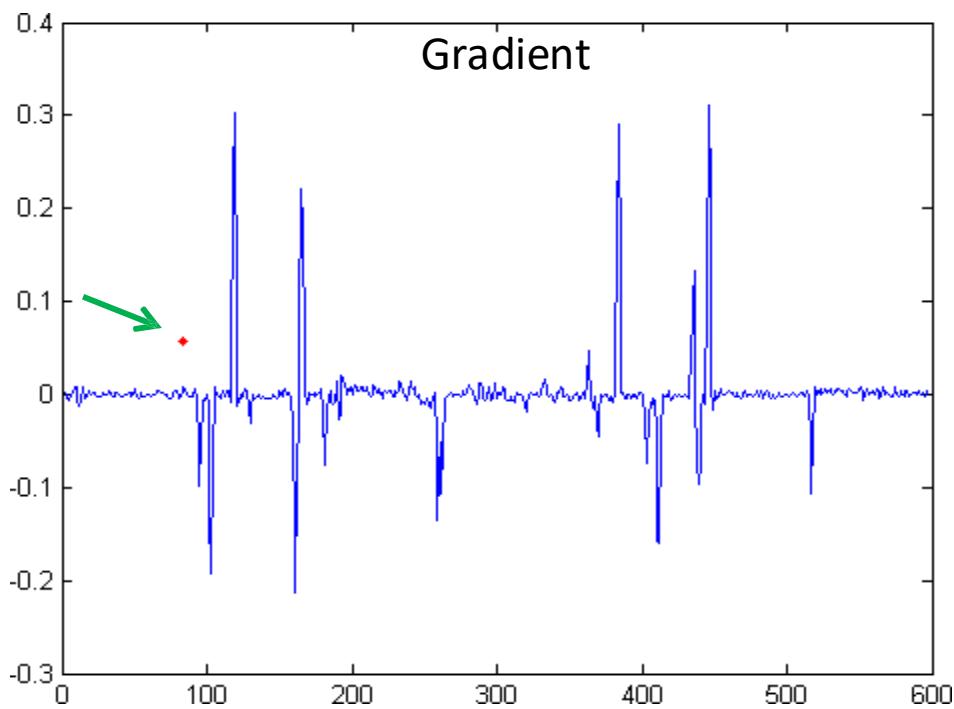
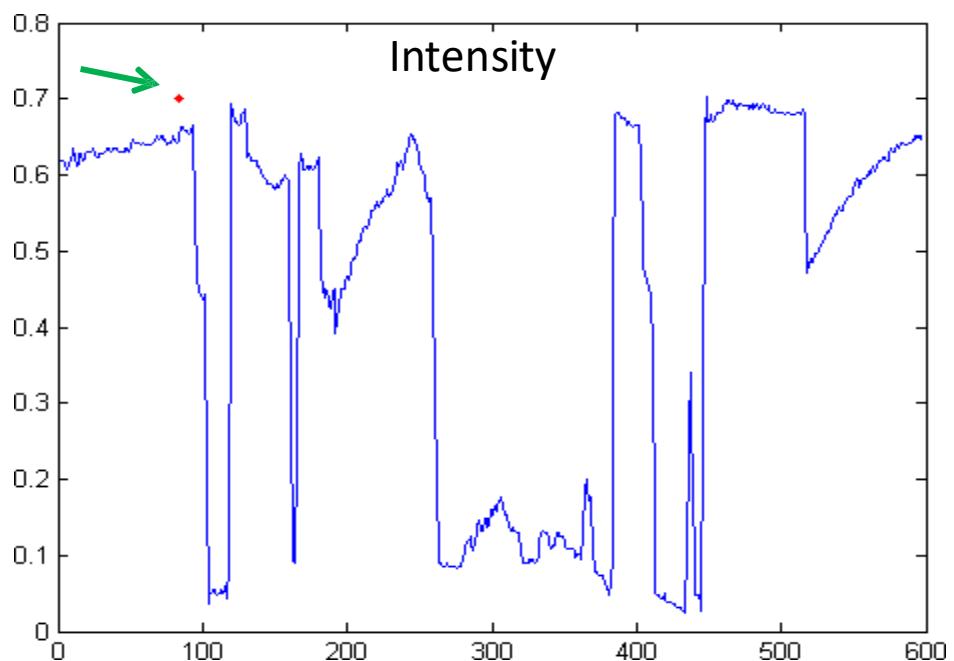
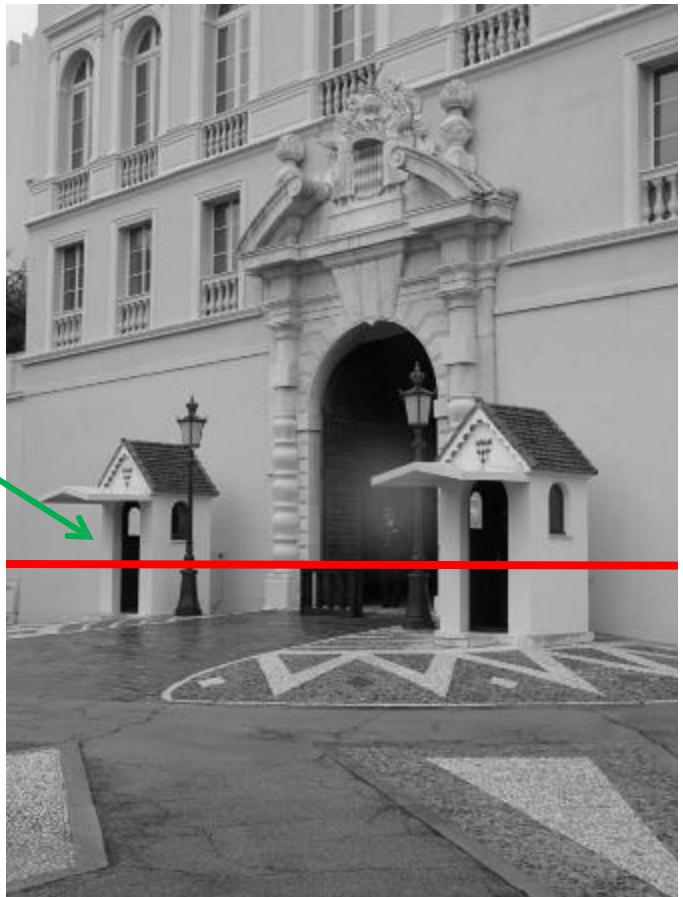
Intensity profile



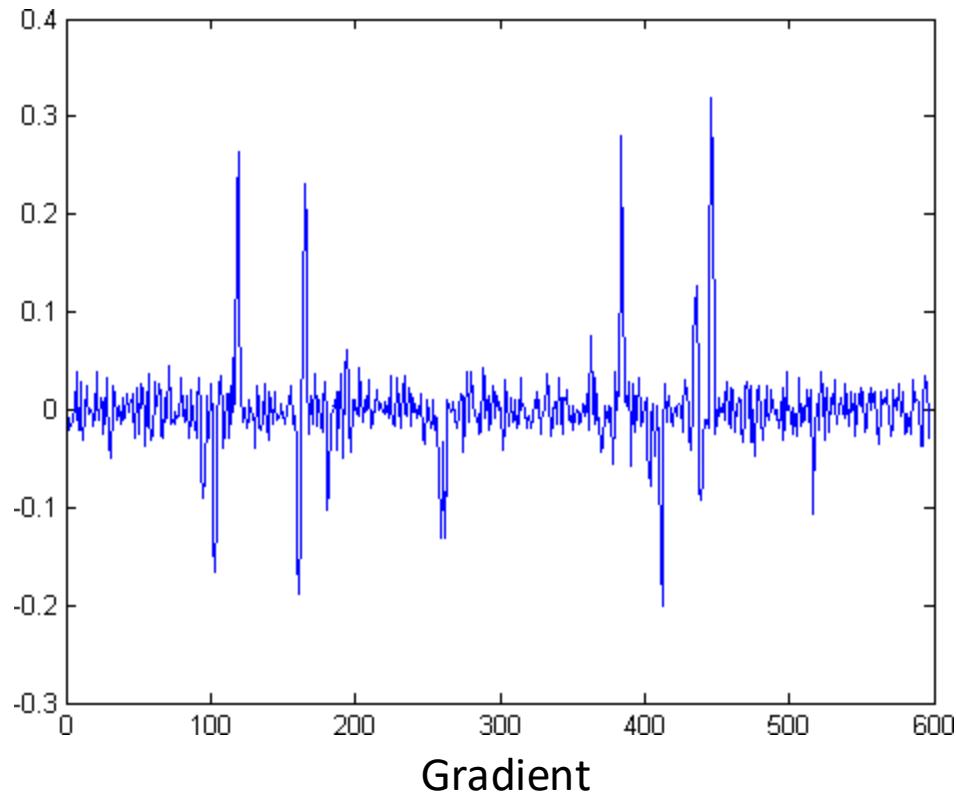
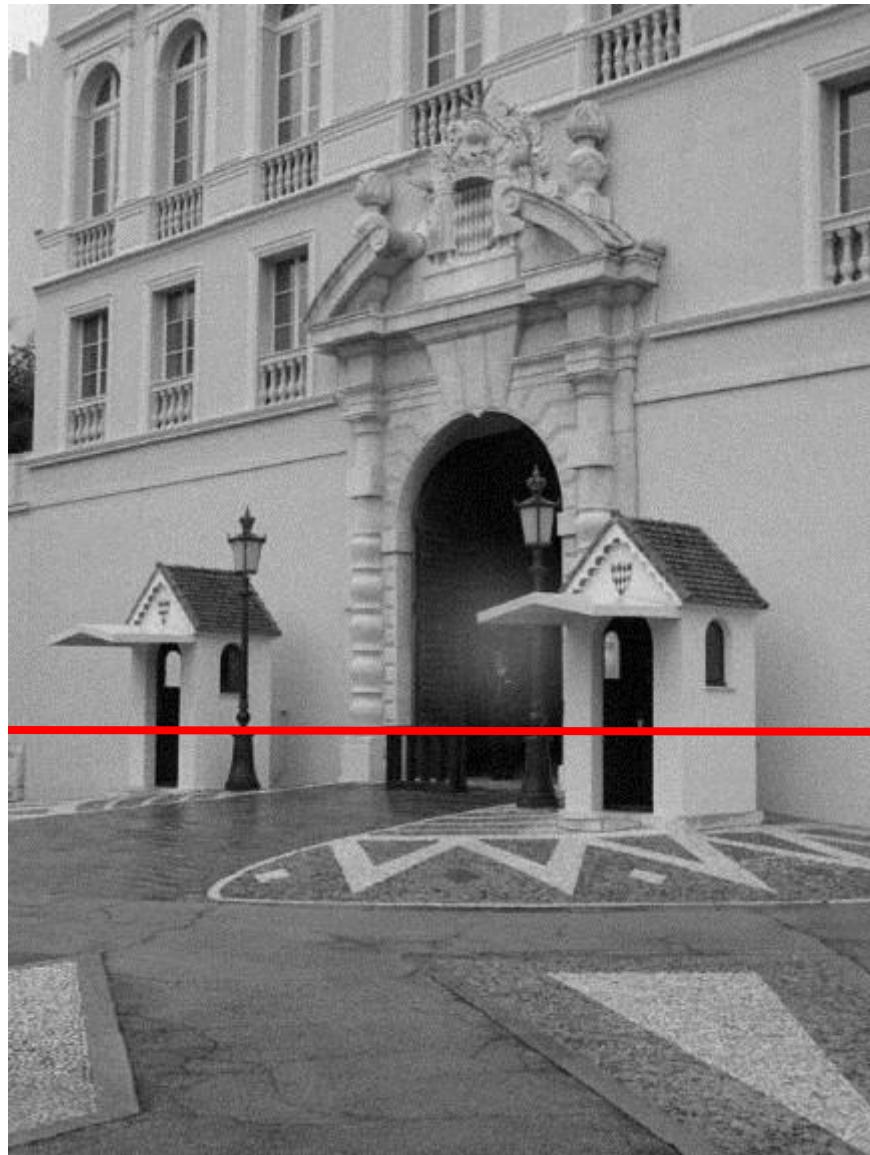
Intensity profile



Intensity profile

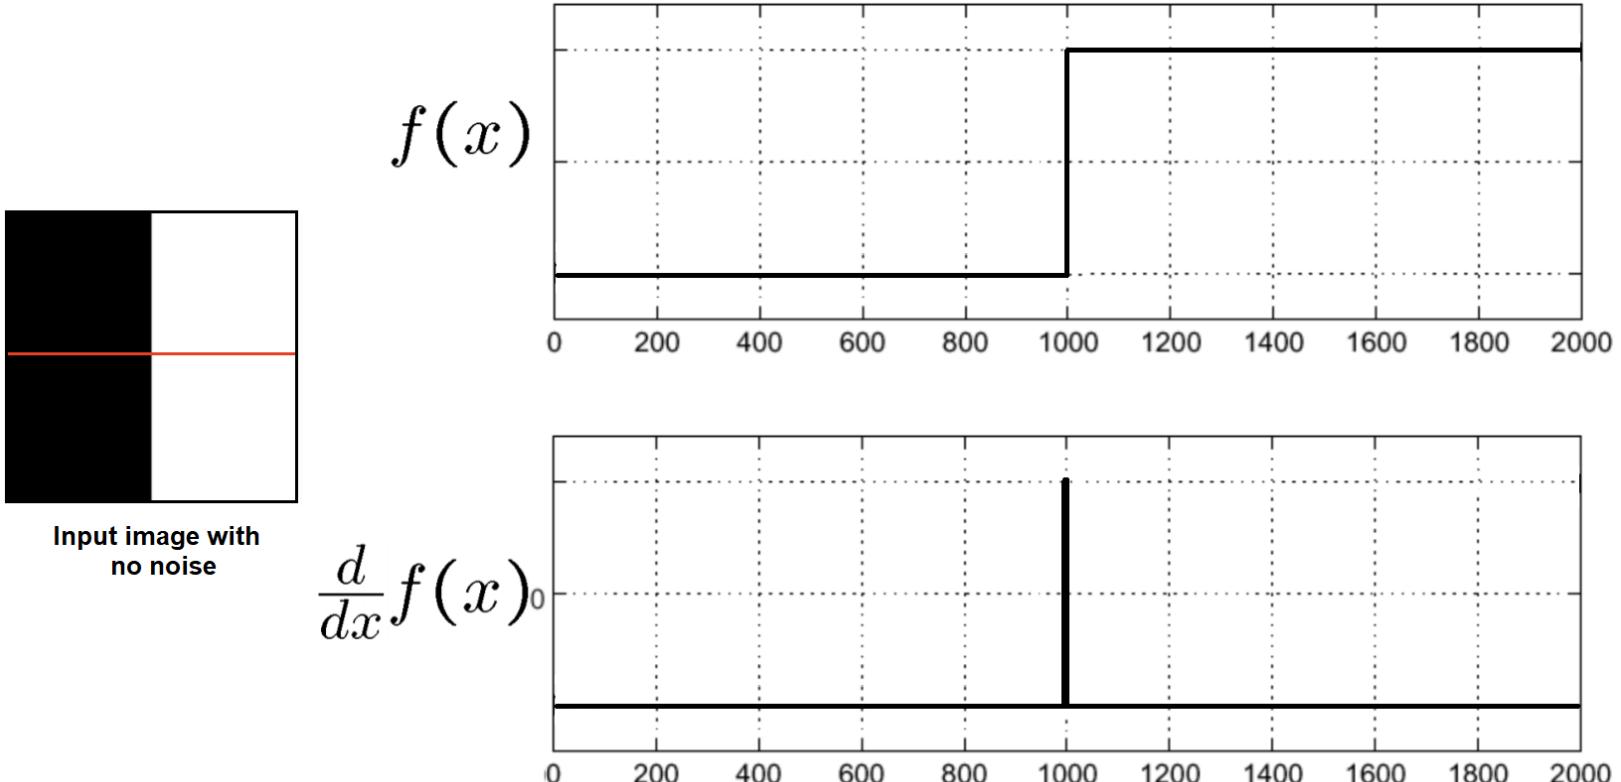


With a little Gaussian noise



Derivative with no noise

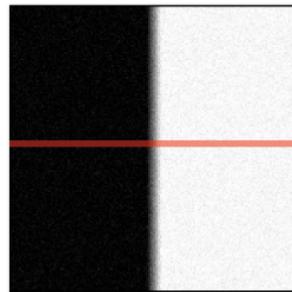
- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



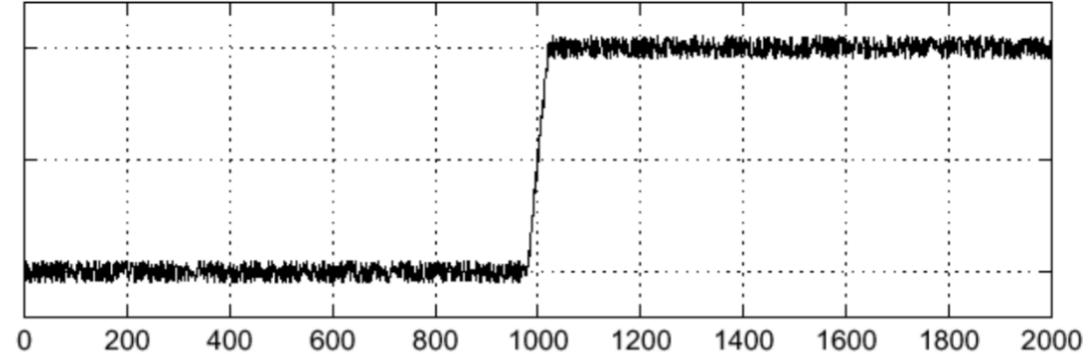
Where is the edge?

Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal

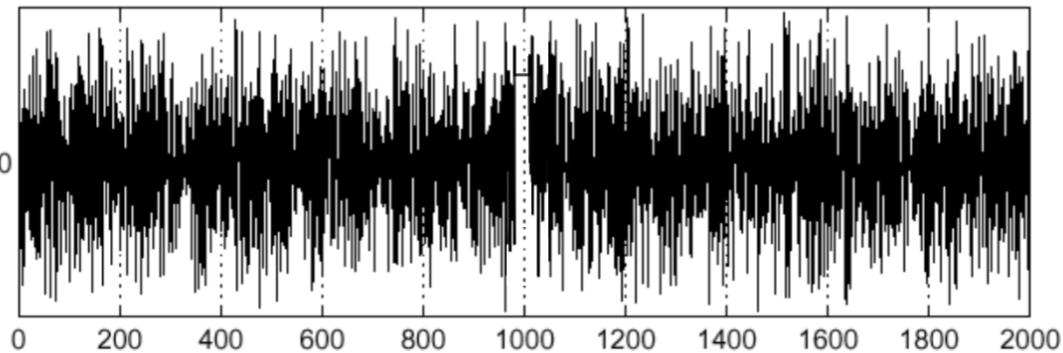


$f(x)$



Noisy input image

$$\frac{d}{dx}f(x)$$

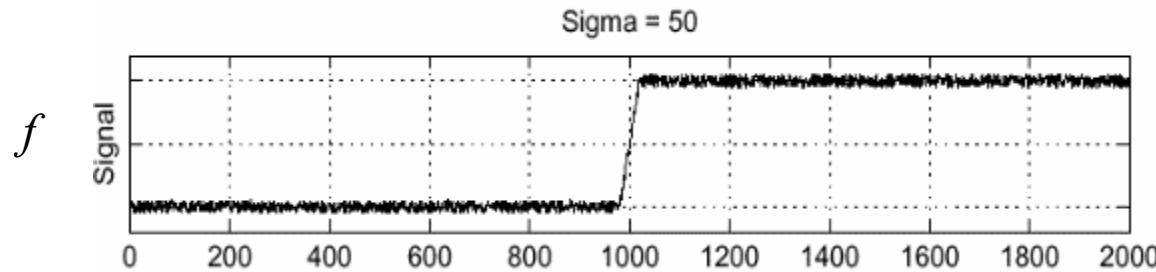


Where is the edge?

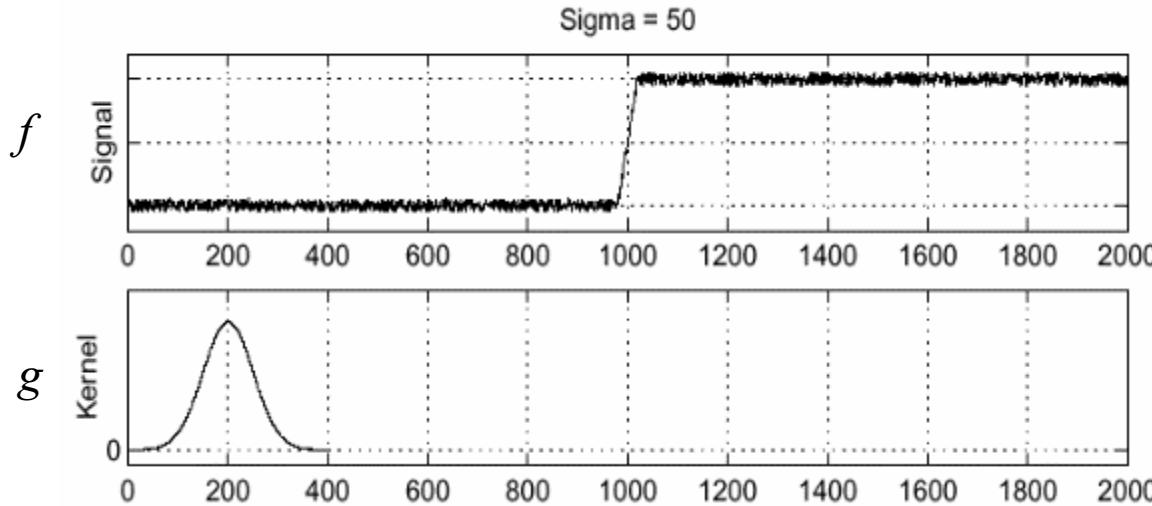
The larger the noise the stronger the response

- What can we do about it?

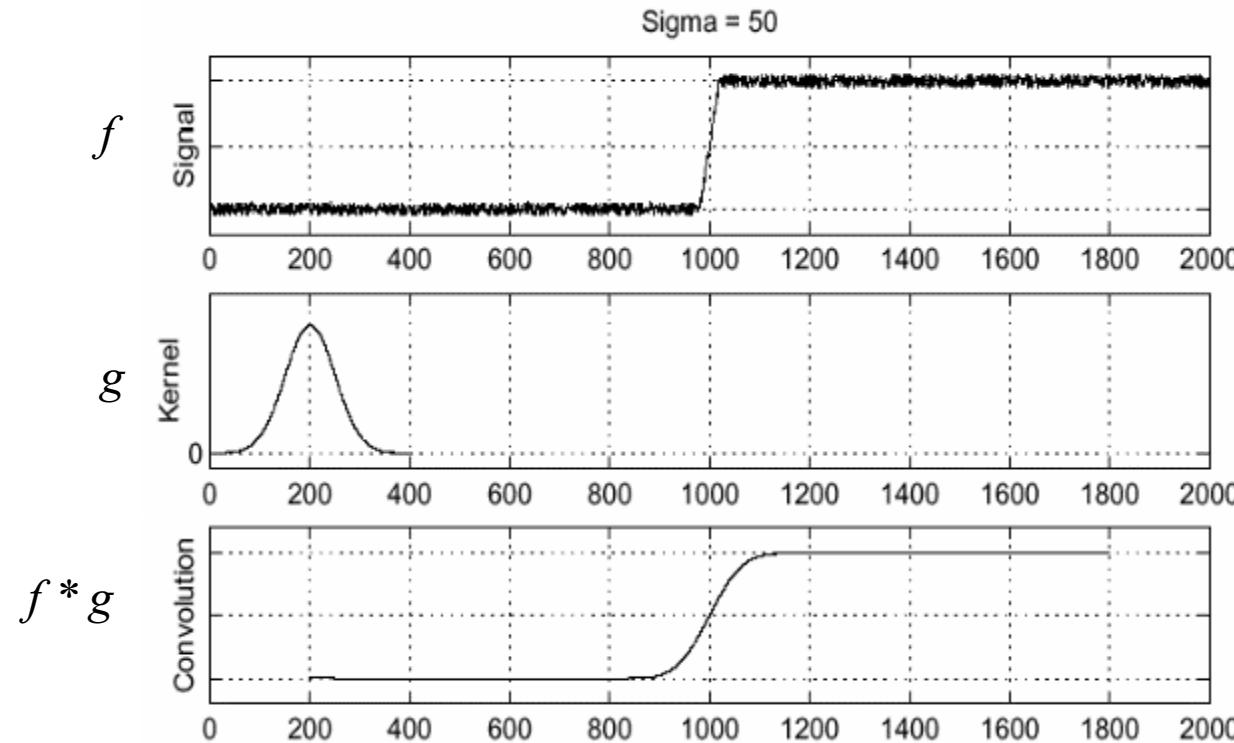
Solution: smooth first



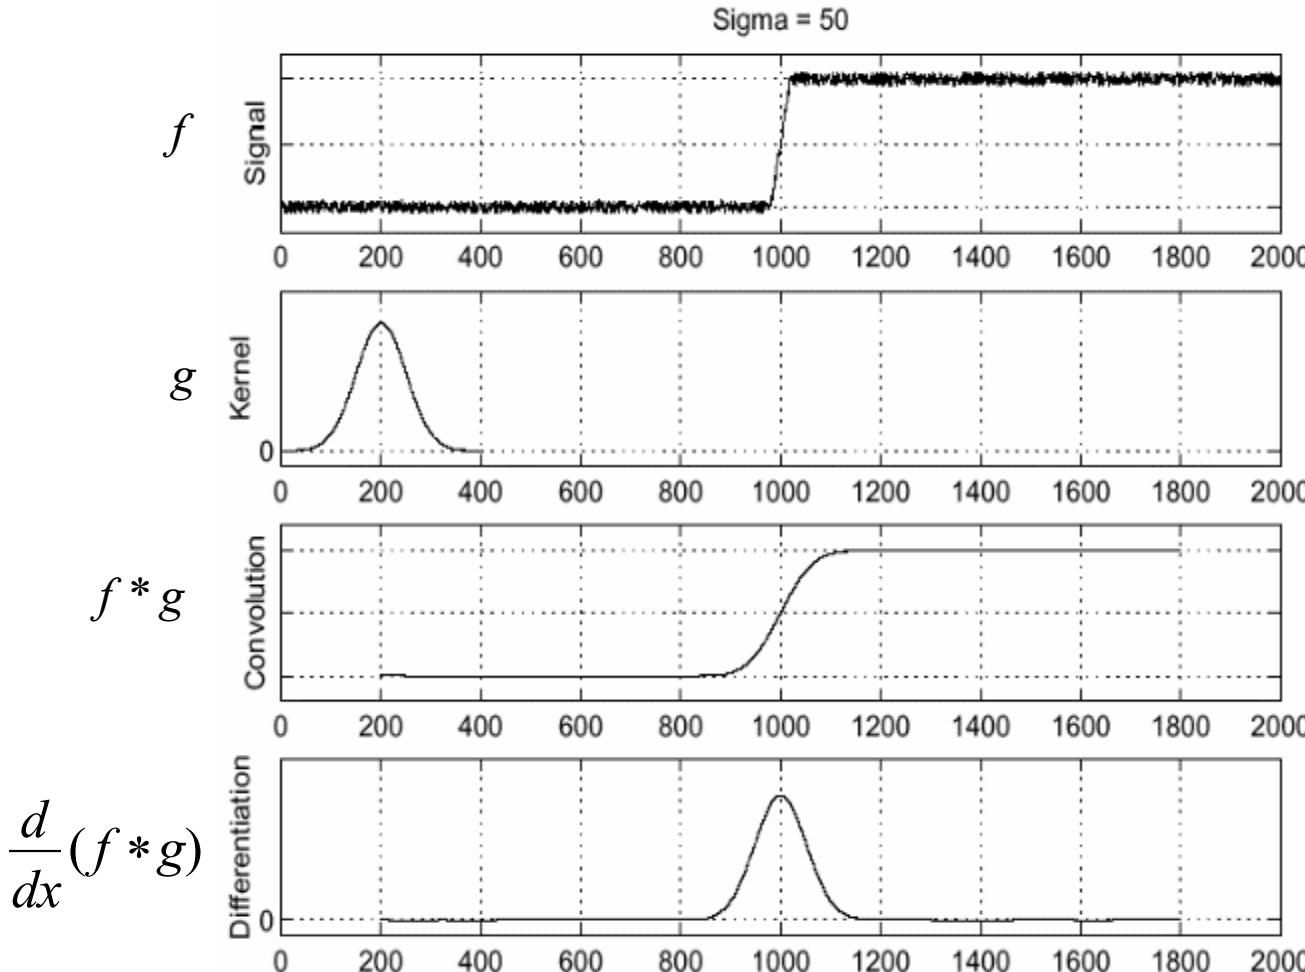
Solution: smooth first



Solution: smooth first



Solution: smooth first



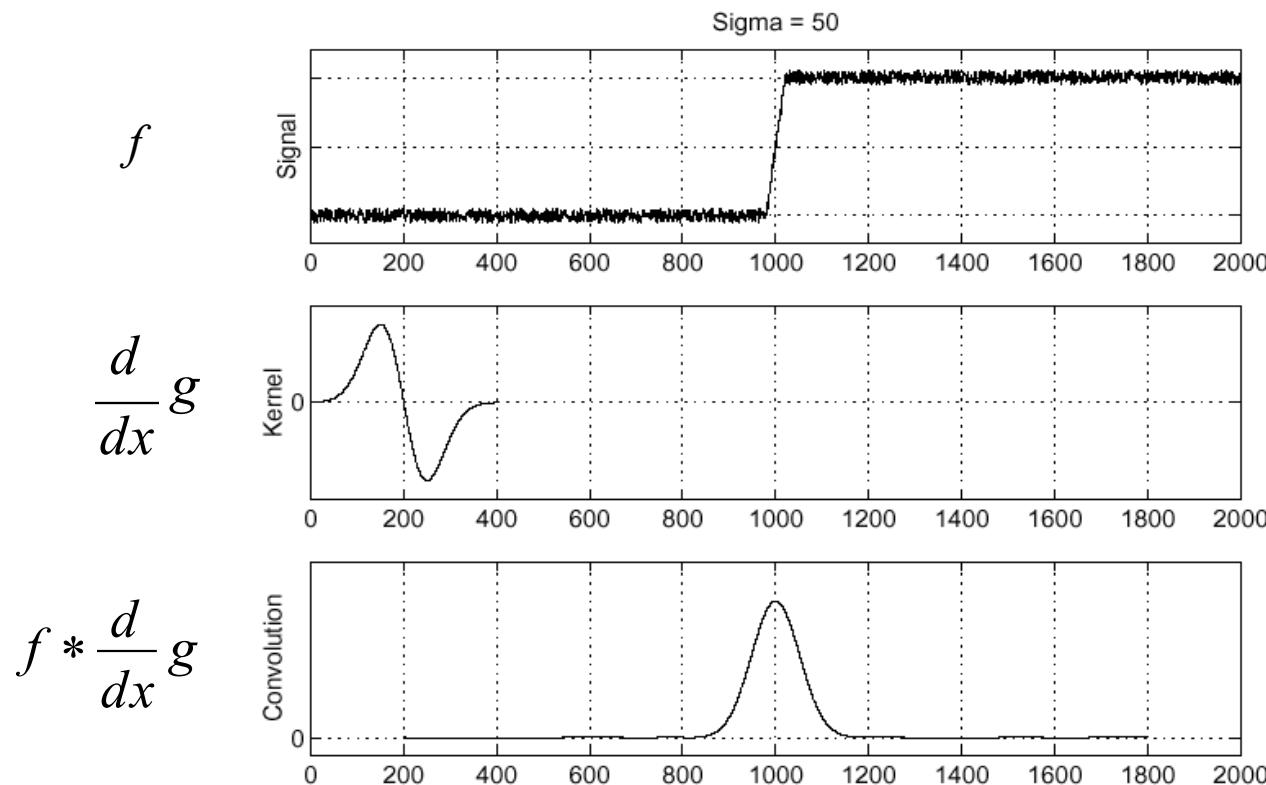
- To find edges, look for peaks in

$$\frac{d}{dx}(f * g)$$

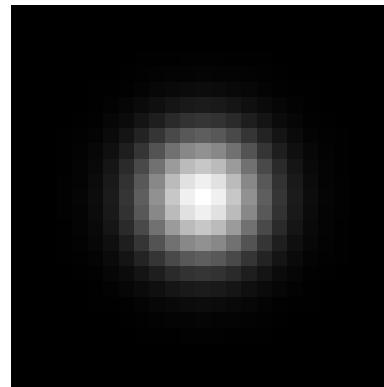
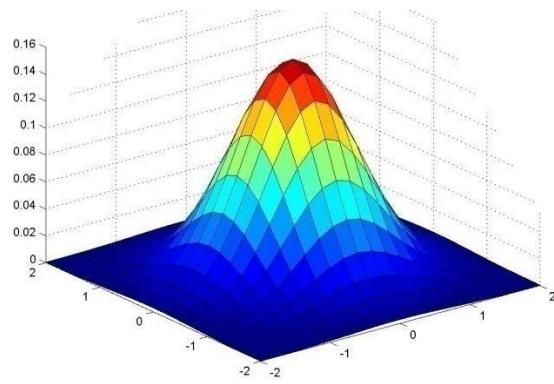
Derivative theorem of convolution

- Differentiation in convolution, and convolution is associative:
- This saves us one operation:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

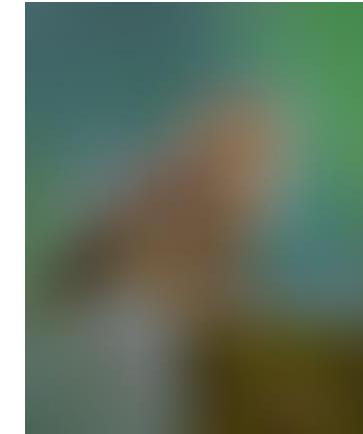
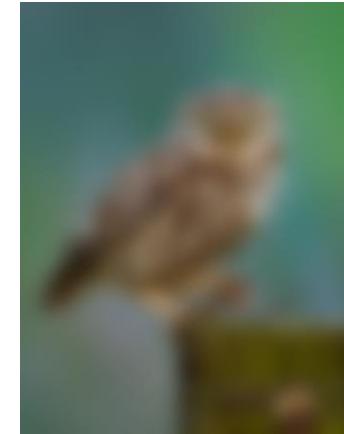


Gaussian Kernel

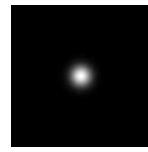


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Gaussian filters



$\sigma = 1$ pixel



$\sigma = 5$ pixels

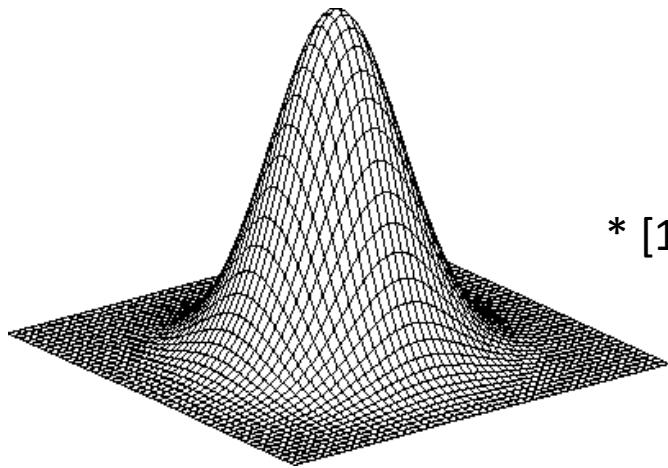


$\sigma = 10$ pixels



$\sigma = 30$ pixels

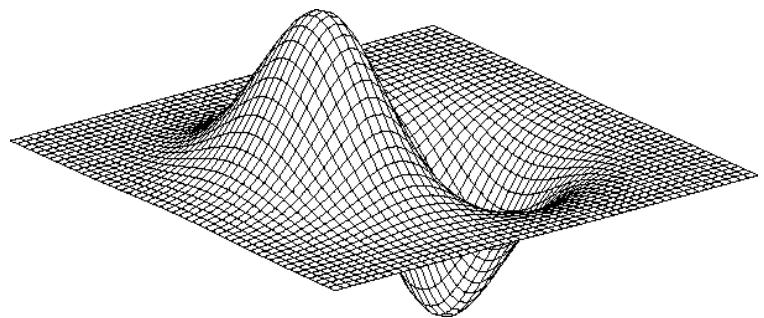
Derivative of Gaussian filter



Gaussian

$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

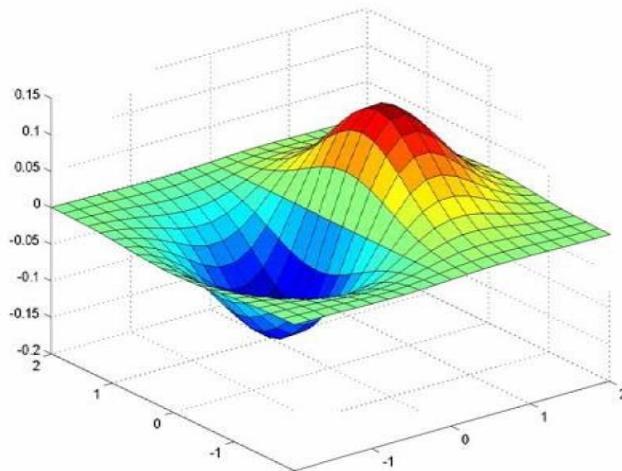
* [1 0 -1] =



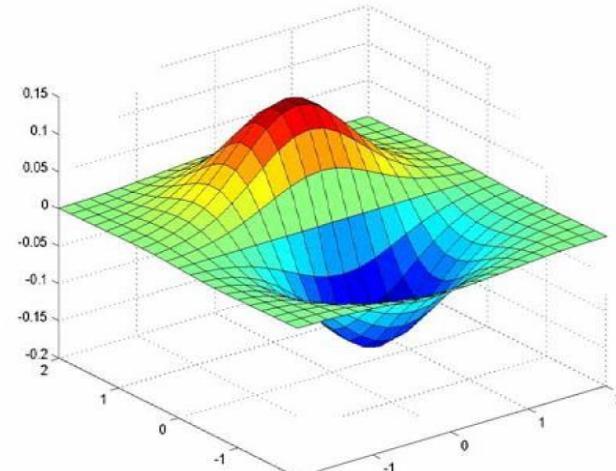
derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

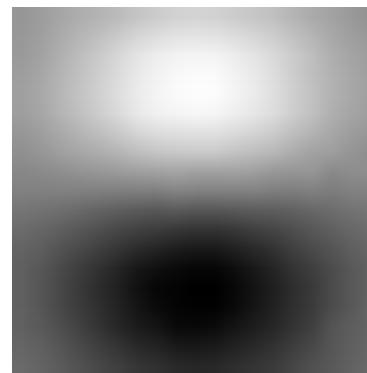
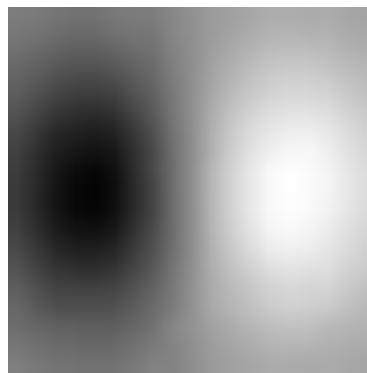
Derivative of Gaussian filter



x-direction



y-direction



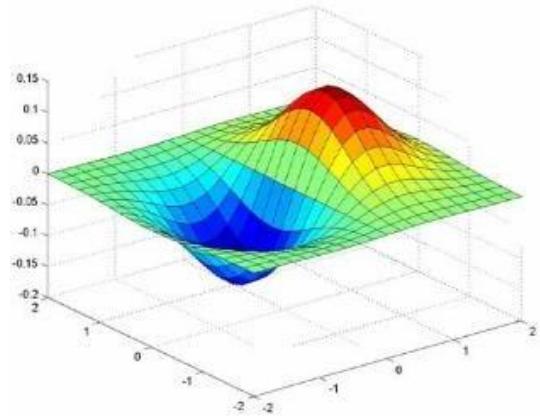
- Which one finds horizontal and vertical edges ?

Example



input image (“Lena”)

Compute Gradients (DoG)

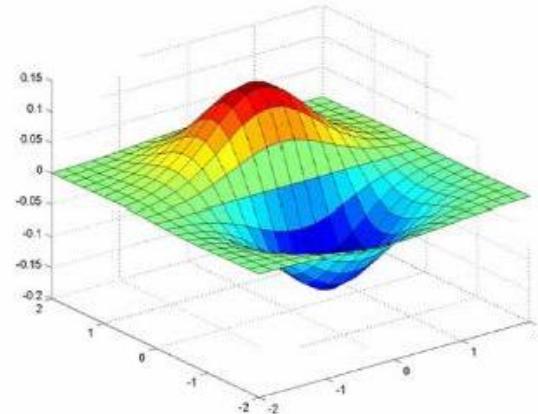
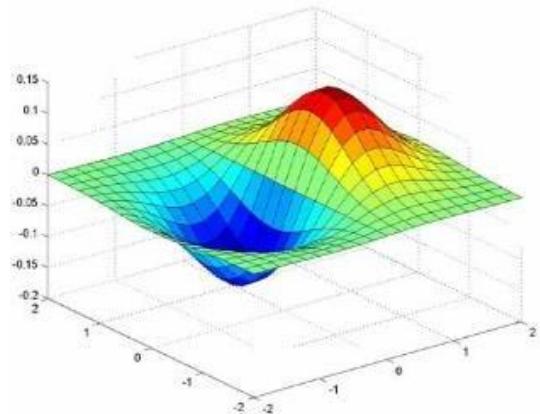


Input Image



X-Derivative of Gaussian

Compute Gradients (DoG)



Input Image

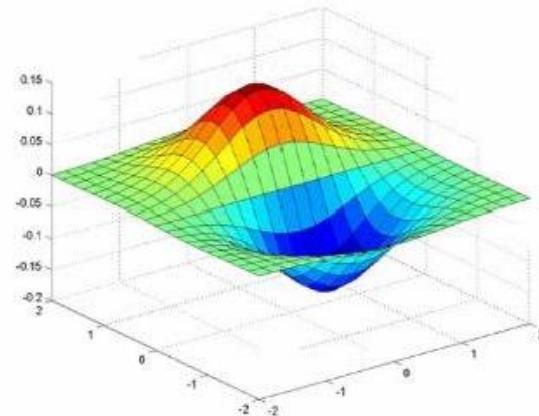
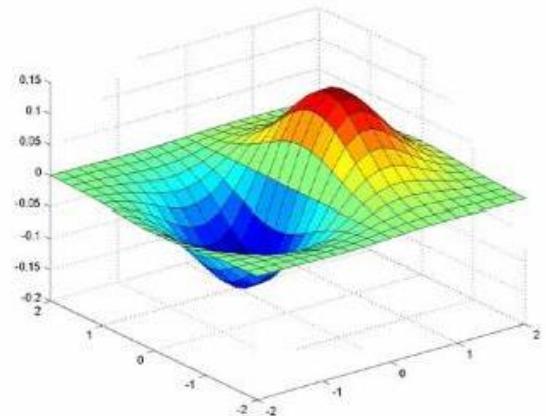


X-Derivative of Gaussian



Y-Derivative of Gaussian

Compute Gradients (DoG)



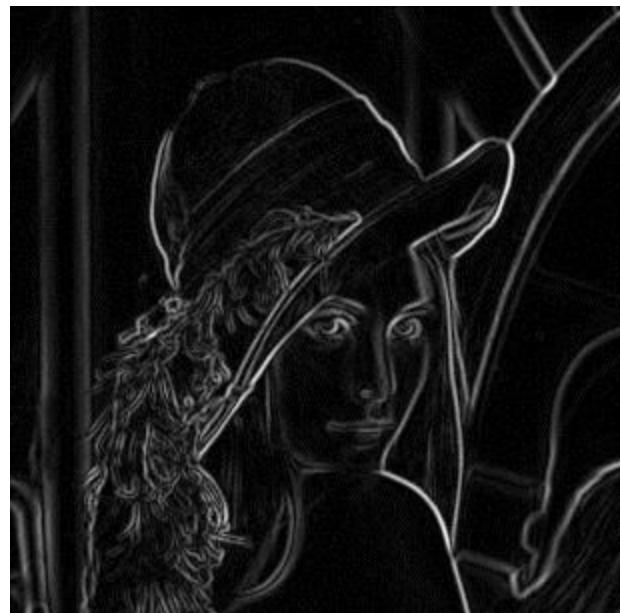
Input Image



X-Derivative of Gaussian

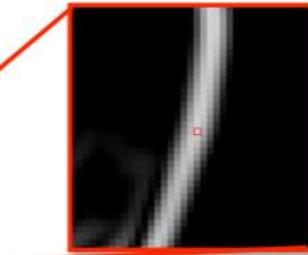


Y-Derivative of Gaussian



Gradient Magnitude

Designing an edge detector



where is the edge?

Designing an edge detector

- Criteria for a good edge detector:
 - **Good detection:**
 - find all real edges, ignoring noise or other artifacts

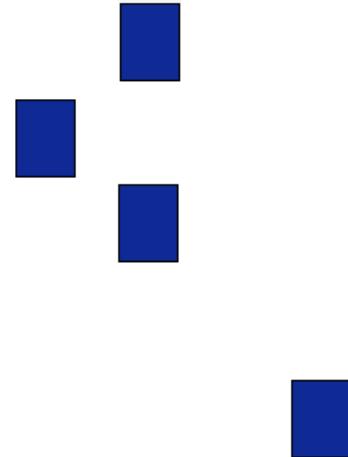
Designing an edge detector

- Criteria for a good edge detector:
 - **Good detection:**
 - find all real edges, ignoring noise or other artifacts
 - **Good localization**
 - detect edges as close as possible to the true edges
 - return one point only for each true edge point

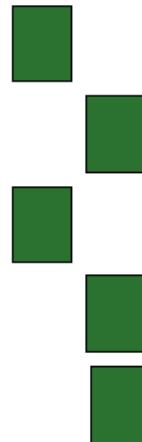
Designing an edge detector



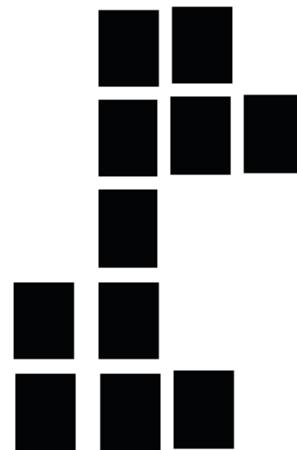
True Edge



Poor robustness to noise



Poor localization

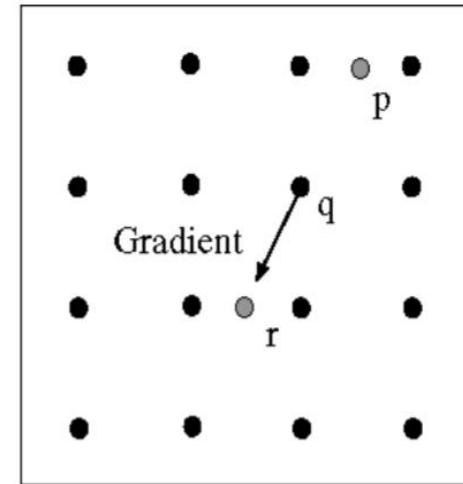
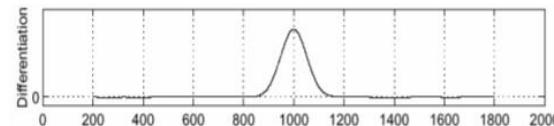
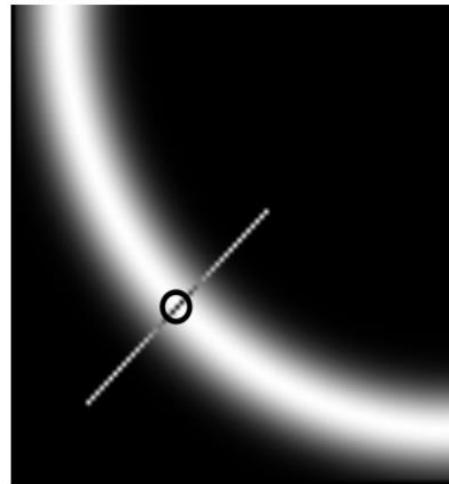


Too many responses

Non-Maxima Suppression

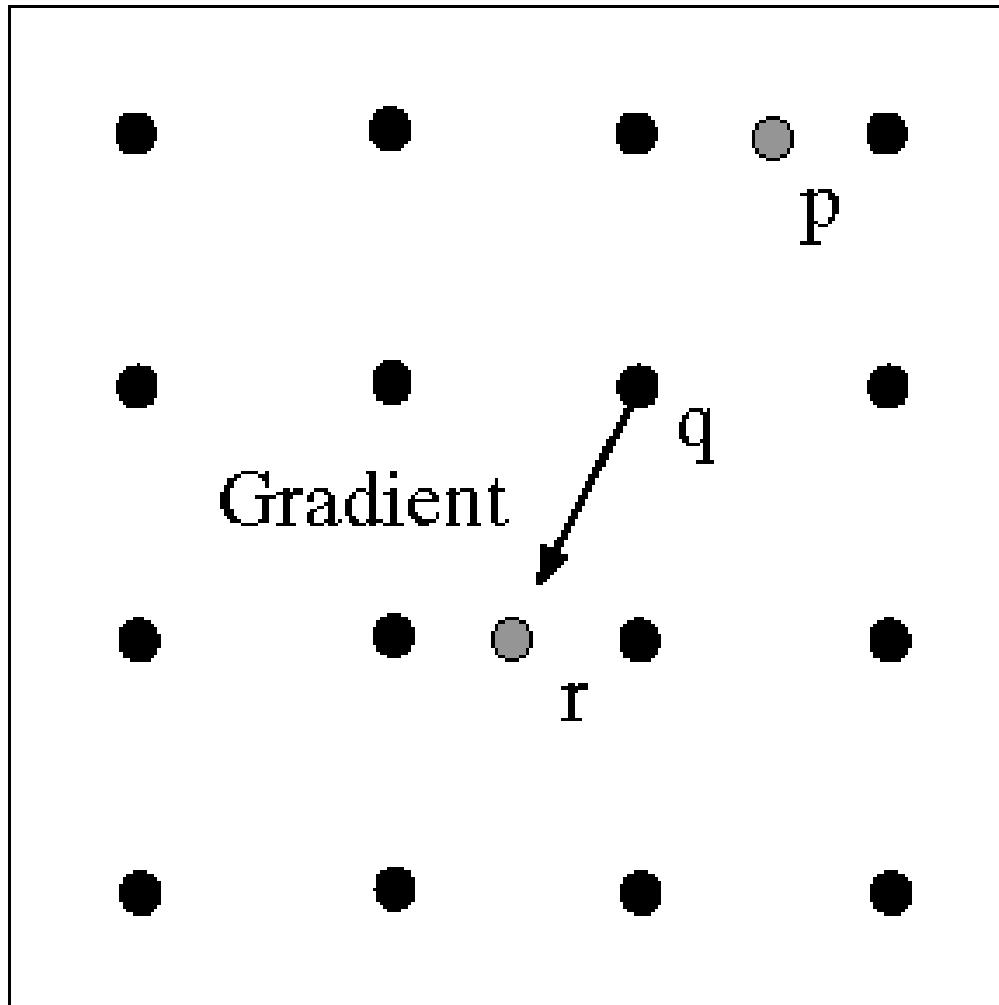


where is the edge

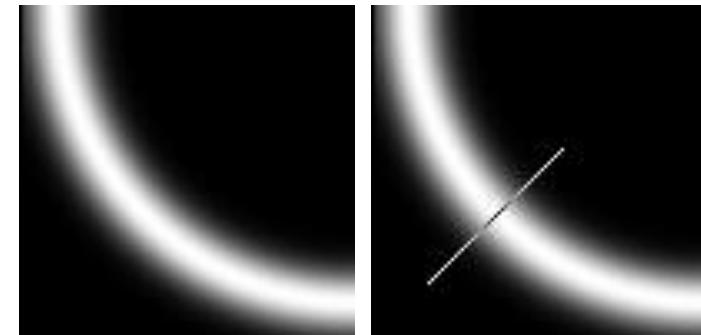


- Get orientation $\theta = \arctan(g_y, g_x)$
- Check if pixel is local maximum along gradient direction:
- Could require checking interpolated pixels p and r

Non-maximum suppression for each orientation



At q , we have a maximum if the value is larger than those at both p and at r . Interpolate to get these values.



Before Non-max Suppression



After non-max suppression



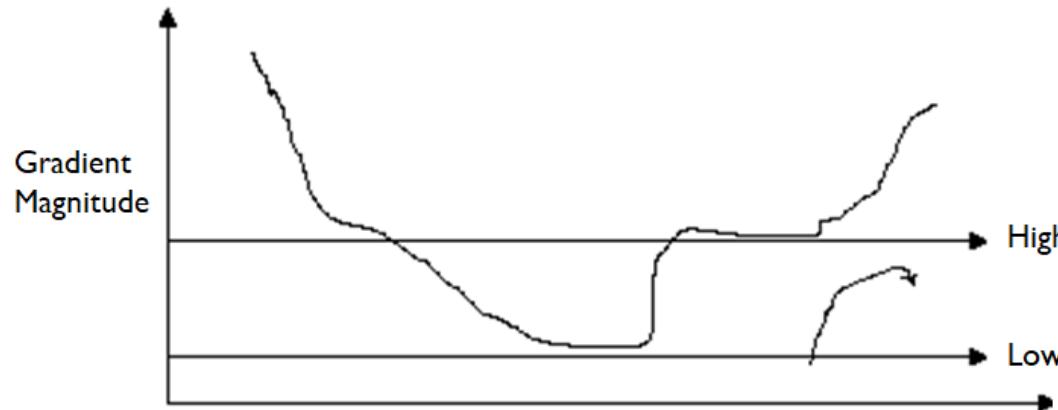
Hysteresis thresholding

- Check for well-connected edges
- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



Hysteresis thresholding

- Check for well-connected edges. How?
- use a high threshold to start edge curves and a low threshold to continue them.
- How does it work?
 - If gradient at pixel > 'High' => 'edge pixel'
 - If gradient at pixel < 'Low' => 'non-edge pixel'
 - If gradient at pixel > 'Low' and < 'High' => 'edge pixel' iff it is connected to an 'edge pixel' directly or via pixels between 'Low' and 'High'



Final Canny Edges



Canny edge detector

- The most widely used edge detector

A computational approach to edge detection

[J Canny - IEEE Transactions on pattern analysis and machine ...](#), 1986 - ieeexplore.ieee.org

Abstract: This paper describes a computational approach to edge detection. The success of the approach depends on the definition of a comprehensive set of goals for the computation of edge points. These goals must be precise enough to delimit the desired behavior of the detector while making minimal assumptions about the form of the solution. We define detection and localization criteria for a class of edges, and present mathematical forms for ...

Cited by 27743 Related articles All 27 versions Import into BibTeX Save More

J. Canny, [*A Computational Approach To Edge Detection*](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

Canny edge detector

- Probably the most widely used edge detector in computer vision (Canny 1986)

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width

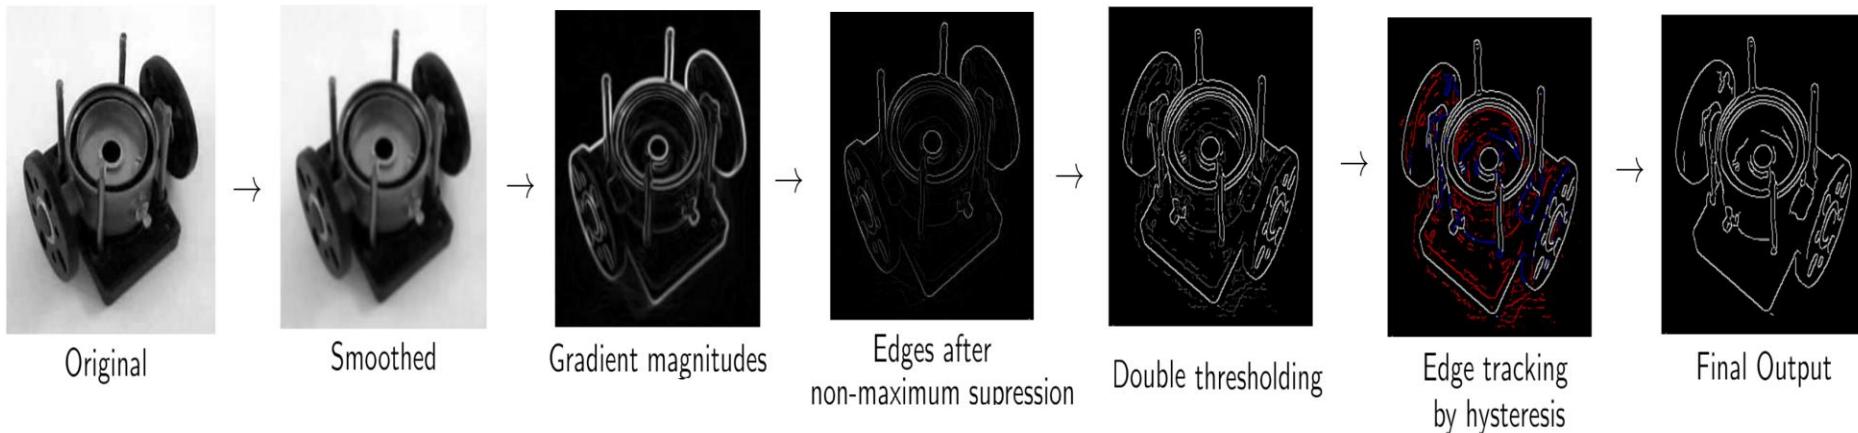
Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width
4. Thresholding and linking (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
 2. Find magnitude and orientation of gradient
 3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width
 4. Thresholding and linking (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
-
- MATLAB: `edge(image, 'canny')`

Canny edge Pipeline: Example



Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$



Canny with $\sigma = 2$

The choice of σ depends on desired behavior

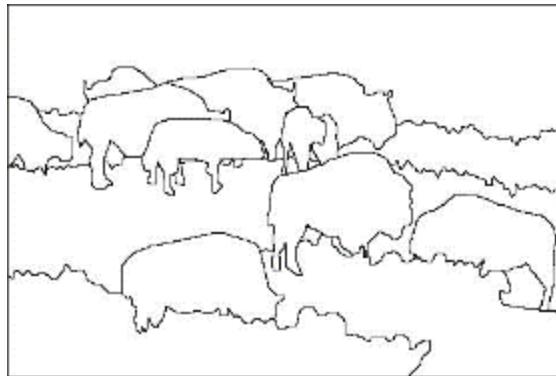
- large σ detects large scale edges
- small σ detects fine features

Learning to detect boundaries

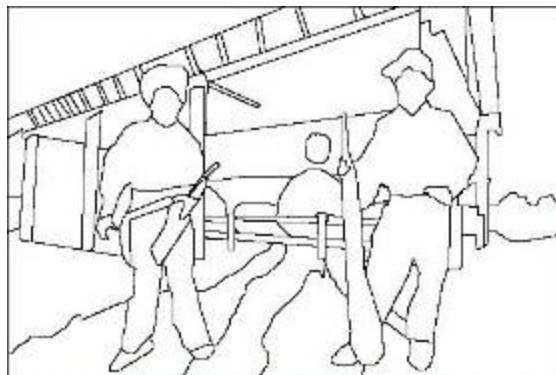
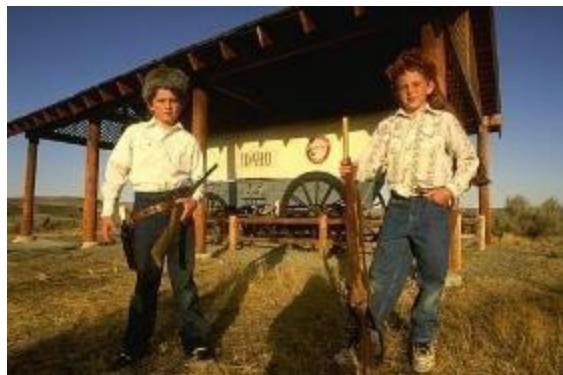
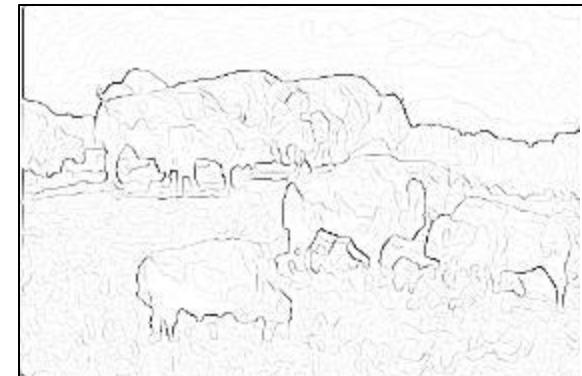
image



human segmentation



gradient magnitude

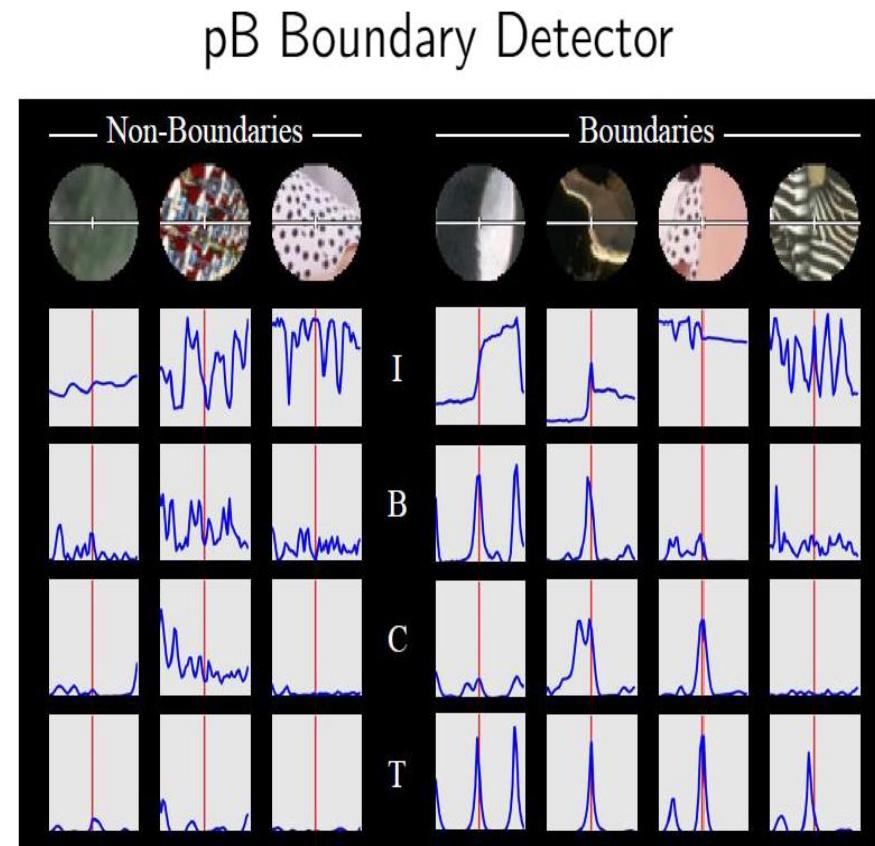
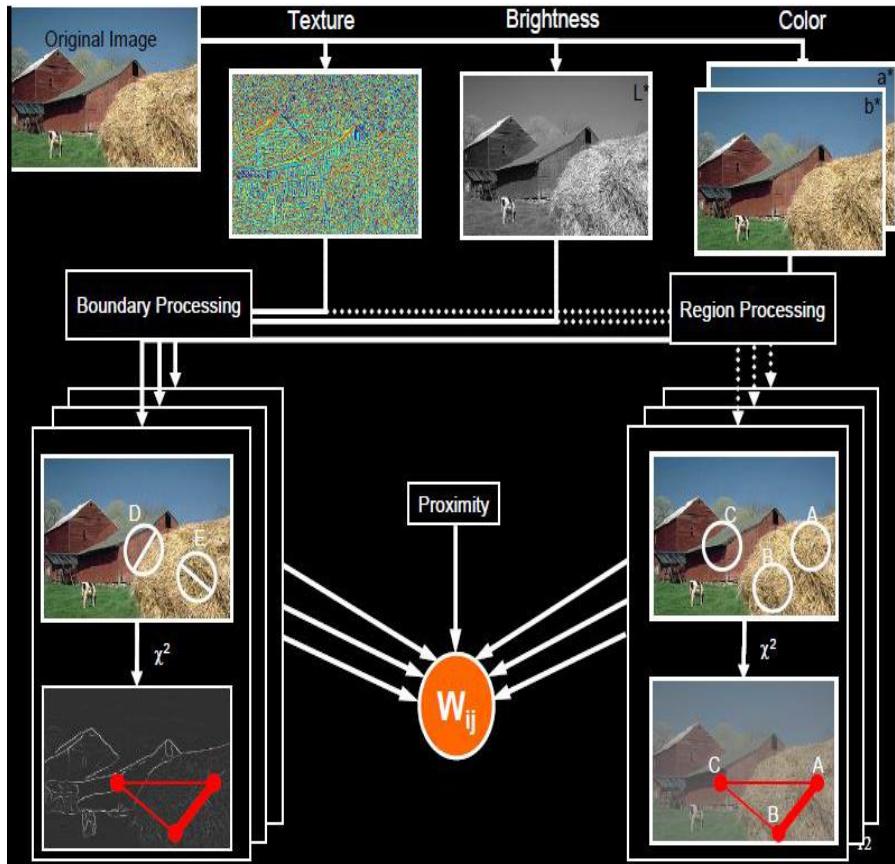


- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

More Recent Methods in Edge Detection

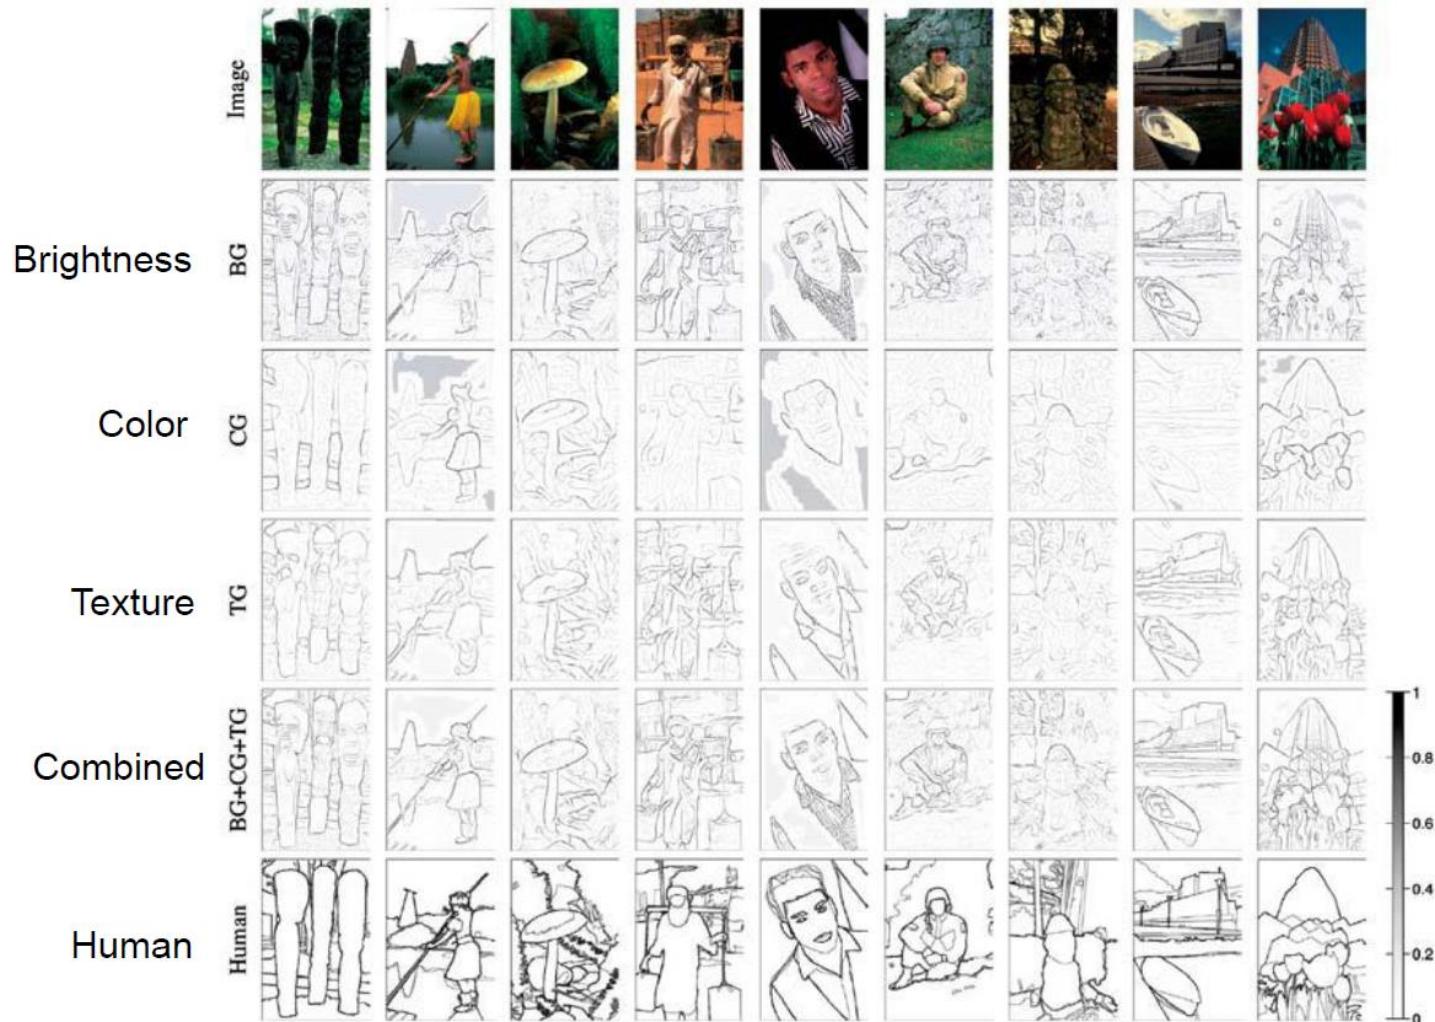
- Learning to Detect Natural Image Boundaries Using Local Brightness, Color, and Texture Cues (Martin et al, 2004)



I-Intensity, B-Brightness gradient, C-Color gradient, T-Texture gradient

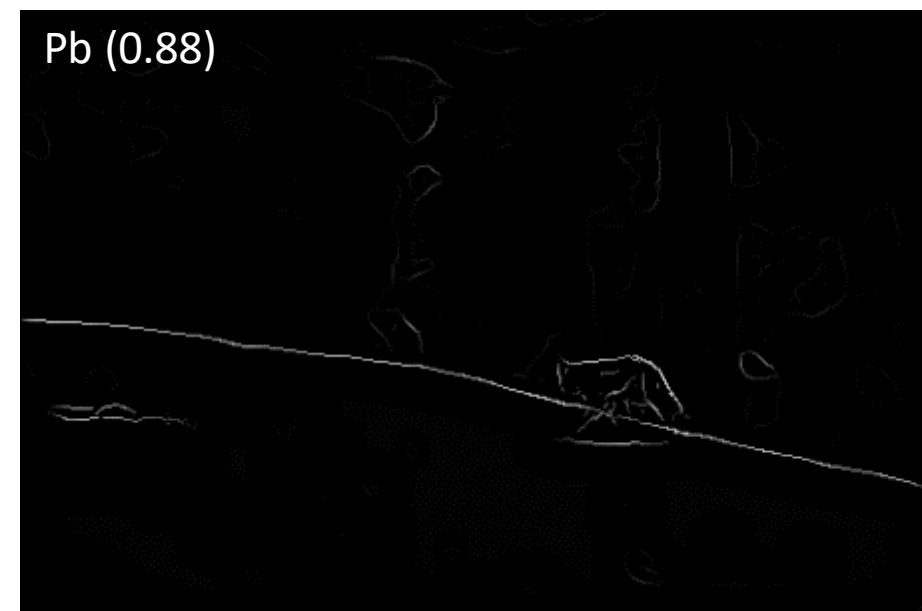
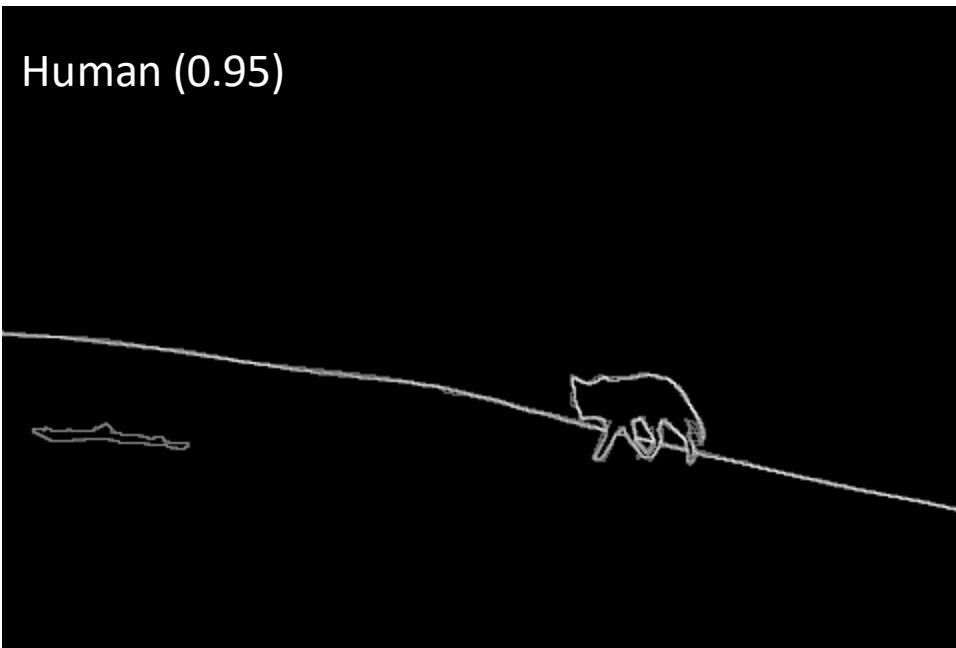
More Recent Methods in Edge Detection

- Learning to Detect Natural Image Boundaries Using Local Brightness, Color, and Texture Cues (Martin et al, 2004)



Source: Derek Hoiem

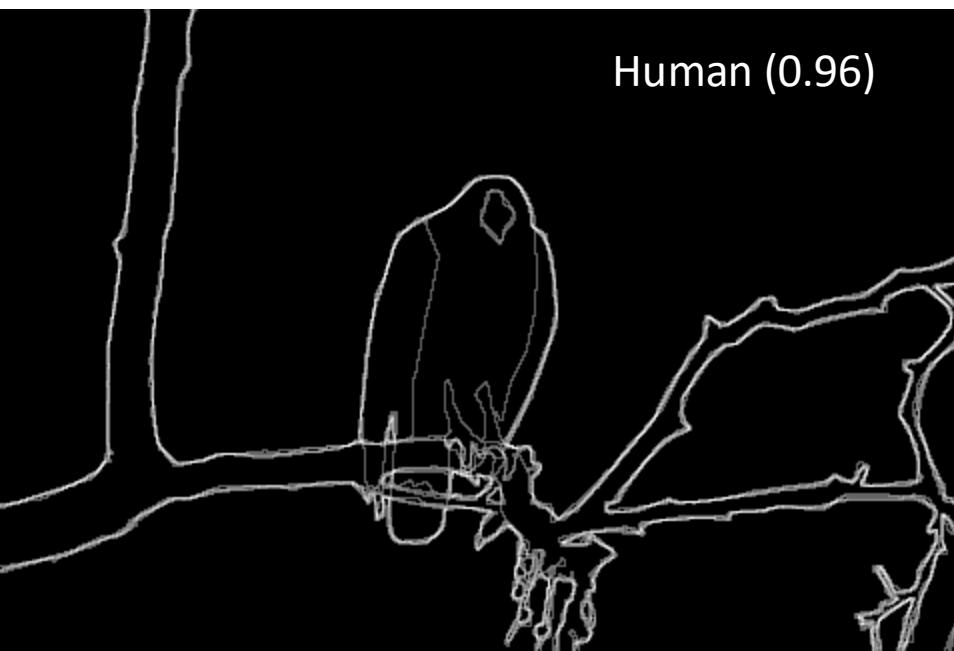
Results



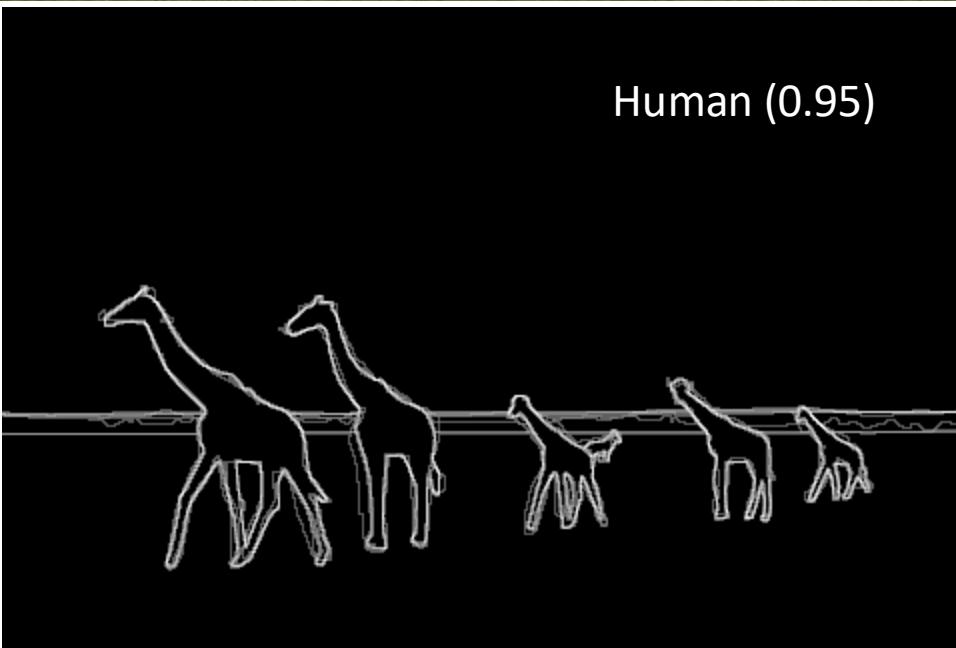
Results



Human (0.96)



Results

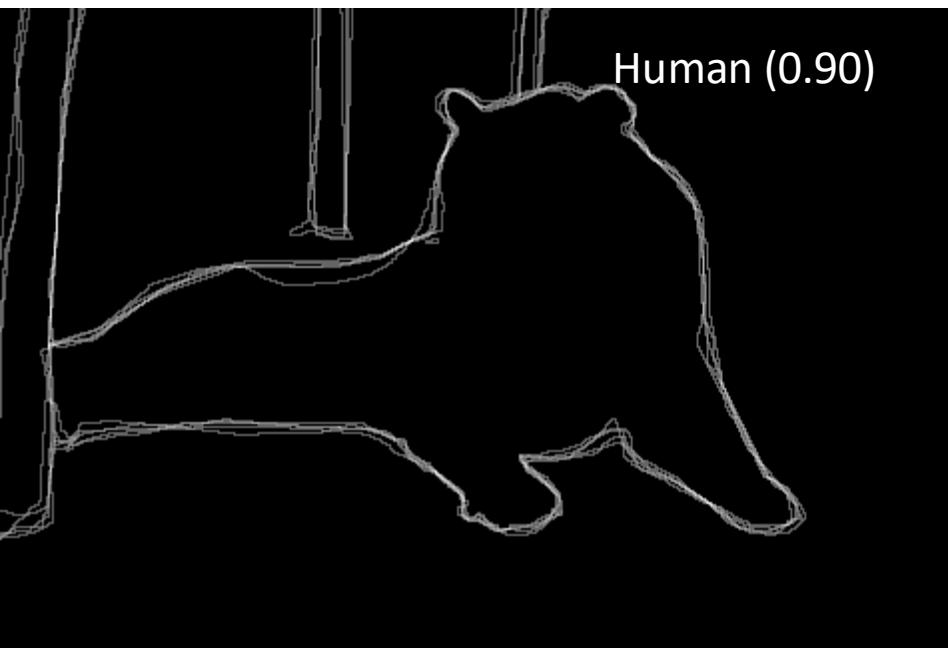


Human (0.95)



Pb (0.63)

Results



For more:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html>

State of edge detection

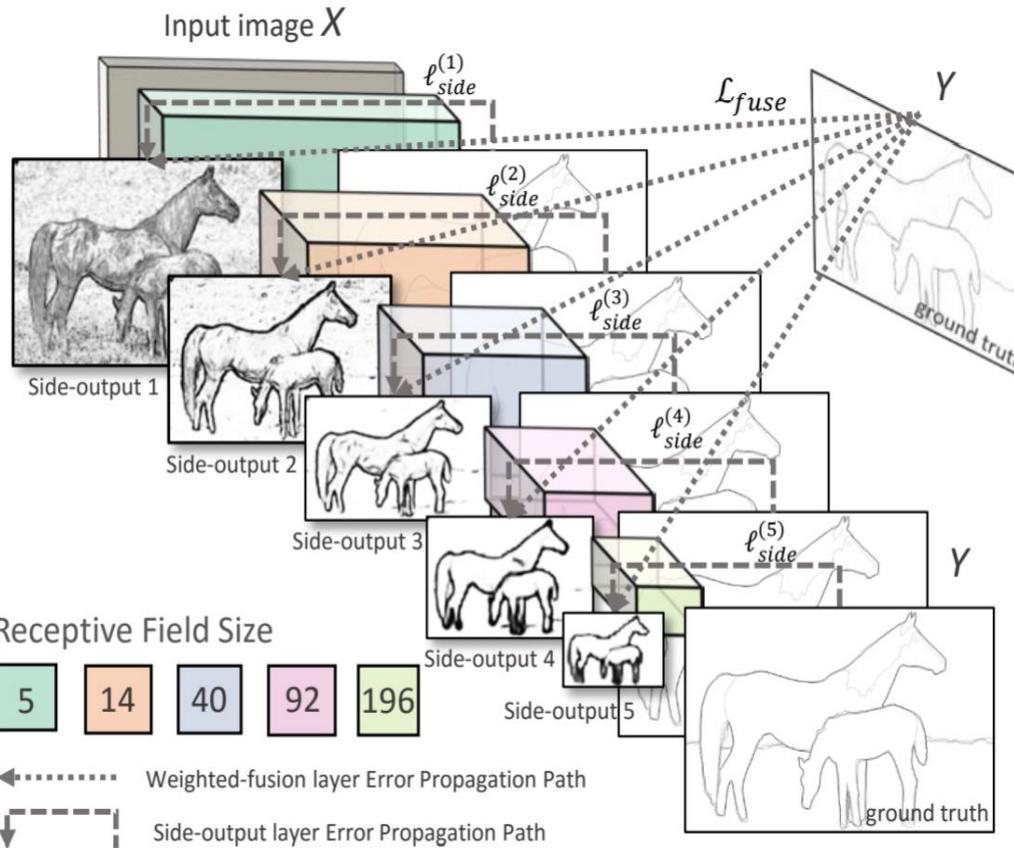
- Local edge detection is mostly solved
 - Intensity gradient, color, texture
- Often used in combination with object detectors or region classifiers
- Deep learning approach is more common nowadays

More Recent Methods in Edge Detection

- Structured Forests for Fast Edge Detection (Dollar et al, 2013)
- Crisp Boundary Detection using Pointwise Mutual Information (Isola et al, 2014)
- Holistically Nested Edge Detection (Xie et al, 2015)

More Recent Methods in Edge Detection

- Holistically Nested Edge Detection (Xie et al, 2015)



	ODS	OIS	AP	FPS
Human	.80	.80	-	-
Canny	.600	.640	.580	15
Felz-Hutt [9]	.610	.640	.560	10
BEL [5]	.660*	-	-	1/10
gPb-owt-ucm [1]	.726	.757	.696	1/240
Sketch Tokens [24]	.727	.746	.780	1
SCG [31]	.739	.758	.773	1/280
SE-Var [6]	.746	.767	.803	2.5
OEF [13]	.749	.772	.817	-
DeepNets [21]	.738	.759	.758	1/5†
N4-Fields [10]	.753	.769	.784	1/6†
DeepEdge [2]	.753	.772	.807	1/10 ³ †
CSCNN [19]	.756	.775	.798	-
DeepContour [34]	.756	.773	.797	1/30†
HED (ours)	.782	.804	.833	2.5†, 1/12

Finding straight lines



Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x, y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x, y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and theta of edgelets using eig of x, y 2nd moment matrix of their points

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x, y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and theta of edgelets using eig of x, y 2nd moment matrix of their points

$$\mathbf{M} = \begin{bmatrix} \sum(x - \mu_x)^2 & \sum(x - \mu_x)(y - \mu_y) \\ \sum(x - \mu_x)(y - \mu_y) & \sum(y - \mu_y)^2 \end{bmatrix} \quad [v, \lambda] = \text{eig}(\mathbf{M})$$

Larger eigenvector
↓

$$\theta = \text{atan } 2(v(2,2), v(1,2))$$
$$conf = \lambda_2 / \lambda_1$$

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x, y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and theta of edgelets using eig of x, y 2nd moment matrix of their points

$$\mathbf{M} = \begin{bmatrix} \sum(x - \mu_x)^2 & \sum(x - \mu_x)(y - \mu_y) \\ \sum(x - \mu_x)(y - \mu_y) & \sum(y - \mu_y)^2 \end{bmatrix} \quad [v, \lambda] = \text{eig}(\mathbf{M})$$

Larger eigenvector
↓

$$\theta = \text{atan} 2(v(2,2), v(1,2))$$
$$conf = \lambda_2 / \lambda_1$$

5. Threshold on straightness, store segment

Canny lines → ... → straight edges

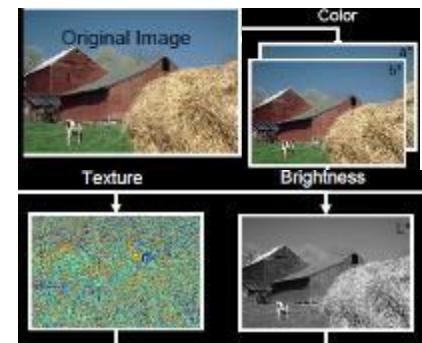


Things to remember

- Canny edge detector =
smooth → derivative → thin → threshold → link



- Pb: learns weighting of gradient, color, texture differences



- Straight line detector =
canny + gradient orientations → orientation binning
→ linking → check for straightness

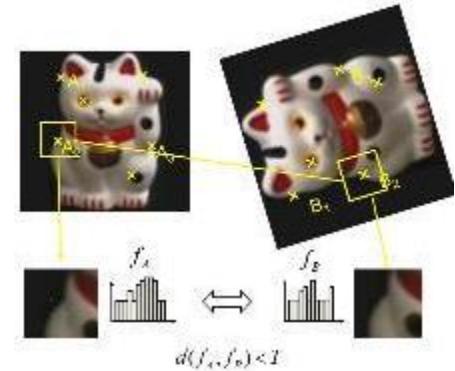


Acknowledgements

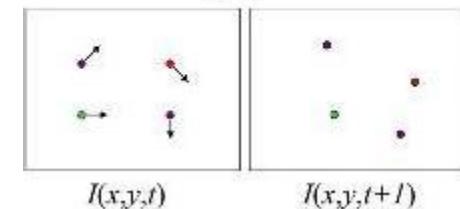
- Thanks to the following researchers for making their teaching/research material online
 - Forsyth
 - Steve Seitz
 - Noah Snavely
 - J.B. Huang
 - Derek Hoiem
 - D. Lowe
 - A. Bobick
 - S. Lazebnik
 - K. Grauman
 - R. Zaleski

Next classes: Correspondence and Alignment

- Detecting interest points



- Tracking points



- Object/image alignment and registration
 - Image stitching

