

# Interest Points







Computer Vision

Adduru U G Sankararao, IIIT Sri City

# What have we learned so far?

- Light and color
  - What an image records
- Filtering in spatial domain
  - Filtering = weighted sum of neighboring pixels
  - Smoothing, sharpening, measuring texture
- Filtering in frequency domain
  - Filtering = change frequency of the input image
  - Denoising, sampling, image compression
- Image pyramid (Gaussian and Laplacian)
  - Multi-scale analysis
- Edge detection
  - Canny edge = smooth  $\rightarrow$  derivative  $\rightarrow$  thin  $\rightarrow$  threshold  $\rightarrow$  link
  - Finding straight lines



# Today's class

- What is interest point?
- Blob detection
- Corner detection
- Handling scale and orientation

# Why extract features?



- Motivation: panorama stitching
  - We have two images – how do we combine them?





# Why extract features?


- Motivation: panorama stitching
  - We have two images – how do we combine them?



Step 1: extract features

# Why extract features?

- Motivation: panorama stitching
  - We have two images – how do we combine them?



Step 1: extract features

Step 2: match features

# Why extract features?

- Motivation: panorama stitching
  - We have two images – how do we combine them?



- Step 1: extract features
- Step 2: match features
- Step 3: align images



# Applications

- Key Points/Interest Points are used for:
  - Image alignment
  - 3D reconstruction
  - Motion tracking
  - Object recognition
  - Robot navigation
  - Indexing and database retrieval





# Desired Properties of local features


- repeatable: in a transformed image, the same feature is detected at a transformed position
  - distinctive: different image features can be discriminated by their local appearance
  - localized: relatively small regions, robust to occlusion
- 
- elongated: edges, ridges
  - + isotropic: blobs, extremal regions
  - + points: corners and junctions

# Overview of Keypoint Matching




1. Find a set of distinctive key-points

# Overview of Keypoint Matching



1. Find a set of distinctive keypoints
2. Define a region around each keypoint

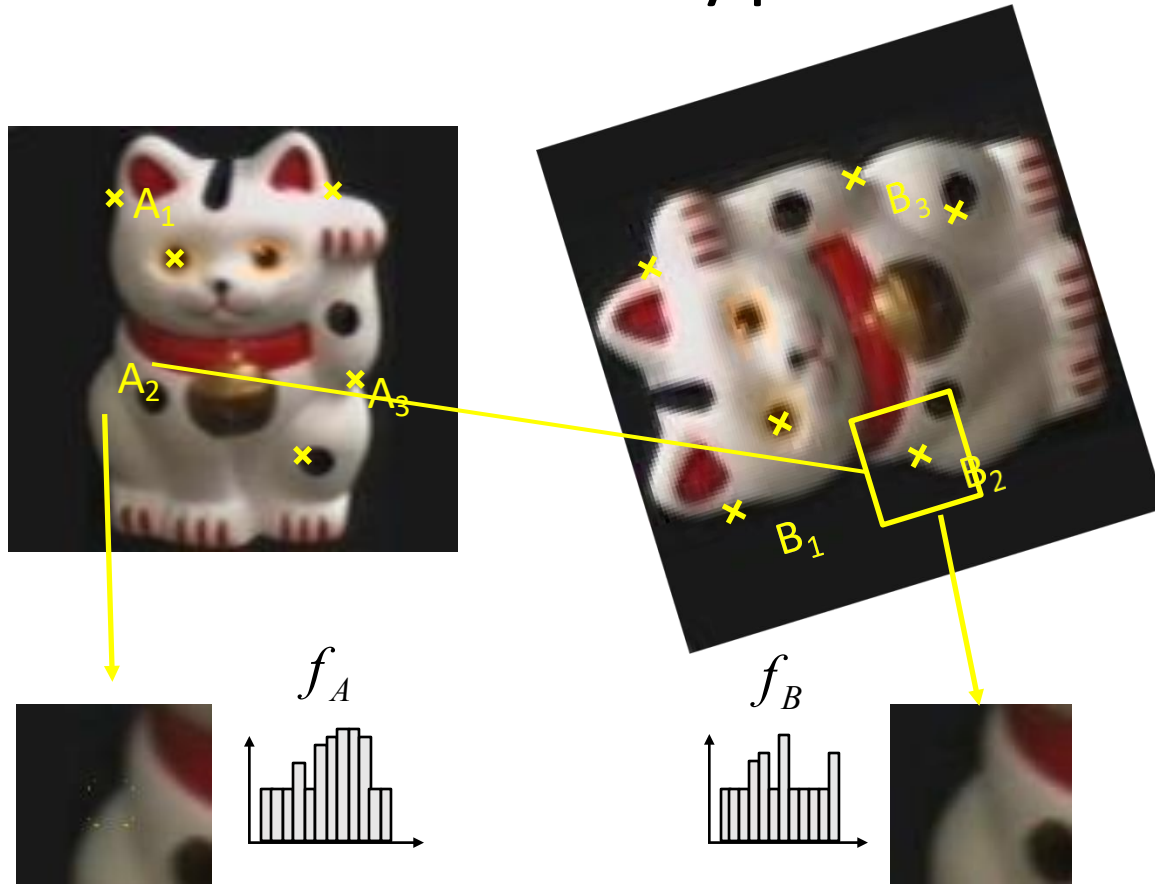
# Overview of Keypoint Matching



1. Find a set of distinctive key-points
2. Define a region around each keypoint
3. Extract and normalize the region content

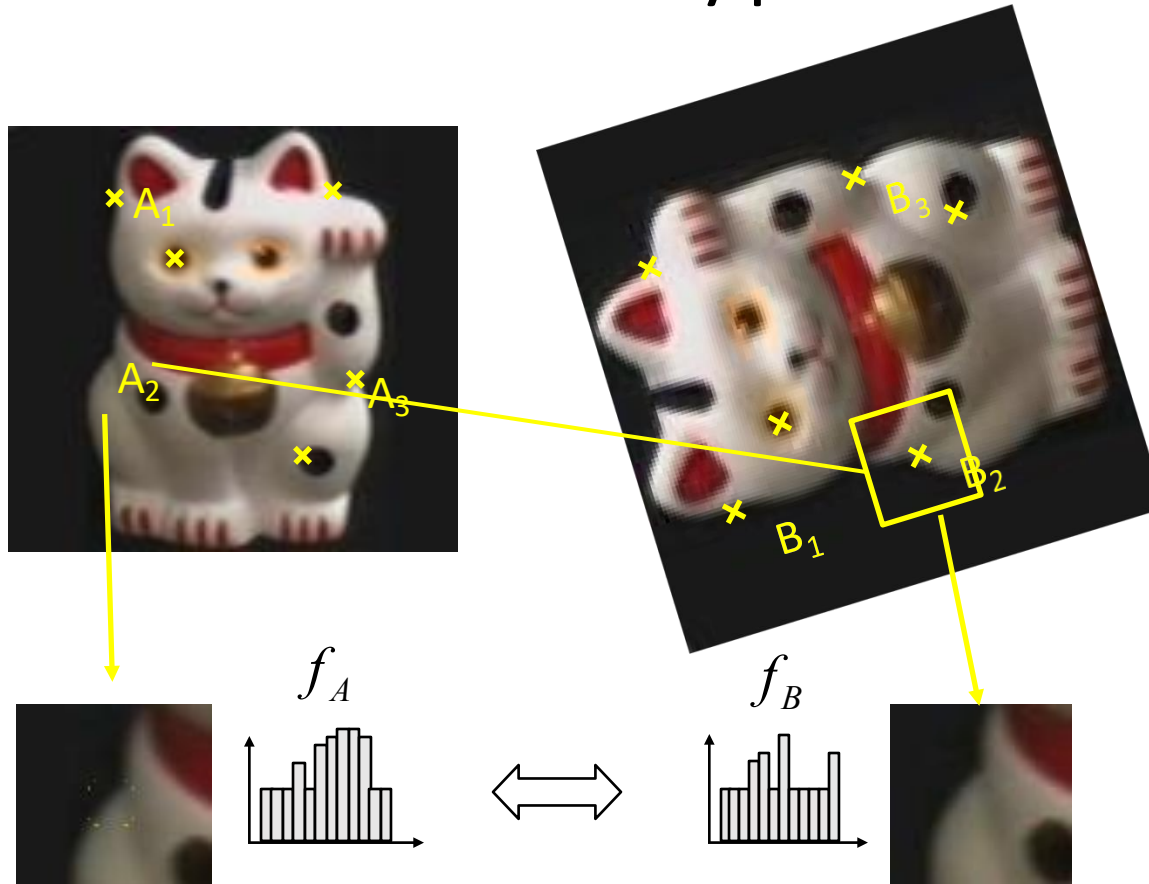


# Overview of Keypoint Matching



1. Find a set of distinctive key-points
2. Define a region around each keypoint
3. Extract and normalize the region content
4. Compute a local descriptor (feature vector) from the normalized region

# Overview of Keypoint Matching



$$d(f_A, f_B) < T$$


1. Find a set of distinctive key-points
2. Define a region around each keypoint
3. Extract and normalize the region content
4. Compute a local descriptor (feature vector) from the normalized region
5. Match local descriptors

# Goals for Keypoints




Detect points that are *repeatable* and *distinctive*

# Laplacian of Gaussian



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_{\sigma}(u, v)$$

$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$



# Laplacian of Gaussian

- Discrete approximation of the second derivative:

Example of a  $3 \times 3$   
Laplacian of Gaussian  
filter:

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

$$\frac{\partial^2 f}{\partial x^2} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

$$\frac{\partial^2 f}{\partial y^2} = f(x, y+1) + f(x, y-1) - 2f(x, y)$$

Substituting in the LoG equation:  $\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$

How did we obtain this  
filter?

$$\nabla^2 f = f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)$$

- Converting this equation to a filter results in the given LoG matrix.
-

# Laplacian of Gaussian



*Original Image.*



*Laplacian of Gaussian*

**What else can LoG do? → Blob detection**



# Blob detection

- Blobs: regions in an image that differ in properties (like brightness, colour, or texture) compared to surrounding areas.
- A blob is simply a group of connected pixels that share some common property (e.g., intensity  $>$  threshold).
- A *blob* could be circular, oval, or just any connected region of pixels.
- Blobs often correspond to objects or features in the image.



# Basic idea

- Convolve the image with a “blob filter” at multiple scales and look for extrema of filter response in the resulting *scale space*





T. Lindeberg. [Feature detection with automatic scale selection.](#)  
*IJCV* 30(2), pp 77-116, 1998.




# LoG as Blob filter

- Laplacian of Gaussian: Circularly symmetric operator, can be used for blob detection in 2D




$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Recall: Edge detection




# Edge detection, Take 2

$f$




Edge

$\frac{d^2}{dx^2} g$



Second derivative  
of Gaussian  
(Laplacian)


$f * \frac{d^2}{dx^2} g$



Edge = zero crossing  
of second derivative

# From edges to blobs

- Edge = ripple
- Blob = superposition of two ripples




**Spatial selection:** the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob



# Scale selection

- The scale of the blob is found by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:




# Scale normalization






- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by  $\sigma$
- Laplacian is the second Gaussian derivative, so it must be multiplied by  $\sigma^2$

# Effect of scale normalization






Original signal



Unnormalized Laplacian response



Scale-normalized Laplacian response





↑  
maximum

# Blob detection in 2D

- *Scale-normalized* Laplacian of Gaussian:

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$



# Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



# Scale-space blob detector: Example




# Scale-space blob detector: Example




$\sigma = 2$

# Scale-space blob detector: Example




sigma = 2.5018

# Scale-space blob detector: Example



sigma = 3.1296


# Scale-space blob detector: Example



sigma = 3.9149




# Scale-space blob detector: Example



sigma = 4.8972

# Scale-space blob detector: Example



sigma = 6.126

# Scale-space blob detector: Example



sigma = 7.6631

# Scale-space blob detector: Example



sigma = 9.5859

# Scale-space blob detector: Example



sigma = 11.9912


# Scale-space blob detector: Example



sigma = 15



# Scale-space blob detector: Example



# From Blobs to Corners

- In the following image, what are some interesting features to choose?




# From Blobs to Corners


- Look for image regions that are unusual. How to define "unusual"?
- Texture-less patches are nearly impossible to localize.
- Patches with large contrast changes (gradients) are easier to localize.
- But straight line segments at a single orientation suffer from the aperture problem (we'll see next slide), i.e., it is only possible to align the patches along the direction normal to the edge direction.
- Gradients in at least two (significantly) different orientations are the easiest, e.g., corners.

# Corner Detection: Basic Idea


- Consider a small window of pixels.
- How does the window change when you shift it?



“flat” region:  
no change in all  
directions





“edge”:  
no change along the  
edge direction



“corner”:  
significant change in  
all directions

# Computing Autocorrelation


- In the previous slide, how to quantify the "significant" change of the window?
- Answer: **Autocorrelation function** (ACF) or Sum of squared differences (SSD).



# Computing Autocorrelation

Consider shifting the window  $W$  by  $(u, v) = (\Delta u, \Delta v)$

- how do the pixels in  $W$  change?
- compare each pixel before and after by finding the ACF
- this defines an ACF “error”  $E(u, v)$ :






# Computing Autocorrelation


Consider shifting the window  $W$  by  $(u, v) = (\Delta u, \Delta v)$

- how do the pixels in  $W$  change?
- compare each pixel before and after by finding ACF
- this defines an ACF “error”  $E(u, v)$ :




$$E(u, v) = \sum_{(x, y) \in W} w(x, y) (I(x + u, y + v) - I(x, y))^2$$

Window function  $w(x, y) =$



1 in window, 0 outside

or



Gaussian

# Small motion assumption

Taylor Series expansion of  $I$ :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

# Small motion assumption

Taylor Series expansion of  $I$ :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion  $(u, v)$  is small, then first order approximation is good

$$I(x + u, y + v) \approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v$$

# Small motion assumption

Taylor Series expansion of  $I$ :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion  $(u, v)$  is small, then first order approximation is good

$$\begin{aligned} I(x+u, y+v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

shorthand:  $I_x = \frac{\partial I}{\partial x}$

# Small motion assumption

Taylor Series expansion of  $I$ :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion  $(u, v)$  is small, then first order approximation is good

$$\begin{aligned} I(x+u, y+v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$


shorthand:  $I_x = \frac{\partial I}{\partial x}$

Plugging this into the formula on the previous slide...

# Corner detection: the math

Using the small motion assumption,  
replace  $I$  with a linear approximation

(Shorthand:  $I_x = \frac{\partial I}{\partial x}$  )




$$E(u, v) = \sum_{(x, y) \in W} w(x, y) (I(x + u, y + v) - I(x, y))^2$$



# Computing Autocorrelation

Using the small motion assumption,  
replace  $I$  with a linear approximation

$$\text{(Shorthand: } I_x = \frac{\partial I}{\partial x} \text{ )}$$




$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} (I(x+u, y+v) - I(x, y))^2 \\ &\approx \sum_{(x,y) \in W} (I(x, y) + I_x(x, y)u + I_y(x, y)v - I(x, y))^2 \end{aligned}$$

# Computing Autocorrelation

Using the small motion assumption,  
replace  $I$  with a linear approximation


$$\text{(Shorthand: } I_x = \frac{\partial I}{\partial x} \text{ )}$$



$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} (I(x+u, y+v) - I(x, y))^2 \\ &\approx \sum_{(x,y) \in W} (I(x, y) + I_x(x, y)u + I_y(x, y)v - I(x, y))^2 \\ &\approx \sum_{(x,y) \in W} (I_x(x, y)u + I_y(x, y)v)^2 \end{aligned}$$


# Computing Autocorrelation

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} (I_x(x, y)u + I_y(x, y)v)^2 \\ &\approx \sum_{(x,y) \in W} (I_x^2 u^2 + 2I_x I_y uv + I_y^2 v^2) \end{aligned}$$



# Computing Autocorrelation

$$\begin{aligned}
 E(u, v) &\approx \sum_{(x,y) \in W} w(x,y) (I_x(x, y)u + I_y(x, y)v)^2 \\
 &\approx \sum_{(x,y) \in W} w(x,y) (I_x^2 u^2 + 2I_x I_y uv + I_y^2 v^2)
 \end{aligned}$$



$$E(u, v) = [u \ v] H [u \ v]^T$$

$$H = \sum \sum w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = w * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Thus,  $E(u,v)$  is locally approximated as a *quadratic form*


# Computing Autocorrelation

- The weighted summations have been replaced with discrete convolutions with the weighting kernel  $w$ .
- The eigenvalues of  $\mathbf{H}$  reveal the amount of intensity change in the two principal orthogonal gradient directions in the window.

$$H = \mathbf{U} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \mathbf{U}^T \quad \text{with} \quad H\mathbf{u}_i = \lambda\mathbf{u}_i$$

# Computing Autocorrelation

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} (I_x(x, y)u + I_y(x, y)v)^2 \\ &\approx \sum_{(x,y) \in W} (I_x^2 u^2 + 2I_x I_y uv + I_y^2 v^2) \\ &\approx Au^2 + 2Buv + Cv^2 \end{aligned}$$



$$A = \sum_{(x,y) \in W} I_x^2 \quad B = \sum_{(x,y) \in W} I_x I_y \quad C = \sum_{(x,y) \in W} I_y^2$$

- Thus,  $E(u,v)$  is locally approximated as a *quadratic form*

# The second moment matrix

The surface  $E(u,v)$  is locally approximated by a quadratic form.

$$E(u, v) \approx Au^2 + 2Buv + Cv^2$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$

# The second moment matrix

The surface  $E(u,v)$  is locally approximated by a quadratic form.

$$\begin{aligned} E(u, v) &\approx Au^2 + 2Buv + Cv^2 \\ &\approx \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



# The second moment matrix

The surface  $E(u,v)$  is locally approximated by a quadratic form.

$$\begin{aligned} E(u, v) &\approx Au^2 + 2Buv + Cv^2 \\ &\approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$

# The second moment matrix


The surface  $E(u,v)$  is locally approximated by a quadratic form.

$$E(u, v) \approx Au^2 + 2Buv + Cv^2$$
$$\approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$




Let's try to understand its shape.

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$




Horizontal edge:  $I_x = 0$

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Horizontal edge:  $I_x = 0$


$$H = \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix}$$

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$


$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Horizontal edge:  $I_x = 0$

$$H = \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix}$$




$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$




Vertical edge:  $I_y = 0$

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Vertical edge:  $I_y = 0$


$$H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$$

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$A = \sum_{(x,y) \in W} I_x^2$$


$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$



Vertical edge:  $I_y = 0$

$$H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$$






# Interpretation

- The shape of  $H$  tells us something about the distribution of gradients around a pixel
- We can visualize  $H$  as an ellipse with axis lengths determined by the eigenvalues of  $H$  and orientation determined by the eigenvectors of  $H$

Ellipse equation:

$$\begin{bmatrix} u & v \end{bmatrix} H \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$



# Quick eigenvalue/eigenvector review

The **eigenvectors** of a matrix **A** are the vectors **x** that satisfy:

$$Ax = \lambda x$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to **x**

# Quick eigenvalue/eigenvector review

The **eigenvectors** of a matrix **A** are the vectors **x** that satisfy:

$$Ax = \lambda x$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to **x**

- The eigenvalues are found by solving:

$$\det(A - \lambda I) = 0$$

# Quick eigenvalue/eigenvector review

The **eigenvectors** of a matrix **A** are the vectors **x** that satisfy:

$$Ax = \lambda x$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to **x**

- The eigenvalues are found by solving:

$$\det(A - \lambda I) = 0$$

- In our case, **A** = **H** is a 2x2 matrix, so we have

$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$$

# Quick eigenvalue/eigenvector review

The **eigenvectors** of a matrix **A** are the vectors **x** that satisfy:

$$Ax = \lambda x$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to **x**

- The eigenvalues are found by solving:

$$\det(A - \lambda I) = 0$$

- In our case, **A** = **H** is a 2x2 matrix, so we have

- The solution:  $\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$

$$\lambda_{\pm} = \frac{1}{2} \left[ (h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

# Quick eigenvalue/eigenvector review

The **eigenvectors** of a matrix **A** are the vectors **x** that satisfy:

$$Ax = \lambda x$$

The scalar  $\lambda$  is the **eigenvalue** corresponding to **x**

- The eigenvalues are found by solving:

$$\det(A - \lambda I) = 0$$

- In our case, **A** = **H** is a 2x2 matrix, so we have

- The solution: 
$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$$


$$\lambda_{\pm} = \frac{1}{2} \left[ (h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

Once you know  $\lambda$ , you find **x** by solving

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$


# Interpreting the eigenvalues

*How do the eigenvalues determine if an image point is a corner?*



# Interpreting the eigenvalues


*How do the eigenvalues determine if an image point is a corner?*



“edge”:

$$\lambda_1 \gg \lambda_2$$


$$\lambda_2 \gg \lambda_1$$



“corner”:

$\lambda_1$  and  $\lambda_2$  are large,

$$\lambda_1 \sim \lambda_2;$$



“flat” region

$\lambda_1$  and  $\lambda_2$  are small;



# Harris Corner Detector

Here's what you do

- Compute the gradient at each point in the image
- Compute  $\mathbf{H}$  for each image window to get its cornerness scores.
- Compute the eigen values or compute the following function


$$M_c = \lambda_1 \lambda_2 - \kappa (\lambda_1 + \lambda_2)^2 = \det(\mathbf{H}) - k \operatorname{trace}^2(\mathbf{H})$$

- Find points whose surrounding window gave larger cornerness response ( $M_c > \text{threshold}$ )
- Take points of local maxima, perform non-maximum suppression.

# Harris detector example




Computer cornerness scores (red high, blue low)






Computer cornerness scores (red high, blue low)



Threshold ( $M_c > \text{value}$ )



Find local maxima of  $M_c$  (Nonmax Suppression)



# Harris features (in red)



# Harris Corner Detector: Variants

- Harris and Stephens '88 is rotationally invariant and downweights edge-like features where  $\lambda_1 \gg \lambda_0$ .

$$\det(\mathbf{A}) - \alpha \text{trace}(\mathbf{A})^2 = \lambda_0 \lambda_1 - \alpha (\lambda_0 + \lambda_1)^2$$


- Triggs '04 suggested  $\lambda_0 - \alpha \lambda_1$ .
- Brown et al, '05 use harmonic mean:

$$\frac{\det(\mathbf{A})}{\text{trace}(\mathbf{A})} = \frac{\lambda_0 \lambda_1}{\lambda_0 + \lambda_1}$$


which is smoother when  $\lambda_0 \approx \lambda_1$



# Harris Detector – Responses [Harris88]



*Effect:* A very precise corner detector.




# Harris Detector – Responses [Harris88]




# Harris Detector: Invariance Properties

- Scaling




# Harris Detector: Invariance Properties

- Scaling



# Harris Detector: Invariance Properties


- Scaling



*Not invariant to scaling*

# Harris Detector: Invariance Properties


- Rotation



Ellipse rotates but its shape (i.e. eigenvalues) remains the same

# Harris Detector: Invariance Properties

- Rotation



Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response is invariant to image rotation



# Harris Detector: Invariance Properties

- Photometric change: Affine intensity change:  $I \rightarrow aI + b$ 
  - ✓ Only derivatives are used =>  
invariance to intensity shift  $I \rightarrow I + b$





# Harris Detector: Invariance Properties

- Photometric change: Affine intensity change:  $I \rightarrow aI + b$ 
  - ✓ Only derivatives are used =>  
invariance to intensity shift  $I \rightarrow I + b$
  - ✓ Intensity scale:  $I \rightarrow aI$



# Harris Detector: Invariance Properties

- Photometric change: Affine intensity change:  $I \rightarrow aI + b$ 
  - ✓ Only derivatives are used =>  
invariance to intensity shift  $I \rightarrow I + b$
  - ✓ Intensity scale:  $I \rightarrow aI$



*Partially invariant to affine intensity change*

# Things to remember

- Keypoint detection: repeatable and distinctive
  - Corners,
  - Invariant to scale, rotation, etc.
- Harris Corner Detection
  - Rotation Invariant
  - Partial Intensity Change Invariant
  - *Not Invariant to Scale*

