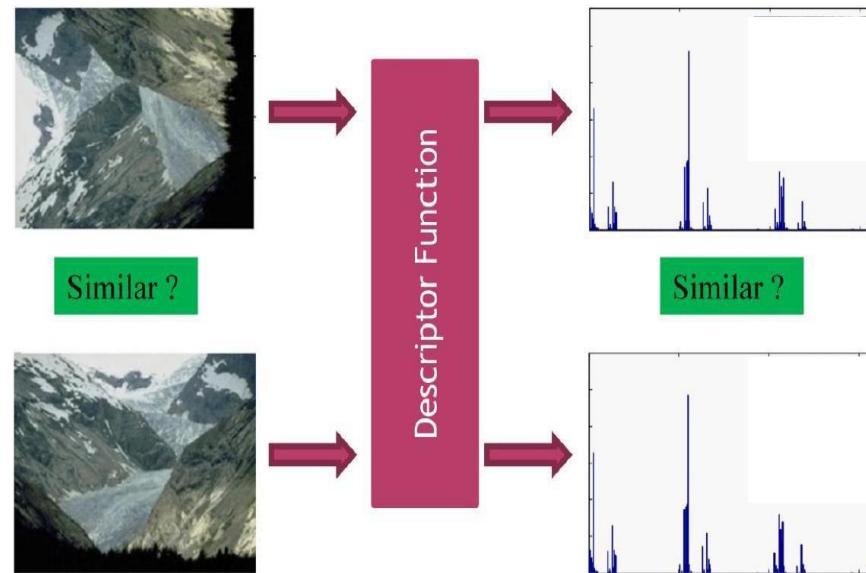


# Feature Detection and Description

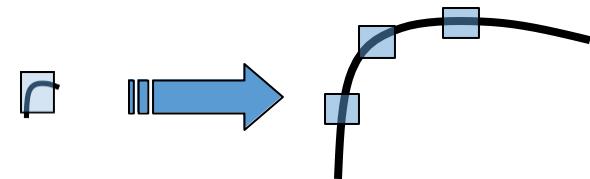


Computer Vision

Adduru U G Sankararao, IIIT Sri City

# Previous Class

- Keypoint detection: repeatable and distinctive
  - Corners,
  - Invariant to scale, rotation, etc.
- Harris Corner Detection
  - Rotation Invariant
  - Partial Intensity Change Invariant
  - *Not Invariant to Scale*

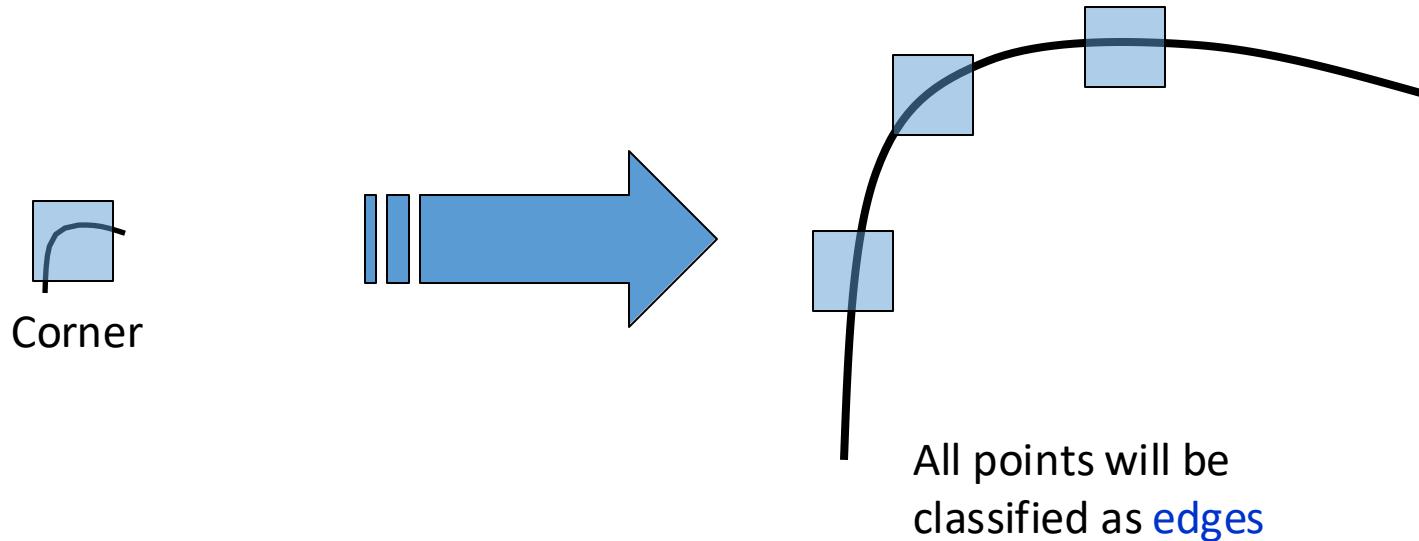


# This lecture

- Scale Space, Image Pyramids, Filter Banks
- Scale Invariant Feature Transform (SIFT)

# Harris Detector: Invariance Properties

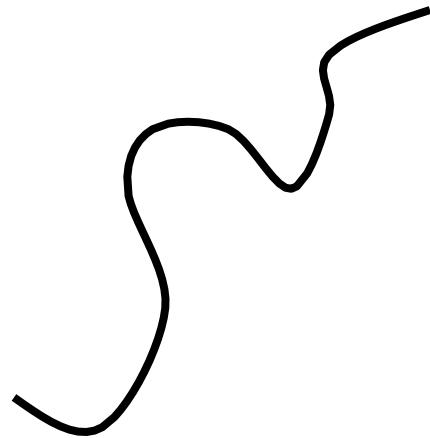
- Scaling



*Not invariant to scaling*

# Scale invariant detection

Suppose you're looking for corners

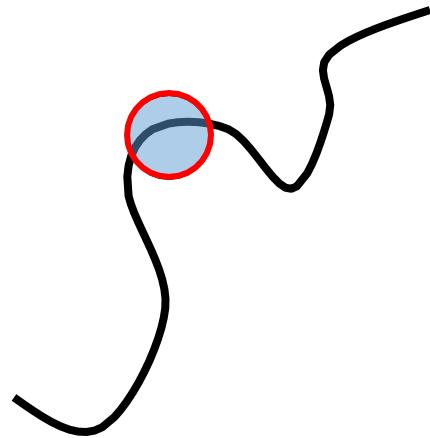


Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Scale invariant detection

Suppose you're looking for corners

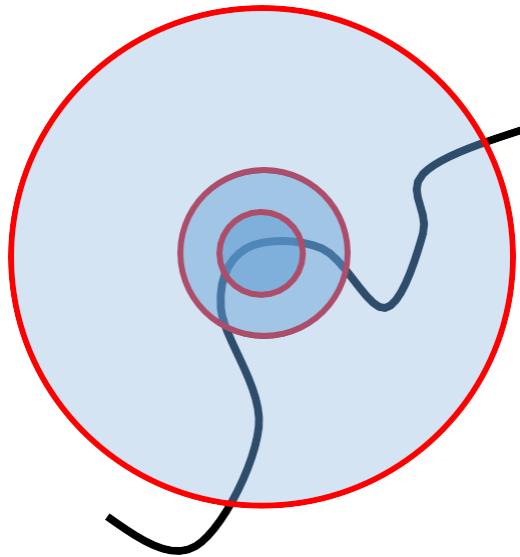


Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Scale invariant detection

Suppose you're looking for corners



Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Automatic Scale Selection

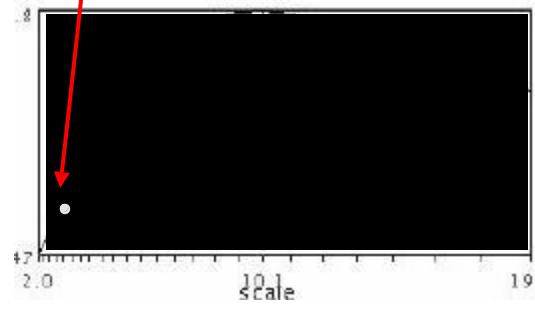
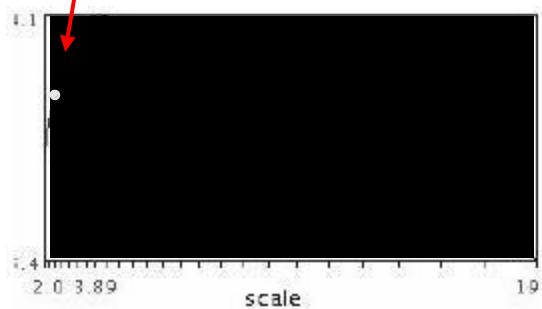


$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

How to find corresponding patch sizes?

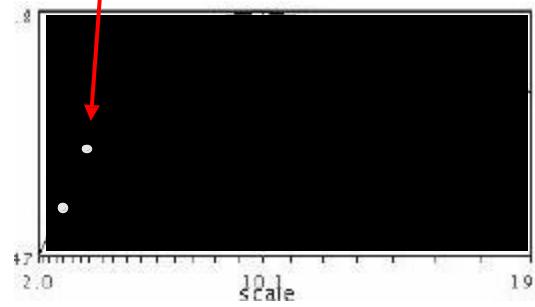
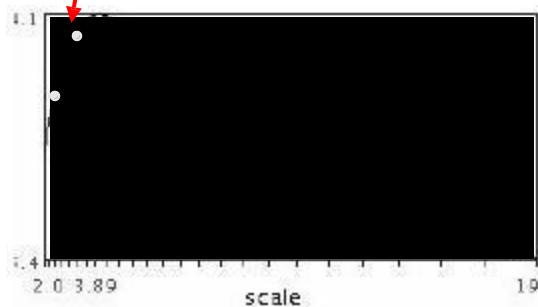
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



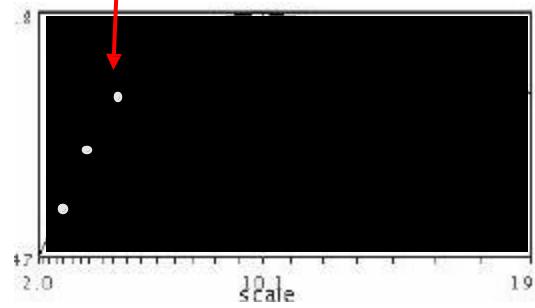
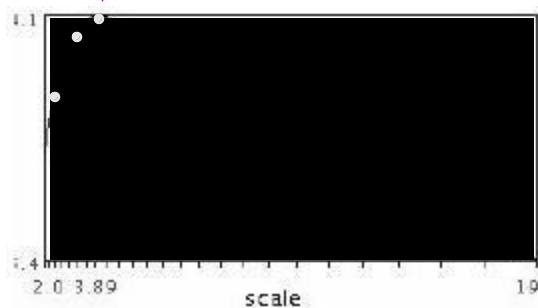
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



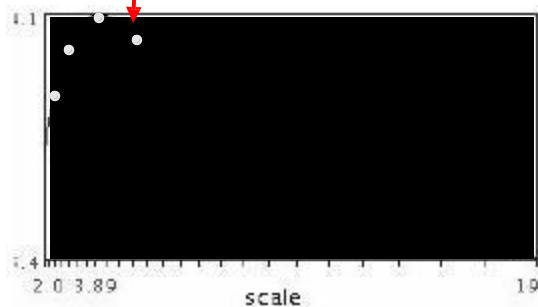
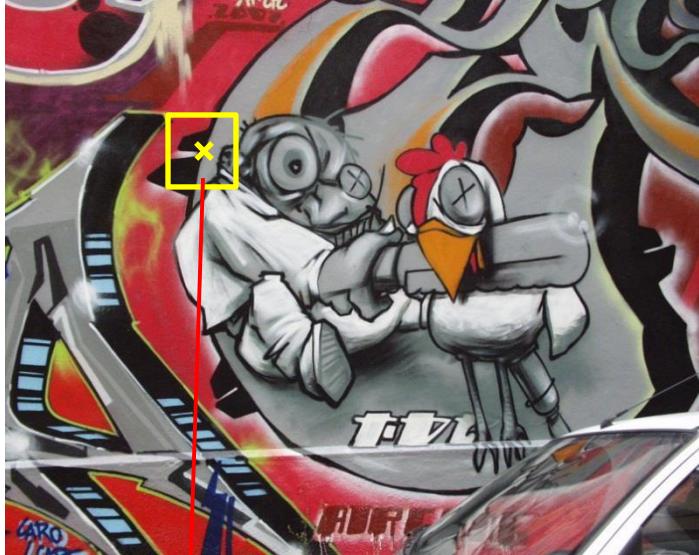
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)

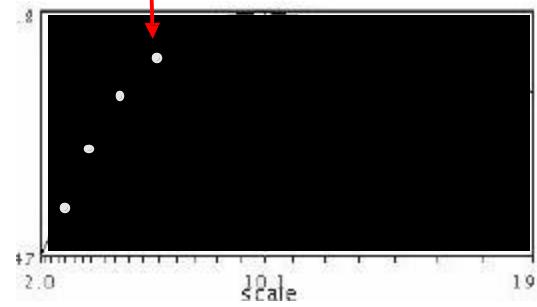


# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



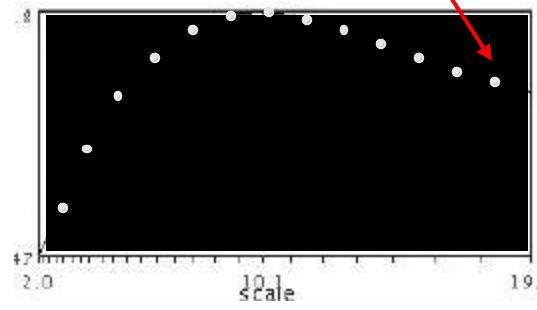
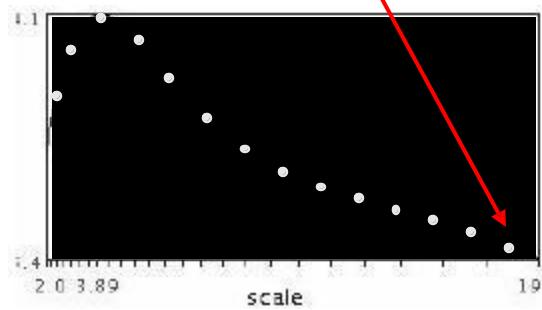
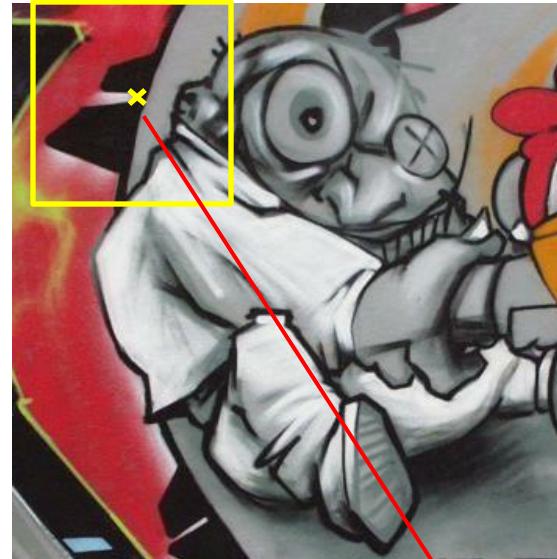
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

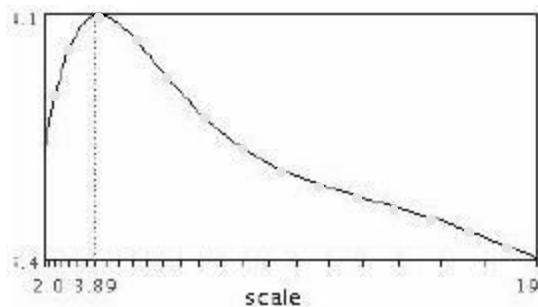
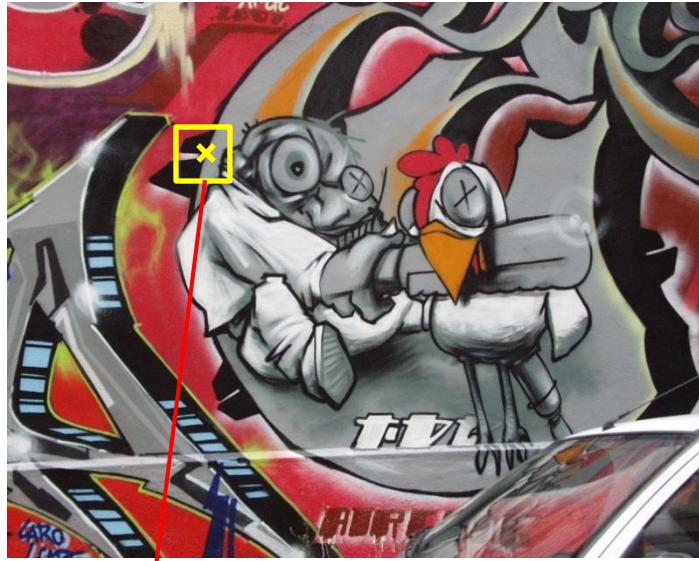
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)

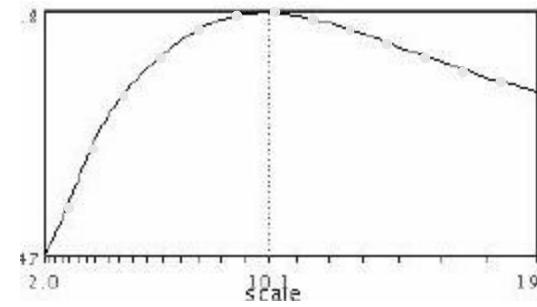


# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma'))$$

K. Grauman, B. Leibe

Is there a better way to do this?

# Automatic Scale Selection:Implementation

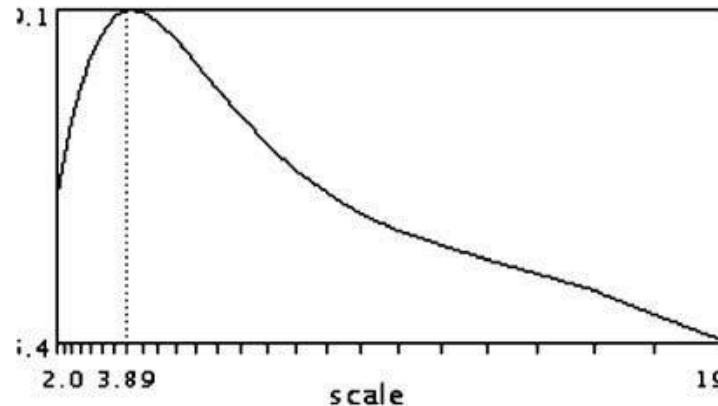
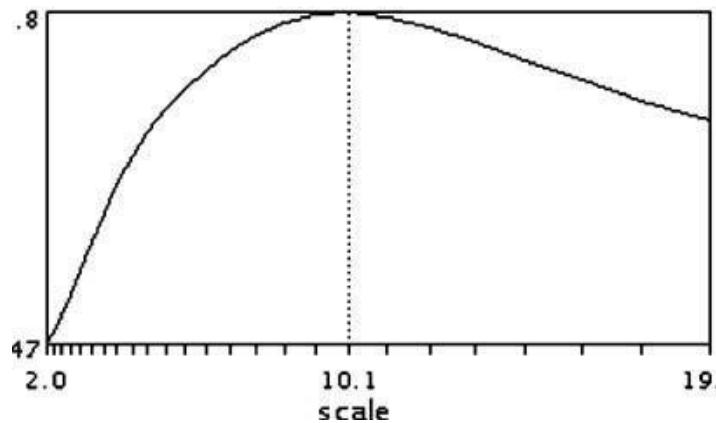
- Instead of computing  $f$  for larger and larger windows, we can implement using a fixed window size with a Gaussian pyramid



(sometimes need to create in-between levels, e.g. a  $\frac{3}{4}$ -size image)

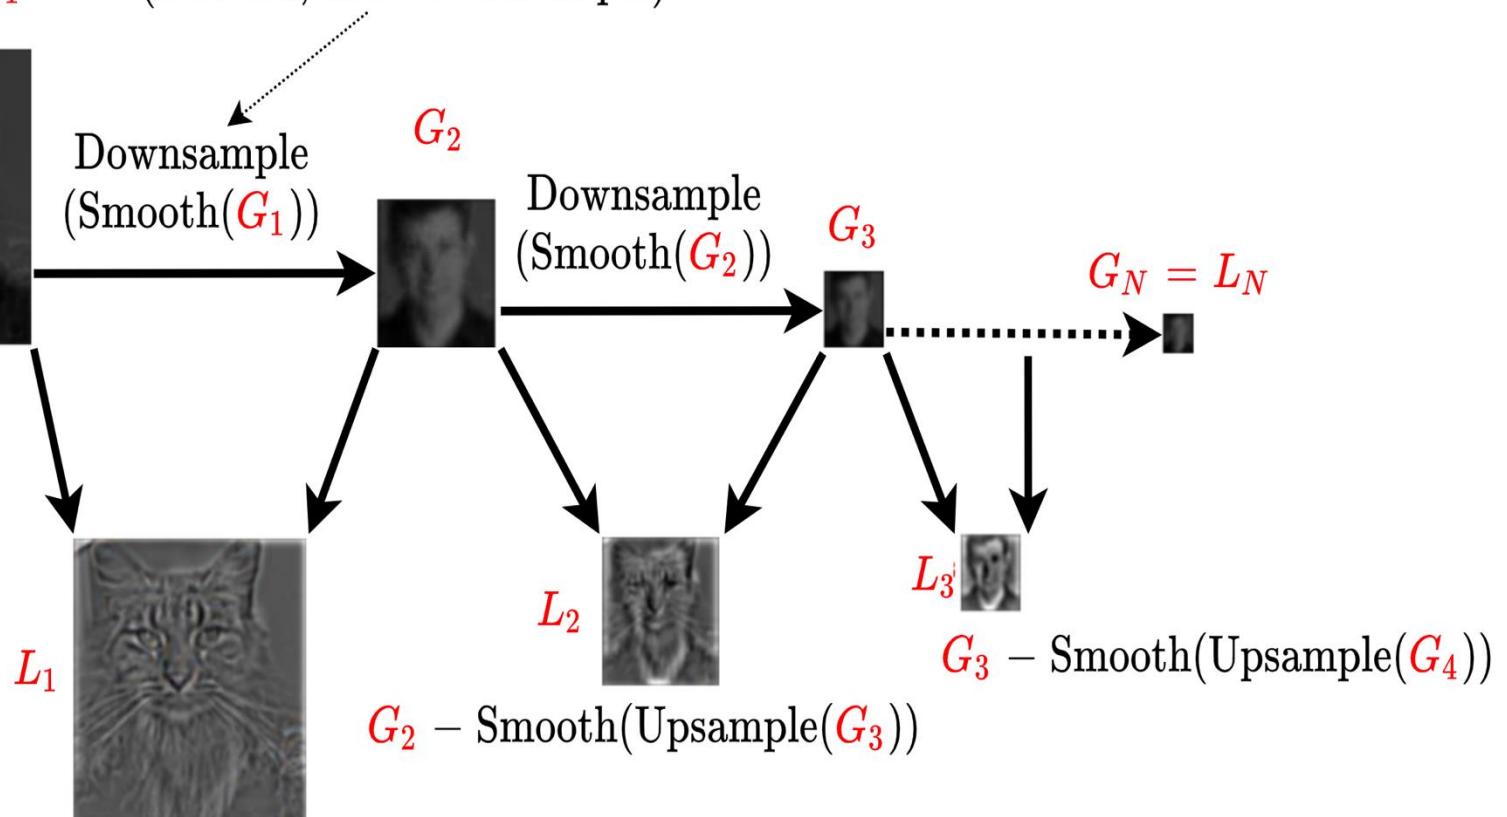
# Keypoint detection with scale selection

- We want to extract keypoints with characteristic scale that is *covariant* with the image transformation



# Gaussian and Laplacian Pyramid

Image =  $G_1$  (Smooth, then Downsample)

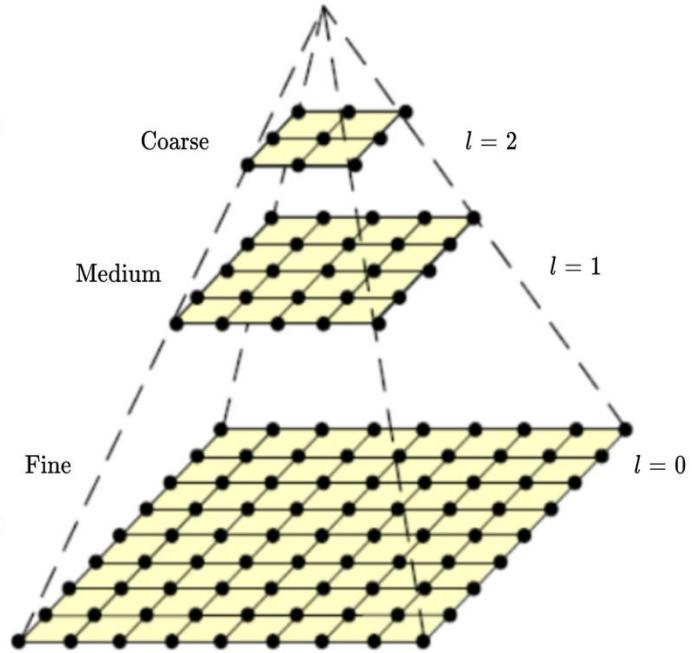


$G_1 - \text{Smooth}(\text{Upsample}(G_2))$

Credit: Derek Hoiem

# Images Pyramids: Uses

- Compression
- Object detection
  - Scale search
  - Features
- Detecting stable interest points
- Registration: Coarse-to-fine  
Image Registration

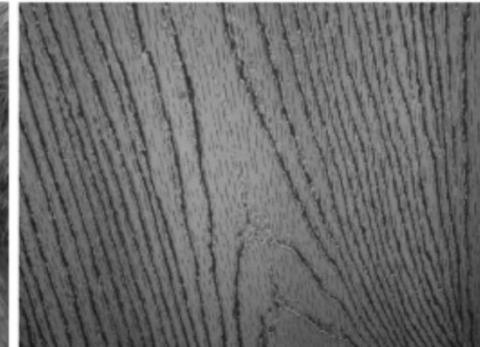


- *Compute Gaussian pyramid.*
- *Align with coarse pyramid.*
- *Successively align with finer pyramids.*
  - *Search smaller range.*

# Textures in Images

- Regular or stochastic patterns caused by bumps, grooves and/or markings.
- Gives us information about spatial arrangement of colours or intensities in an image.

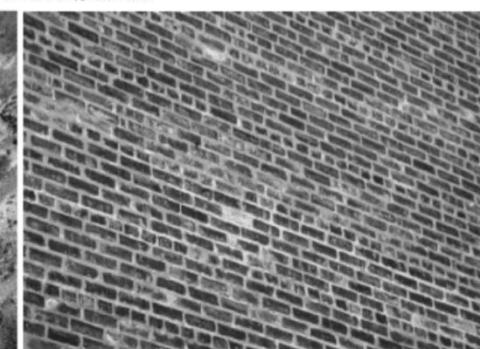
Different Materials



Different Orientation



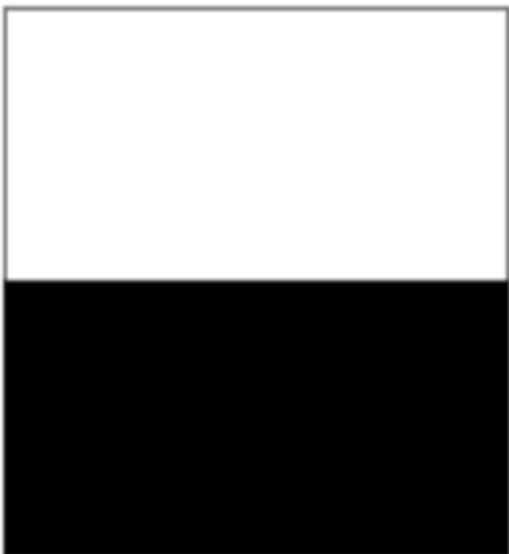
Different Scales



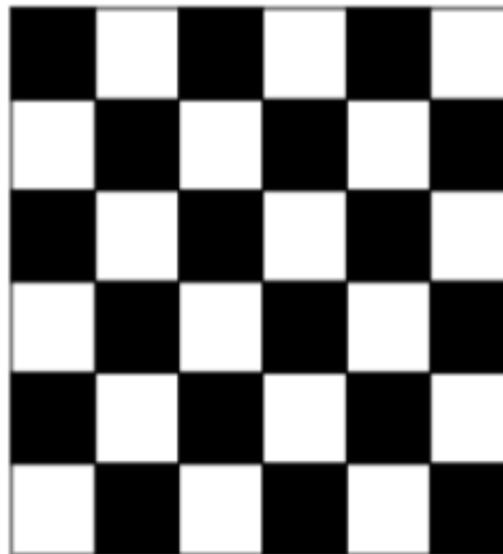
# Textures in Images

- Conveys more information that can be exploited to match regions of interest in images.

Histogram conveys 50% white pixels and 50% black pixels



(Block Pattern)



(Checkerboard Pattern)



(Striped Pattern)

Drastically different textures

# Textures in Images

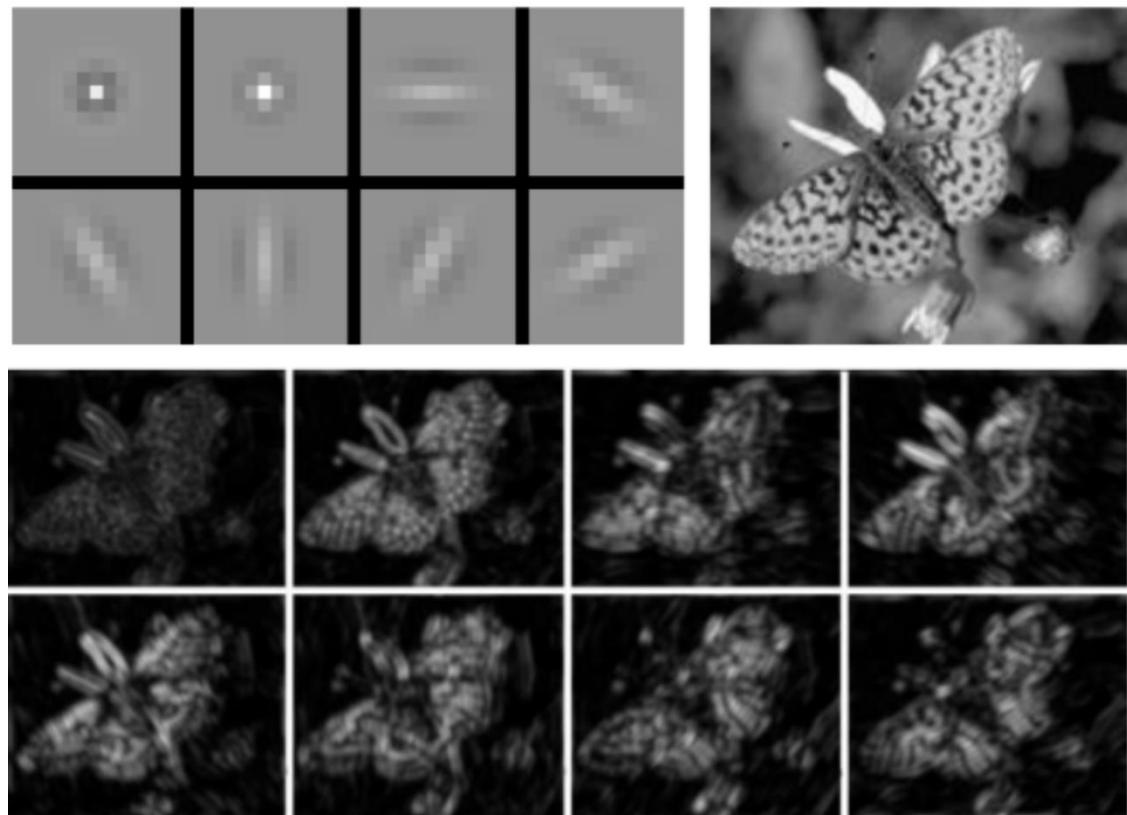
- Conveys more information that can be exploited to match regions of interest in images.

## How to represent textures?

- Compute responses of blobs and edges at various orientations and scales.
- Ways to process:
  - Record simple statistics (e.g., mean, std.) of absolute filter responses.
  - Take vectors of filter responses at each pixel and cluster them.

# Filter Banks

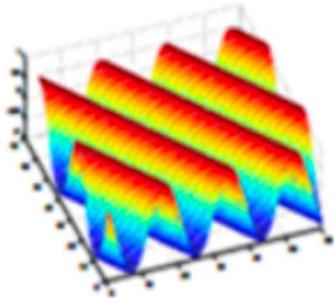
- An array of bandpass filters that separates the input signal into multiple components, each one carrying a single frequency sub-band of the original signal.
- Process image with each filter and keep responses (or squared/abs responses).



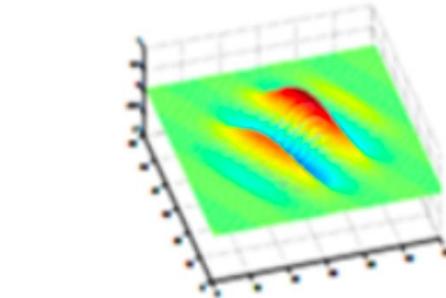
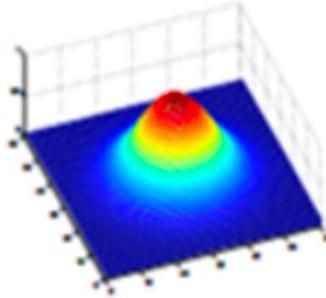
# Gabor Filters

- Special classes of bandpass filters (i.e., they allow a certain 'band' of frequencies and reject the others).
- A Gabor filter can be viewed as a sinusoidal signal of particular frequency and orientation, modulated by a Gaussian wave.

A 2-D Gaussian



A sinusoid oriented  $30^\circ$  with x-axis

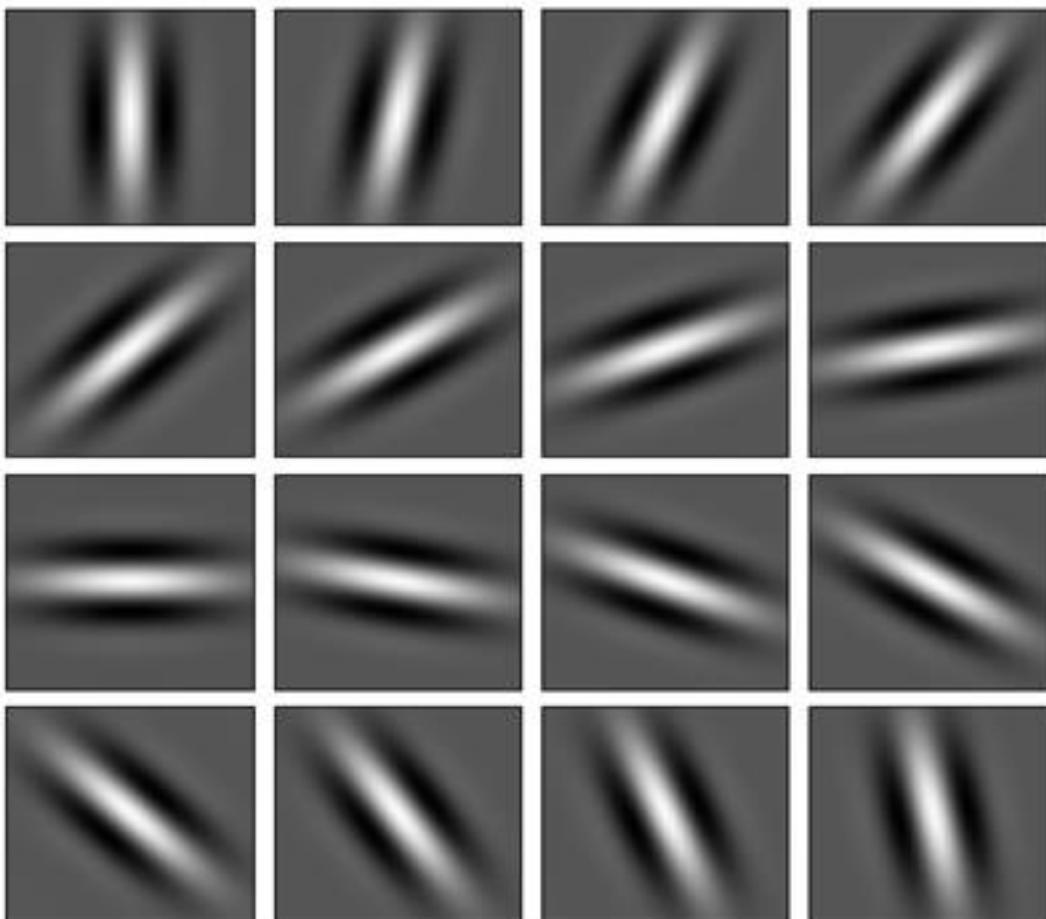


A corresponding 2-D Gabor Filter

*A 2-D Gabor filter obtained by modulating the sine wave with a Gaussian*

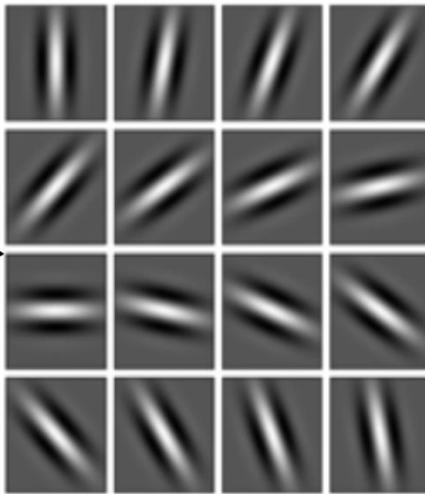
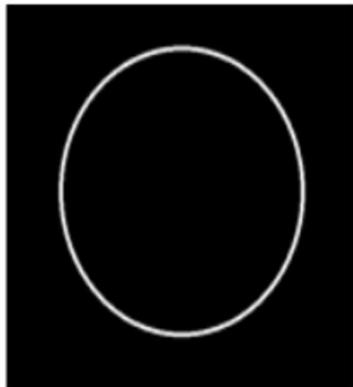
# Gabor Filter Banks

- Bank of 16 Gabor filters at an orientation of 11.25° (i.e. if the first filter is at 00.00, then the second will be at 11.25, the third will be at 22.50, and so on.)

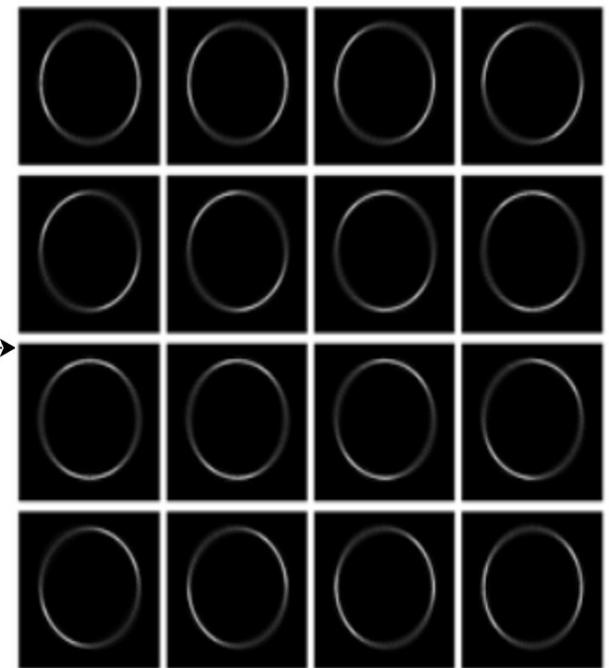


# Gabor Filter Banks

- Bank of 16 Gabor filters at an orientation of 11.25° (i.e. if the first filter is at 00.00, then the second will be at 11.25, the third will be at 22.50, and so on.)



Input Image of a Circle



A bank of 16 Gabor filters

Output Image of the circle after passed through individual Gabor filters

# Scale Invariant Feature Transform (SIFT)

- David G. Lowe, [Distinctive Image Features from Scale-invariant Keypoints](#), IJCV 2004: Over 50000 citations
- Transforms image data into scale-invariant coordinates
- Fundamental to many core vision problems/applications:
  - Recognition, Motion tracking, Multiview geometry

# Widely Used Local Descriptors

- SIFT – Scale Invariant Feature Transform

Distinctive image **features** from **scale-invariant** keypoints

DG Lowe - International journal of computer vision, 2004 - Springer

... the assigned orientation, **scale**, and location for each **feature**, thereby providing **invariance** to these ... **Invariant** Feature **Transform** (SIFT), as it **transforms** image data into **scale-invariant** coordinates relative ... that densely cover the image over the full range of **scales** and locations ...

☆ 99 [Cited by 50252](#) Related articles All 179 versions

- LBP – Local Binary Pattern

Multiresolution gray-scale and rotation invariant texture classification with **local binary patterns**

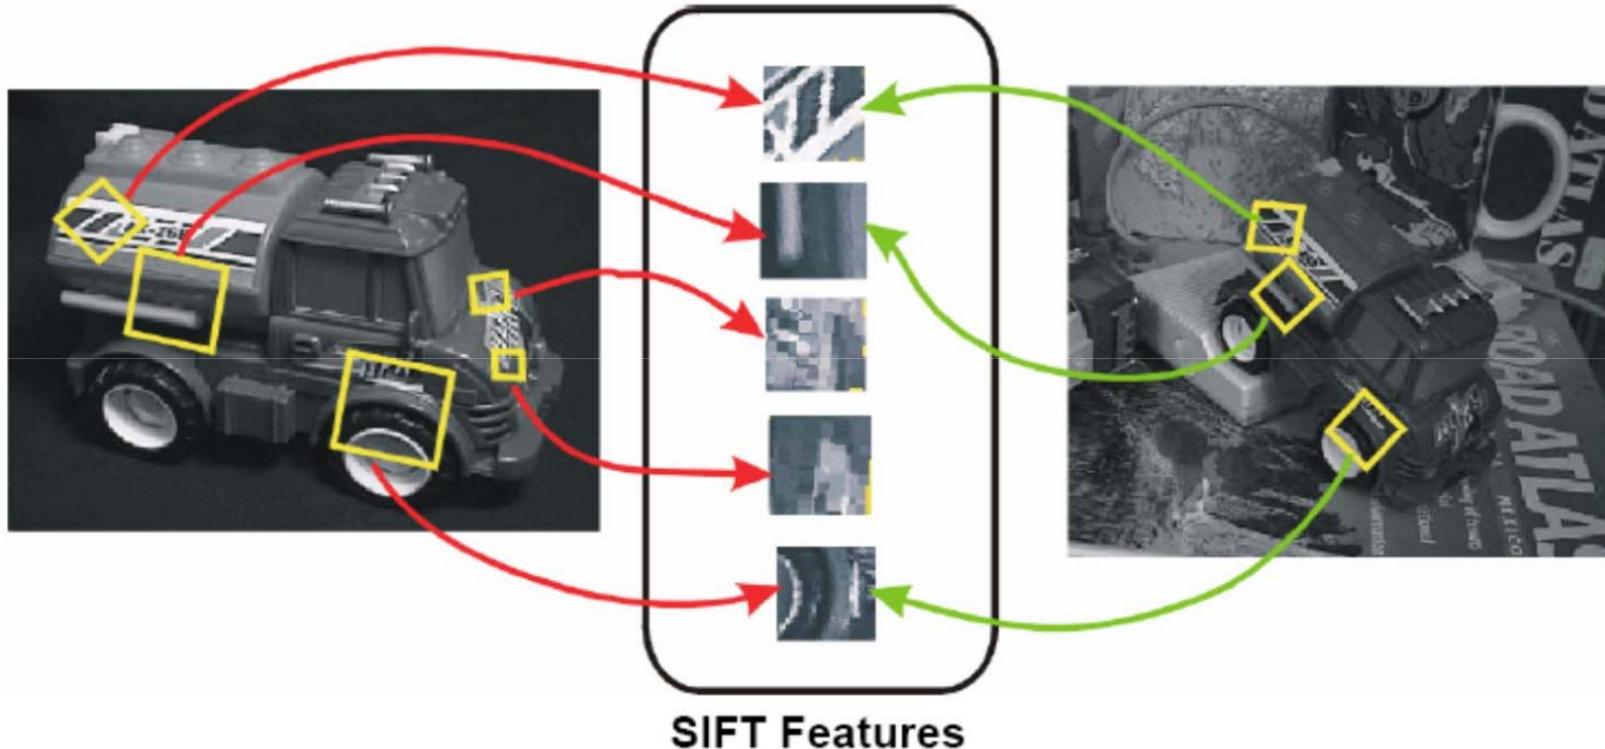
T Ojala, M Pietikainen... - ... Transactions on **pattern** ..., 2002 - ieeexplore.ieee.org

Presents a theoretically very simple, yet efficient, multiresolution approach to gray-scale and rotation invariant texture classification based on **local binary patterns** and nonparametric discrimination of sample and prototype distributions. The method is based on recognizing ...

☆ 99 [Cited by 12251](#) Related articles All 16 versions

# SIFT: Invariant Local Features

- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, shear.

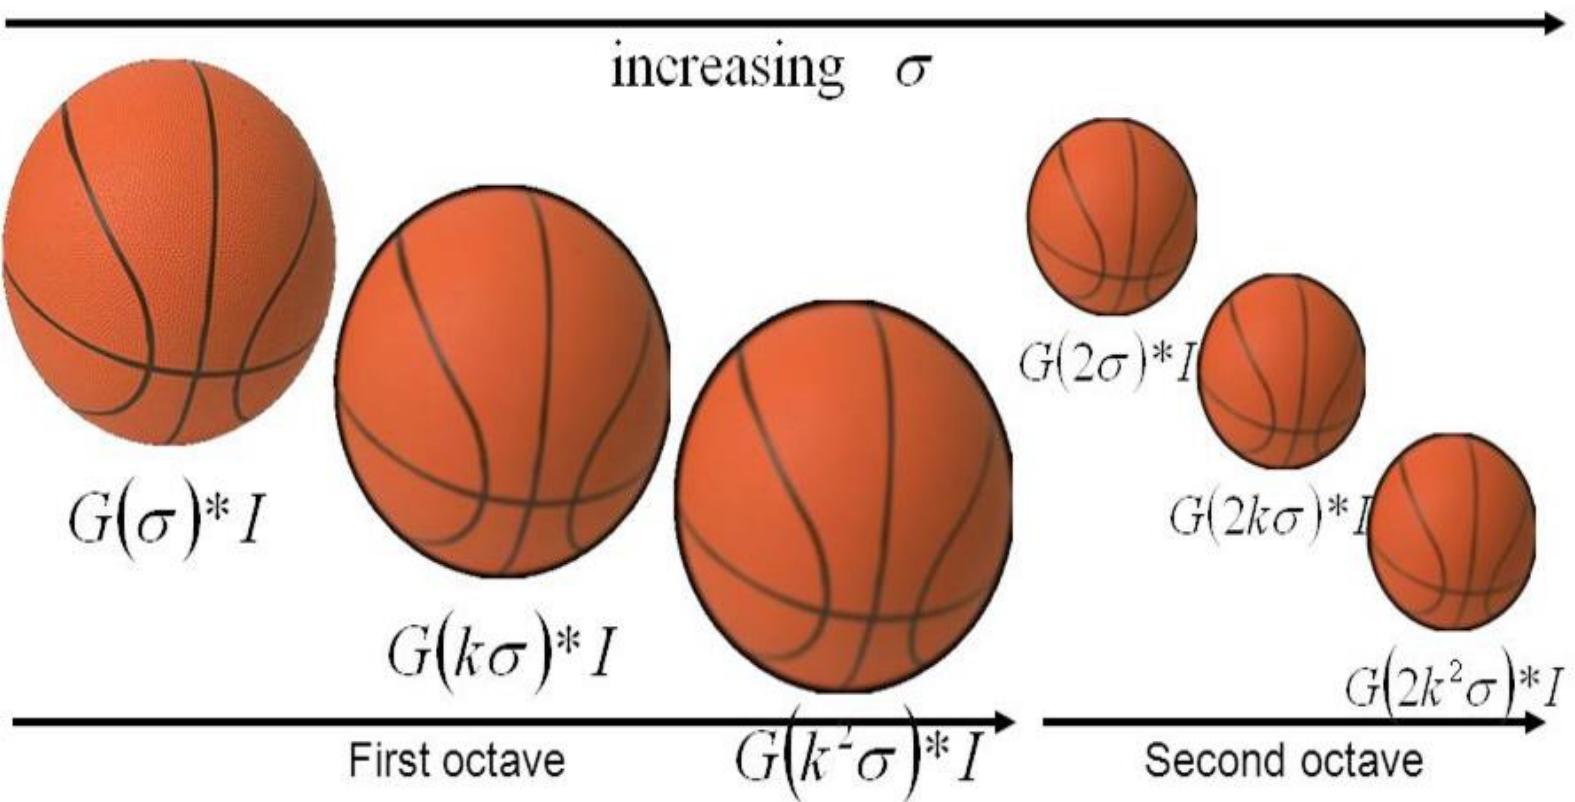


# SIFT:

- Step 1: Scale-space Extrema Detection - Detect interest points (invariant to scale and orientation) using DOG.
- Step 2: Keypoint Localization - Determine location and scale at each candidate location, and select them based on stability.
- Step 3: Orientation Estimation - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- Step 4: Keypoint Descriptor - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion.

# SIFT: Scale Space Extrema Detection

- Constructing Scale-space



# SIFT: Scale Space Extrema Detection

- Difference of Gaussians

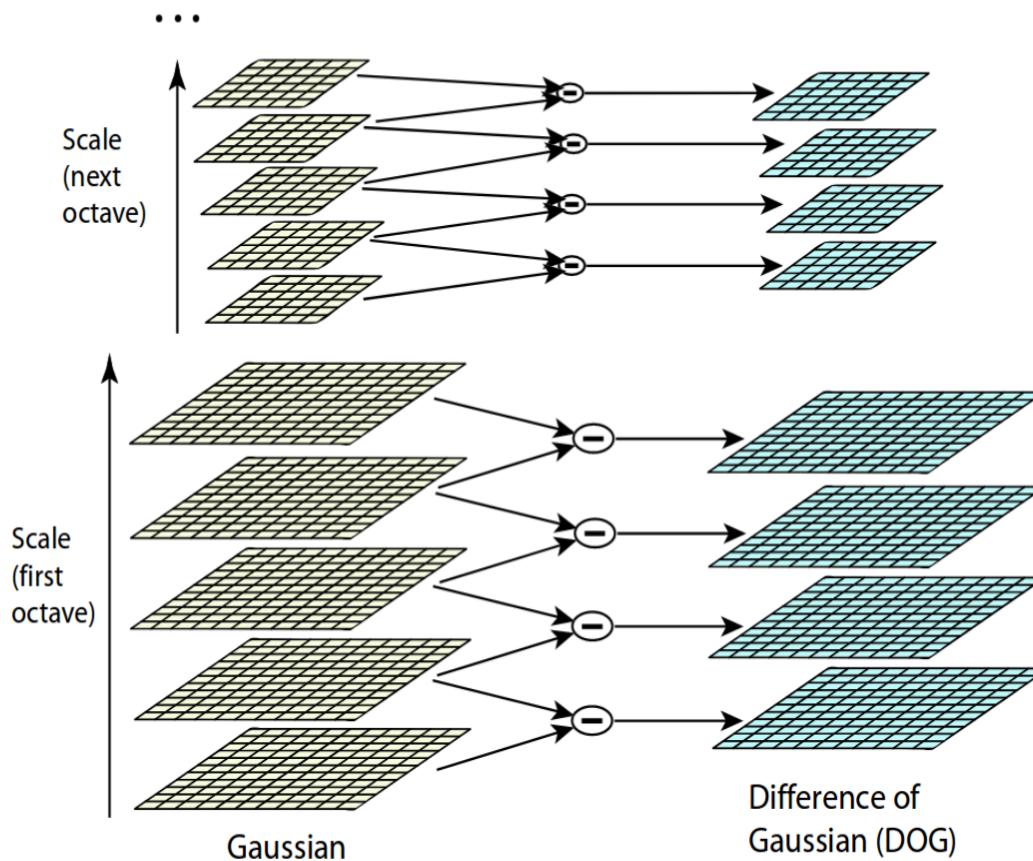
An octave corresponds to doubling the value of  $\sigma$

$$D(x, y, \sigma) = \hat{I}(x, y, \sigma) - \hat{I}(x, y, k\sigma)$$

where  $\hat{I}(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$

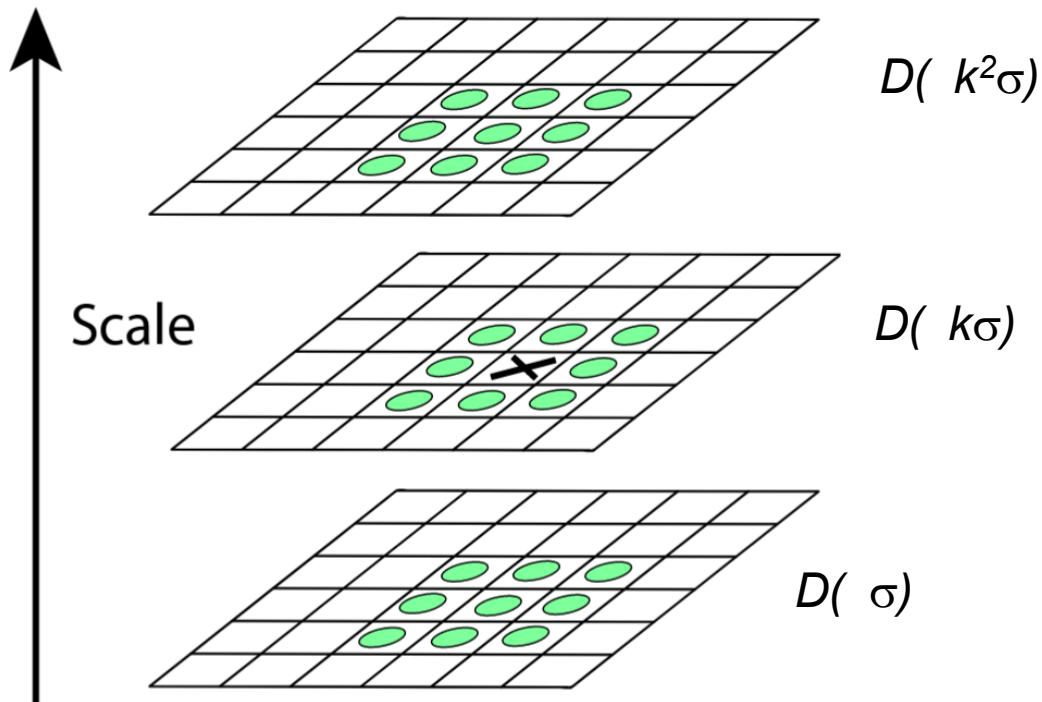


*Original image*



# SIFT: Scale Space Extrema Detection

- Compare a pixel (X) with 26 pixels in current and adjacent scales (Green Circles)
- Select a pixel (X) if it is larger/smaller than all 26 pixels



# Difference-of-Gaussian (DoG)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

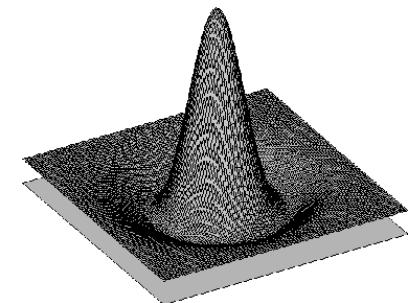
(Difference of Gaussians)



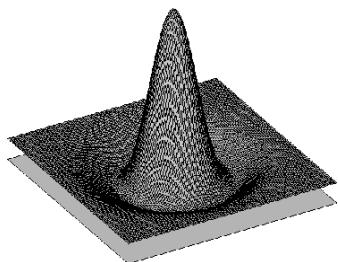
-



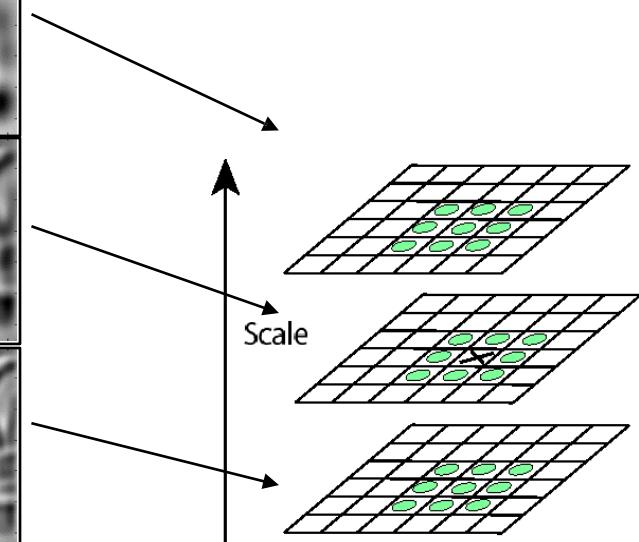
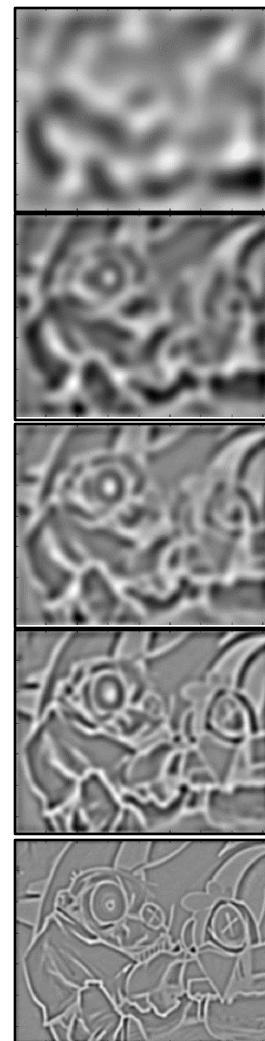
=



# Find local maxima in position-scale space of Difference-of-Gaussian

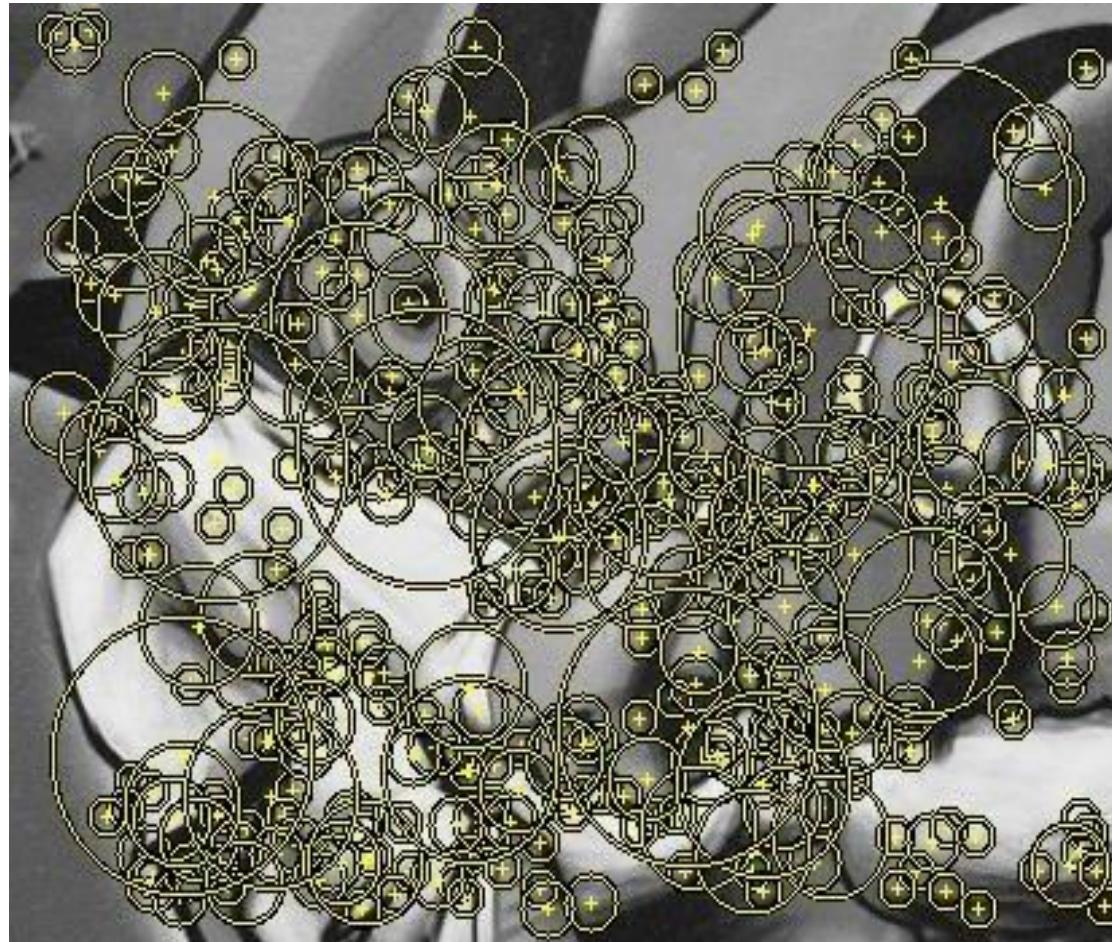


$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow$$



⇒ List of  
( $x, y, s$ )

# Results: Difference-of-Gaussian



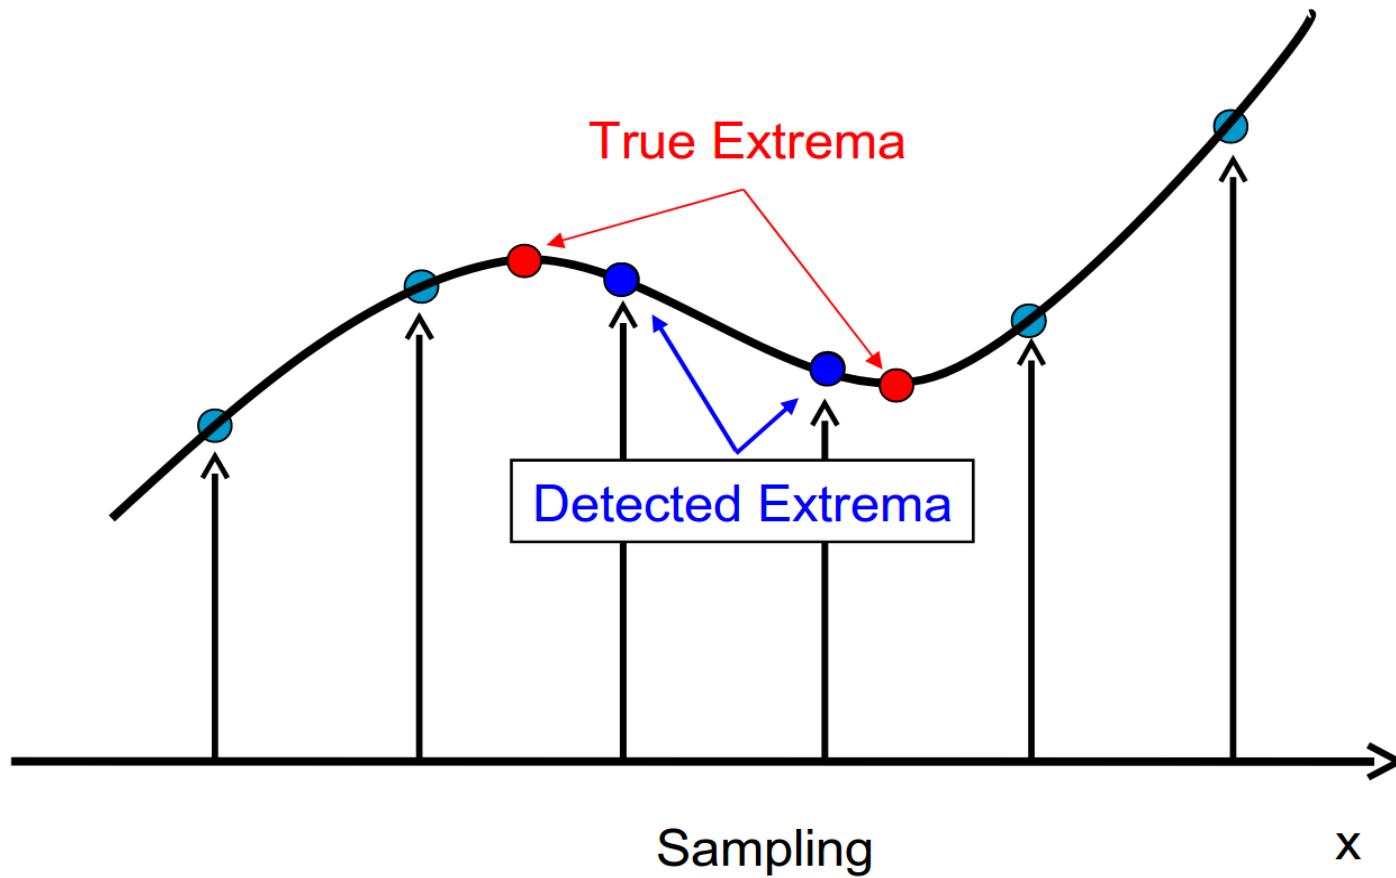
K. Grauman, B. Leibe

# SIFT Algorithm Stages

- Step 1: Scale-space Extrema Detection - Detect interest points (invariant to scale and orientation) using DOG.
- Step 2: Keypoint Localization - Determine location and scale at each candidate location, and select them based on stability.
- Step 3: Orientation Estimation - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- Step 4: Keypoint Descriptor - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion.

# SIFT: Keypoint Localization

- The problem:



# SIFT: Keypoint Localization

- Use Taylor series expansion of the scale-space function:

$$D(\mathbf{s}_0 + \Delta\mathbf{s}) = D(\mathbf{s}_0) + \frac{\partial D}{\partial \mathbf{s}} \Big|_{\mathbf{s}_0} \Delta\mathbf{s} + \frac{1}{2} \Delta\mathbf{s}^T \frac{\partial^2 D}{\partial \mathbf{s}^2} \Big|_{\mathbf{s}_0} \Delta\mathbf{s}$$

where  $\mathbf{s}_0 = (x_0, y_0, \sigma_0)^T$  and  $\Delta\mathbf{s} = (\delta x, \delta y, \delta \sigma)^T$

- The location of the extremum,  $\hat{\mathbf{s}}$ , is determined by taking the derivative of this function with respect to  $\mathbf{s}$  and setting it to zero:

$$\hat{\mathbf{s}} = - \left( \frac{\partial^2 D}{\partial \mathbf{s}^2} \Big|_{\mathbf{s}_0} \right)^{-1} \frac{\partial D}{\partial \mathbf{s}} \Big|_{\mathbf{s}_0}$$

- Next, reject low contrast points and points that lie on the edges
- **Low contrast points elimination:** Reject keypoint if  $D(\hat{\mathbf{s}})$  is smaller than 0.03 (assuming image values are normalized in  $[0,1]$ )

# SIFT: Keypoint Localization

- Reject key points with strong edge response in one direction only
- Edge elimination: Similar to Harris corner detector. SIFT instead uses **Hessian**

- Compute Hessian of D (principal curvature)

$$H = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$$

$\alpha$  : largest eigenvalue( $\lambda_{max}$ )

$\beta$  : smallest eigenvalue( $\lambda_{min}$ )

$$Tr(H) = D_{xx} + D_{yy} = \alpha + \beta$$

$$Det(H) = D_{xx}D_{yy} - D_{xy}^2 = \alpha\beta$$

- Evaluate ratio

$$\frac{Tr(H)^2}{Det(H)} = \frac{(\alpha + \beta)^2}{\alpha\beta} = \frac{(r\beta + \beta)^2}{r\beta^2}$$

$$\frac{Tr(H)^2}{Det(H)} = \frac{(r + 1)^2}{r} \text{ where } r = \frac{\alpha}{\beta}$$

- The ratio is minimum when  $r = 1$  (eigen values are equal)
- Reject keypoint if the ratio is  $>$  a threshold ( $r = 10$  for SIFT)

# SIFT Algorithm Stages

- Step 1: Scale-space Extrema Detection - Detect interest points (invariant to scale and orientation) using DOG.
- Step 2: Keypoint Localization - Determine location and scale at each candidate location, and select them based on stability.
- Step 3: Orientation Estimation - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- Step 4: Keypoint Descriptor - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion.

# SIFT: Orientation Estimation

- Why? To achieve **rotation invariance**
- Use scale of point to choose correct image:

$$\hat{I}(x, y) = G(x, y, \sigma) * I(x, y)$$

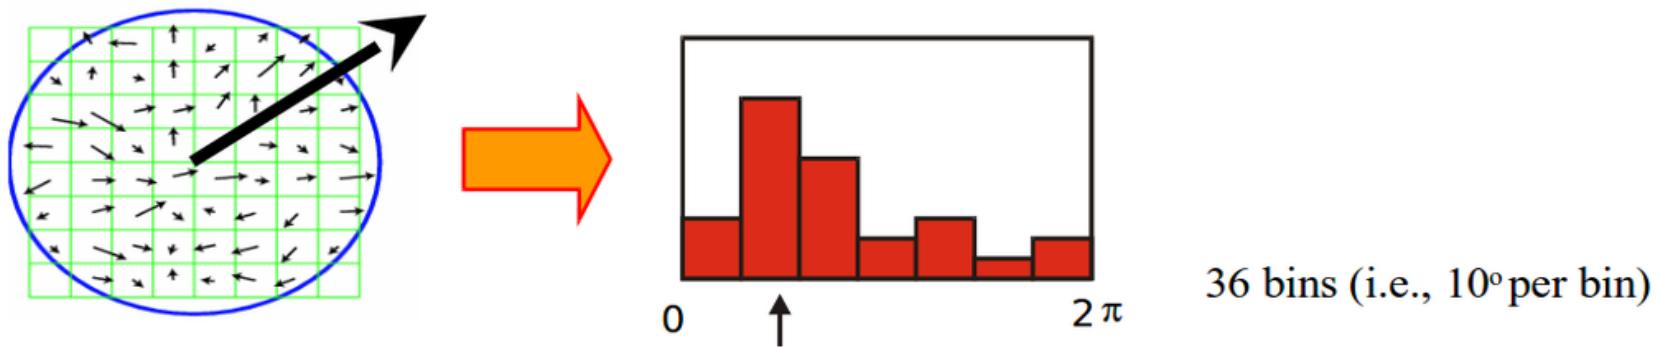
- Compute gradient magnitude and orientation using finite differences:

$$m(x, y) = \sqrt{(\hat{I}(x+1, y) - \hat{I}(x-1, y))^2 + (\hat{I}(x, y+1) - \hat{I}(x, y-1))^2}$$

$$\theta(x, y) = \tan^{-1} \left( \frac{(\hat{I}(x, y+1) - \hat{I}(x, y-1))}{(\hat{I}(x+1, y) - \hat{I}(x-1, y))} \right)$$

# SIFT: Orientation Estimation

- Create histogram of gradient directions, within a region around the keypoint, at selected scale:



- Histogram entries are weighted by:
  - gradient magnitude, and
  - a Gaussian function with equal to 1.5 times scale of the keypoint
- Select **the peak as direction of keypoint**
- Introduce additional key points at same location if another peak is within 80% of max peak of histogram with different direction

# SIFT:



*From 233x189 original image to 832 DoG Extrema*

# SIFT:



*From 832 DoG Extrema to 729 keypoints after low contrast threshold*

# SIFT:

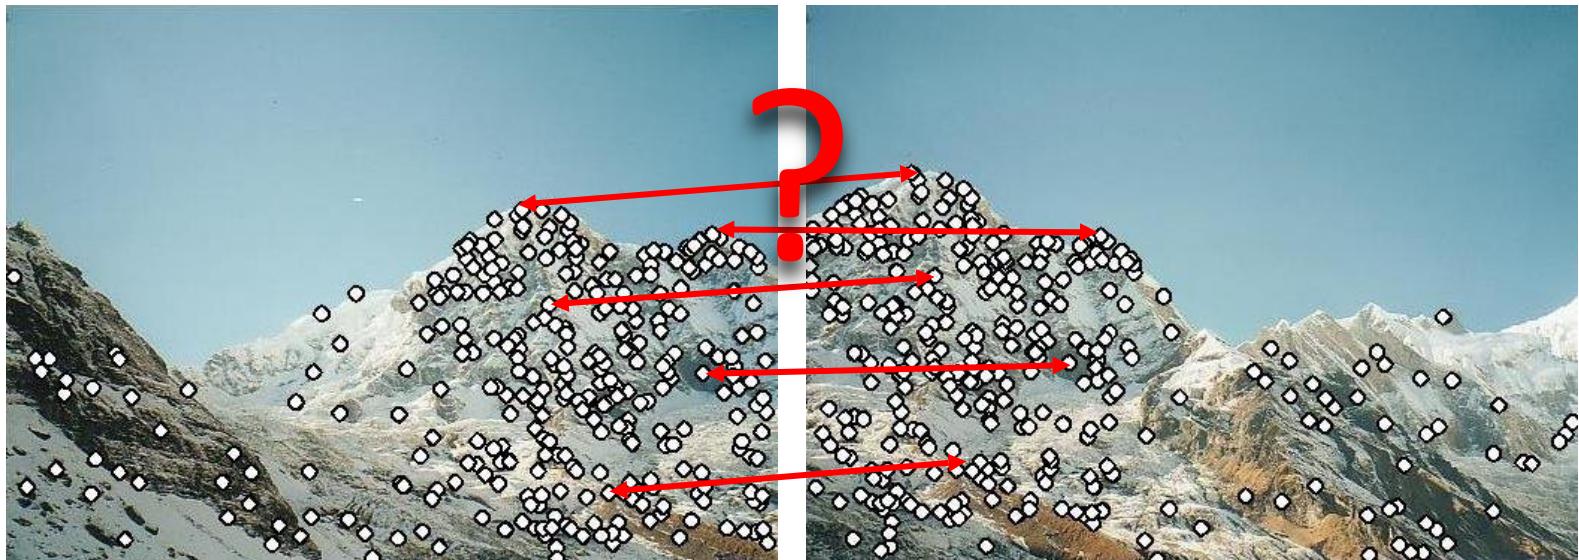


*From 729 keypoints to 536 keypoints after testing ratio based on Hessian*

# Feature descriptors

We know how to detect good points

Next question: **How to match them?**

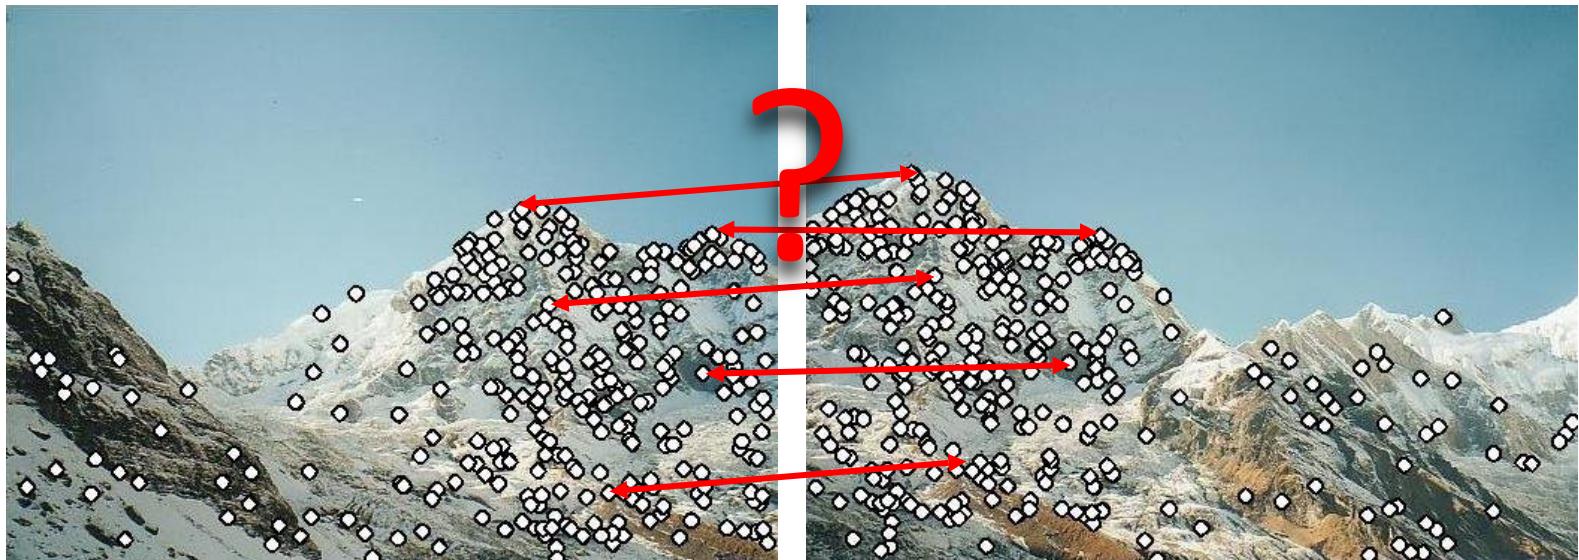


**Answer:** Come up with a *descriptor* for each point,  
find similar descriptors between the two images

# Feature descriptors

We know how to detect good points

Next question: **How to match them?**



Lots of possibilities (this is a popular research area)

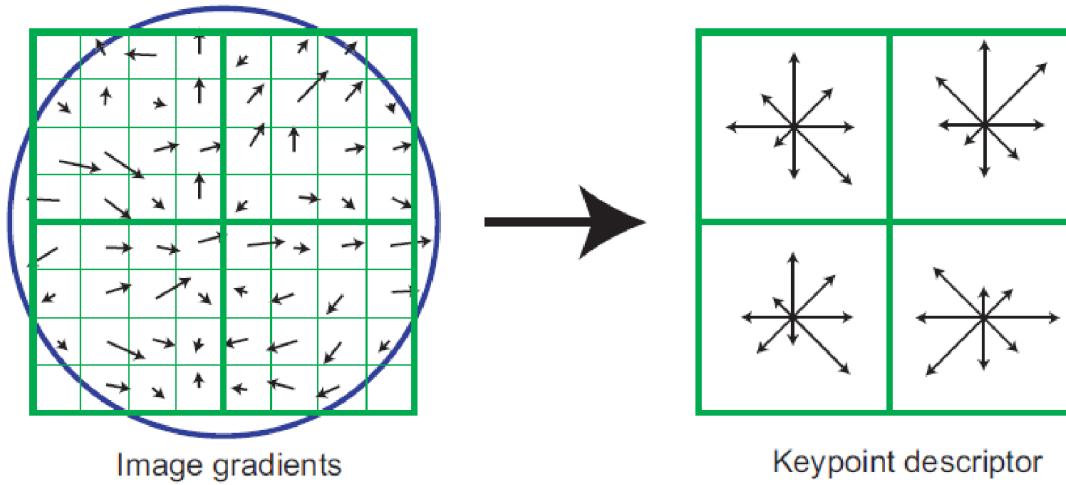
- Simple option: match square windows around the point using descriptors

# SIFT Algorithm Stages

- Step 1: Scale-space Extrema Detection - Detect interest points (invariant to scale and orientation) using DOG.
- Step 2: Keypoint Localization - Determine location and scale at each candidate location, and select them based on stability.
- Step 3: Orientation Estimation - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- Step 4: Keypoint Descriptor - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion.

# SIFT: Keypoint Descriptor

- Around the detected keypoint, compute gradient at each pixel in a 16x16 window, using the appropriate level of the Gaussian pyramid at which the keypoint was detected.



- Downweight gradients by a Gaussian fall-off function (blue circle) to reduce the influence of gradients far from the center.
- In each 4x4 quadrant, compute a gradient orientation histogram using 8 orientation histogram bins.

# SIFT: Keypoint Descriptor

- The resulting 128 non-negative values form a raw version of the SIFT descriptor vector.
- To reduce the effects of contrast or gain (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length.
- To further make the descriptor robust to other **photometric variations**, values are clipped to 0.2 and the resulting vector is once again renormalized to unit length.

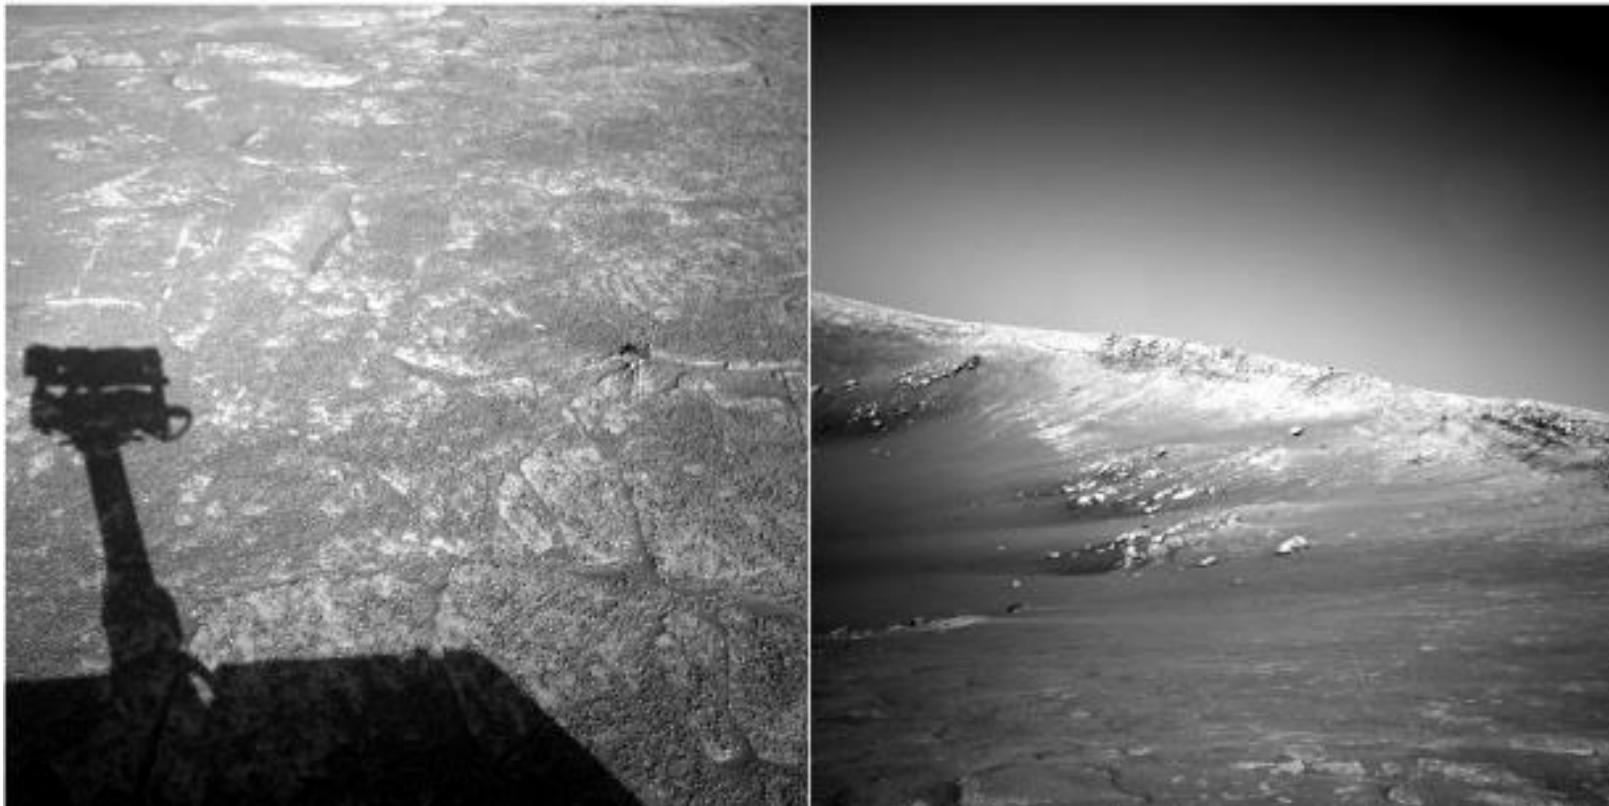
# SIFT:

- Extraordinarily robust feature detection
- Changes in viewpoint: up to about 60 degree out of plane rotation
- Changes in illumination: sometimes even day vs night (below)
- Fast and efficient → can run in real-time



*Credit: Raquel Urtasun, Szeliski*

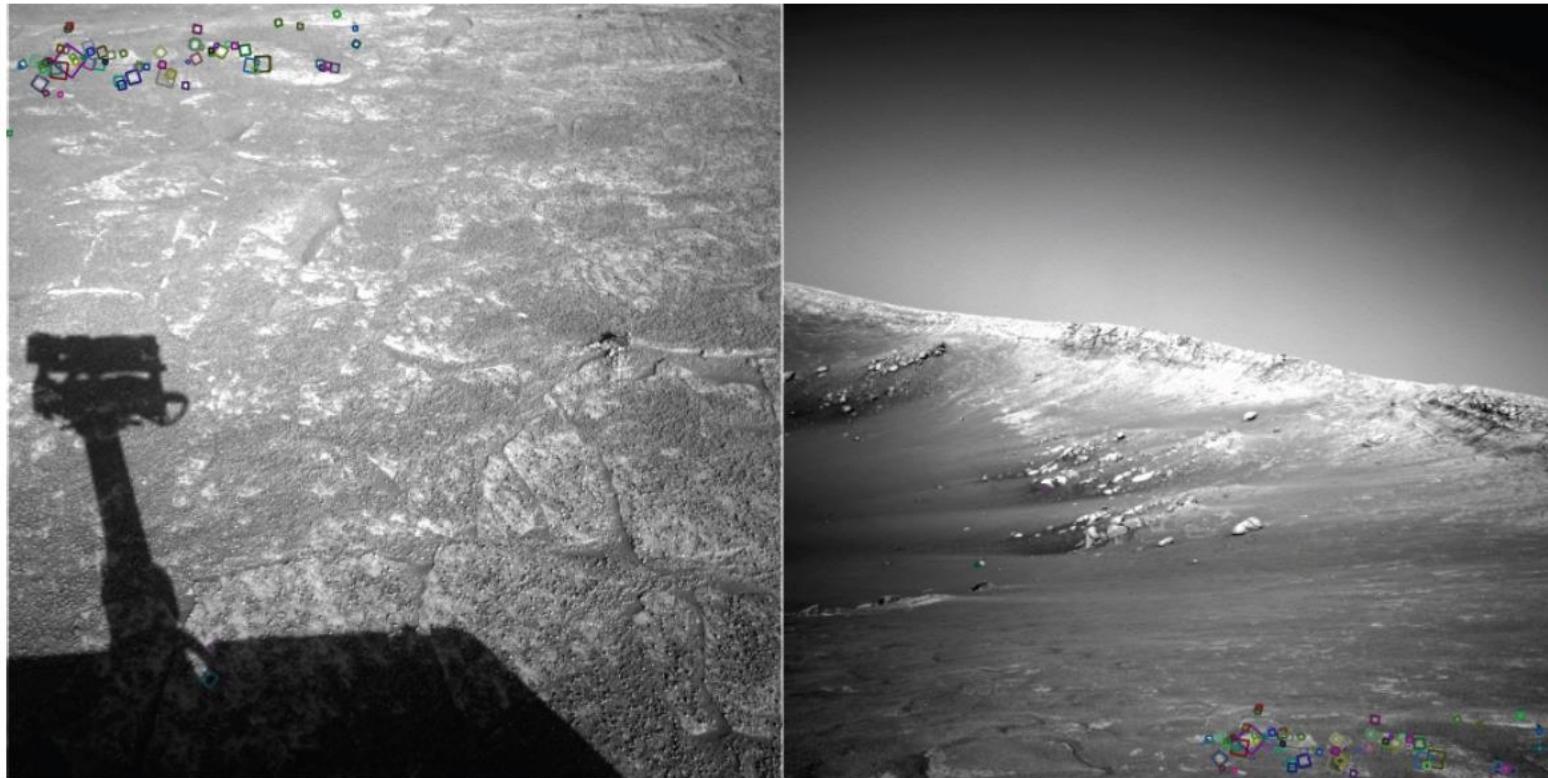
# SIFT: Example



*Mars Rover images*

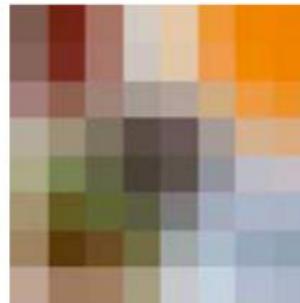
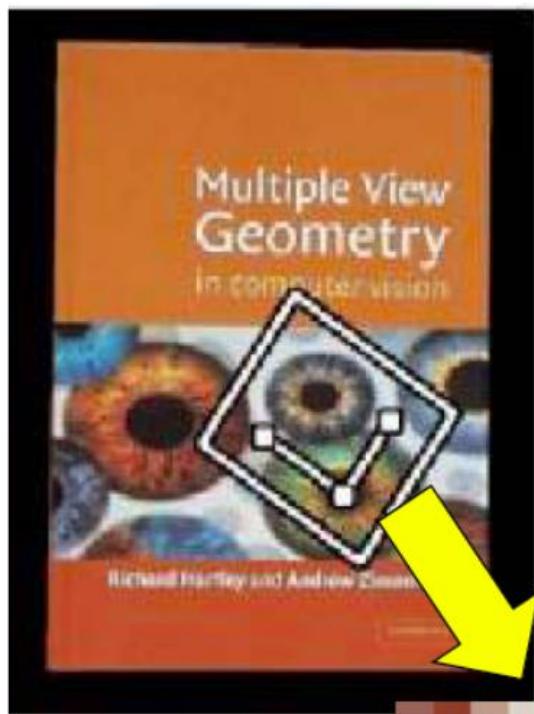
# SIFT: Example

*Maybe, look for tiny squares. . . ?*



*Mars Rover images with SIFT features*

# SIFT: Invariances(Geometric Transformations)



e.g. scale,  
translation,  
rotation

# SIFT: Invariances(Photometric Transformations)

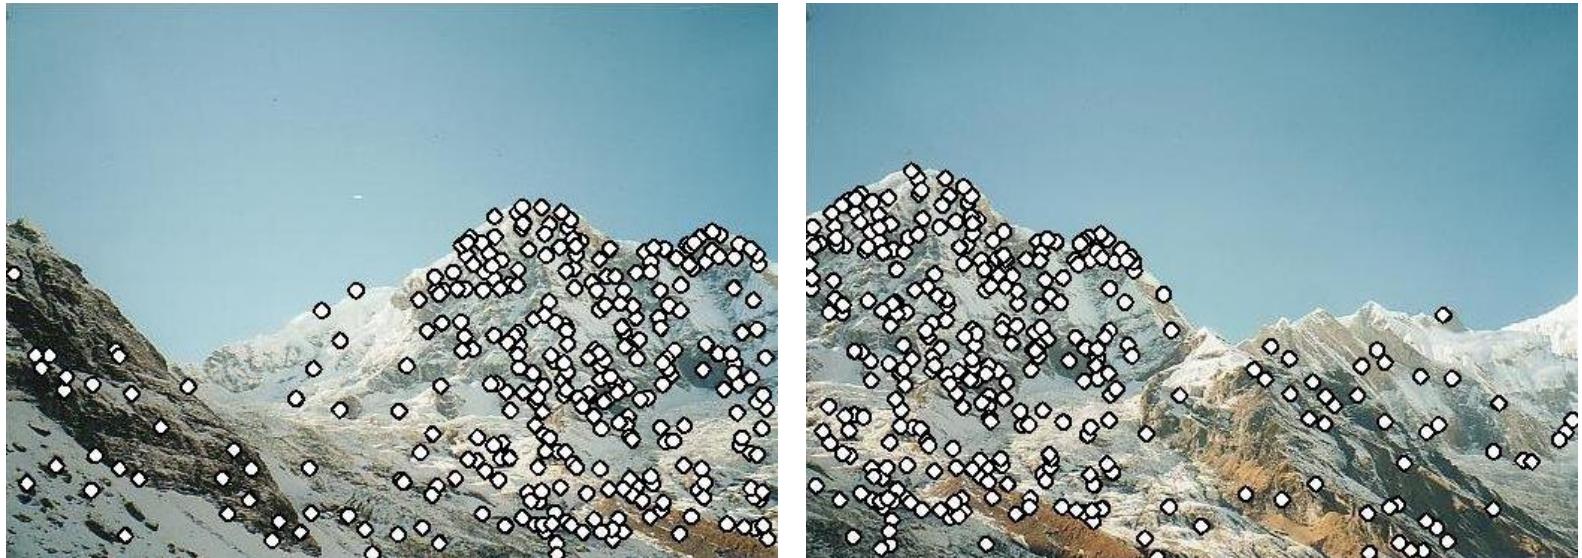


Credit: Raquel Urtasun, Szeliski

# SIFT application: Image Stitching

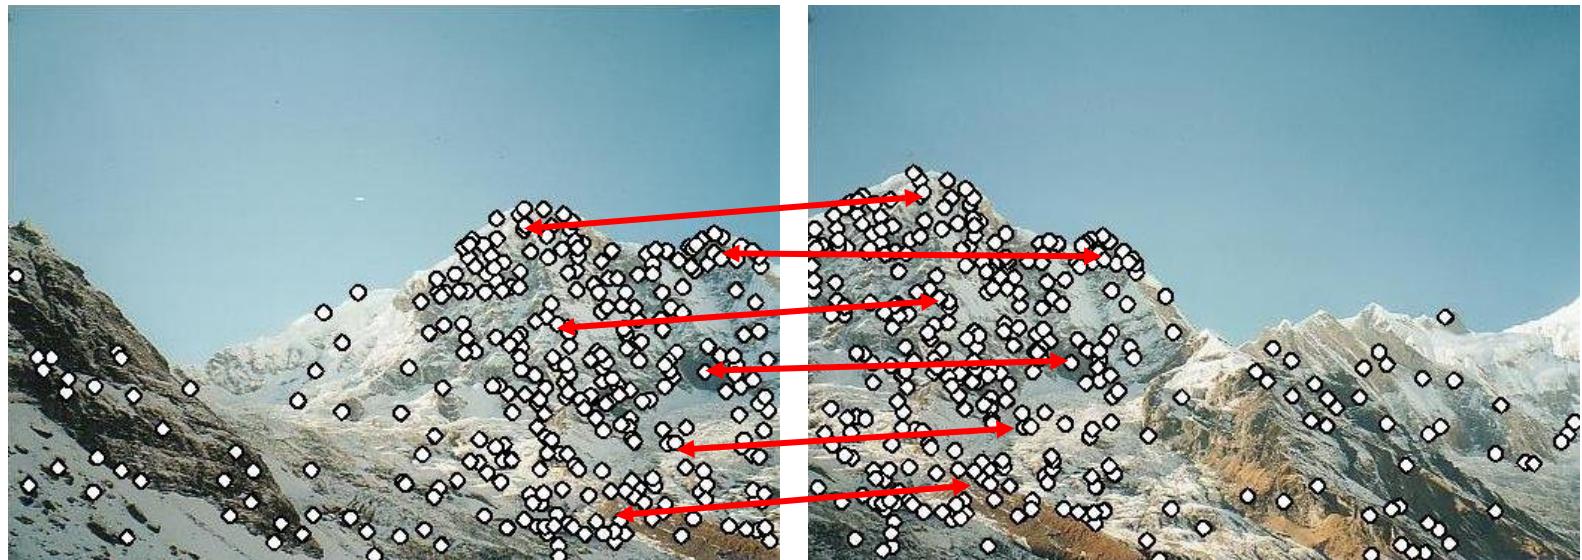


# SIFT application: Image Stitching



Step 1: extract features, describe features

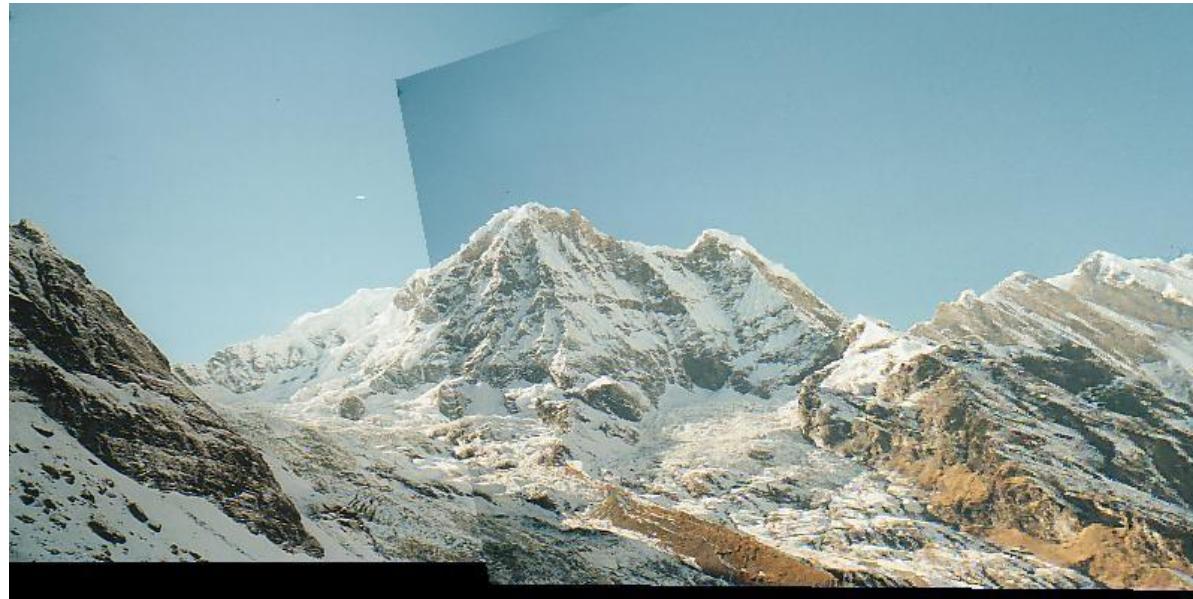
# SIFT application: Image Stitching



Step 1: extract features, describe features

Step 2: match features

# SIFT application: Image Stitching

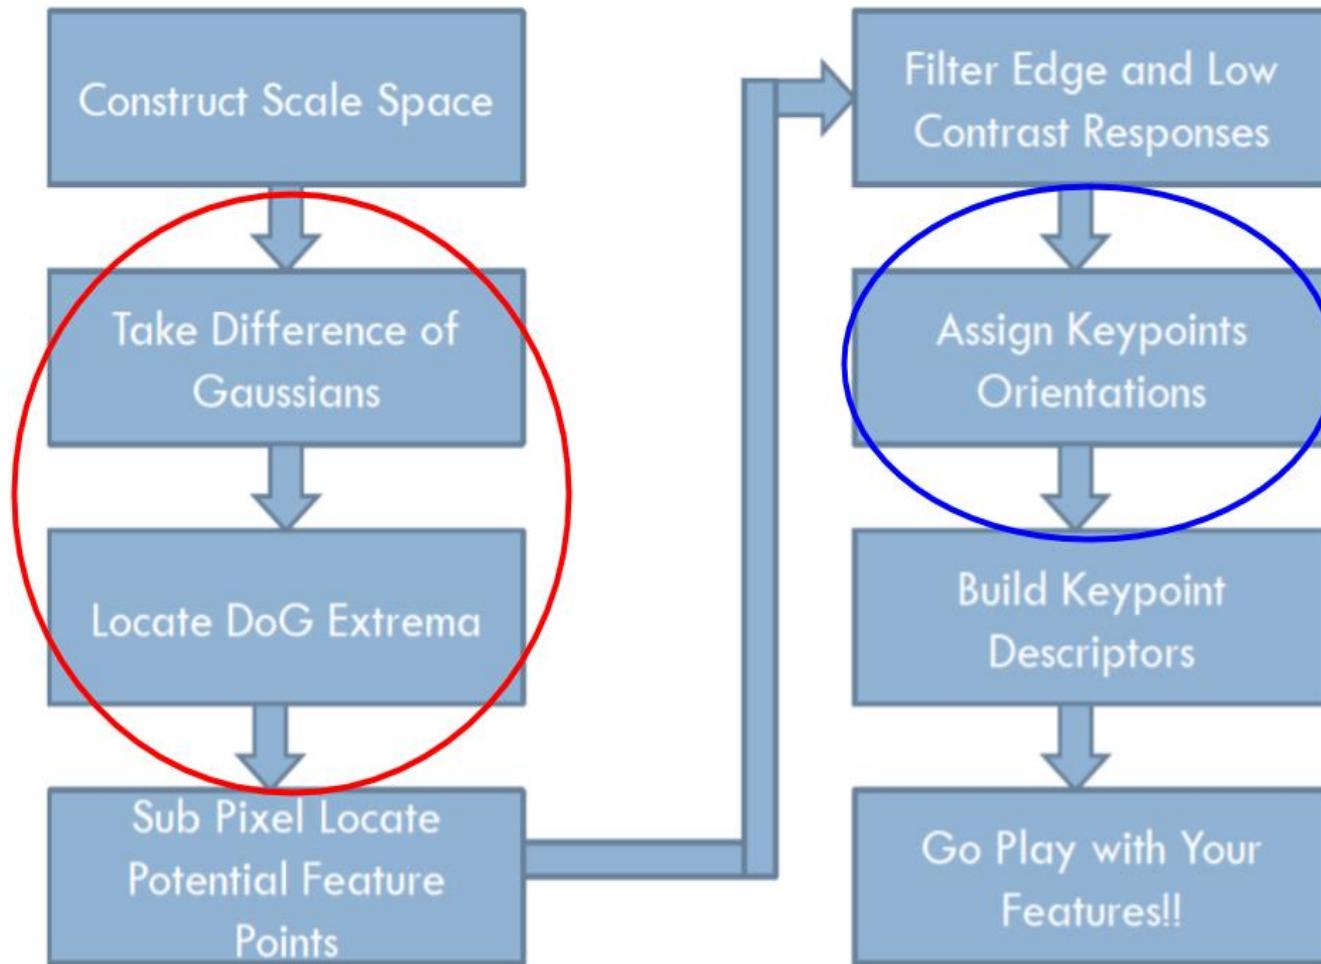


Step 1: extract features, describe features

Step 2: match features

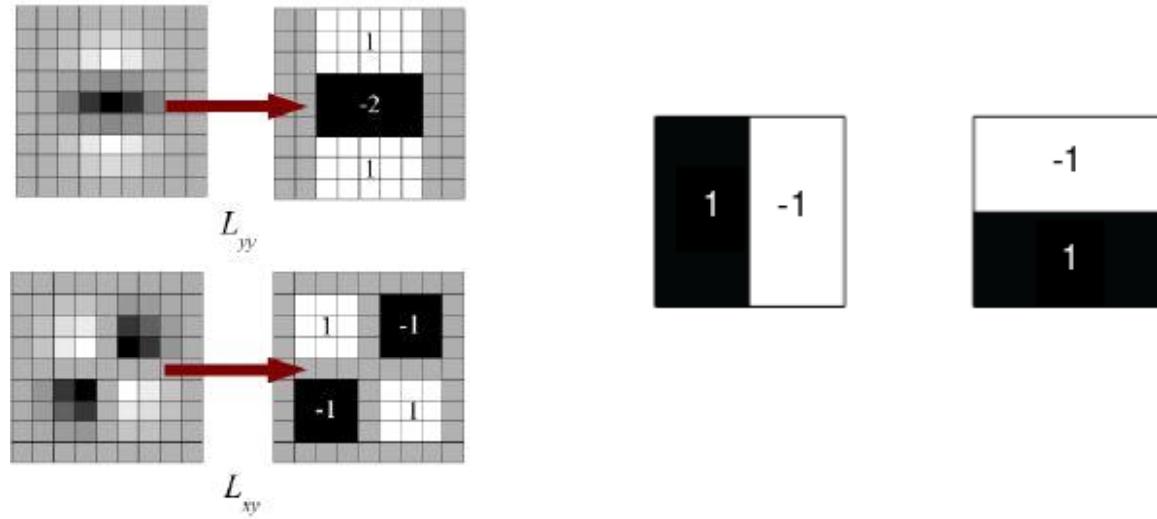
Step 3: align/fitting images

# SIFT



# *SURF: Speeded Up Robust Features* *(ECCV 2006, 26k+ citations)*

- Uses box filters instead of Gaussians to approximate Laplacians
- Uses Haar wavelets to get keypoint orientations.
- Haar wavelets are simple filters which can be used to find gradients in the x and y directions
- SURF is good at handling blur and rotation variations
- SURF is not as good as SIFT on invariance to illumination & viewpoint changes
- SURF is nearly 3 times faster than SIFT

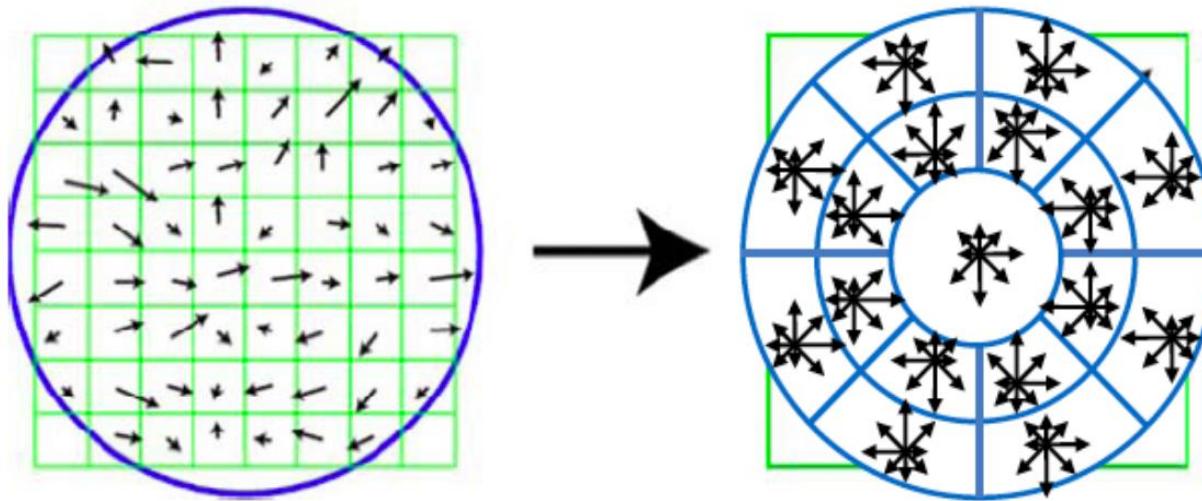


For more information:

<https://medium.com/data-breach/introduction-to-surf-speeded-up-robust-features-c7396d6e7c4e>  
<http://www.vision.ee.ethz.ch/~surf/>

# Gradient Location-Orientation Histogram (GLOH)

- Variant of SIFT that uses a log-polar binning structure instead of four quadrants.
- Uses 17 spatial bins and 16 orientation bins.
- The 272D histogram is then projected onto a 128D descriptor using PCA trained on a large dataset.



Credit: Matthew Brown, Kristen Grauman, Raquel Urtasun