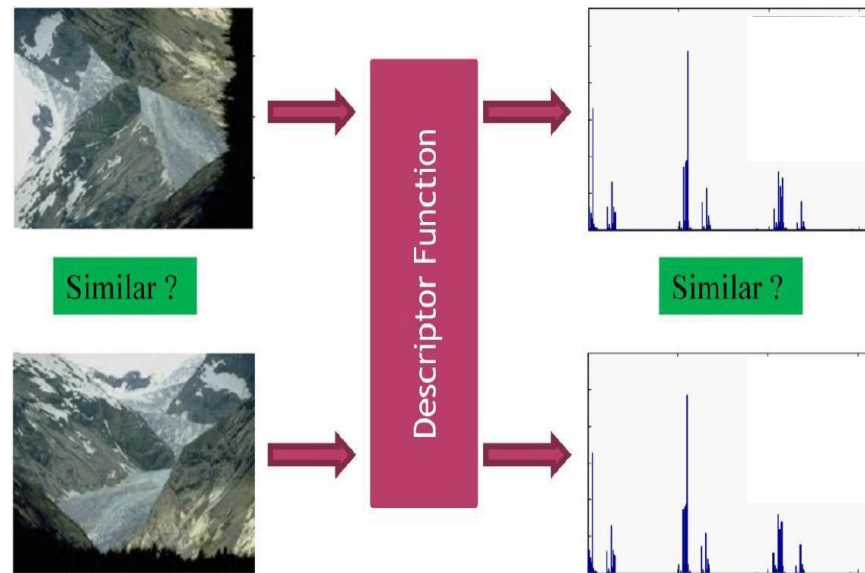


Feature Matching





Computer Vision

Adduru U G Sankararao, IIIT Sri City


Previous Class

- Local Descriptors:
 - Discriminative
 - Robust
 - Compact
 - Efficient
- Scale Invariant Feature Transform
 - Utilizes gradient information
 - Gradient direction binning
- Next:
 - Feature Matching
 - Evaluation Metrics



Image/Region Matching

- Automatically recognize whether two images/regions contain the similar content.
- Comparing the image pixels as they are, will not work.



Image/Region Matching

- Pixel-based distances on high-dimensional data (and images especially) can be very unintuitive.

original

shifted

messed up

darkened



Challenges

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter

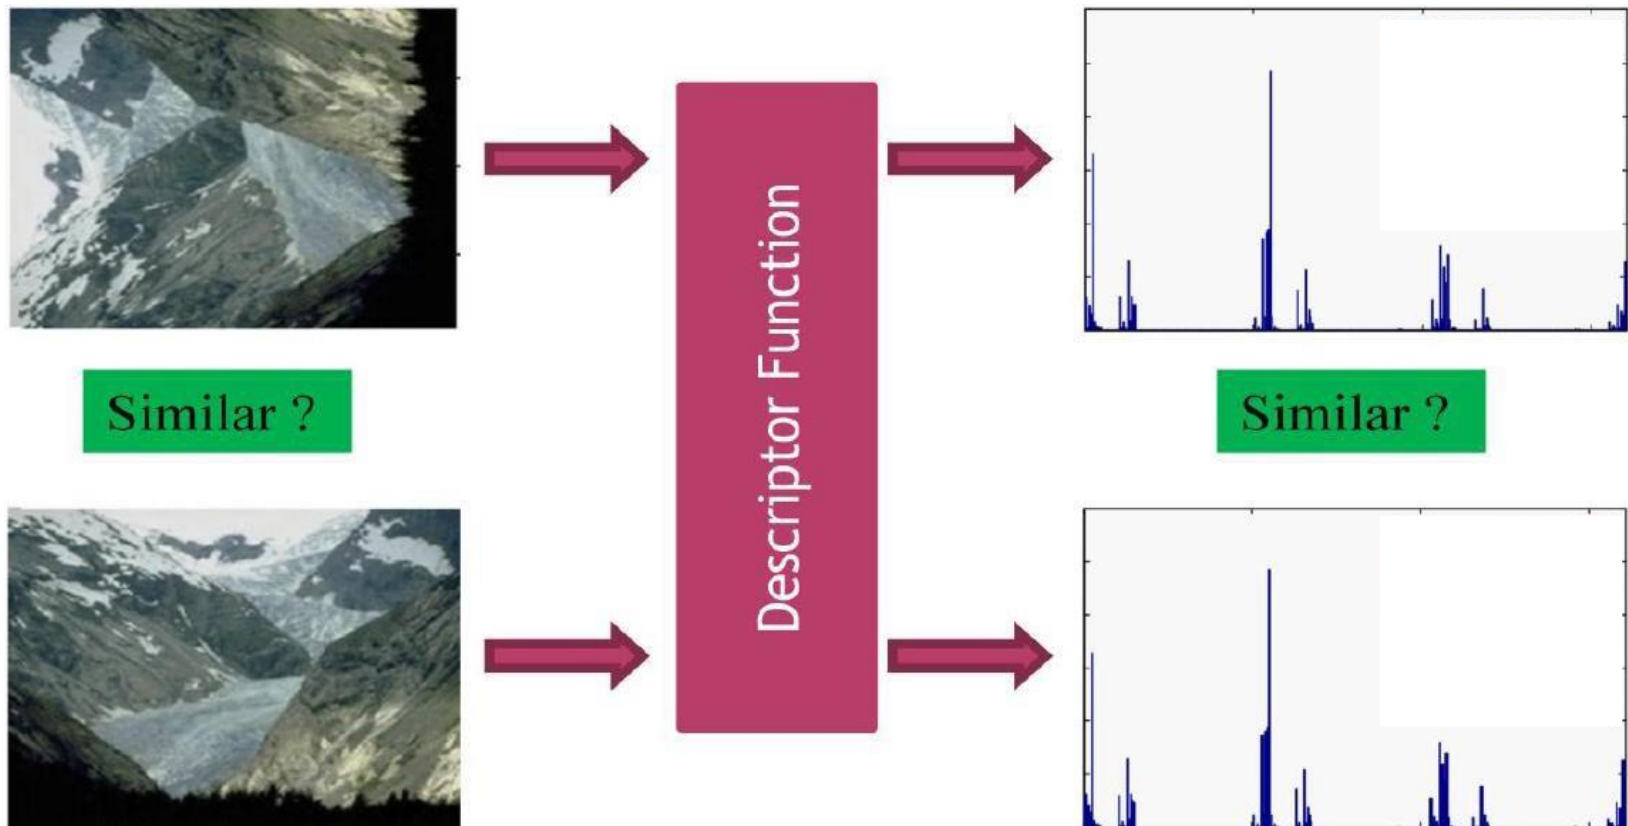


Intra-class variation



Solution

- Descriptors allow certain differences between the images.



Comparing using descriptor function
(images are taken from Corel-database and RSHD descriptor is used)

Feature matching

Given a feature in I_1 , how to find the best match in I_2 ?

1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min distance

How to define the difference between two features f_1, f_2 ?


- Simple approach: L_2 distance, $||f_1 - f_2||$

Techniques:


- Brute Force matching (L_2 , `cv::BFMatcher` in OpenCV)
- Nearest Neighbour matching
- FLANN-based matching

Brute Force Matcher (BFMatcher)

- Compares each descriptor in one image with all descriptors in the other image.
- Finds the closest match based on a distance metric (eg: L2 norm (Euclidean))
- Simple, easy to use, but Slow when there are lots of features.
- Can give good scores to ambiguous (incorrect) matches




I_1




I_2

Nearest Neighbours based matching

- Instead of just taking the single closest match, it finds k best matches for each descriptor, (Usually $k = 2$).
- Then apply Lowe's ratio test to keep only reliable matches:
- Better approach: **ratio distance** = $||f_1 - f_2|| / ||f_1 - f_2'||||$
 - f_2 is best SSD match to f_1 in I_2
 - f_2' is 2nd best SSD match to f_1 in I_2
 - Keep a match if the ratio is less than a threshold
 - gives large values for ambiguous matches, helps remove ambiguous matches



I_1




I_2

Feature matching example




Feature matching example




51 matches

Feature matching example




Feature matching example




58 matches

Feature matching:



Correct matches: Inliers



Incorrect matches: Outliers

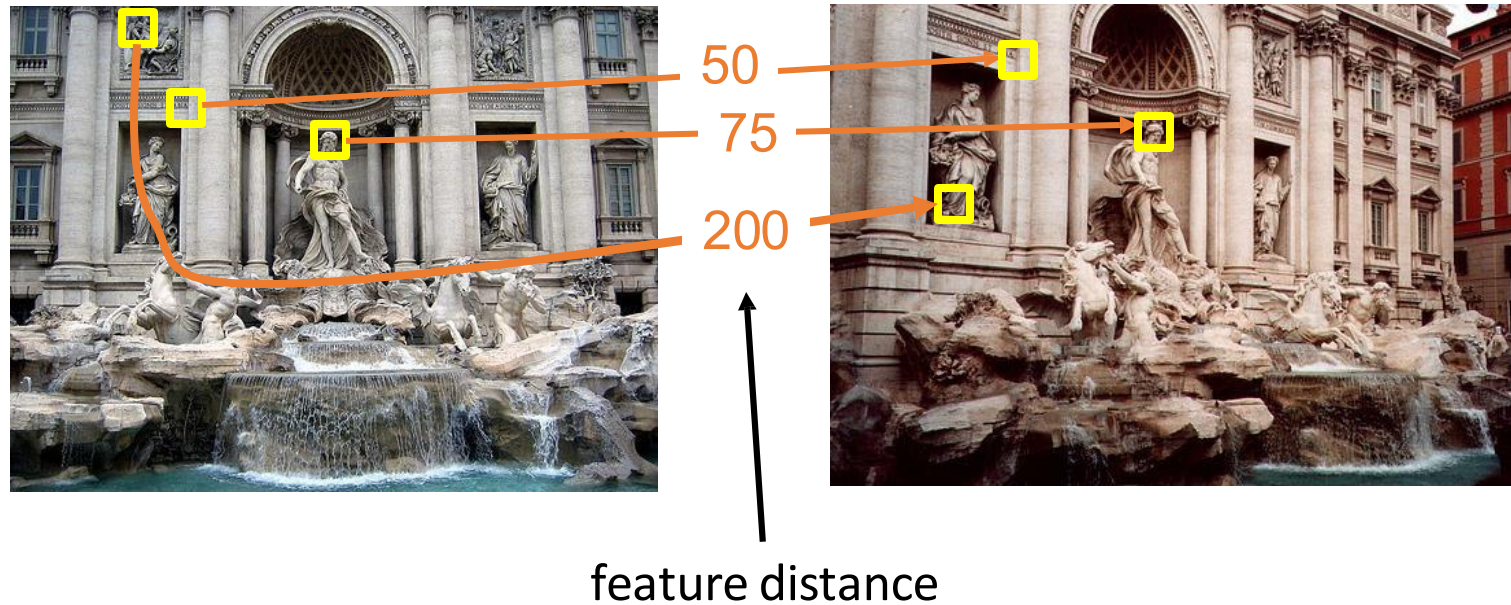
Need to
remove
outliers

Evaluating the results

How can we measure the performance of a feature matcher?

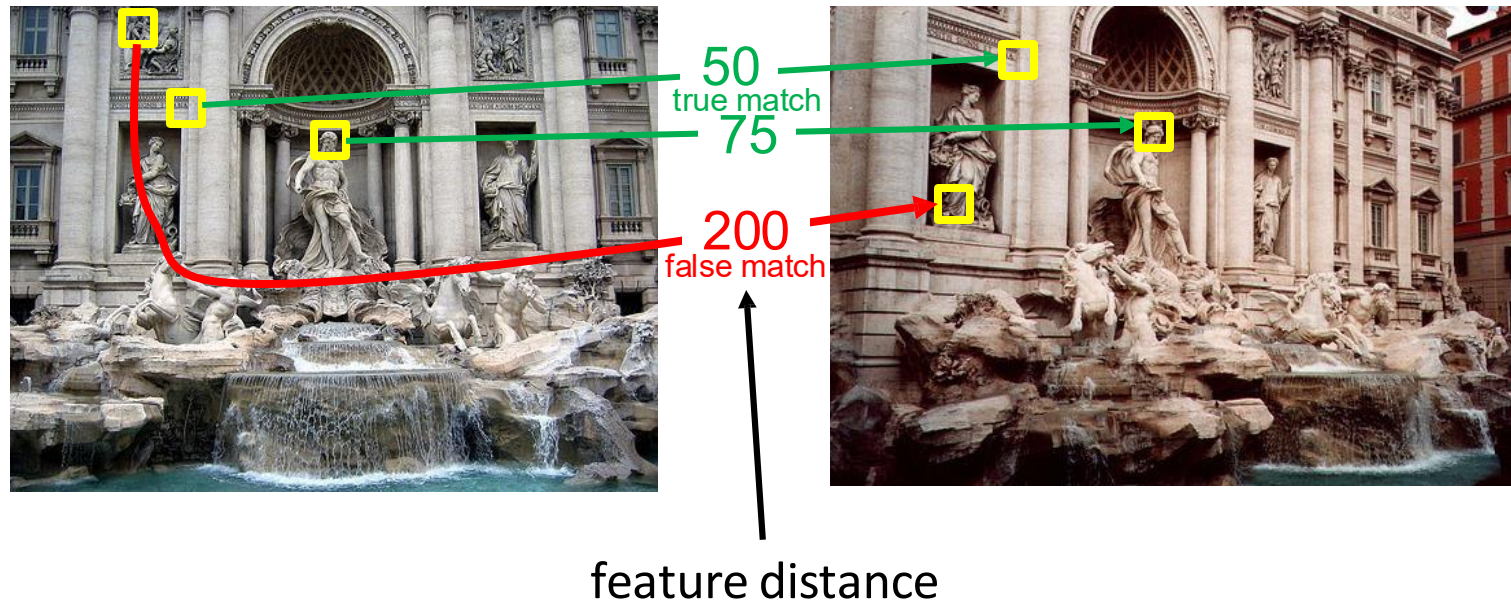
Evaluating the results

How can we measure the performance of a feature matcher?



True/false positives

How can we measure the performance of a feature matcher?



The **distance threshold** affects performance

- True positives = # of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

Evaluation Metrics

	True matches	True non-matches	
Predicted matches	TP = 18	FP = 4	P' = 22
Predicted non-matches	FN = 2	TN = 76	N' = 78
	P = 20	N = 80	Total = 100

PPV = 0.82

TPR = 0.90	FPR = 0.05
------------	------------

ACC = 0.94

- true positive rate (TPR),

$$\text{TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}} = \frac{\text{TP}}{P};$$

PPV = Precision

- false positive rate (FPR),

$$\text{FPR} = \frac{\text{FP}}{\text{FP} + \text{TN}} = \frac{\text{FP}}{N};$$

TPR. = Recall

- positive predictive value (PPV),

$$\text{PPV} = \frac{\text{TP}}{\text{TP} + \text{FP}} = \frac{\text{TP}}{P'};$$

F1-Score = HM (Pre, Rec)


- accuracy (ACC),

$$\text{ACC} = \frac{\text{TP} + \text{TN}}{P + N}.$$

Evaluating the results

How can we measure the performance of a feature matcher?

$$\frac{\text{\# true positives}}{\text{\# correctly matched features (positives)}} \\ \text{“recall”}$$



$$\frac{\text{\# false positives}}{\text{\# incorrectly matched features (negatives)}} \\ 1 - \text{“precision”}$$


Evaluating the results

How can we measure the performance of a feature matcher?

Aera under ROC



AUC more → better

$$\frac{\text{\# true positives}}{\text{\# correctly matched features (positives)}} \\ \text{"recall"}$$



$$\frac{\text{\# false positives}}{\text{\# incorrectly matched features (negatives)}} \\ 1 - \text{"precision"}$$

Evaluating Results



- As the threshold θ is increased, the number of true positives (TP) and false positives (FP) increases.

Image matching

dense registration*

[Lucas and Kanade 1981]



- for each location in an image, find a displacement with respect to another reference image
- appropriate for small displacements, *e.g.* stereopsis or optical flow

dense registration*

[Lucas and Kanade 1981]



- for each location in an image, find a displacement with respect to another reference image
- appropriate for small displacements, *e.g.* stereopsis or optical flow

dense registration*

[Lucas and Kanade 1981]



- for each location in an image, find a displacement with respect to another reference image
- appropriate for small displacements, *e.g.* stereopsis or optical flow


dense registration*

[Lucas and Kanade 1981]



- for each location in an image, find a displacement with respect to another reference image
- appropriate for small displacements, *e.g.* stereopsis or optical flow


one dimension*



- assuming $g(x) = f(x + t)$ and t is small,

$$\frac{df}{dx}(x) \approx \frac{f(x + t) - f(x)}{t} = \frac{g(x) - f(x)}{t}$$


one dimension*



- assuming $g(x) = f(x + t)$ and t is small,

$$\frac{df}{dx}(x) \approx \frac{f(x + t) - f(x)}{t} = \frac{g(x) - f(x)}{t}$$

one dimension*



- assuming $g(x) = f(x + t)$ and t is small,

$$\frac{df}{dx}(x) \approx \frac{f(x + t) - f(x)}{t} = \frac{g(x) - f(x)}{t}$$

two dimensions: least squares*

- again, assume an image patch defined by window w ; what is the error between the patch shifted by \mathbf{t} in reference image f and a patch at the origin in shifted image g ?

$$\begin{aligned} E(\mathbf{t}) &= \sum_{\mathbf{x}} w(\mathbf{x})(f(\mathbf{x} + \mathbf{t}) - g(\mathbf{x}))^2 \\ &\approx \sum_{\mathbf{x}} w(\mathbf{x})(f(\mathbf{x}) + \mathbf{t}^\top \nabla f(\mathbf{x}) - g(\mathbf{x}))^2 \end{aligned}$$

- error minimized when gradient vanishes

$$\mathbf{0} = \frac{\partial E}{\partial \mathbf{t}} = \sum_{\mathbf{x}} w(\mathbf{x}) 2 \nabla f(\mathbf{x}) (f(\mathbf{x}) + \mathbf{t}^\top \nabla f(\mathbf{x}) - g(\mathbf{x}))$$


- least-squares solution

$$\left(w * (\nabla f)(\nabla f)^\top \right) \mathbf{t} = w * ((\nabla f)(g - f))$$

the aperture problem*



the aperture problem*



wide-baseline matching


- in dense registration, we started from a local “template matching” process and found an efficient solution based on a Taylor approximation
- both make sense for small displacements
- in wide-baseline matching, every part of one image may appear anywhere in the other
- we start by **pairwise matching of local descriptors** without any order and then attempt to enforce some **geometric consistency according to a rigid motion model**

wide-baseline matching




- a region in one image may appear anywhere in the other

wide-baseline matching



- features detected independently in each image

wide-baseline matching



- tentative correspondences by pairwise descriptor matching

wide-baseline matching



- subset of correspondences that are 'inlier' to a rigid transformation

descriptor extraction

for each detected feature in each image

- construct a local histogram of gradient orientations
- find one or more dominant orientations corresponding to peaks in the histogram
- resample local patch at given location, scale, affine shape and orientation
- extract one descriptor for each dominant orientation

descriptor matching




descriptor matching



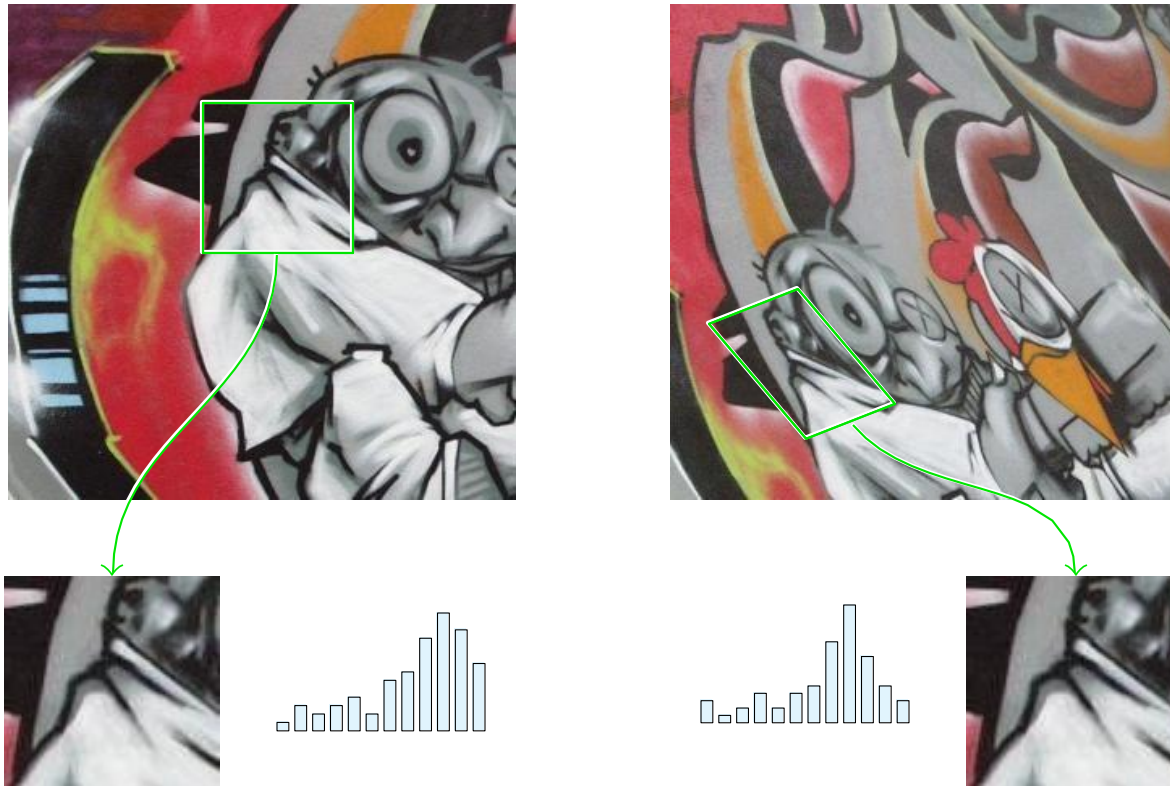
- detect features

descriptor matching



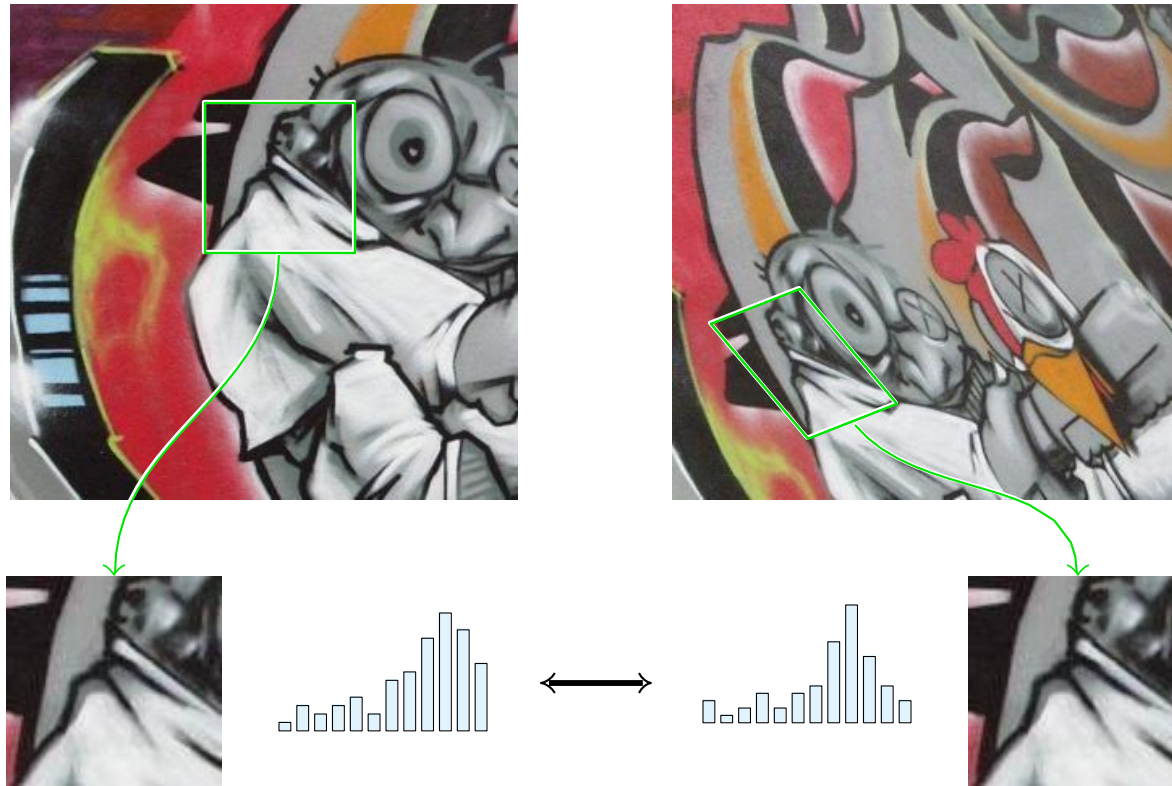
- detect features - find dominant orientation, resample patches

descriptor matching



- detect features - find dominant orientation, resample patches - extract descriptors

descriptor matching




- detect features - find dominant orientation, resample patches - extract descriptors - match pairwise

descriptor matching

- for each descriptor in one image, find its two nearest neighbors in the other
- if ratio of distance of first to distance of second is small, make a correspondence
- this yields a list of **tentative** correspondences

ratio test



- ratio of first to second nearest neighbor distance can determine the probability of a true correspondence