


Data Mining: Your Journey from Raw Data to Actionable Insight

A comprehensive course in extracting knowledge, identifying patterns, and making predictions from complex datasets.

The Challenge and The Opportunity

We operate in a world of unprecedented data volume. The key to competitive advantage lies not just in collecting this data, but in extracting meaningful knowledge from it. This course provides the techniques to transform massive datasets into strategic assets.



Extract Knowledge: Learn both predictive and descriptive methods to uncover hidden insights.



Build Models: Master techniques for classification, clustering, and association rule mining.



Evaluate Performance: Develop the skills to quantify the effectiveness and accuracy of your models.

A Structured Path to Expertise

This course is structured as a three-part journey, designed to build your skills logically from the ground up—from fundamental principles to advanced applications.



1. THE FOUNDATION

Mastering the art of data preparation.



2. THE CORE TOOLKIT

Building your arsenal of essential mining techniques.



3. THE FRONTIER

Applying your knowledge to modern, complex data challenges.


The Foundation: Laying the Groundwork with Data Preprocessing

High-quality insights begin with high-quality data. In this foundational module, you will learn the essential techniques to clean, structure, and refine raw data, ensuring it's ready for powerful analysis.



Data Summarization

Measuring central tendency and dispersion.




Data Cleaning

Handling missing values and noisy data.



Data Integration & Transformation

Combining and standardizing diverse data sources.




Data Reduction

Techniques for simplifying datasets without losing critical information (e.g., dimensionality and numerosity reduction).

The Core Toolkit: Mastering Essential Mining Techniques

With a solid foundation in place, we move to the core of data mining. You will master a powerful toolkit of algorithms for discovering patterns, predicting outcomes, and organizing data.



Toolkit Part 1: Uncovering Hidden Relationships

Association Rule & Correlation Analysis

What items are frequently purchased together? What events are correlated? This module teaches you how to discover these valuable relationships in large transactional datasets.

Key Concepts Covered

- Frequent Itemset Mining: The core of finding patterns.
- The Apriori Algorithm: A classic, efficient method for generating rules.
- From Association to Correlation: Moving beyond simple co-occurrence to understand true statistical relationships.



Toolkit Part 2: Predicting Future Outcomes

Supervised Learning: Classification & Prediction

Learn to build models that predict a specific outcome based on labeled historical data. This is the engine behind everything from spam detection to medical diagnosis.


Methods You Will Master




Probabilistic Models:
Naïve Bayesian
Classification




Neural Networks:
Classification by
Backpropagation



Optimal Margin
Classifiers: Support
Vector Machines
(SVMs)



Intuitive Models:
Decision Tree
Induction



Human-Readable
Models: Rule-Based
Classification
(IF-THEN rules)


Toolkit Part 3: Discovering Natural Structures

Unsupervised Learning: Cluster Analysis

What if you don't have labeled data? Unsupervised learning allows you to explore datasets and identify natural groupings or clusters on your own. This is key for customer segmentation, anomaly detection, and exploratory analysis.

Clustering Methodologies Covered

- Partitioning Methods: k-Means and k-Medoids
- Hierarchical Methods: Agglomerative and Divisive Clustering
- Density-Based Methods: Finding clusters of arbitrary shape
- Grid-Based Methods: Efficiently clustering large datasets



The Frontier: Applying Your Skills to Advanced Challenges

The world of data is not static. This final module takes your foundational and core skills and applies them to the complex, evolving data sources that define modern analytics.




**Mining Data
Streams**



**Analyzing Time-
Series Data**



**Extracting Insight
from Text**



**Understanding the
World Wide Web**

A Closer Look at the Frontier



Data Streams

- Techniques for classifying and clustering dynamic, evolving data streams in real-time.



Time-Series Data

- Methods for trend analysis and similarity search in sequential data.



Text Mining

- Dimensionality reduction for text, information retrieval, and other text mining approaches.



Web Mining

- Analyzing page layout, link structures, web document classification, and web usage patterns.


Your New Capabilities

Upon completing this course, you will be able to:

- ✓ **Expertly Prepare Data:** Master data cleaning, aggregation, and transformation for any analytical task.
- ✓ **Discover Association Rules:** Systematically find valuable patterns in transactional data and quantify their strength.
- ✓ **Build Predictive Models:** Apply and compare probabilistic and discriminative supervised learning techniques to solve real-world problems.
- ✓ **Implement Rule-Based Systems:** Construct and interpret classification models based on intuitive IF-THEN logic.
- ✓ **Deploy Clustering Techniques:** Understand and apply a wide variety of unsupervised learning methods to uncover hidden data structures.
- ✓ **Tackle Advanced Data:** Confidently mine data from streams, time-series, unstructured text, and the web.

Essential Reading & Resources

Core Textbook




Title: Data Mining: Concepts and Techniques, 3rd Edition

Authors: Han J, Kamber M, Pei J.


Description: The foundational text for this course, providing comprehensive coverage of all core concepts.

Recommended Further Reading




Title: Introduction to Data Mining

Authors: Pang-Ning Tan, Michael Steinbach, et al.



Title: Mining of Massive Datasets

Authors: Jure Leskovec, Anand Rajaraman, Jeffrey David Ullman.



This course is more than a collection of algorithms.

It is a structured methodology for turning data into your most valuable strategic asset.